

CONTENTS

1	Math and Machine Learning Basics	6
1.1	Linear Algebra (Quick Review) (Ch. 2)	7
1.1.1	Example: Principal Component Analysis	9
1.2	Probability & Information Theory (Quick Review) (Ch. 3)	11
1.3	Numerical Computation (Ch. 4)	13
1.4	Machine Learning Basics (Ch. 5)	16
1.4.1	Estimators, Bias and Variance (5.4)	16
1.4.2	Maximum Likelihood Estimation (5.5)	18
1.4.3	Bayesian Statistics (5.6)	19
1.4.4	Supervised Learning Algorithms (5.7)	21
2	Deep Networks: Modern Practices	22
2.1	Deep Feedforward Networks (Ch. 6)	23
2.1.1	Back-Propagation (6.5)	24
2.2	Regularization for Deep Learning (Ch. 7)	25
2.3	Optimization for Training Deep Models (Ch. 8)	27
2.4	Convolutional Neural Networks (Ch. 9)	31
2.5	Sequence Modeling (RNNs) (Ch. 10)	34
2.5.1	Review: The Basics of RNNs	34
2.5.2	RNNs as Directed Graphical Models	39
2.5.3	Challenge of Long-Term Deps. (10.7)	41
2.5.4	LSTMs and Other Gated RNNs (10.10)	42
2.6	Applications (Ch. 12)	43
2.6.1	Natural Language Processing (12.4)	43
2.6.2	Neural Language Models (12.4.2)	44
3	Deep Learning Research	45
3.1	Linear Factor Models (Ch. 13)	46

3.2	Autoencoders (Ch. 14)	49
3.3	Representation Learning (Ch. 15)	50
3.4	Structured Probabilistic Models for DL (Ch. 16)	51
3.4.1	Sampling from Graphical Models	53
3.4.2	Inference and Approximate Inference	53
3.5	Monte Carlo Methods (Ch. 17)	55
3.6	Approximate Inference (Ch. 19)	57
4	Papers and Tutorials	59
4.1	WaveNet	61
4.2	Neural Style	65
4.3	Neural Conversation Model	67
4.4	NMT By Jointly Learning to Align & Translate	69
4.4.1	Detailed Model Architecture	70
4.5	Effective Approaches to Attention-Based NMT	72
4.6	Using Large Vocabularies for NMT	74
4.7	Candidate Sampling – TensorFlow	77
4.8	Attention Terminology	79
4.9	TextRank	81
4.9.1	Keyword Extraction	83
4.9.2	Sentence Extraction	84
4.10	Simple Baseline for Sentence Embeddings	85
4.11	Survey of Text Clustering Algorithms	87
4.11.1	Feature Selection and Transformation Methods	87
4.11.2	Distance-based Clustering Algorithms	90
4.11.3	Probabilistic Document Clustering and Topic Models	91
4.11.4	Online Clustering with Text Streams	93
4.11.5	Semi-Supervised Clustering	94
4.12	Deep Sentence Embedding Using LSTMs	95
4.13	Clustering Massive Text Streams	98

4.14	Supervised Universal Sentence Representations (InferSent)	100
4.15	Dist. Rep. of Sentences from Unlabeled Data (FastSent)	101
4.16	Latent Dirichlet Allocation	103
4.17	Conditional Random Fields	106
4.18	Attention Is All You Need	108
4.19	Hierarchical Attention Networks	111
4.20	Joint Event Extraction via RNNs	114
4.21	Event Extraction via Bidi-LSTM Tensor NNs	116
4.22	Reasoning with Neural Tensor Networks	118
4.23	Language to Logical Form with Neural Attention	119
4.24	Seq2SQL: Generating Structured Queries from NL using RL	121
4.25	SLING: A Framework for Frame Semantic Parsing	124
4.26	Poincaré Embeddings for Learning Hierarchical Representations	126
4.27	Enriching Word Vectors with Subword Information (FastText)	128
4.28	DeepWalk: Online Learning of Social Representations	130
4.29	Review of Relational Machine Learning for Knowledge Graphs	132
4.30	Fast Top-K Search in Knowledge Graphs	135
4.31	Dynamic Recurrent Acyclic Graphical Neural Networks (DRAGNN)	137
4.32	Neural Architecture Search with Reinforcement Learning	140
4.33	Joint Extraction of Events and Entities within a Document Context	142
4.34	Globally Normalized Transition-Based Neural Networks	145
4.35	An Introduction to Conditional Random Fields	148
4.35.1	Inference	152
4.36	Co-sampling: Training Robust Networks for Extremely Noisy Supervision	155
5	NLP with Deep Learning	156
5.1	Word Vector Representations (Lec 2)	157
5.2	GloVe (Lec 3)	160
6	Speech and Language Processing	163

6.1	Introduction (Ch. 1 2nd Ed.)	164
6.2	Morphology (Ch. 3 2nd Ed.)	165
6.3	N-Grams (Ch. 6 2nd Ed.)	166
6.4	Naive Bayes and Sentiment (Ch. 6 3rd Ed.)	168
6.5	Hidden Markov Models (Ch. 9 3rd Ed.)	170
6.6	POS Tagging (Ch. 10 3rd Ed.)	173
6.7	Formal Grammars (Ch. 11 3rd Ed.)	176
6.8	Vector Semantics (Ch. 15)	177
6.9	Semantics with Dense Vectors (Ch. 16)	180
6.10	Information Extraction (Ch. 21 3rd Ed)	183
7	Probabilistic Graphical Models	186
7.1	Foundations (Ch. 2)	187
7.1.1	Appendix	188
7.2	The Bayesian Network Representation (Ch. 3)	190
7.3	Undirected Graphical Models (Ch. 4)	194
7.3.1	Exercises	200
7.4	Local Probabilistic Models (Ch. 5)	201
7.5	Template-Based Representations (Ch. 6)	202
7.6	Variable Elimination (Ch. 9)	204
7.7	Clique Trees (Ch. 10)	208
7.8	Inference as Optimization (Ch. 11)	212
8	Information Theory, Inference, and Learning Algorithms	213
8.1	Introduction to Information Theory (Ch. 1)	214
8.2	Probability, Entropy, and Inference (Ch. 2)	216
8.2.1	More About Inference (Ch. 3 Summary)	217
8.3	The Source Coding Theorem (Ch. 4)	219
8.3.1	Data Compression and Typicality	221
8.3.2	Further Analysis and Q&A	223

8.4	Monte Carlo Methods (Ch. 29)	225
8.5	Variational Methods (Ch. 33)	227
9	Bayesian Data Analysis	229
9.1	Probability and Inference (Ch. 1)	230
9.2	Single-Parameter Models (Ch. 2)	231
9.3	Asymptotics and Connections to Non-Bayesian Approaches (Ch. 4)	234
10	Blogs	237
10.1	Conv Nets: A Modular Perspective	238
10.2	Understanding Convolutions	239
10.3	Deep Reinforcement Learning	241
10.4	Deep Learning for Chatbots (WildML)	243
10.5	Attentional Interfaces – Neural Perspective	245

MATH AND MACHINE LEARNING BASICS

CONTENTS

1.1	Linear Algebra (Quick Review) (Ch. 2)	7
1.1.1	Example: Principal Component Analysis	9
1.2	Probability & Information Theory (Quick Review) (Ch. 3)	11
1.3	Numerical Computation (Ch. 4)	13
1.4	Machine Learning Basics (Ch. 5)	16
1.4.1	Estimators, Bias and Variance (5.4)	16
1.4.2	Maximum Likelihood Estimation (5.5)	18
1.4.3	Bayesian Statistics (5.6)	19
1.4.4	Supervised Learning Algorithms (5.7)	21

Linear Algebra (Quick Review) (Ch. 2)

Table of Contents Local

Written by Brandon McKinzie

- For A^{-1} to exist, $Ax = b$ must have exactly one solution for every value of b . Determining whether a solution exists $\forall b \in \mathbb{R}^m$ means requiring that the **column space** (range) of A be all of \mathbb{R}^m . It is helpful to see Ax expanded out explicitly in this way:

Unless stated otherwise, assume $A \in \mathbb{R}^{m \times n}$

$$Ax = \sum_i x_i A_{:,i} = x_1 \begin{pmatrix} A_{1,1} \\ \vdots \\ A_{m,1} \end{pmatrix} + \cdots + x_m \begin{pmatrix} A_{1,m} \\ \vdots \\ A_{m,m} \end{pmatrix} \quad (2.27)$$

- Necessary: A must have at least m columns ($n \geq m$). (“wide”).
- Necessary *and* sufficient: matrix must contain at least one set of m linearly independent columns.
- Invertibility: In addition to above, need matrix to be *square* (re: at most m columns \wedge at least m columns).

- A square matrix with linearly dependent columns is known as **singular**. A (necessarily square) matrix is singular if and only if one or more eigenvalues are zero.

- A **norm** is any function f that satisfies the following properties:

$$\|x\|_\infty = \max_i |x_i|$$

$$f(x) = 0 \Rightarrow x = \mathbf{0} \quad (1)$$

$$f(x + y) \leq f(x) + f(y) \quad (2)$$

$$\forall \alpha \in \mathbb{R}, f(\alpha x) = |\alpha| f(x) \quad (3)$$

- An **orthogonal matrix** is a square matrix whose rows are mutually orthonormal and whose columns are mutually orthonormal:

$$A^T A = A A^T = I \quad (2.37)$$

$$A^{-1} = A^T \quad (2.38)$$

Note that orthonorm cols implies orthonorm rows (if square). To prove, consider the relationship between $A^T A$ and $A A^T$

- Suppose square matrix $A \in \mathbb{R}^{n \times n}$ has n linearly independent eigenvectors $\{v^{(1)}, \dots, v^{(n)}\}$. The **eigendecomposition** of A is then given by¹

$$A = V \text{diag}(\lambda) V^{-1} \quad (2.40)$$

In the special case where A is real-symmetric, $A = Q \Lambda Q^T$. **Interpretation:** Ax can be decomposed into the following three steps:

All real-symmetric A have an eigendecomposition, **but it might not be unique!**

¹This appear to imply that unless the columns of V are also normalized, can't guarantee that its inverse equals its transpose? (since that is the only difference between it and an orthogonal matrix)

- 1) **Change of basis:** The vector $(Q^T \mathbf{x})$ can be thought of as how \mathbf{x} would appear in the basis of eigenvectors of \mathbf{A} .
- 2) **Scale:** Next, we scale each component $(Q^T \mathbf{x})_i$ by an amount λ_i , yielding the new vector $(\Lambda(Q^T \mathbf{x}))$.
- 3) **Change of basis:** Finally, we rotate this new vector back from the eigen-basis into its original basis, yielding the transformed result of $Q\Lambda Q^T \mathbf{x}$.

A common convention to sort the entries of Λ in descending order.

- **Positive definite:** all λ are positive; **positive semidefinite:** all λ are positive or zero.
 - PSD: $\forall \mathbf{x}, \mathbf{x}^T \mathbf{A} \mathbf{x} \geq 0$
 - PD: $\mathbf{x}^T \mathbf{A} \mathbf{x} = 0 \Rightarrow \mathbf{x} = \mathbf{0}$.²
- Any real matrix $\mathbf{A} \in \mathbb{R}^{m \times n}$ has a **singular value decomposition** of the form,

$$\mathbf{A} = \mathbf{U} \mathbf{D} \mathbf{V}^T \quad (10)$$

$$\mathbf{U} \in \mathbb{R}^{m \times m} \quad (7)$$

$$\mathbf{D} \in \mathbb{R}^{m \times n} \quad (8)$$

$$\mathbf{V} \in \mathbb{R}^{n \times n} \quad (9)$$

where both \mathbf{U} and \mathbf{V} are orthogonal matrices, and \mathbf{D} is diagonal.

- The **singular values** are the diagonal entries D_{ii} .
- The **left(right)-singular vectors** are the columns of $\mathbf{U}(\mathbf{V})$.
- Eigenvectors of $\mathbf{A} \mathbf{A}^T$ are the L-S vectors. Eigenvectors of $\mathbf{A}^T \mathbf{A}$ are the R-S vectors. The eigenvalues of both $\mathbf{A} \mathbf{A}^T$ and $\mathbf{A}^T \mathbf{A}$ are given by the singular values squared.
- The Moore-Penrose **pseudoinverse**, denoted \mathbf{A}^+ , enables us to find an “inverse” of sorts for a (possibly) non-square matrix \mathbf{A} . Most algorithms compute \mathbf{A}^+ via

$$\mathbf{A}^+ = \mathbf{V} \mathbf{D}^+ \mathbf{U}^T \quad (11)$$

\mathbf{A}^+ is useful, e.g., when we want to solve $\mathbf{A} \mathbf{x} = \mathbf{y}$ by left-multiplying each side to obtain $\mathbf{x} = \mathbf{B} \mathbf{y}$. It is far more likely for solution(s) to exist when \mathbf{A} is wider than it is tall.

- The **determinant** of a matrix is $\det(\mathbf{A}) = \prod_i \lambda_i$. Conceptually, $|\det(\mathbf{A})|$ tells how much [multiplication by] \mathbf{A} expands/contracts space. If $\det(\mathbf{A}) = 1$, the transformation preserves volume.

²I proved this and it made me happy inside. Check it out. Let \mathbf{A} be positive definite. Then

$$\mathbf{x}^T \mathbf{A} \mathbf{x} = \mathbf{x}^T \mathbf{Q} \Lambda \mathbf{Q}^T \mathbf{x} \quad (4)$$

$$= \sum_i (\mathbf{Q}^T \mathbf{x})_i \lambda_i (\mathbf{Q}^T \mathbf{x})_i \quad (5)$$

$$= \sum_i \lambda_i (\mathbf{Q}^T \mathbf{x})_i^2 \quad (6)$$

Since all terms in the summation are non-negative and all $\lambda_i > 0$, we have that $\mathbf{x}^T \mathbf{A} \mathbf{x} = 0$ if and only if $(\mathbf{Q}^T \mathbf{x})_i = 0 = \mathbf{q}^{(i)} \cdot \mathbf{x}$ for all i . Since the set of eigenvectors $\{\mathbf{q}^{(i)}\}$ form an orthonormal basis, we have that \mathbf{x} must be the zero vector.

EXAMPLE: PRINCIPAL COMPONENT ANALYSIS

Task. Say we want to apply lossy compression (less memory, but may lose precision) to a collection of m points $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$. We will do this by converting each $\mathbf{x}^{(i)} \in \mathbb{R}^n$ to some $\mathbf{c}^{(i)} \in \mathbb{R}^l$ ($l < n$), i.e. finding functions f and g such that:

$$f(\mathbf{x}) = \mathbf{c} \quad \text{and} \quad \mathbf{x} \approx g(f(\mathbf{x})) \quad (12)$$

Decoding function (g). As is, we still have a rather general task to solve. *PCA* is defined by choosing $g(\mathbf{c}) = \mathbf{D}\mathbf{c}$, with $\mathbf{D} \in \mathbb{R}^{n \times l}$, where all columns of \mathbf{D} are both (1) orthogonal and (2) unit norm.

Encoding function (f). Now we need a way of mapping \mathbf{x} to \mathbf{c} such that $g(\mathbf{c})$ will give us back a vector optimally close to \mathbf{x} . We've already defined g , so this amounts to finding the optimal \mathbf{c}^* such that:

$$\mathbf{c}^* = \arg \min_{\mathbf{c}} \|\mathbf{x} - g(\mathbf{c})\|_2^2 \quad (13)$$

$$(\mathbf{x} - g(\mathbf{c}))^T (\mathbf{x} - g(\mathbf{c})) = \mathbf{x}^T \mathbf{x} - 2\mathbf{x}^T g(\mathbf{c}) + g(\mathbf{c})^T g(\mathbf{c}) \quad (14)$$

$$\mathbf{c}^* = \arg \min_{\mathbf{c}} \left[-2\mathbf{x}^T \mathbf{D}\mathbf{c} + \mathbf{c}^T \mathbf{c} \right] \quad (15)$$

$$= \mathbf{D}^T \mathbf{x} = f(\mathbf{x}) \quad (16)$$

which means the *PCA reconstruction operation* is defined as $r(\mathbf{x}) = \mathbf{D}\mathbf{D}^T \mathbf{x}$.

Optimal \mathbf{D} . It is important to notice that we've been able to determine e.g. the optimal \mathbf{c}^* for some \mathbf{x} because each \mathbf{x} has a (allowably) different \mathbf{c}^* . However, we use *the same* matrix \mathbf{D} for all our samples $\mathbf{x}^{(i)}$, and thus must optimize it over all points in our collection. With that out of the way, we just do what we always do: minimize over the L^2 distance between points and their reconstruction. Formally, we minimize the Frobenius norm of the matrix of errors:

$$\mathbf{D}^* = \arg \min_{\mathbf{D}} \sqrt{\sum_{i,j} \left(\mathbf{x}_j^{(i)} - r(\mathbf{x}^{(i)})_j \right)^2} \quad s.t. \quad \mathbf{D}^T \mathbf{D} = \mathbf{I} \quad (17)$$

Consider the case of $l = 1$ which means $\mathbf{D} = \mathbf{d} \in \mathbb{R}^n$. In this case, after [insert math here], we obtain

$$\mathbf{d}^* = \arg \max_{\mathbf{d}} \text{Tr} \left(\mathbf{d}^T \mathbf{X}^T \mathbf{X} \mathbf{d} \right) \quad s.t. \quad \mathbf{d}^T \mathbf{d} = 1 \quad (18)$$

where, as usual, $\mathbf{X} \in \mathbb{R}^{m,n}$. It should be clear that the optimal \mathbf{d} is just the largest eigenvector of $\mathbf{X}^T \mathbf{X}$.

Probability & Information Theory (Quick Review) (Ch. 3)

Table of Contents Local

Written by Brandon McKinzie

Expectation. For some function $f(x)$, $\mathbb{E}_{x \sim P}[f(x)]$ is the mean value that f takes on when x is drawn from P . The formula for discrete and continuous variables, respectively is as follows:

$$\mathbb{E}_{x \sim P}[f(x)] = \sum_x P(x)f(x) \quad (3.9)$$

$$\mathbb{E}_{x \sim P}[f(x)] = \int p(x)f(x)dx \quad (3.10)$$

Variance. A measure of how much the values of a function of a random variable x vary as we sample different values of x from its distribution.

$$\text{Var}[f(x)] = \mathbb{E}[(f(x) - \mathbb{E}[f(x)])^2] \quad (3.11)$$

Covariance. Gives some sense of how much two values are *linearly* related to each other, as well as the *scale* of these variables.

$$\text{Cov}[f(x), g(x)] = \mathbb{E}[(f(x) - \mathbb{E}[f(x)])(g(x) - \mathbb{E}[g(x)])] \quad (3.13)$$

→ Large $|\text{Cov}[f, g]|$ means the function values change a lot and both functions are far from their means at the same time.

→ **Correlation** normalizes the contribution of each variable in order to measure only how much the variables are related.

Covariance Matrix of a random vector $\mathbf{x} \in \mathbb{R}^n$ is an $n \times n$ matrix, such that

$$\text{Cov}[\mathbf{x}]_{i,j} = \text{Cov}[x_i, x_j] \quad (3.14)$$

and if we want the “sample” covariance matrix taken over m data point samples, then

$$\Sigma := \frac{1}{m} \sum_{k=1}^m (x_k - \bar{x})(x_k - \bar{x})^T \quad (19)$$

where m is the number of data points.

Measure Theory.

- A set of points that is negligibly small is said to have **measure zero**. In practical terms, think of such a set as occupying no volume in the space we are measuring (interested in).
- A property that holds **almost everywhere** holds throughout all space except for on a set of measure zero.

In \mathbb{R}^2 , a line has measure zero.

Functions of RVs.

- **Common mistake:** Suppose $\mathbf{y} = g(\mathbf{x})$, and g is invertible/continuous/differentiable. It is NOT true that $p_y(\mathbf{y}) = p_x(g^{-1}(\mathbf{y}))$. This fails to account for the distortion of [probability] space introduced by g . Rather,

$$p_x(\mathbf{x}) = p_y(g(\mathbf{x})) \left| \frac{\partial g(\mathbf{x})}{\partial \mathbf{x}} \right| \quad (3.47)$$

Information Theory. Denote the **self-information** of an event $\mathbf{x} = x$ to be

$$I(x) \triangleq -\log P(x) \quad (20)$$

where \log is always assumed to be the natural logarithm. We can quantify the amount of uncertainty in an entire probability distribution using the **Shannon entropy**,

$$H(\mathbf{x}) = \mathbb{E}_{\mathbf{x} \sim P} [I(x)] = -\mathbb{E}_{\mathbf{x} \sim P} [\log P(x)] \quad (21)$$

which gives the expected amount of information in an event drawn from that distribution. Taking it a step further, say we have two separate probability distributions $P(\mathbf{x})$ and $Q(\mathbf{x})$. We can measure how different these distributions are with the **Kullback-Leibler (KL) divergence**:

$$D_{KL}(P||Q) \triangleq \mathbb{E}_{\mathbf{x} \sim P} \left[\log \frac{P(x)}{Q(x)} \right] = \mathbb{E}_{\mathbf{x} \sim P} [\log P(x) - \log Q(x)] \quad (22)$$

Note that the expectation is taken over P , thus making D_{KL} not symmetric (and thus not a true distance measure), since $D_{KL}(P||Q) \neq D_{KL}(Q||P)$. Finally, a closely related quantity is the **cross-entropy**, $H(P, Q)$, defined as:

$$H(P, Q) \triangleq H(P) + D_{KL}(P||Q) \quad (23)$$

$$= -\mathbb{E}_{\mathbf{x} \sim P} [\log Q(x)] \quad (24)$$

Numerical Computation (Ch. 4)

Table of Contents Local

Written by Brandon McKinzie

Some terminology. **Underflow** is when numbers near zero are rounded to zero. Similarly, **overflow** is when large [magnitude] numbers are approximated as $\pm\infty$. **Conditioning** refers to how rapidly a function changes w.r.t. small changes in its inputs. Consider the function $f(\mathbf{x}) = \mathbf{A}^{-1}\mathbf{x}$. When \mathbf{A} has an eigenvalue decomposition, its *condition number* is

$$\max_{i,j} \left| \frac{\lambda_i}{\lambda_j} \right| \quad (4.2)$$

which is the ratio of the magnitude of the largest and smallest eigenvalue. When this is large, matrix inversion is sensitive to error in the input [of $f(\mathbf{x})$].

Gradient-based optimization. Recall from basic calculus that the **directional derivative** of $f(\mathbf{x})$ in direction $\hat{\mathbf{u}}$ (a unit vector) is defined as the slope of the function f in direction $\hat{\mathbf{u}}$. By definition of the derivative, this is given by (with $\mathbf{v} := \mathbf{x} + \alpha\hat{\mathbf{u}}$)

$$\lim_{\alpha \rightarrow 0} \frac{f(\mathbf{x} + \alpha\hat{\mathbf{u}}) - f(\mathbf{x})}{\alpha} = \left. \frac{\partial f(\mathbf{x} + \alpha\hat{\mathbf{u}})}{\partial \alpha} \right|_{\alpha=0} \quad (25)$$

$$= \sum_i \left. \frac{\partial f(\mathbf{v})}{\partial v_i} \frac{\partial v_i}{\partial \alpha} \right|_{\alpha=0} \quad (26)$$

$$= \sum_i \left. (\nabla_{\mathbf{v}} f(\mathbf{v}))_i u_i \right|_{\alpha=0} \quad (27)$$

$$= \left. \hat{\mathbf{u}}^T \nabla_{\mathbf{v}} f(\mathbf{v}) \right|_{\alpha=0} \quad (28)$$

$$= \hat{\mathbf{u}}^T \nabla_{\mathbf{x}} f(\mathbf{x}) \quad (29)$$

where it's important to recognize the distinction between $\lim_{\alpha \rightarrow 0}$ and *setting* α to zero, which is denoted by $|_{\alpha=0}$. If we want to *find* the direction $\hat{\mathbf{u}}$ such that this directional derivative is a minimum, i.e.

$$\hat{\mathbf{u}}^* = \arg \min_{\hat{\mathbf{u}}, \hat{\mathbf{u}}^T \hat{\mathbf{u}}=1} \hat{\mathbf{u}}^T \nabla_{\mathbf{x}} f(\mathbf{x}) \quad (30)$$

$$= \arg \min_{\hat{\mathbf{u}}, \hat{\mathbf{u}}^T \hat{\mathbf{u}}=1} \|\hat{\mathbf{u}}\|_2 \|\nabla_{\mathbf{x}} f(\mathbf{x})\|_2 \cos(\theta) \quad (31)$$

$$= \cos(\theta) \quad (32)$$

and we see that $\hat{\mathbf{u}}$ points in the opposite direction as the gradient.

Jacobian and Hessian Matrices. For when we want partial derivatives of some function f whose input and output are both vectors. The **Jacobian matrix** contains all such partial derivatives. Sometimes we want to know about second derivatives too, since this tells us whether a gradient step will cause as much of an improvement as we would expect based on the gradient alone. The **Hessian matrix** $\mathbf{H}(f)(\mathbf{x})$ is defined such that

$$\begin{aligned} f : \mathbb{R}^m &\rightarrow \mathbb{R}^n \\ \mathbf{J} &\in \mathbb{R}^{n \times m} \text{ where} \\ J_{i,j} &= \frac{\partial}{\partial x_j} f(\mathbf{x})_i \end{aligned}$$

The Hessian is the Jacobian of the gradient.

$$\mathbf{H}(f)(\mathbf{x})_{i,j} = \frac{\partial^2}{\partial x_i \partial x_j} f(\mathbf{x}) \quad (4.6)$$

The second derivative in a specific direction $\hat{\mathbf{d}}$ is given by $\hat{\mathbf{d}}^T \mathbf{H} \hat{\mathbf{d}}$. It tells us how well we can expect a gradient descent step to perform. How so? Well, it shows up in the second-order approximation to the function $f(\mathbf{x})$ about our current spot, which we can denote $\mathbf{x}^{(0)}$. The standard gradient descent step will move us from $\mathbf{x}^{(0)} \rightarrow \mathbf{x}^{(0)} - \epsilon \mathbf{g}$, where \mathbf{g} is the gradient evaluated at $\mathbf{x}^{(0)}$. Plugging this in to the 2nd order approximation shows us how \mathbf{H} can give information related to how “good” of a step that really was. Mathematically,

$$f(\mathbf{x}) \approx f(\mathbf{x}^{(0)}) + (\mathbf{x} - \mathbf{x}^{(0)})^T \mathbf{g} + \frac{1}{2} (\mathbf{x} - \mathbf{x}^{(0)})^T \mathbf{H} (\mathbf{x} - \mathbf{x}^{(0)}) \quad (4.8)$$

$$f(\mathbf{x}^{(0)} - \epsilon \mathbf{g}) \approx f(\mathbf{x}^{(0)}) - \epsilon \mathbf{g}^T \mathbf{g} + \frac{1}{2} \epsilon^2 \mathbf{g}^T \mathbf{H} \mathbf{g} \quad (4.9)$$

If $\mathbf{g}^T \mathbf{H} \mathbf{g}$ is positive, then we can easily solve for the optimal $\epsilon = \epsilon^*$ that decreases the Taylor series approximation as

$$\epsilon^* = \frac{\mathbf{g}^T \mathbf{g}}{\mathbf{g}^T \mathbf{H} \mathbf{g}} \quad (4.10)$$

which can be as low as $1/\lambda_{max}$ (the worst case), and as high as $1/\lambda_{min}$ with the λ being the eigenvalues of the Hessian. The best (and perhaps only) way to take what we learned about the “second derivative test” in single-variable calculus and apply it to the multidimensional case with \mathbf{H} is by using the *eigendecomposition of \mathbf{H}* . Why? Because we can examine the eigenvalues of the Hessian to determine whether the critical point $\mathbf{x}^{(0)}$ is a local maximum, local minimum, or saddle point⁴. If all eigenvalues are positive (remember that this is equivalent to saying that the Hessian is **positive definite!**), the point is a local minimum.

The condition number of the Hessian at a given point can give us an idea about how much the second derivatives (along different directions) differ from each other

³ In the same manner that I derived equation 29, we can derive the second derivative in a specified direction $\hat{\mathbf{d}}$:

$$\frac{\partial^2}{\partial \alpha^2} f(\mathbf{x} + \alpha \hat{\mathbf{d}}) \Big|_{\alpha=0} = \frac{\partial}{\partial \alpha} \hat{\mathbf{d}}^T \nabla_{\mathbf{v}} f(\mathbf{v}) \Big|_{\alpha=0} \quad (33)$$

$$= \sum_i d_i \frac{\partial}{\partial \alpha} \frac{\partial f(\mathbf{v})}{\partial v_i} \Big|_{\alpha=0} \quad (34)$$

$$= \sum_i d_i \frac{\partial}{\partial v_i} \frac{\partial f(\mathbf{v})}{\partial \alpha} \Big|_{\alpha=0} \quad (35)$$

$$= \sum_i \sum_j d_i \frac{\partial^2 f(\mathbf{v})}{\partial v_i \partial v_j} d_j \Big|_{\alpha=0} \quad (36)$$

$$= \hat{\mathbf{d}}^T \mathbf{H} \hat{\mathbf{d}} \quad (37)$$

⁴Emphasis on “values” in “eigenvalues” because it’s important not to get tripped up here about what the

Constrained optimization: minimizing/maximizing a function $f(\mathbf{x})$ constrained to only values of \mathbf{x} in some set \mathbb{S} . One way of approaching such a problem is to re-design the unconstrained optimization problem such that the re-designed problem's solution satisfies the constraints. For example, to minimize $f(\mathbf{x})$ for $\mathbf{x} \in \mathbb{R}^2$ with constraint $\|\mathbf{x}\|_2 = 1$, we can minimize $g(\theta) = f([\cos \theta, \sin \theta]^T)$ wrt θ , then return $[\cos \theta, \sin \theta]^T$ as the solution to the original problem.

The **Karush-Kuhn-Tucker** (KKT) approach, a generalization of **Lagrange multipliers**, provides a general approach for re-designing the optimization problem, with procedure as follows:

1. Find m functions $g^{(i)}$ and n functions $h^{(j)}$ such that your set of allowed values \mathbb{S} can be written

$$\mathbb{S} = \{\mathbf{x} \mid \forall i, g^{(i)}(\mathbf{x}) = 0 \text{ and } \forall j, h^{(j)}(\mathbf{x}) \leq 0\} \quad (38)$$

The equations involving $g^{(i)}$ are called the **equality constraints** and the inequalities involving $h^{(j)}$ are called the **inequality constraints**.

2. Introduce new variables λ_i (for the equality constraints) and α_j (for the inequality constraints). These are called the KKT multipliers. The generalized Lagrangian is then defined as

$$\mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}, \boldsymbol{\alpha}) = f(\mathbf{x}) + \sum_i \lambda_i g^{(i)}(\mathbf{x}) + \sum_j \alpha_j h^{(j)}(\mathbf{x}) \quad (39)$$

3. Solve the re-designed unconstrained optimization problem:

$$\min_{\mathbf{x}} \max_{\boldsymbol{\lambda}} \max_{\boldsymbol{\alpha}, \boldsymbol{\alpha} \geq 0} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}, \boldsymbol{\alpha}) \quad (40)$$

which has the same optimal objective function value and set of optimal points \mathbf{x} as the original constrained problem, $\min_{\mathbf{x} \in \mathbb{S}} f(\mathbf{x})$. Any time the constraints are satisfied, the expression $\max_{\boldsymbol{\lambda}} \max_{\boldsymbol{\alpha}, \boldsymbol{\alpha} \geq 0} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}, \boldsymbol{\alpha})$ evaluates to $f(\mathbf{x})$, and any time a constraint is violated, the same expression evaluates to ∞ .

eigenvectors of the Hessian mean. The reason for the decomposition is that it gives us an orthonormal basis (out of which we can get any direction) and therefore the magnitude of the second derivative along each of these directions as the eigenvalues.

Machine Learning Basics (Ch. 5)

Table of Contents Local

Written by Brandon McKinzie

Capacity, Overfitting, and Underfitting. Difference between ML and optimization is that, in addition to wanting low training error, we want **generalization error** (test error) to be low as well. The ideal model is an oracle that simply knows the true probability distribution $p(\mathbf{x}, y)$ that generates the data. The error incurred by such an oracle, due things like inherently stochastic mappings from \mathbf{x} to y or other variables, is called the **Bayes error**. The **no free lunch theorem** states that, averaged over all possible data-generating distributions, every classification algorithm has the same error rate when classifying previously unobserved points. Therefore, the goal of ML research is to understand what kinds of distributions are relevant to the “real world” that an AI agent experiences, and what kinds of ML algorithms perform well on data drawn from the relevant data-generating distributions.

ESTIMATORS, BIAS AND VARIANCE (5.4)

Point Estimation: attempt to provide “best” prediction of some quantity, such as some parameter or even a whole function. Formally, a point estimator or *statistic* is any function of the data:

$$\hat{\theta}_m = g\left(x^{(1)}, \dots, x^{(m)}\right) \quad (5.19)$$

where, since the data is drawn from a random process, $\hat{\theta}$ is a random variable. **Function estimation** is identical in form, where we want to estimate some $f(x)$ with \hat{f} , a point estimator in *function space*.

Bias. Defined below, where the expectation is taken over the data-generating distribution⁵. Bias measures the expected deviation from the true value of the func/param.

$$\text{bias} \left[\hat{\theta}_m \right] = \mathbb{E} \left[\hat{\theta}_m \right] - \theta \quad (5.20)$$

TODO: Figure out how to derive $\mathbb{E} \left[\hat{\theta}_m^2 \right]$ for Gaussian distribution [helpful link].

⁵May want to double-check this, but I’m fairly certain this is what the book meant when it said “data,” based on later examples.

Bias-Variance Tradeoff.

→ **Conceptual Info.** Two sources of error for an estimator are (1) bias and (2) variance, which are both defined as deviations from a certain value. Bias gives deviation from the *true* value, while variance gives the [expected] deviation from this *expected* value.

→ **Summary of main formulas.**

$$\text{bias} [\hat{\theta}_m] = \mathbb{E} [\hat{\theta}_m] - \theta \quad (41)$$

$$\text{Var} [\hat{\theta}_m] = \mathbb{E} \left[\left(\hat{\theta}_m - \mathbb{E} [\hat{\theta}_m] \right)^2 \right] \quad (42)$$

→ **MSE decomposition.** The MSE of the estimates is given by⁶

$$\text{MSE} = \mathbb{E} \left[(\hat{\theta}_m - \theta)^2 \right] \quad (5.53)$$

$$= \text{Bias}(\hat{\theta})^2 + \text{Var} [\hat{\theta}_m] \quad (5.54)$$

and desirable estimators are those with low MSE.

Consistency. As the number of training data points increases, we want the estimators to converge to the true values. Specifically, below are the definitions for *weak* and *strong* consistency, respectively.

$$\begin{aligned} \text{plim}_{m \rightarrow \infty} \hat{\theta}_m &= \theta \\ p \left(\lim_{m \rightarrow \infty} \hat{\theta}_m = \theta \right) &= 1 \end{aligned} \quad (5.55)$$

where the symbol plim means $P(|\hat{\theta}_m - \theta| > \epsilon) \rightarrow 0$ as $m \rightarrow \infty$.

⁶Derivation:

$$\text{MSE} = \mathbb{E} [\hat{\theta}^2 + \theta^2 - 2\theta\hat{\theta}] \quad (43)$$

$$= \mathbb{E} [\hat{\theta}^2] + \theta^2 - 2\theta\mathbb{E} [\hat{\theta}] \quad (44)$$

$$= (\mathbb{E} [\hat{\theta}]^2 - \mathbb{E} [\hat{\theta}]^2) + \mathbb{E} [\hat{\theta}^2] + \theta^2 - 2\theta\mathbb{E} [\hat{\theta}] \quad (45)$$

$$= \left(\mathbb{E} [\hat{\theta}]^2 + \theta^2 - 2\theta\mathbb{E} [\hat{\theta}] \right) + \left(\mathbb{E} [\hat{\theta}^2] - \mathbb{E} [\hat{\theta}]^2 \right) \quad (46)$$

$$= \text{Bias}(\hat{\theta})^2 + \text{Var} [\hat{\theta}_m] \quad (47)$$

MAXIMUM LIKELIHOOD ESTIMATION (5.5)

Consider set of m examples $\mathbb{X} = \{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ drawn independently from the true (but unknown) $p_{data}(\mathbf{x})$. Let $p_{model}(\mathbf{x}; \boldsymbol{\theta})$ be parametric family of probability distributions over the same space indexed by $\boldsymbol{\theta}$. The maximum likelihood estimator for $\boldsymbol{\theta}$ can be expressed as

$$\boldsymbol{\theta}_{ML} = \arg \max_{\boldsymbol{\theta}} \mathbb{E}_{\mathbf{x} \sim \hat{p}_{data}} [\log p_{model}(\mathbf{x}; \boldsymbol{\theta})] \quad (5.59)$$

where we've chosen to express with log for underflow/gradient reasons. One interpretation of ML is to view it as minimizing the dissimilarity, as measured by the KL divergence⁷, between \hat{p}_{data} and p_{model} .

Any loss consisting of a negative log-likelihood is a **cross-entropy** between the \hat{p}_{data} distribution and the p_{model} distribution.

Thoughts: Let's look at D_{KL} in some more detail. First, I'll rewrite it with the explicit definition of $\mathbb{E}_{\mathbf{x} \sim \hat{p}_{data}} [\log (\hat{p}_{data}(\mathbf{x}))]$:

$$D_{KL}(\hat{p}_{data} || p_{model}) = \mathbb{E}_{\mathbf{x} \sim \hat{p}_{data}} [\log (\hat{p}_{data}(\mathbf{x})) - \log (p_{model}(\mathbf{x}))] \quad (48)$$

$$= \left(\frac{1}{N} \left(\sum_{i=1}^N \log (\text{Counts}(\mathbf{x}_i)) \right) - \log N \right) - \mathbb{E}_{\mathbf{x} \sim \hat{p}_{data}} [\log (p_{model}(\mathbf{x}))] \quad (49)$$

Note also that our goal is to find parameters $\boldsymbol{\theta}$ such that D_{KL} is minimized. It is for *this* reason, that we wish to optimize over $\boldsymbol{\theta}$, that minimizing D_{KL} amounts to maximizing the quantity, $\mathbb{E}_{\mathbf{x} \sim \hat{p}_{data}} [\log (p_{model}(\mathbf{x}))]$. Sure, I can agree this is *true*, but **why is our goal to minimize D_{KL} , as opposed to minimizing $|D_{KL}|$?** I'm assuming it is because optimizing w.r.t. an absolute value is challenging numerically.

Conditional Log-Likelihood and MSE. We can readily generalize $\boldsymbol{\theta}_{ML}$ to estimate a conditional probability $p(\mathbf{y} | \mathbf{x}; \boldsymbol{\theta})$ in order to predict \mathbf{y} given \mathbf{x} , since

We are assuming the examples are i.i.d. here.

$$\boldsymbol{\theta}_{ML} = \arg \max_{\boldsymbol{\theta}} \sum_{i=1}^m \log P(\mathbf{y}^{(i)} | \mathbf{x}^{(i)}; \boldsymbol{\theta}) \quad (5.63)$$

where $\mathbf{x}^{(i)}$ are fed as *inputs* to the model; this is why we can formulate MLE as a conditional probability.

⁷ The KL divergence is given by

$$D_{KL}(\hat{p}_{data} || p_{model}) = \mathbb{E}_{\mathbf{x} \sim \hat{p}_{data}} [\log \hat{p}_{data}(\mathbf{x}) - \log p_{model}(\mathbf{x})] \quad (5.60)$$

BAYESIAN STATISTICS (5.6)

Distinction between frequentist and bayesian approach:

- **Frequentist:** Estimate θ \rightarrow make predictions thereafter based on this estimate.
- **Bayesian:** Consider all possible values of θ when making predictions.

The prior. Before observing the data, we represent our knowledge of θ using the **prior probability distribution** $p(\theta)$. Unlike maximum likelihood, which makes predictions using a *point estimate* of θ (a single value), the Bayesian approach uses Bayes' rule to make predictions using the *full distribution* over θ . In other words, rather than focusing on the most accurate value estimate of θ , we instead focus on pinning down a range of possible θ values and how likely we believe each of these values to be.

It is common to choose a high-entropy prior, e.g. uniform.

So what happens to θ after we observe the data? We update it using Bayes' rule⁸:

$$p(\theta | x^{(1)}, \dots, x^{(m)}) = \frac{p(x^{(1)}, \dots, x^{(m)} | \theta)p(\theta)}{p(x^{(1)}, \dots, x^{(m)})} \quad (50)$$

Note that we still haven't mentioned how to actually make *predictions*. Since we no longer have just one value for θ , but rather we have a posterior *distribution* $p(\theta | x^{(1)}, \dots, x^{(m)})$, we must integrate over this to get the predicted likelihood of the next sample $x^{(m+1)}$:

$$p(x^{(m+1)} | x^{(1)}, \dots, x^{(m)}) = \int p(x^{(m+1)} | \theta)p(\theta | x^{(1)}, \dots, x^{(m)})d\theta \quad (51)$$

$$= \mathbb{E}_{\theta \sim p(\theta | x^{(1)}, \dots, x^{(m)})} [p(x^{(m+1)} | \theta)] \quad (52)$$

Linear Regression: MLE vs. Bayesian. Both want to model the conditional distribution $p(y | \mathbf{x})$ (the conditional likelihood). To derive the standard linear regression algorithm, we *define*

$$p(y | \mathbf{x}) = \mathcal{N}(y; \hat{y}(\mathbf{x}; \mathbf{w}), \sigma^2) \quad (53)$$

$$\hat{y}(\mathbf{x}; \mathbf{w}) = \mathbf{w}^T \mathbf{x} \quad (54)$$

Assume σ^2 is some fixed constant chosen by the user.

⁸In practice, we typically compute the denominator by simply normalizing the probability distribution, i.e. it is effectively the partition function.

- **Maximum Likelihood Approach:** We can use the definition above (and the i.i.d. assumption) to evaluate the conditional log-likelihood as

$$\sum_{i=1}^m \log p(y^{(i)} | \mathbf{x}^{(i)}; \boldsymbol{\theta}) = -m \log \sigma - \frac{m}{2} \log(2\pi) - \sum_{i=1}^m \frac{\|\hat{\mathbf{y}}^{(i)} - \mathbf{y}^{(i)}\|^2}{2\sigma^2} \quad (5.65)$$

where only the last term has any dependence on \mathbf{w} . Therefore, to obtain \mathbf{w}_{ML} we take the derivative of the last term w.r.t. \mathbf{w} , set that to zero, and solve for \mathbf{w} . We see that finding the \mathbf{w} that maximizes the conditional log-likelihood is equivalent to finding the \mathbf{w} that minimizes the training MSE.

Recall that the training MSE is $\frac{1}{m} \sum_{i=1}^m \|\hat{\mathbf{y}}^{(i)} - \mathbf{y}^{(i)}\|^2$

- **Bayesian Approach:** Our conditional likelihood is already given in equation 53. Next, we must define a prior distribution over \mathbf{w} . As is common, we choose a Gaussian prior to express our high degree of uncertainty about $\boldsymbol{\theta}$ (implying we'll choose a relatively large variance):

$$p(\mathbf{w}) := \mathcal{N}(\mathbf{w}; \boldsymbol{\mu}_0, \boldsymbol{\Lambda}_0) \quad (55)$$

Typically assume $\boldsymbol{\Lambda}_0 = \text{diag}(\boldsymbol{\lambda}_0)$

We can then compute [the unnormalized] $p(\mathbf{w} | \mathbf{X}, \mathbf{y}) \propto p(\mathbf{y} | \mathbf{X}, \mathbf{w})p(\mathbf{w})$ [and then normalize it].

Maximum A Posteriori (MAP) Estimation. Often we either prefer a point estimate for θ , or we find out that computing the posterior distribution is intractable and a point estimate offers a tractable estimation. The obvious way of obtaining this while still taking the Bayesian route is to just argmax the posterior and use that as your point estimate:

$$\boldsymbol{\theta}_{MAP} = \arg \max_{\boldsymbol{\theta}} p(\boldsymbol{\theta} | \mathbf{x}) = \arg \max_{\boldsymbol{\theta}} \log p(\mathbf{x} | \boldsymbol{\theta}) + \log p(\boldsymbol{\theta}) \quad (56)$$

where the second form shows how this is basically maximum likelihood with incorporation of the prior. We don't want just any $\boldsymbol{\theta}$ that maximizes the likelihood of our data if there is virtually no chance of that value of $\boldsymbol{\theta}$ in the first place.

Logistic Regression. We’ve already seen that linear regression corresponds to the family

$$p(y \mid \mathbf{x}) = \mathcal{N}(y; \boldsymbol{\theta}^T \mathbf{x}, \mathbf{I}) \quad (5.80)$$

which we can generalize to the binary **classification** scenario by interpreting as the probability of class 1. One way of doing this while ensuring the output is between 0 and 1 is to use the logistic sigmoid function:

$$p(y = 1 \mid \mathbf{x}; \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}^T \mathbf{x}) \quad (5.81)$$

Equation 5.81 is the definition of logistic regression

Unfortunately, there is no closed-form solution for $\boldsymbol{\theta}$, so we must search via maximizing the log-likelihood.

Support Vector Machines. Driving by a linear function $\mathbf{w}^T \mathbf{x} + \mathbf{b}$ like logistic regression, but instead of outputting probabilities it outputs a class identity, which depends on the sign of $\mathbf{w}^T \mathbf{x} + \mathbf{b}$. SVMs make use of the **kernel trick**, the “trick” being that we can rewrite $\mathbf{w}^T \mathbf{x} + \mathbf{b}$ completely in terms of dot products between examples. The general form of our prediction function becomes

$$f(\mathbf{x}) = b + \sum_i \alpha_i k(\mathbf{x}, \mathbf{x}^{(i)}) \quad (5.83)$$

If our kernel function is just $k(x, x^{(i)}) = x^T x^{(i)}$ then we’ve just rewritten \mathbf{w} in the form $\mathbf{w} \rightarrow \mathbf{X}^T \boldsymbol{\alpha}$

where the *kernel* [function] takes the general form $k(\mathbf{x}, \mathbf{x}^{(i)}) = \phi(\mathbf{x}) \cdot \phi(\mathbf{x}^{(i)})$. A major drawback to kernel machines (methods) in general is that the cost of evaluating the decision function $f(\mathbf{x})$ is linear in the number of training examples. SVMs, however, are able to mitigate this by learning an α with mostly zeros. The training examples with *nonzero* α_i are known as **support vectors**.

DEEP NETWORKS: MODERN PRACTICES

CONTENTS

2.1	Deep Feedforward Networks (Ch. 6)	23
2.1.1	Back-Propagation (6.5)	24
2.2	Regularization for Deep Learning (Ch. 7)	25
2.3	Optimization for Training Deep Models (Ch. 8)	27
2.4	Convolutional Neural Networks (Ch. 9)	31
2.5	Sequence Modeling (RNNs) (Ch. 10)	34
2.5.1	Review: The Basics of RNNs	34
2.5.2	RNNs as Directed Graphical Models	39
2.5.3	Challenge of Long-Term Deps. (10.7)	41
2.5.4	LSTMs and Other Gated RNNs (10.10)	42
2.6	Applications (Ch. 12)	43
2.6.1	Natural Language Processing (12.4)	43
2.6.2	Neural Language Models (12.4.2)	44

Deep Feedforward Networks (Ch. 6)

Table of Contents Local

Written by Brandon McKinzie

The strategy/purpose of [feedforward] deep learning is to *learn the set of features/representation describing \mathbf{x}* with a mapping ϕ before applying a linear model. In this approach, we have a model

$$y = f(\mathbf{x}; \boldsymbol{\theta}, \mathbf{w}) = \phi(\mathbf{x}; \boldsymbol{\theta})^T \mathbf{w}$$

with ϕ defining a hidden layer.

ReLU and their generalizations. Some nice properties of ReLUs are...

- Derivatives through a ReLU remain large and consistent whenever the unit is active.
- Second derivative is 0 a.e. and the derivative is 1 everywhere the unit is active, meaning the gradient direction is more useful for learning than it would be with activation functions that introduce 2nd-order effects (see equation 4.9)

Recall the ReLU activation function:
 $g(z) = \max\{0, z\}$
 a.e. is short for "almost everywhere"

Generalizing to aid gradients when $z < 0$. Three such generalizations are based on using a nonzero slope α_i when $z_i < 0$:

$$h_i = g(\mathbf{z}, \boldsymbol{\alpha})_i = \max(0, z_i) + \alpha_i \min(0, z_i) \quad (57)$$

- Absolute value rectification: fix $\alpha_i = -1$ to obtain $g(z) = |z|$.
- Leaky ReLU: fix α_i to a small value like 0.01.
- Parametric ReLU (PReLU): treats α_i like a learnable parameter.

Logistic sigmoid and hyperbolic tangent. Sigmoid activations on hidden units is a bad idea, since they're only sensitive to their inputs near zero, with small gradients everywhere else. If sigmoid activations must be used, tanh is probably a better substitute, since it resembles the identity (i.e. a linear function) near zero.

The chain rule. Suppose $z = f(\mathbf{y})$ where $\mathbf{y} = g(\mathbf{x})$ (see margin for dimensions). Then⁹,

$$\frac{\partial z}{\partial x_i} = (\nabla_{\mathbf{x}} z)_i = \sum_{j=1}^n \frac{\partial z}{\partial y_j} \frac{\partial y_j}{\partial x_i} = \sum_{j=1}^n (\nabla_{\mathbf{y}} z)_j \frac{\partial y_j}{\partial x_i} = \sum_{j=1}^n (\nabla_{\mathbf{y}} z)_j (\nabla_{\mathbf{x}} y_j)_i \quad (6.45)$$

$$\rightarrow \nabla_{\mathbf{x}} z = \left(\frac{\partial \mathbf{y}}{\partial \mathbf{x}} \right)^T \nabla_{\mathbf{y}} z = \mathbf{J}_{\mathbf{y}=g(\mathbf{x})}^T \nabla_{\mathbf{y}} z \quad (6.46)$$

$$\mathbf{x} \in \mathbb{R}^m$$

$$\mathbf{y} \in \mathbb{R}^n$$

$$z : \mathbb{R}^n \rightarrow \mathbb{R}$$

$$g : \mathbb{R}^m \rightarrow \mathbb{R}^n$$

From this we see that the gradient of a variable x can be obtained by multiplying a Jacobian matrix $\frac{\partial \mathbf{y}}{\partial \mathbf{x}}$ by a gradient $\nabla_{\mathbf{y}} z$.

⁹Note that we can view $z = f(\mathbf{y})$ as a multi-variable function of the dimensions of \mathbf{y} ,

$$z = f(y_1, y_2, \dots, y_n)$$

Regularization for Deep Learning (Ch. 7)

Recall the definition of regularization: “any modification we make to a learning algorithm that is intended to reduce its generalization error but not its training error.”

Limiting Model Capacity. Recall that **Capacity** [of a model] is the ability to fit a wide variety of functions. Low cap models may struggle to fit training set, while high cap models may overfit by simply memorizing the training set. We can limit model capacity by adding a parameter norm penalty $\Omega(\theta)$ to the objective function J :

$$\tilde{J}(\theta; X, y) = J(\theta; X, y) + \alpha\Omega(\theta) \quad \text{where} \quad \alpha \in [0, \infty) \quad (7.1)$$

where we typically choose Ω to only penalize the *weights* and leave biases unregularized.

L2-Regularization. Defined as setting $\Omega(\theta) = \frac{1}{2}\|w\|_2^2$. Assume that $J(w)$ is quadratic, with minimum at w^* . Since quadratic, we can approximate J with a second-order expansion about w^* .

$$\hat{J}(w) = J(w^*) + \frac{1}{2}(w - w^*)^T H(w - w^*) \quad (7.6)$$

$$\nabla_w \hat{J}(w) = H(w - w^*) \quad (7.7)$$

where $H_{ij} = \frac{\partial^2 J}{\partial w_i \partial w_j} \big|_{w^*}$. If we add in the [derivative of] the weight decay and set to zero, we obtain the solution

$$\tilde{w} = (H + \alpha I)^{-1} H w^* \quad (7.10)$$

$$= Q(\Lambda + \alpha I)^{-1} \Lambda Q^T w^* \quad (7.13)$$

which shows that the effect of regularization is to rescale the i eigenvectors of H by $\frac{\lambda_i}{\lambda_i + \alpha}$. This means that eigenvectors with $\lambda_i \gg \alpha$ are relatively unchanged, but the eigenvectors with $\lambda_i \ll \alpha$ are shrunk to nearly zero. In other words, only directions along which the parameters contribute significantly to reducing the objective function are preserved relatively intact.

Sparse Representations. Weight decay acts by placing a penalty directly on the model parameters. Another strategy is to place a penalty on the *activations* of the units, encouraging their activations to be sparse. It's important to distinguish the difference between sparse parameters and sparse *representations*. In the former, if we take the example of some $\mathbf{y} = \mathbf{B}\mathbf{h}$, there are many zero entries in some parameter matrix \mathbf{B} while, in the latter, there are many zero entries in the representation vector \mathbf{h} . The modification to the loss function, analogous to 7.1, takes the form

$$\tilde{J}(\boldsymbol{\theta}; \mathbf{X}, \mathbf{y}) = J(\boldsymbol{\theta}; \mathbf{X}, \mathbf{y}) + \alpha \Omega(\mathbf{h}) \quad \text{where} \quad \alpha \in [0, \infty) \quad (7.48)$$

Adversarial Training Even for networks that perform at human level accuracy have a nearly 100 percent error rate on examples that are intentionally constructed to search for an input \mathbf{x}' near a data point \mathbf{x} such that the model output for \mathbf{x}' is very different than the output for \mathbf{x} .

In many cases, \mathbf{x}' can be so similar to \mathbf{x} that a human cannot tell the difference!

$$\mathbf{x}' \leftarrow \mathbf{x} + \epsilon \cdot \text{sign}(\nabla_{\mathbf{x}} J(\boldsymbol{\theta}; \mathbf{x}, \mathbf{y})) \quad (58)$$

In the context of regularization, one can reduce the error rate on the original i.i.d. test set via **adversarial training** – training on adversarially perturbed training examples.

Optimization for Training Deep Models (Ch. 8)

Table of Contents Local

Written by Brandon McKinzie

Empirical Risk Minimization. The ultimate goal of any machine learning algorithm is to reduce the expected generalization error, also called the **risk**:

$$J^*(\theta) = \mathbb{E}_{(\mathbf{x}, y) \sim p_{data}} [L(f(\mathbf{x}; \theta), y)] \quad (59) \quad \text{risk definitions}$$

with emphasis that the risk is over the *true* underlying data distribution p_{data} . If we knew p_{data} , this would be an optimization problem. Since we don't, and only have a set of training samples, it is a machine learning problem. However, we can still just minimize the **empirical risk**, replacing p_{data} in the equation above with \hat{p}_{data} ¹⁰.

So, how is minimizing the empirical risk any different than familiar gradient descent approaches? Aren't they designed to do just that? Well, sort of, but it's technically not the same. When we say "minimize the empirical risk" in the context of optimization, we mean this very literally. Gradient descent methods emphatically do *not* just go and *set* the weights to values such that the empirical risk reaches its lowest possible value – that's not machine learning. Furthermore, many useful loss function such as 0-1 loss¹¹ do not have useful derivatives. ERM ≠ GD

Surrogate Loss Functions and Early Stopping. In cases such as 0-1 loss, where minimization is intractable, one typically optimizes a **surrogate loss function** instead, such as the negative log-likelihood of the correct class. Also, an important difference between pure optimization and our training algorithms is that the latter usually don't halt at a local minimum. Instead, we usually must define some early stopping condition to terminate training before overfitting begins to occur.

We want to minimize the risk, but we don't have access to p_{data} , so ...

We want to minimize the empirical risk, but it's prone to overfitting and our loss function's derivative may be zero/undefined, so ...

We minimize a surrogate loss function iteratively over minibatches until early stopping is triggered.

¹⁰This amount to a simple average over the loss function at each training point.

¹¹The 0-1 loss function is defined as

$$L(\hat{y}, y) = I(\hat{y} \neq y) \quad (60)$$

Batch and Minibatch Algorithms. Computing $\nabla_{\theta} J(\theta)$ as an expectation over the entire training set is expensive, so we typically compute the expectation over a small subset of the examples. Recall that the standard deviation, or standard error $SE(\mu_m)$, of the mean taken over some subset of $m \leq n$ samples, $\mu_m = \frac{1}{m} \sum_{i \sim \text{Rand}(0, n, \text{size}=m)} x^{(i)}$, is given by σ/\sqrt{m} , where σ is the true [sample] standard deviation of the full n data samples. In other words, to improve such a gradient by a factor of 10 requires 100 times more samples-per-batch (and thus 100 times more computation). For this reason, most optimization algorithms actually *converge* much faster if they can rapidly compute approximate estimates of the gradient (re: smaller batches) rather than slowly computing the exact gradient.

The key points to consider when choosing your batch size:

1. Larger batches = more accurate estimates of the gradient, but with less than linear returns.
2. If examples in the batch are processed in parallel (as is typical), then memory roughly scales with batch size.
3. Small batches can offer a regularizing effect. Generalization error is often best for a batch size of 1. However, this requires a low learning rate to maintain stability and thus a longer overall training runtime.

Also, note that online SGD, where we never reuse data points, but simply update parameters as new data comes in, gives an unbiased estimator of the exact gradient of the generalization error (the risk). Once data samples are reused (e.g. when training with multiple epochs), the gradient estimates become biased. The interesting point here is that the availability of increasingly massive datasets is making single-epoch¹² training more common. In such cases, *overfitting is no longer an issue*, but rather underfitting and computational efficiency.

Ill-conditioning of the Hessian matrix \mathbf{H} can cause SGD to get “stuck” in the sense that even very small steps increase the cost function. Recall that a second-order Taylor series expansion of the cost function predicts that an SGD step of $-\epsilon \mathbf{g}$ will add

$$\frac{1}{2} \epsilon^2 \mathbf{g}^T \mathbf{H} \mathbf{g} - \epsilon \mathbf{g}^T \mathbf{g} \quad (61)$$

to the cost. If \mathbf{H} has a large condition number (re: if \mathbf{H} is ill-conditioned), then the range of possible values, $[1/\lambda_{\max}, 1/\lambda_{\min}]$, for $\mathbf{g}^T \mathbf{H} \mathbf{g}$ can become very large. In particular, if $\mathbf{g}^T \mathbf{H} \mathbf{g}$ exceeds $\epsilon \mathbf{g}^T \mathbf{g}$, then the SGD step will *increase* the cost!

¹²Or even less, i.e. not using all of the training data.

Training algorithms. Below, I list some popular training algorithms and their update equations.

- **SGD.**

$$\mathbf{g} \leftarrow \frac{1}{m} \nabla_{\boldsymbol{\theta}} \sum_i^m L\left(f(\mathbf{x}^{(i)}; \boldsymbol{\theta}), \mathbf{y}^{(i)}\right) \quad (62)$$

$$\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} - \epsilon \mathbf{g} \quad (63)$$

- **Momentum.**

$$\mathbf{g} \leftarrow \frac{1}{m} \nabla_{\boldsymbol{\theta}} \sum_i^m L\left(f(\mathbf{x}^{(i)}; \boldsymbol{\theta}), \mathbf{y}^{(i)}\right) \quad (64)$$

$$\mathbf{v} \leftarrow \alpha \mathbf{v} - \epsilon \mathbf{g} \quad (65)$$

$$\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \mathbf{v} \quad (66)$$

- **Nesterov Momentum.** Gradient computations instead evaluated after current velocity is applied.

$$\mathbf{g} \leftarrow \frac{1}{m} \nabla_{\boldsymbol{\theta}} \sum_i^m L\left(f(\mathbf{x}^{(i)}; \boldsymbol{\theta} + \alpha \mathbf{v}), \mathbf{y}^{(i)}\right) \quad (67)$$

$$\mathbf{v} \leftarrow \alpha \mathbf{v} - \epsilon \mathbf{g} \quad (68)$$

$$\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \mathbf{v} \quad (69)$$

- **AdaGrad.** Different learning rate for each model parameter. Individually adapts the learning rates of all model parameters by scaling them inversely proportional to the square root of the sum of all historical squared values of the gradient. Empirically, can result in premature and excessive decrease in the effective learning rate.

$$\mathbf{g} \leftarrow \frac{1}{m} \nabla_{\boldsymbol{\theta}} \sum_i^m L\left(f(\mathbf{x}^{(i)}; \boldsymbol{\theta}), \mathbf{y}^{(i)}\right) \quad (70)$$

$$\mathbf{r} \leftarrow \mathbf{r} + \mathbf{g} \odot \mathbf{g} \quad (71)$$

$$\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} - \frac{\epsilon}{\delta + \sqrt{\mathbf{r}}} \odot \mathbf{g} \quad (72)$$

where the gradient accumulation variable \mathbf{r} is initialized to the zero vector, and the fraction and square root in the last equation is applied element-wise.

- **RMSProp.** Modifies AdaGrad by changing the gradient accumulation into an exponentially weighted moving average.

$$\mathbf{g} \leftarrow \frac{1}{m} \nabla_{\boldsymbol{\theta}} \sum_i^m L\left(f(\mathbf{x}^{(i)}; \boldsymbol{\theta}), \mathbf{y}^{(i)}\right) \quad (73)$$

$$\mathbf{r} \leftarrow \rho \mathbf{r} + (1 - \rho) \mathbf{g} \odot \mathbf{g} \quad (74)$$

$$\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} - \frac{\epsilon}{\delta + \sqrt{\mathbf{r}}} \odot \mathbf{g} \quad (75)$$

It is also common to modify RMSProp to use Nesterov momentum.

- **Adam.** So-named to mean “adaptive moments.” We now call \mathbf{r} the 2nd moment (variance) variable, and introduce \mathbf{s} as the 1st moment (mean) variable, where the moments are for the [true] gradient; the new variables act as estimates of the moments [since we estimate the gradient with a simple average over a minibatch]. Note that these moments are uncentered.

$$\mathbf{g} \leftarrow \frac{1}{m} \nabla_{\boldsymbol{\theta}} \sum_i^m L\left(f(\mathbf{x}^{(i)}; \boldsymbol{\theta}), \mathbf{y}^{(i)}\right) \quad (76)$$

$$\mathbf{s} \leftarrow \frac{1}{1 - \rho_1^{t-1}} [\rho_1 \mathbf{s} + (1 - \rho_1) \mathbf{g}] \quad (77)$$

$$\mathbf{r} \leftarrow \frac{1}{1 - \rho_2^{t-1}} [\rho_2 \mathbf{r} + (1 - \rho_2) \mathbf{g} \odot \mathbf{g}] \quad (78)$$

$$\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} - \frac{\epsilon}{\delta + \sqrt{\mathbf{r}}} \odot \mathbf{s} \quad (79)$$

where the factors proportional to the ρ values serve to correct for bias in the moment estimators.


Convolutional Neural Networks (Ch. 9)

We use a 2-D image I as our input (and therefore require a 2-D kernel K). Note that most neural networks do not technically implement convolution¹³, but instead implement a related function called the *cross-correlation*, defined as

$$S(i, j) = (I * K)(i, j) = \sum_m \sum_n I(i + m, j + n) K(m, n) \quad (9.6)$$

Convolution leverages the following three important ideas:

- **Sparse interactions[/connectivity/weights].** Individual input units only interact/connect with a subset of the output units. Accomplished by making the kernel smaller than the input. It's important to recognize that the receptive field of the units in the deeper layers of a convolutional network is *larger* than the receptive field of the units in the shallow layers, as seen below.



- **Parameter sharing.**
- **Equivariance** (to translation). Changes in inputs [to a function] cause output to change in the same way. Specifically, f is equivariant to g if $f(g(x)) = g(f(x))$. For convolution, g would be some function that translates the input.

¹³Technically the convolution output is defined as

$$S(i, j) = (I * K)(i, j) = \sum_m \sum_n I(m, n) K(i - m, j - n) \quad (9.4)$$

$$= (K * I)(i, j) = \sum_m \sum_n I(i - m, j - n) K(m, n) \quad (9.5)$$

where 9.5 can be asserted due to commutativity of convolution.

Pooling. Helps make the representation approximately **invariant** to small translations of the input. The use of pooling can be viewed as adding an infinitely strong prior¹⁴ that the function the layer learns must be invariant to small translations.

Additional common tricks¹⁵.

- **Local Response Normalization (LRN)**¹⁶. Purpose is to aid generalization ability. Let $a_{x,y}^i$ denote the activity of a neuron computed by applying kernel i at position (x, y) and then applying the ReLU nonlinearity. The response-normalized activity $b_{x,y}^i$ is given by the expression

$$b_{x,y}^i = \frac{a_{x,y}^i}{\left(k + \alpha \sum_j \left(a_{x,y}^j\right)^2\right)^\beta} \quad (80)$$

where j runs from $[i - n/2]_+$ to $\min(N - 1, i + n/2)$, and N is the total number of kernels in the given layer¹⁷. Authors used $k = 2$, $n = 5$, $\alpha = 10^{-4}$, $\beta = 0.75$.

- **Batch Normalization**¹⁸. BN Allows us to use much higher learning rates and be less careful about initialization. Algorithm defined in image below, where each element of the batch, $x_i \equiv x_i^{(k)}$ (where we drop the k for notational simplicity), represents the k th activation output from the previous layer [for the i th sample in the batch] and about to be fed as input to the current layer.

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_{1\dots m}\}$;	
Parameters to be learned: γ, β	
Output: $\{y_i = \text{BN}_{\gamma,\beta}(x_i)\}$	
$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i$	// mini-batch mean
$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2$	// mini-batch variance
$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}}$	// normalize
$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma,\beta}(x_i)$	// scale and shift

Note that one can model each layer's activations as arising from some distribution. When we feed data to a network, we model the data as coming from some data-generating

¹⁴Where the distribution of this prior is over all possible functions learnable by the model.

¹⁵Collected on my own. In other words, not from the deep learning book, but rather a bunch of disconnected resources over time.

¹⁶From section 3.3 of Krizhevsky et al. (2012). AlexNet paper.

¹⁷In other words, the summation is over the adjacent kernel maps, with [total] window size n (manually chosen). The min/max just says $n/2$ to the left (right) unless that would be past the leftmost (rightmost) kernel map.

¹⁸From "Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift" by Ioffe et al.

distribution. Similarly, we can model the activations that occur when feeding the data through as coming from some activation-layer-generating distribution. The problem with this model is that the process of updating our weights during training changes the distribution of activations for each layer, which can make the learning task more difficult. Batch normalization's goal is to reduce this *internal covariate shift*, and is motivated by the practice of normalizing our data to have zero mean and unit variance.

Intuition of some math.

- **Q:** How to intuitively understand the commutativity of convolution?
 - **A:** You must first realize that, *independent of which formula we're thinking of*, as one index into either the image or kernel increases (decreases), the other decreases (increases); they increment in opposite directions. The difference between the two formulas (9.4 and 9.5 in textbook), is just that we start at different ends of the summations (when you actually substitute in the numbers). Note that this doesn't require any symmetry on the boundaries of the summation about zero; it truly is a property of the convolution.
- **Q:** What do the authors mean by "we have flipped the kernel"?
 - **A:** Not much, and it's poor wording. They didn't *do* anything, that is just part of the definition of the convolution. They literally just mean that the convolution has the property that as one index increases, the other decreases (see previous answer). The cross-correlation, however, has the property that as one index increases, the other increases, too.

Sequence Modeling (RNNs) (Ch. 10)

Table of Contents Local

Written by Brandon McKinzie

REVIEW: THE BASICS OF RNNs

Notation/Architecture Used.

- **U**: input \rightarrow hidden.
- **W**: hidden \rightarrow hidden.
- **V**: hidden \rightarrow output.
- **Activations**: tanh [hidden] and softmax [after output].
- **Misc. Details**: $\mathbf{x}^{(t)}$ is a *vector* of inputs fed at time t . Recall that RNNs can be unfolded for any desired number of steps τ . For example, if $\tau = 3$, the general functional representation output of an RNN is $\mathbf{s}^{(3)} = f(\mathbf{s}^{(2)}; \boldsymbol{\theta}) = f(f(\mathbf{s}^{(1)}; \boldsymbol{\theta}); \boldsymbol{\theta})$. Typical RNNs read information out of the state \mathbf{h} to make predictions.

Shape of $\mathbf{x}^{(t)}$ fixed, e.g. vocab length.Black square on recurrent connection \equiv interaction w/delay of a single time step.

Forward Propagation & Loss. Specify initial state $\mathbf{h}^{(0)}$. Then, for each time step from $t = 1$ to $t = \tau$, feed input sequence $\mathbf{x}^{(t)}$ and compute the output sequence $\mathbf{o}^{(t)}$. To determine the loss at each time-step, $L^{(t)}$, we compare $\text{softmax}(\mathbf{o}^{(t)})$ with (one-hot) $\mathbf{y}^{(t)}$.

$$\mathbf{h}^{(t)} = \tanh(\mathbf{a}^{(t)}) \quad \text{where} \quad \mathbf{a}^{(t)} = \mathbf{b} + \mathbf{W}\mathbf{h}^{(t-1)} + \mathbf{U}\mathbf{x}^{(t)} \quad (10.9/8)$$

$$\hat{\mathbf{y}}^{(t)} = \text{softmax}(\mathbf{o}^{(t)}) \quad \text{where} \quad \mathbf{o}^{(t)} = \mathbf{c} + \mathbf{V}\mathbf{h}^{(t)} \quad (10.11/10)$$

Note that this is an example of an RNN that maps input seqs to output seqs of the same length¹⁹. We can then compute, e.g., the log-likelihood loss $L = \sum_t L^{(t)}$ over all time steps as:

$$L = - \sum_t \log \left(p_{\text{model}} \left[\mathbf{y}^{(t)} \mid \{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(t)}\} \right] \right) \quad (10.12/13/14)$$

Convince yourself this is identical to cross-entropy.

¹⁹Where “same length” is related to the number of timesteps (i.e. τ input steps means τ output steps), not anything about the actual shapes/sizes of each individual input/output.

where $y^{(t)}$ is the **ground-truth** (one-hot vector) at time t , whose probability of occurring is given by the corresponding element of $\hat{\mathbf{y}}^{(t)}$

Back-Propagation Through Time.

1. **Internal-Node Gradients.** In what follows, when considering what is included in the chain rule(s) for gradients with respect to a node \mathbf{N} , just need to consider paths from it [through its **descendents**] to loss node(s).

- **Output nodes.** For any given time t , the node $\mathbf{o}^{(t)}$ has only one direct descendant, the loss node $L^{(t)}$. Since no other loss nodes can be reached from $\mathbf{o}^{(t)}$, it is the only one we need consider in the gradient.

$$\begin{aligned}
\left(\nabla_{\mathbf{o}^{(t)}} L\right)_i &= \frac{\partial L}{\partial \mathbf{o}_i^{(t)}} \\
&= \frac{\partial L}{\partial L^{(t)}} \cdot \frac{\partial L^{(t)}}{\partial \mathbf{o}_i^{(t)}} \\
&= (1) \cdot \frac{\partial L^{(t)}}{\partial \mathbf{o}_i^{(t)}} \\
&= \frac{\partial}{\partial \mathbf{o}_i^{(t)}} \left\{ -\log \left(\hat{\mathbf{y}}_{y^{(t)}}^{(t)} \right) \right\}
\end{aligned} \tag{81}$$

Ground-truth $y^{(t)}$ here is a **scalar**, interpreted as the index of the correct label of output vector.

$$\begin{aligned}
&= -\frac{\partial}{\partial \mathbf{o}_i^{(t)}} \left\{ \log \left(\frac{e^{\mathbf{o}_{y^{(t)}}^{(t)}}}{\sum_j e^{\mathbf{o}_j^{(t)}}} \right) \right\} \\
&= -\frac{\partial}{\partial \mathbf{o}_i^{(t)}} \left\{ \mathbf{o}_{y^{(t)}}^{(t)} - \log \left(\sum_j e^{\mathbf{o}_j^{(t)}} \right) \right\} \\
&= -\left\{ \mathbf{1}_{i,y^{(t)}} - \frac{\partial}{\partial \mathbf{o}_i^{(t)}} \log \left(\sum_j e^{\mathbf{o}_j^{(t)}} \right) \right\} \\
&= -\left\{ \mathbf{1}_{i,y^{(t)}} - \frac{1}{\sum_j e^{\mathbf{o}_j^{(t)}}} \frac{\partial \sum_j e^{\mathbf{o}_j^{(t)}}}{\partial \mathbf{o}_i^{(t)}} \right\} \\
&= -\left\{ \mathbf{1}_{i,y^{(t)}} - \frac{e^{\mathbf{o}_i^{(t)}}}{\sum_j e^{\mathbf{o}_j^{(t)}}} \right\} \\
&= -\left\{ \mathbf{1}_{i,y^{(t)}} - \hat{\mathbf{y}}_i^{(t)} \right\} \\
&= \hat{\mathbf{y}}_i^{(t)} - \mathbf{1}_{i,y^{(t)}}
\end{aligned} \tag{10.18}$$

$$\mathbf{1}_{i,y^{(t)}} = \begin{cases} 1 & y^{(t)} = i \\ 0 & \text{otherwise} \end{cases}$$

which leaves all entries of $\mathbf{o}^{(t)}$ unchanged *except* for the entry corresponding to the true label, which will become negative in the gradient. All this means is, since we

want to increase the probability of this entry, driving this value up will *decrease* the loss (hence negative) and driving any other entries up will *increase* the loss proportional to its current estimated probability (driving up an [incorrect] entry that is already high is “worse” than driving up a small [incorrect entry]).

- **Hidden nodes.** First, consider the simplest hidden node to take the gradient of, the last one, $\mathbf{h}^{(\tau)}$ (simplest because only one descendant [path] reaching any loss node(s)).

$$\begin{aligned}
\left(\nabla_{\mathbf{h}^{(\tau)}} L\right)_i &= \frac{\partial L}{\partial L^{(\tau)}} \sum_{k=1}^{n_{out}} \frac{\partial L^{(\tau)}}{\partial \mathbf{o}_k^{(\tau)}} \frac{\partial \mathbf{o}_k^{(\tau)}}{\partial \mathbf{h}_i^{(\tau)}} \\
&= \sum_{k=1}^{n_{out}} \left(\nabla_{\mathbf{o}^{(\tau)}} L\right)_k \frac{\partial \mathbf{o}_k^{(\tau)}}{\partial \mathbf{h}_i^{(\tau)}} \\
&= \sum_{k=1}^{n_{out}} \left(\nabla_{\mathbf{o}^{(\tau)}} L\right)_k \frac{\partial}{\partial \mathbf{h}_i^{(\tau)}} \left\{ c_k + \sum_{j=1}^{n_{hid}} V_{kj} \mathbf{h}_j^{(\tau)} \right\} \\
&= \sum_{k=1}^{n_{out}} \left(\nabla_{\mathbf{o}^{(\tau)}} L\right)_k V_{ki} \\
&= \sum_{k=1}^{n_{out}} (V^T)_{ik} \left(\nabla_{\mathbf{o}^{(\tau)}} L\right)_k \\
&= \left(V^T \nabla_{\mathbf{o}^{(\tau)}} L\right)_i
\end{aligned} \tag{10.19}$$

Before proceeding, **notice the following useful pattern:** If two nodes a and b , each containing n_a and n_b neurons, are fully connected by parameter matrix $W_{n_b \times n_a}$ and directed like $a \rightarrow b \rightarrow L$, then²⁰ $\nabla_a L = W^T \nabla_b L$. Using this result, we can then iterate and take gradients back in time from $t = \tau - 1$ to $t = 1$ as follows:

$$\nabla_{\mathbf{h}^{(t)}} L = \left(\frac{\partial \mathbf{h}^{(t+1)}}{\partial \mathbf{h}^{(t)}} \right)^T \left(\nabla_{\mathbf{h}^{(t+1)}} L \right) + \left(\frac{\partial \mathbf{o}^{(t)}}{\partial \mathbf{h}^{(t)}} \right)^T \left(\nabla_{\mathbf{o}^{(t)}} L \right) \tag{10.20}$$

$$\begin{aligned}
&= W^T \left(\nabla_{\mathbf{h}^{(t+1)}} L \right) \text{diag}(1 - \tanh^2(\mathbf{a}^{(t+1)})) + V^T \left(\nabla_{\mathbf{o}^{(t)}} L \right) \\
&= W^T \left(\nabla_{\mathbf{h}^{(t+1)}} L \right) \text{diag}(1 - (\mathbf{h}^{(t+1)})^2) + V^T \left(\nabla_{\mathbf{o}^{(t)}} L \right)
\end{aligned} \tag{10.21}$$

$$\begin{aligned}
\frac{d}{dx} \tanh(x) &= 1 - \tanh^2(x) \\
(\text{diag}(\mathbf{a}))_{ii} &\triangleq a_i
\end{aligned}$$

2. Parameter Gradients. Now we can compute the gradients for the parameter matrices/vectors, where it is crucial to remember that a given parameter matrix (e.g. U) is shared across *all* time steps t . We can treat tensor derivatives in the same form as

²⁰More generally,

$$\nabla_a L = \left(\frac{\partial \mathbf{b}}{\partial \mathbf{a}} \right)^T \nabla_b L$$

which is a good example of how vector derivatives map into a matrix. For example, let $\mathbf{a} \in \mathbb{R}^{n_a}$ and $\mathbf{b} \in \mathbb{R}^{n_b}$. Then

$$\frac{\partial \mathbf{b}}{\partial \mathbf{a}} \in \mathbb{R}^{n_b \times n_a}$$

previously done with vectors after a quick abstraction: For any tensor \mathbf{X} of arbitrary rank (e.g. if rank-4 then index like \mathbf{X}_{ijkl}), use single variable (e.g. i) to represent the complete tuple of indices²¹.

- **Bias parameters [vectors]**. These are nothing new, since just vectors.

$$\begin{aligned} (\nabla_{\mathbf{c}} L) &= \sum_t \left(\frac{\partial \mathbf{o}^{(t)}}{\partial \mathbf{c}^{(t)}} \right)^T (\nabla_{\mathbf{o}^{(t)}} L) \\ &= \sum_t (\nabla_{\mathbf{o}^{(t)}} L) \end{aligned} \quad (10.22)$$

$$\begin{aligned} (\nabla_{\mathbf{c}} L) &= \sum_t \left(\frac{\partial \mathbf{h}^{(t)}}{\partial \mathbf{b}^{(t)}} \right)^T (\nabla_{\mathbf{h}^{(t)}} L) \\ &= \sum_t \text{diag} \left(1 - (\mathbf{h}^{(t)})^2 \right) (\nabla_{\mathbf{h}^{(t)}} L) \end{aligned} \quad (10.23)$$

- **V** ($n_{out} \times n_{hid}$).

$$\nabla_{\mathbf{V}} L = \sum_t \nabla_{\mathbf{V}} L^{(t)} \quad (82a)$$

$$= \sum_t \nabla_{\mathbf{V}} L^{(t)}(\mathbf{o}_1^{(t)}, \dots, \mathbf{o}_{n_{out}}^{(t)}) \quad (82b)$$

$$= \sum_t \sum_i^{n_{out}} (\nabla_{\mathbf{o}^{(t)}} L)_i \nabla_{\mathbf{V}} \mathbf{o}_i^{(t)} \quad (82c)$$

$$= \sum_t \sum_i^{n_{out}} (\nabla_{\mathbf{o}^{(t)}} L)_i \nabla_{\mathbf{V}} \left\{ c_i + \sum_{j=1}^{n_{hid}} V_{ij} \mathbf{h}_j^{(t)} \right\} \quad (82d)$$

$$= \sum_t \sum_i^{n_{out}} (\nabla_{\mathbf{o}^{(t)}} L)_i \begin{bmatrix} 0 & 0 & \dots & 0 \\ \vdots & \vdots & \dots & \vdots \\ \mathbf{h}_1^{(t)} & \mathbf{h}_2^{(t)} & \dots & \mathbf{h}_{n_{hid}}^{(t)} \\ \vdots & \vdots & \dots & \vdots \\ 0 & 0 & \dots & 0 \end{bmatrix} \quad (82e)$$

$$= \sum_t (\nabla_{\mathbf{o}^{(t)}} L) (\mathbf{h}^{(t)})^T \quad (82f)$$

²¹More details on tensor derivatives: Consider the chain defined by $\mathbf{Y} = g(\mathbf{X})$, and $z = f(\mathbf{Y})$, where z is some vector. Then

$$\nabla_{\mathbf{X}} z = \sum_j (\nabla_{\mathbf{X}} Y_j) \frac{\partial z}{\partial Y_j}$$

where if 82e confuses you, see the footnote²².

- \mathbf{W} ($n_{hid} \times n_{hid}$). This one is a bit odd, since \mathbf{W} is, in a sense, even more “shared” across time steps than \mathbf{V} ²³. The authors here define/choose, when evaluating $\nabla_{\mathbf{W}} h_i^{(t)}$ to only concern themselves with $\mathbf{W} := \mathbf{W}^{(t)}$, i.e. the direct connections to \mathbf{h} at time t .

$$\nabla_{\mathbf{W}} L = \sum_t^\tau \nabla_{\mathbf{W}} L^{(t)} \quad (84a)$$

$$= \sum_t^\tau \sum_i^{n_{hid}} \left(\nabla_{\mathbf{h}^{(t)}} L \right)_i \nabla_{\mathbf{W}^{(t)}} \mathbf{h}_i^{(t)} \quad (10.25)$$

$$= \sum_t^\tau \sum_i^{n_{hid}} \left(\nabla_{\mathbf{h}^{(t)}} L \right)_i \left(\text{diag} \left(1 - (\mathbf{h}^{(t)})^2 \right) \begin{bmatrix} 0 & 0 & \dots & 0 \\ \vdots & \vdots & \dots & \vdots \\ \mathbf{h}_1^{(t-1)} & \mathbf{h}_2^{(t-1)} & \dots & \mathbf{h}_{n_{hid}}^{(t-1)} \\ \vdots & \vdots & \dots & \vdots \\ 0 & 0 & \dots & 0 \end{bmatrix} \right) \quad (84b)$$

$$= \sum_t^\tau \text{diag} \left(1 - (\mathbf{h}^{(t)})^2 \right) \left(\nabla_{\mathbf{h}^{(t)}} L \right) (\mathbf{h}^{(t-1)})^T \quad (10.26)$$

- \mathbf{U} ($n_{hid} \times n_{in}$). Very similar to the previous calculation.

$$\nabla_{\mathbf{U}} L = \sum_t^\tau \nabla_{\mathbf{U}} L^{(t)} \quad (85a)$$

$$= \sum_t^\tau \sum_i^{n_{hid}} \left(\nabla_{\mathbf{h}^{(t)}} L \right)_i \nabla_{\mathbf{U}^{(t)}} \mathbf{h}_i^{(t)} \quad (10.27)$$

$$= \sum_t^\tau \text{diag} \left(1 - (\mathbf{h}^{(t)})^2 \right) \left(\nabla_{\mathbf{h}^{(t)}} L \right) (\mathbf{x}^{(t)})^T \quad (10.28)$$

²² The general lesson learned here is that, for some matrix $\mathbf{W} \in \mathbb{R}^{a \times b}$ and vector $\mathbf{x} \in \mathbb{R}^b$,

$$\sum_i \nabla_{\mathbf{W}} [(\mathbf{W}\mathbf{x})_i] = \begin{bmatrix} \mathbf{x}^T \\ \mathbf{x}^T \\ \vdots \\ \mathbf{x}^T \end{bmatrix} \quad (83)$$

where, of course, the output has the same dimensions as \mathbf{W} .

²³Specifically, $\mathbf{h}^{(t)}$ is both

- An explicit function of the parameter matrix $\mathbf{W}^{(t)}$ directly feeding into it.
- An implicit function of all other $\mathbf{W}^{t=i}$ that came before.

This is different than before, where we had $\mathbf{o}^{(t)}$ not implicitly depending on earlier $\mathbf{V}^{(t=i)}$. In other words, $\mathbf{h}^{(t)}$ is a descendant of all earlier (and current) \mathbf{W} .

RNNs AS DIRECTED GRAPHICAL MODELS

The advantage of RNNs is their efficient parameterization of the joint distribution over $\mathbf{y}^{(i)}$ via parameter sharing. This introduces a built-in assumption that we can model the effect of $y^{(i)}$ in the distant past on the current $y^{(t)}$ *via its effect on \mathbf{h}* . We are also assuming that the conditional probability distribution over the variables at $t + 1$ given the variables at time t is **stationary**. Next, we want to know how to draw *samples* from such a model. Specifically, how to sample from the conditional distribution ($y^{(t)}$ given $y^{(t-1)}$) at each time step.

Say we want to model a sequence of scalar random variables $\mathbb{Y} \triangleq \{y^{(1)}, \dots, y^{(\tau)}\}$ for some sequence length τ . Without making independence assumptions just yet, we can parameterize the joint distribution $P(\mathbb{Y})$ with basic definitions of probability:

$$P(\mathbb{Y}) \triangleq P(y^{(1)}, \dots, y^{(\tau)}) = \prod_{t=1}^{\tau} P(y^{(t)} \mid y^{(t-1)}, \dots, y^{(1)}) \quad (86)$$

where I've drawn an example of the *complete graph* for $\tau = 4$ below.



The complete graph can represent the direct dependencies between any pairs of y values.

If each value y could take on the same fixed set of k values, we would need to learn k^4 parameters to represent the joint distribution $P(\mathbb{Y})$. This is clearly inefficient, since the number of parameters needed scales like $\mathcal{O}(k^\tau)$. If we relax the restriction that each $y^{(i)}$ must depend *directly* on all past $y^{(j)}$, we can considerably reduce the number of parameters needed to compute the probability of some particular sequence.

We could include latent variables \mathbf{h} at each timestep that capture the dependencies, reminiscent of a classic RNN:



Since in the RNN case all factors $P(h^{(t)} | h^{(t-1)})$ are deterministic, we don't need any additional parameters to compute this probability²⁴, other than the single m^2 parameters needed to convert any $h^{(t)}$ to the next $h^{(t+1)}$ (which is shared across all transitions). Now, the number of parameters needed as a function of sequence length is constant, and as a function of k is just $\mathcal{O}(k)$.

Finally, to view the RNN as a graphical model, we must describe how to sample from it, namely how to sample a sequence \mathbf{y} from $P(\mathbb{Y})$, if parameterized by our graphical model above. In the general case where we don't know the value of τ for our sequence \mathbf{y} , one approach is to have a EOS symbol that, if found during sampling, means we should stop there. Also, in the typical case where we actually want to model $P(y | x)$ for input sequence x , we can reinterpret the parameters θ of our graphical model as a function of \mathbf{x} the input sequence. In other words, the graphical model interpretation becomes a function of \mathbf{x} , where \mathbf{x} determines the exact values of the probabilities the graphical model takes on – an “instance” of the graphical model.

Bidirectional RNNs. In many applications, it is desirable to output a prediction of $\mathbf{y}^{(t)}$ that may depend on *the whole sequence*. For example, in speech recognition, the interpretation of words/sentences can also depend on what is *about* to be said. Below is a typical bidirectional RNN, where the inputs $\mathbf{x}^{(t)}$ are fed both to a “forward” RNN (\mathbf{h}) and a “backward” RNN (\mathbf{g}).



Notice how the output units $\mathbf{o}^{(t)}$ have the nice property of depending on both the past and future while being most sensitive to input values around time t .

²⁴Don't forget that, in a neural net, a variable $y^{(t)}$ is represented by a *layer*, which itself is composed of k nodes, each associated with one of the k unique values that $y^{(t)}$ could be.

Encoder-Decoder Seq2Seq Architectures (10.4) Here we discuss how an RNN can be trained to map an input sequence to output sequence which is not necessarily the same length. (Not really much of a discussion...figure below says everything.)



CHALLENGE OF LONG-TERM DEPS. (10.7)

Gradients propagated over many stages either vanish (usually) or explode. We saw how this could occur when we took parameter gradients earlier, and for weight matrices \mathbf{W} further along from the loss node, the expression for $\nabla_{\mathbf{W}} L$ contained multiplicative Jacobian factors. Consider the (linear activation) repeated function composition of an RNN's hidden state in 10.36. We can rewrite it as a power method (10.37), and if \mathbf{W} admits an eigendecomposition (remember \mathbf{W} is necessarily square here), we can further simplify as seen in 10.38.

$$\mathbf{h}^{(t)} = \mathbf{W}^T \mathbf{h}^{(t-1)} \quad (10.36)$$

$$= (\mathbf{W}^t)^T \mathbf{h}^{(0)} \quad (10.37)$$

$$= \mathbf{Q}^T \mathbf{\Lambda}^t \mathbf{Q} \mathbf{h}^{(0)} \quad (10.38)$$

Q: Explain interp. of mult. \mathbf{h} by \mathbf{Q} as opposed to the usual \mathbf{Q}^T explained in the linear algebra review.

Any component of $\mathbf{h}^{(0)}$ that isn't aligned with the largest eigenvector will eventually be discarded.²⁵

If, however, we have a non-recurrent network such that the state elements are repeatedly multiplied by different $w^{(t)}$ at each time step, the situation is different. Suppose the different $w^{(t)}$ are i.i.d. with mean 0 and variance v . The variance of the product is easily seen to

²⁵Make sure to think about this from the right perspective. The largest value of $t = \tau$ in the RNNs we've seen would correspond with either (1) the largest output sequence or (2) the largest input sequence (if fixed-vector output). After we extract the output from a given forward pass, we reset the clock and either back-propagate errors (if training) or get ready to feed another sequence.

be $\mathcal{O}(v^n)^{26}$. To obtain some desired variance v^* we may choose the individual weights with variance $v = \sqrt[n]{v^*}$.

LSTMs AND OTHER GATED RNNs (10.10)

While leaky units have connection weights that are either manually chosen constants or are trainable parameters, gated RNNs generalize this to connection weights that may change *at each time step*. Furthermore, gated RNNs can learn to both accumulate and *forget*, while leaky units are designed for just accumulation²⁷

LSTM (10.10.1). The idea is we want self-loops to produce paths where the gradient can flow for long durations. The self-loop weights are **gated**, meaning they are controlled by another hidden unit, interpreted as being conditioned on *context*. Listed below are the main components of the LSTM architecture.

- **Forget gate** $f_i^{(t)} = \sigma \left(b_i^f + \sum_j U_{i,j}^f x_j^{(t)} + \sum_j W_{i,j}^f h_j^{(t-1)} \right)$.
- **Internal state** $s_i^{(t)} = f_i^{(t)} \odot s_i^{(t-1)} + g_i^{(t)} \odot \sigma \left(b_i + \sum_j U_{i,j} x_j^{(t)} + \sum_j W_{i,j} h_j^{(t-1)} \right)$.
- **External input gate** $g_i^{(t)} = \sigma \left(b_i^g + \sum_j U_{i,j}^g x_j^{(t)} + \sum_j W_{i,j}^g h_j^{(t-1)} \right)$.
- **Output gate** $q_i^{(t)} = \sigma \left(b_i^o + \sum_j U_{i,j}^o x_j^{(t)} + \sum_j W_{i,j}^o h_j^{(t-1)} \right)$.

The subscript, i , identifies the cell. The superscript, t , denotes the time.

The final hidden state can then be computed via

$$h_i^{(t)} = \tanh(s_i^{(t)}) \odot q_i^{(t)} \quad (89)$$

²⁶Quick sketch of (my) proof:

$$\text{Var} [w^{(i)}] = v = \mathbb{E} [(w^{(i)})^2] - \cancel{\mathbb{E} [w^{(i)}]^2} \quad (87)$$

$$\text{Var} \left[\prod_t^n w^{(t)} \right] = \mathbb{E} \left[\left(\prod_t^n w^{(t)} \right)^2 \right] = \prod_t^n \mathbb{E} [(w^{(t)})^2] = v^n \quad (88)$$

²⁷Q: Isn't choosing to update with higher relative weight on the present the same as forgetting? A: Sort of. It's like "soft forgetting" and will inevitably erase more/less than desired (smeary). In this context, "forget" means to set the weight of a specific past cell to zero.

Applications (Ch. 12)

Table of Contents Local

Written by Brandon McKinzie

NATURAL LANGUAGE PROCESSING (12.4)

Begins on pg. 448

n-grams. A **language model** defines a probability distribution over sequences of [discrete] tokens (words/characters/etc). Early models were based on the *n-gram*: a [fixed-length] sequence of n tokens. Such models define the conditional distribution for the n th token, given the $(n - 1)$ previous tokens:

$$P(x_t \mid x_{t-(n-1)}, \dots, x_{t-1})$$

where x_i denotes the token at step/index/position i in the sequence.

To define distributions over longer sequences, we can just use Bayes rule over the shorter distributions, as usual. For example, say we want to find the [joint] distribution for some τ -gram ($\tau > n$), and we have access to an n -gram model and a [perhaps different] model for the initial sequence $P(x_1, \dots, x_{n-1})$. We compute the τ distribution simply as follows:

$$P(x_1, \dots, x_\tau) = P(x_1, \dots, x_{n-1}) \prod_{t=n}^{\tau} P(x_t \mid x_{t-1}, \dots, x_{t-(n-2)}, x_{t-(n-1)}) \quad (12.5)$$

where it's important to see that each factor in the product is a distribution over a length- n sequence. Since we need that initial factor, it is common to train both an n -gram model and an $n - 1$ -gram model simultaneously.

Let's do a specific example for a trigram ($n = 3$).

- **Assumptions [for this trigram model example]:**
 - For any $n \geq 3$, $P(x_n \mid x_1, \dots, x_{n-1}) = P(x_n \mid x_{n-2}, x_{n-1})$.
 - When we get to computing the full joint distribution over some sequence of arbitrary length, we assume we have access to both P_3 and P_2 , the joint distributions over all subsequences of length 3 and 2, respectively.
- **Example sequence:** We want to know how to use a trigram model on the sequence ['THE', 'DOG', 'RAN', 'AWAY'].

- **Derivation:** We can use the built-in model assumption to derive the following formula.

$$\begin{aligned}
P(\text{THE DOG RAN AWAY}) &= P_3(\text{AWAY} \mid \text{THE DOG RAN}) P_3(\text{THE DOG RAN}) \\
&= P_3(\text{AWAY} \mid \text{DOG RAN}) P_3(\text{THE DOG RAN}) \\
&= \frac{P_3(\text{DOG RAN AWAY})}{P_2(\text{DOG RAN})} P_3(\text{THE DOG RAN}) \\
&= P_3(\text{THE DOG RAN}) P_3(\text{DOG RAN AWAY}) / P_2(\text{DOG RAN})
\end{aligned} \tag{12.7}$$

Limitations of n-gram. The last example illustrates some potential problems one may encounter that arise [if using MLE] when the full joint we seek is nonzero, but (a) some P_n factor is zero, or (b) P_{n-1} is zero. Some methods of dealing with this are as follows.

Recall that, in MLE, the P_n and P_{n-1} are usually approximated via counting occurrences in the training set

- **Smoothing:** shifting probability mass from the observed tuples to unobserved ones that are similar.
- **Back-off methods:** look up the lower-order (lower values of n) n -grams if the frequency of the context $x_{t-1}, \dots, x_{t-(n-1)}$ is too small to use the higher-order model.

In addition, n -gram models are vulnerable to the curse of dimensionality, since most n -grams won't occur in the training set²⁸, even for modest n .

NEURAL LANGUAGE MODELS (12.4.2)

Designed to overcome curse of dimensionality by using a distributed representation of words. Recognize that any model trained on sentences of length n and then told to generalize to new sentences [also of length n] must deal with a space²⁹ of possible sentences that is exponential in n . Such word representations (i.e. viewing words as existing in some high-dimensional space) are often called **word embeddings**. The idea is to map the words (or sentences) from the raw high-dimensional [vocab sized] space to a smaller feature space, where similar words are closer to one another. Using distributed representations may also be used with graphical models (think Bayes' nets) in the form multiple *latent variables*.

²⁸For a given vocabulary, which usually has much more than n possible words, consider how many possible sequences of length n .

²⁹Ok I tried re-wording that from the book's confusing wording but that was also a bit confusing. Let me break it down. Say you train on a thousand sentences each of length 5. For a given vocabulary of size VOCAB_SIZE, the number of possible sequences of length 5 is $(\text{VOCAB_SIZE})^5$, which can be quite a lot more than a thousand (not to mention the possibility of duplicate training examples). To the naive model, all points in this high-dimensional space are basically the same. A neural language model, however, tries to arrange the space of possibilities in a meaningful way, so that an unforeseen sample at test time can be said "similar" as some previously seen training example. It does this by *embedding* words/sentences in a lower-dimensional feature space.

DEEP LEARNING RESEARCH

CONTENTS

3.1	Linear Factor Models (Ch. 13)	46
3.2	Autoencoders (Ch. 14)	49
3.3	Representation Learning (Ch. 15)	50
3.4	Structured Probabilistic Models for DL (Ch. 16)	51
3.4.1	Sampling from Graphical Models	53
3.4.2	Inference and Approximate Inference	53
3.5	Monte Carlo Methods (Ch. 17)	55
3.6	Approximate Inference (Ch. 19)	57

Linear Factor Models (Ch. 13)

Table of Contents Local

Written by Brandon McKinzie

Overview. Much research is in building a *probabilistic model*³⁰ of the input, $p_{\text{model}}(x)$. Why? Because then we can perform *inference* to predict stuff about our environment given any of the other variables. We call the other variables **latent variables**, h , with

$$p_{\text{model}}(x) = \sum_h \Pr(h) \Pr(x | h) = \mathbb{E}_h [p_{\text{model}}(x | h)] \quad (90)$$

So what? Well, the latent variables provide another means of *data representation*, which can be useful. **Linear factor models** (LFM) are some of the simplest probabilistic models with latent variables.

A linear factor model is defined by the use of a stochastic linear decoder function that generates \mathbf{x} by adding noise to a linear transformation of \mathbf{h} .

Note that \mathbf{h} is a *vector* of arbitrary size, where we assume $p(\mathbf{h})$ is a **factorial distribution**: $p(\mathbf{h}) = \prod_i p(h_i)$. This roughly means we assume the elements of \mathbf{h} are mutually independent³¹. The LFM describes the data-generation process as follows:

1. Sample the explanatory factors: $\mathbf{h} \sim p(\mathbf{h})$.
2. Sample the real-valued observable variables given the factors:

$$\mathbf{x} = \mathbf{W}\mathbf{h} + \mathbf{b} + \text{noise} \quad (91)$$

Probabilistic PCA and Factor Analysis.

- **Factor analysis:**

$$\mathbf{h} \sim \mathcal{N}(\mathbf{h}; \mathbf{0}, \mathbf{I}) \quad (92)$$

$$\text{noise} \sim \mathcal{N}(\mathbf{0}, \boldsymbol{\psi} \equiv \text{diag}(\boldsymbol{\sigma}^2)) \quad (93)$$

$$\mathbf{x} \sim \mathcal{N}(\mathbf{x}; \mathbf{b}, \mathbf{W}\mathbf{W}^T + \boldsymbol{\psi}) \quad (94)$$

where the last relation can be shown by recalling that a linear combination of Gaussian variables is itself Gaussian, and showing that $\mathbb{E}_h[\mathbf{x}] = \mathbf{b}$, and $\text{Cov}(\mathbf{x}) = \mathbf{W}\mathbf{W}^T + \boldsymbol{\psi}$.

³⁰Whereas, before, we've been building *functions* of the input (deterministic).

³¹Note that, technically, this assumption isn't strictly the definition of mutual independence, which requires that every *subset* (i.e. not just the full set) of $\{h_i \in \mathbf{h}\}$ follow this factorial property.

It is worth emphasizing the interpretation of $\boldsymbol{\psi}$ as the matrix of **conditional variances** σ_i^2 . *Huh?* Let's take a step back. The fact that we were able to separate the distributions in the above relations for \mathbf{h} and noise is from a built-in assumption that $\Pr(x_i|\mathbf{h}, x_{j \neq i}) = \Pr(x_i|\mathbf{h})$ ³².

The Big Idea

The latent variable \mathbf{h} is a big deal because it **captures the dependencies** between the elements of \mathbf{x} . *How do I know?* Because of our assumption that the x_i are conditionally independent given \mathbf{h} . If, once we specify \mathbf{h} , all the elements of \mathbf{x} become independent, then any information about their interrelationship is hiding somewhere in \mathbf{h} .

Detailed walkthrough of Factor Analysis (a.k.a me slowly reviewing, months after taking this note):

- **Goal.** Analyze and understand the motivations behind how Factor Analysis defines the data-generation process under the framework of LFMs (defined in steps 1 and 2 earlier). Assume \mathbf{h} has dimension n .
- **Prior.** Defines $p(\mathbf{h}) := \mathcal{N}(\mathbf{h}; \mathbf{0}, \mathbf{I})$, the unit-variance Gaussian. Explicitly,

$$p(\mathbf{h}) := \frac{1}{(2\pi)^{n/2}} e^{-\frac{1}{2} \sum_i h_i^2}$$

- **Noise.** Assumed to be drawn from a Gaussian with diagonal covariance matrix $\boldsymbol{\psi} := \text{diag}(\boldsymbol{\sigma}^2)$. Explicitly,

$$p(\text{noise} = \mathbf{a}) := \frac{1}{(2\pi)^{n/2} \prod_i \sigma_i} e^{-\frac{1}{2} \sum_i a_i^2 / \sigma_i^2}$$

- **Deriving distribution of \mathbf{x} .** We use the fact that any linear combination of Gaussians is itself Gaussian. Thus, deriving $p(\mathbf{x})$ is reduced to computing it's mean and covariance matrix.

$$\boldsymbol{\mu}_x = \mathbb{E}_{\mathbf{h}} [\mathbf{W}\mathbf{h} + \mathbf{b}] \tag{95}$$

$$= \int p(\mathbf{h})(\mathbf{W}\mathbf{h} + \mathbf{b})d\mathbf{h} \tag{96}$$

$$= \mathbf{b} + \int \frac{1}{(2\pi)^{n/2}} e^{-\frac{1}{2} \sum_i h_i^2} \mathbf{W}\mathbf{h} d\mathbf{h} \tag{97}$$

$$= \mathbf{b} \tag{98}$$

$$\text{Cov}(\mathbf{x}) = \mathbb{E} [(\mathbf{x} - \mathbb{E}[\mathbf{x}])(\mathbf{x} - \mathbb{E}[\mathbf{x}])^T] \tag{99}$$

$$= \mathbb{E} [(\mathbf{W}\mathbf{h} + \text{noise})(\mathbf{h}^T \mathbf{W}^T + \text{noise}^T)] \tag{100}$$

$$= \mathbb{E} [(\mathbf{W}\mathbf{h}\mathbf{h}^T \mathbf{W}^T)] + \boldsymbol{\psi} \tag{101}$$

$$= \mathbf{W}\mathbf{W}^T + \boldsymbol{\psi} \tag{102}$$

where we compute the expectation of \mathbf{x} over \mathbf{h} because \mathbf{x} is defined as a function of \mathbf{h} , and noise is always expectation zero.

³²Due to $\langle \text{MATH} \rangle$, this introduces a constraint that knowing the value of some element x_j doesn't alter the probability $\Pr(x_i = W_i \cdot h + b_i + \text{noise})$. Given how we've defined the variable \mathbf{h} , this means that knowing noise_j provides no clues about noise_i. Mathematically, the noise must have a diagonal covariance matrix.

- **Thoughts.** Not really seeing why this is useful/noteworthy. Feels very contrived (many assumptions) and restrictive – it only applies if the dependencies between each x_i can be modeled with a random variable \mathbf{h} sampled from a unit variance Gaussian.
- **Probabilistic PCA:** Just factor analysis with $\boldsymbol{\psi} = \sigma^2 \mathbf{I}$. So zero-mean spherical Gaussian noise. It becomes regular PCA as $\sigma \rightarrow 0$. Here we can use an iterative EM algorithm for estimating the parameters \mathbf{W} .

Autoencoders (Ch. 14)

Table of Contents Local

Written by Brandon McKinzie

Introduction. An autoencoder learns to copy its input to its output, via an encoder function $\mathbf{h} = f(\mathbf{x})$ and a decoder function $\mathbf{r} = g(\mathbf{h})$. Modern autoencoders generalize this to allow for stochastic mappings $p_{\text{encoder}}(\mathbf{h} \mid \mathbf{x})$ and $p_{\text{decoder}}(\mathbf{x} \mid \mathbf{h})$. \mathbf{r} for “reconstruction”

Undercomplete Autoencoders. Constrain dimension of \mathbf{h} to be smaller than that of \mathbf{x} . The learning process minimizes some $L(\mathbf{x}, g(f(\mathbf{x})))$, where the loss function could be e.g. mean squared error. Be careful not to have too many learnable parameters in the functions g and f (thus increasing model capacity), since that defeats the purpose of using an undercomplete autoencoder in the first place.

Regularized Autoencoders. We can remove the undercomplete constraint/necessity by modifying our loss function. For example, a **sparse autoencoder** one that adds a penalty $\Omega(\mathbf{h})$ to the loss function that encourages the *activations on* (not connections to/from) the hidden layer to be sparse. One way to achieve *actual zeros* in \mathbf{h} is to use rectified linear units for the activations.

Representation Learning (Ch. 15)

Table of Contents Local

Written by Brandon McKinzie

Greedy Layer-Wise Unsupervised Pretraining. Given:

- Unsupervised learning algorithm \mathcal{L} which accepts as input a training set of examples \mathbf{X} , and outputs an encoder/feature function f .
- $f^{(i)}(\tilde{\mathbf{X}})$ denotes the output of the i th layer of f , given as *immediate input* the (possibly transformed) set of examples $\tilde{\mathbf{X}}$.
- Let m denote the number of layers (“stages”) in the encoder function (note that each layer/stage here *must* use a representation learning algorithm for its \mathcal{L} e.g. an RBM, autoencoder, sparse coding model, etc.)

The procedure is as follows:

1. Initialize.

$$f(\cdot) \leftarrow I(\cdot) \quad (103)$$

$$\tilde{\mathbf{X}} = \mathbf{X} \quad (104)$$

2. For each layer (stage) i in $\text{range}(m)$, do:

$$f^{(k)} = \mathcal{L}(\tilde{\mathbf{X}}) \quad (105)$$

$$f(\cdot) \leftarrow f^{(k)}(f(\cdot)) \quad (106)$$

$$\tilde{\mathbf{X}} \leftarrow f^{(k)}(\tilde{\mathbf{X}}) \quad (107)$$

In English: just apply the regular learning/training process for each layer/stage **sequentially and individually**³³.

When this is complete, we can run **fine-tuning**: train all layers together (including any later layers that could not be pretrained) with a supervised learning algorithm. Note that we do indeed allow the pretrained encoding stages to be optimized here (i.e. not fixed).

³³In other words, you proceed one layer at a time *in order*. You don't touch layer i until the weights in layer $i - 1$ have been learned.

Structured Probabilistic Models for DL (Ch. 16)

Table of Contents Local

Written by Brandon McKinzie

Motivation. In addition to classification, we can ask probabilistic models to perform other tasks such as density estimation ($\mathbf{x} \rightarrow p(\mathbf{x})$), denoising, missing value imputation, or sampling. What these [other] tasks have in common is they require a *complete understanding of the input*. Let's start with the most naive approach of modeling $p(\mathbf{x})$, where \mathbf{x} contains n elements, each of which can take on k distinct values: we store a lookup table of all possible \mathbf{x} and the corresponding probability value $p(\mathbf{x})$. This requires k^n parameters³⁴. Instead, we use graphs to describe model structure (direct/indirect interactions) to drastically reduce the number of parameters.

Directed Models. Also called **belief networks** or **Bayesian networks**. Formally, a directed graphical model defined on a set of variables $\{\mathbf{x}\}$ is defined by a DAG, \mathcal{G} , whose vertices are the random variables in the model, and a set of **local conditional probability distributions**, $p(x_i | \text{Pa}_{\mathcal{G}}(x_i))$, where $\text{Pa}_{\mathcal{G}}(x_i)$ gives the parents of x_i in \mathcal{G} . The probability distribution over \mathbf{x} is given by

$$p(\mathbf{x}) = \prod_i p(x_i | \text{Pa}_{\mathcal{G}}(x_i)) \quad (108)$$

Undirected Graphical Models. Also called **Markov Random Fields (MRFs)** or **Markov Networks**. Appropriate for situations where interactions do not have a well-defined direction. Each **clique** \mathcal{C} (any set of nodes that are all [maximally] connected) in \mathcal{G} is associated with a factor $\phi(\mathcal{C})$. The factor $\phi(\mathcal{C})$, also called a **clique potential**, is just a function (not necessarily a probability) that outputs a number when given a possible set of values over the nodes in \mathcal{C} . The output number measures the affinity of the variables in that clique for being in the states specified by the inputs. The set of all factors in \mathcal{G} defines an **unnormalized probability distribution**:

$$\tilde{p}(\mathbf{x}) = \prod_{\mathcal{C} \in \mathcal{G}} \phi(\mathcal{C}) \quad (109)$$

Clique potentials are constrained to be nonnegative.

³⁴Consider the common NLP case where our vector \mathbf{x} contains n word tokens, each of which can take on any symbol in our vocabulary of size v . If we assign $n = 100$ and $v = 100,000$, which are relatively common values for this case, this amounts to $(1e5)^{1e2} = 10^{500}$ parameters.

The Partition Function. To obtain a valid probability distribution, we must normalize the probability distribution:

$$p(\mathbf{x}) = \frac{1}{Z} \tilde{p}(\mathbf{x}) \quad (110)$$

$$Z = \int \tilde{p}(\mathbf{x}) d\mathbf{x} \quad (111)$$

where the normalizing function $Z = Z(\{\phi\})$ is known as the **partition function** (physicists' terminology). It is typically intractable to compute, so we resort to approximations. Note that Z isn't even guaranteed to exist – it's only for those definitions of the clique potentials that cause the integral over $\tilde{p}(\mathbf{x})$ to converge/be defined.

Energy-Based Models (EBMs). A convenient way to enforce $\forall \mathbf{x}, \tilde{p}(\mathbf{x}) > 0$ is to use EBMs, where

$$\tilde{p}(\mathbf{x}) \triangleq \exp(-E(\mathbf{x})) \quad (112)$$

and $E(\mathbf{x})$ is known as the **energy function**³⁵. Many algorithms need to compute not $p_{\text{model}}(\mathbf{x})$ but only $\log \tilde{p}_{\text{model}}(\mathbf{x})$ (unnormalized log probabilities - logits!). For EBMs with latent variables \mathbf{h} , such algorithms are phrased in terms of the **free energy**:

$$\mathcal{F}(\mathbf{x} = x) = -\log \sum_{\mathbf{h}} \exp(-E(\mathbf{x} = x, \mathbf{h} = h)) \quad (113)$$

where we sum over all possible assignments of the latent variables.

Separation and D-Separation. We want to know which subsets of variables are conditionally independent from each other, given the values of other subsets of variables. A set of variables \mathbb{A} is **separated** (if undirected model)/**d-separated** (if directed model) from another set of variables \mathbb{B} given a third set of variables \mathbb{S} if the graph structure implies that \mathbb{A} is independent from \mathbb{B} given \mathbb{S} .

- **Separation.** For *undirected* models. If variables a and b are connected by a path involving only unobserved variables (an **active** path), then a and b are *not* separated. Otherwise, they are separated. Any paths containing at least one observed variable are called **inactive**.
- **D-Separation**³⁶. For *directed* models. Although there are rules that help determine whether a path between a and b is d-separated, it is simplest to just determine whether a is independent from b given any observed variables along the path.

³⁵Physics throwback: this mirrors the Boltzmann factor, $\exp(-\varepsilon/\tau)$, which is proportional to the probability of the system being in quantum energy state ε .

³⁶The D stands for dependence.

SAMPLING FROM GRAPHICAL MODELS

For directed graphical models, we can do **ancestral sampling** to produce a sample \mathbf{x} from the joint distribution represented by the model. Just sort the variables x_i into a topological ordering such that $\forall i, j : j > i \iff x_i \in \text{Pa}_{\mathcal{G}}(x_j)$. To produce the sample, just sequentially sample from the beginning, $x_1 \sim P(x_1)$, $x_2 \sim P(x_2 \mid \text{Pa}_{\mathcal{G}}(x_1))$, etc.

For undirected graphical models, one simple approach is **Gibbs sampling**. Essentially, this involves drawing a conditioned sample from $x_i \sim p(x_i \mid \text{neighbors}(x_i))$ for each x_i . This process is repeated many times, where each subsequent pass uses the previously sampled values in $\text{neighbors}(x_i)$ to obtain an asymptotically converging [to the correct distribution] estimate for a sample from $p(\mathbf{x})$.

INFERENCE AND APPROXIMATE INFERENCE

One of the main tasks with graphical models is predicting the values of some subset of variables given another subset: inference. Although the graph structures we've discussed allow us to represent complicated, high-dimensional distributions with a reasonable number of parameters, the graphs used for deep learning are usually not restrictive enough to allow efficient inference. **Approximate inference** for deep learning usually refers to variational inference, in which we approximate the distribution $p(\mathbf{h} \mid \mathbf{v})$ by seeking an approximate distribution $q(\mathbf{h} \mid \mathbf{v})$ that is as close to the true one as possible.

Example: Restricted Boltzmann Machine. The quintessential example of how graphical models are used for deep learning. The canonical RBM is an energy-based model with **binary** visible and hidden units. Its energy function is

$$E(\mathbf{v}, \mathbf{h}) = -\mathbf{b}^T \mathbf{v} - \mathbf{c}^T \mathbf{h} - \mathbf{v}^T \mathbf{W} \mathbf{h} \quad (114)$$

where \mathbf{b} , \mathbf{c} , and \mathbf{W} are unconstrained, real-valued, learnable parameters. One could interpret the values of the bias parameters as the affinities for the associated variable being its given value, and the value $\mathbf{W}_{i,j}$ as the affinity of v_i being its value and h_j being its value at the same time³⁷.

The restrictions on the RBM structure, namely the fact that there are no intra-layer connections, yields nice properties. Since $\tilde{p}(\mathbf{h}, \mathbf{v})$ can be factored into clique potentials, we can say

³⁷More concretely, remember that \mathbf{v} is a one-hot vector representing some state that can assume $\text{len}(\mathbf{v})$ unique values, and similarly for \mathbf{h} . Then $\mathbf{W}_{i,j}$ gives the affinity for the state associated with v being its i th value and the state associated with h being its j th value.

that:

$$p(\mathbf{h} \mid \mathbf{v}) = \prod_i p(h_i \mid \mathbf{v}) \quad (115)$$

$$p(\mathbf{v} \mid \mathbf{h}) = \prod_i p(v_i \mid \mathbf{h}) \quad (116)$$

Also, due to the restriction of binary variables, each of the conditionals is easy to compute, and can be quickly derived as

$$p(h_i = 1 \mid \mathbf{v}) = \sigma \left(c_i + \mathbf{v}^T \mathbf{W}_{:,i} \right) \quad (117)$$

allowing for efficient block Gibbs sampling.

Monte Carlo Methods (Ch. 17)

Table of Contents Local

Written by Brandon McKinzie

Monte Carlo Sampling (Basics). We can approximate the value of a (usually prohibitively large) sum/integral by viewing it as an *expectation* under some distribution. We can then approximate its value by taking samples from the corresponding probability distribution and taking an empirical average. Mathematically, the basic idea is show below:

$$s = \int p(\mathbf{x})f(\mathbf{x})d\mathbf{x} = \mathbb{E}_p[f(\mathbf{x})] \quad \rightarrow \quad \hat{s}_n = \frac{1}{n} \sum_{i=1, \mathbf{x}^{(i)} \sim p}^n f(\mathbf{x}^{(i)}) \quad (118)$$

As we've seen before, the empirical average is an unbiased³⁸ estimator. Furthermore, the central limit theorem tells us that the distribution of \hat{s}_n converges to a normal distribution with mean s and variance $\text{Var}[f(\mathbf{x})]/n$.

Importance Sampling. What if it's not feasible for us to sample from p ? We can approach this a couple ways, both of which will exploit the following identity:

$$p(\mathbf{x})f(\mathbf{x}) = q(\mathbf{x})\frac{p(\mathbf{x})f(\mathbf{x})}{q(\mathbf{x})} \quad (122)$$

- **Optimal importance sampling.** We can use the aforementioned identity/decomposition to find the **optimal** q^* – optimal in terms of number of samples required to achieve a given level of accuracy. First, we rewrite our estimator \hat{s}_p (they now use subscript to denote the sampling distribution) as \hat{s}_q :

$$\hat{s}_q = \frac{1}{n} \sum_{i=1, \mathbf{x}^{(i)} \sim q}^n \frac{p(\mathbf{x}^{(i)})f(\mathbf{x}^{(i)})}{q(\mathbf{x}^{(i)})} \quad (123)$$

³⁸ Recall that expectations on such an average are still taken over the underlying (assumed) probability distribution:

$$\mathbb{E}_p[\hat{s}_n] = \frac{1}{n} \sum_{i=1}^n \mathbb{E}_p[f(\mathbf{x}^{(i)})] \quad (119)$$

$$= \frac{1}{n} \sum_{i=1}^n s \quad (120)$$

$$= s \quad (121)$$

You should think of the expectation $\mathbb{E}_p[f(\mathbf{x}^{(i)})]$ as the expected value of the *random sample* from the underlying distribution, which of course is s , because we defined it that way.

At first glance, it feels a little wonky, but recognize that we are *sampling from q instead of p* (i.e. if this were an integral, it would be over $q(\mathbf{x})d\mathbf{x}$). The catch is that, now, the variance can be greatly sensitive to the choice of q :

$$\text{Var} [\hat{s}_q] = \text{Var} \left[\frac{p(\mathbf{x})f(\mathbf{x})}{q(\mathbf{x})} \right] / n \quad (124)$$

with the optimal (minimum) value of q at:

$$q^* = \frac{p(\mathbf{x}) | f(\mathbf{x}) |}{Z} \quad (125)$$

- **Biased importance sampling.** Computing the optimal value of q can be as challenging/infeasible as sampling from p . Biased sampling does not require us to find a normalization constant for p or q . Instead, we compute:

$$\hat{s}_{BIS} = \frac{\sum_{i=1}^n \frac{\tilde{p}(\mathbf{x}^{(i)})}{\tilde{q}(\mathbf{x}^{(i)})} f(\mathbf{x}^{(i)})}{\sum_{i=1}^n \frac{\tilde{p}(\mathbf{x}^{(i)})}{\tilde{q}(\mathbf{x}^{(i)})}} \quad (126)$$

where \tilde{p} and \tilde{q} are the unnormalized forms of p and q , and the $\mathbf{x}^{(i)}$ samples are still drawn from [the original/unknown] q . $\mathbb{E} [\hat{s}_{BIS}] \neq s$ except asymptotically when $n \rightarrow \infty$.

Approximate Inference (Ch. 19)

Table of Contents Local

Written by Brandon McKinzie

Overview. Most graphical models with multiple layers of hidden variables have intractable posterior distributions. Many approximate inference approaches make use of the observation that exact inference can be described as an optimization problem. Assume we have a probabilistic model consisting of observed variables \mathbf{v} and latent variables \mathbf{h} . We want to compute $\log p(\mathbf{v}; \boldsymbol{\theta})$, but it's too costly to marginalize out \mathbf{h} . Instead, we compute a lower bound $\mathcal{L}(\mathbf{v}, \boldsymbol{\theta}, q)$ – often called the **evidence lower bound** (ELBO) or **variational free energy** – on $\log p(\mathbf{v}; \boldsymbol{\theta})$:

$$\mathcal{L}(\mathbf{v}, \boldsymbol{\theta}, q) = \log p(\mathbf{v}; \boldsymbol{\theta}) - D_{KL}(q(\mathbf{h} | \mathbf{v}) || p(\mathbf{h} | \mathbf{v}; \boldsymbol{\theta})) \quad (127)$$

$$= -\mathbb{E}_{\mathbf{h} \sim q} [\log p(\mathbf{h}, \mathbf{v})] + H(q) \quad (128)$$

q is an arbitrary probability distribution over \mathbf{h}

where the second form is the more canonical definition³⁹. Note that $\mathcal{L}(\mathbf{v}, \boldsymbol{\theta}, q)$ is a lower-bound on $\log p(\mathbf{v}; \boldsymbol{\theta})$ by definition, since

$$\log p(\mathbf{v}; \boldsymbol{\theta}) - \mathcal{L}(\mathbf{v}, \boldsymbol{\theta}, q) = D_{KL}(q(\mathbf{h} | \mathbf{v}) || p(\mathbf{h} | \mathbf{v}; \boldsymbol{\theta})) \geq 0$$

³⁹This can be derived easily from the first form. Hint:

$$\log \frac{q(\mathbf{h} | \mathbf{v})}{p(\mathbf{h} | \mathbf{v})} = \log \frac{q(\mathbf{h} | \mathbf{v})}{p(\mathbf{h}, \mathbf{v}; \boldsymbol{\theta}) / p(\mathbf{v}; \boldsymbol{\theta})}$$

PAPERS AND TUTORIALS

CONTENTS

4.1	WaveNet	61
4.2	Neural Style	65
4.3	Neural Conversation Model	67
4.4	NMT By Jointly Learning to Align & Translate	69
4.4.1	Detailed Model Architecture	70
4.5	Effective Approaches to Attention-Based NMT	72
4.6	Using Large Vocabularies for NMT	74
4.7	Candidate Sampling – TensorFlow	77
4.8	Attention Terminology	79
4.9	TextRank	81
4.9.1	Keyword Extraction	83
4.9.2	Sentence Extraction	84
4.10	Simple Baseline for Sentence Embeddings	85
4.11	Survey of Text Clustering Algorithms	87
4.11.1	Feature Selection and Transformation Methods	87
4.11.2	Distance-based Clustering Algorithms	90
4.11.3	Probabilistic Document Clustering and Topic Models	91
4.11.4	Online Clustering with Text Streams	93
4.11.5	Semi-Supervised Clustering	94
4.12	Deep Sentence Embedding Using LSTMs	95
4.13	Clustering Massive Text Streams	98

4.14	Supervised Universal Sentence Representations (InferSent)	100
4.15	Dist. Rep. of Sentences from Unlabeled Data (FastSent)	101
4.16	Latent Dirichlet Allocation	103
4.17	Conditional Random Fields	106
4.18	Attention Is All You Need	108
4.19	Hierarchical Attention Networks	111
4.20	Joint Event Extraction via RNNs	114
4.21	Event Extraction via Bidi-LSTM Tensor NNs	116
4.22	Reasoning with Neural Tensor Networks	118
4.23	Language to Logical Form with Neural Attention	119
4.24	Seq2SQL: Generating Structured Queries from NL using RL	121
4.25	SLING: A Framework for Frame Semantic Parsing	124
4.26	Poincaré Embeddings for Learning Hierarchical Representations	126
4.27	Enriching Word Vectors with Subword Information (FastText)	128
4.28	DeepWalk: Online Learning of Social Representations	130
4.29	Review of Relational Machine Learning for Knowledge Graphs	132
4.30	Fast Top-K Search in Knowledge Graphs	135
4.31	Dynamic Recurrent Acyclic Graphical Neural Networks (DRAGNN)	137
4.32	Neural Architecture Search with Reinforcement Learning	140
4.33	Joint Extraction of Events and Entities within a Document Context	142
4.34	Globally Normalized Transition-Based Neural Networks	145
4.35	An Introduction to Conditional Random Fields	148
4.35.1	Inference	152
4.36	Co-sampling: Training Robust Networks for Extremely Noisy Supervision . . .	155

WaveNet

Table of Contents Local

*Written by Brandon McKinzie***Abstract.**

- **What’s WaveNet?** A deep neural network for generating raw audio waveforms.
- **How does it generate them?** **IDK**
- **What’s it good for?** Text-to-speech, generating music, any waveform really.

Introduction.

- Inspired by recent advances in **neural autoregressive generative models**, and based on the PixelCNN architecture.
- Long-range dependencies dealt with via “dilated causal convolutions, which exhibit very large receptive fields.”

WaveNet. The joint probability of a waveform $x = \{x_1, \dots, x_T\}$ is factorised as a product of conditional probabilities,

$$p(x) = \prod_{t=1}^T p(x_t \mid x_1, \dots, x_{t-1}) \quad (129)$$

which are modeled by a stack of convolutional layers (no pooling).

The model outputs a categorical distribution over the next value x_t with a softmax layer and it is optimized to maximize the log-likelihood of the data w.r.t. the parameters.

Main ingredient of WaveNet is *dilated* causal convolutions, illustrated below. Note the absence of recurrent connections, which makes them faster to train than RNNs, but at the cost of requiring many layers, or large filters to increase the receptive field⁴⁰.

⁴⁰Loose interpretation of receptive fields here is that large fields can take into account more info (back in time) as opposed to smaller fields, which can be said to be “short-sighted”



Figure 3: Visualization of a stack of *dilated* causal convolutional layers.

Excellent concise definition from paper:

A dilated convolution (a convolution with holes) is a convolution where the filter is applied over an area larger than its length by skipping input values with a certain step. It is equivalent to a convolution with a larger filter derived from the original filter by dilating it with zeros, but is significantly more efficient. A dilated convolution effectively allows the network to operate on a coarser scale than with a normal convolution. This is similar to pooling or strided convolutions, but here the output has the same size as the input. As a special case, dilated convolution with dilation 1 yields the standard convolution.

Softmax distributions. Chose to model the conditional distributions $p(x_t \mid x_1, \dots, x_{t-1})$ with a softmax layer. To deal with the fact that there are 2^{16} possible values, first apply a “ μ -law companding transformation” to data, and then quantize it to 256 possible values:

$$f(x_t) = \text{sign}(x_t) \frac{\ln(1 + \mu|x_t|)}{\ln(1 + \mu)} \quad (130)$$

which (after plotting in Wolfram) looks identical to the sigmoid function.

Gated activation and res/skip connections. Use the same gated activation unit as PixelCNN:

$$z = \tanh(W_{f,k} * x) \odot \sigma(W_{g,k} * x) \quad (131)$$

where $*$ denotes conv operator, \odot denotes elem-wise mult., k is layer index, f, g denote filter/-gate, and W is learnable conv filter. This is illustrated below, along with the residual/skip connections used to speed up convergence/enable training deeper models.



Conditional Wavenets. Can also model conditional distribution of x given some additional h (e.g. speaker identity).

$$p(x | h) = \prod_{t=1}^T p(x_t | x_1, \dots, x_{t-1}, h) \quad (132)$$

→ **Global conditioning.** Single h that influences output dist. accross all times. Activation becomes:

$$z = \tanh(W_{f,k} * x + V_{f,k}^T h) \odot \sigma(W_{g,k} * x + V_{g,k}^T h) \quad (133)$$

→ **Local conditioning** (confusing). Have a second time-series h_t . They first transform this h_t using a “transposed conv net (learned unsampling) that maps it to a new time-series $y = f(h)$ w/same resolution as x .”

Experiments.

- **Multi-Speaker Speech Generation.** Dataset: multi-speaker corpus of 44 hours of data from 109 different speakers⁴¹. Receptive field of 300 milliseconds.
- **Text-to-Speech.** Single-speaker datasets of 24.6 hours (English) and 34.8 hours (Chinese) speech. Locally conditioned on *linguistic features*. Receptive field of 240 milliseconds. Outperformed both LSTM-RNN and HMM.
- **Music.** Trained the WaveNets to model two music datasets: (1) 200 hours of annotated music audio, and (2) 60 hours of solo piano music from youtube. Larger receptive fields sounded more musical.
- **Speech Recognition.** “With WaveNets we have shown that layers of dilated convolutions allow the receptive field to grow longer in a much cheaper way than using LSTM units.”

⁴¹Speakers encoded as ID in form of a one-hot vector

Conclusion (verbatim): “This paper has presented WaveNet, a deep generative model of audio data that operates directly at the waveform level. WaveNets are autoregressive and combine causal filters with dilated convolutions to allow their receptive fields to grow exponentially with depth, which is important to model the long-range temporal dependencies in audio signals. We have shown how WaveNets can be conditioned on other inputs in a global (e.g. speaker identity) or local way (e.g. linguistic features). When applied to TTS, WaveNets produced samples that outperform the current best TTS systems in subjective naturalness. Finally, WaveNets showed very promising results when applied to music audio modeling and speech recognition.”

Neural Style

Table of Contents Local

Written by Brandon McKinzie

Notation.

- **Content image:** \mathbf{p}
- **Filter responses:** the matrix $P^l \in \mathcal{R}^{N_l \times M_l}$ contains the activations of the filters in layer l , where P_{ij}^l would give the activation of the i th filter at position j in layer l . N_l is number of feature maps, each of size M_l (height \times width of the feature map)⁴².
- **Reconstructed image:** \mathbf{x} (initially random noise). Denote its corresponding filter response matrix at layer l as P^l .

Content Reconstruction.

1. Feed in **content image** \mathbf{p} into pre-trained network, saving any desired filter responses during the forward pass. These are interpreted as the various “encodings” of the image done by the network. Think of them analogously to “ground-truth” labels.
2. Define \mathbf{x} as the **generated image**, which we first initialize to random noise. We will be changing the pixels of \mathbf{x} via gradient descent updates.
3. Define the **loss function**. After each forward pass, evaluate with squared-error loss between the two representations at the layer of interest:

$$\mathcal{L}_{content}(\mathbf{p}, \mathbf{x}, l) = \frac{1}{2} \sum_{i,j} (F_{ij}^l - P_{ij}^l)^2 \quad (1)$$

$$\frac{\partial \mathcal{L}_{content}}{\partial F_{ij}^l} = \begin{cases} (F^l - P^l)_{ij} & F_{ij}^l > 0 \\ 0 & F_{ij}^l < 0 \end{cases} \quad (2)$$

where it appears we are assuming ReLU activations (?).

4. Compute iterative updates to \mathbf{x} via **gradient descent** until it generates the same response in a certain layer of the CNN as the original image \mathbf{p} .

⁴²If not clear, M_l is a scalar, for any given value of l .

Style Representation. On top of the CNN responses in each layer, the authors built a style representation that computes the correlations between the different [aforementioned] filter responses. The correlation matrix for layer l is denoted in the standard way with a Gram matrix $G^l \in \mathcal{R}^{N_l \times N_l}$, with entries

$$G_{ij}^l = \langle F_i^l, F_j^l \rangle = \sum_k F_{ik}^l F_{jk}^l \quad (3)$$

To generate a texture that matches the style of a given image, do the following.

1. Let \mathbf{a} denote the original [style] image, with corresponding A^l . Let \mathbf{x} , initialized to random noise, denote the generated [style] image, with corresponding G^l .
2. The contribution to the loss of layer l , denoted E_l , to the total loss, denoted \mathcal{L}_{style} , is given by

$$E_l = \frac{1}{4N_l^2 M_l^2} \sum_{ij} (G_{ij}^l - A_{ij}^l)^2 \quad (4)$$

$$\mathcal{L}_{style}(\mathbf{a}, \mathbf{x}) = \sum_{l=0}^L w_l E_l \quad (5)$$

$$\frac{\partial E_l}{\partial F_{ij}^l} = \begin{cases} \frac{1}{N_l^2 M_l^2} \left((F^l)^T (G^l - A^l) \right)_{ji} & F_{ij}^l > 0 \\ 0 & F_{ij}^l < 0 \end{cases} \quad (6)$$

where w_l are [as of yet unspecified] weighting factors of the contribution of layer l to the total loss.

Mixing content with style. Essentially a joint minimization that combines the previous two main ideas.

1. Begin with the following images: white noise \mathbf{x} , content image \mathbf{p} , and style image \mathbf{a} .
2. The loss function to minimize is a linear combination of 1 and 5:

$$\mathcal{L}_{total}(\mathbf{p}, \mathbf{a}, \mathbf{x}, l) = \alpha \mathcal{L}_{content}(\mathbf{p}, \mathbf{x}, l) + \beta \mathcal{L}_{style}(\mathbf{a}, \mathbf{x}) \quad (7)$$

Note that we can choose which layers \mathcal{L}_{style} uses by tweaking the layer weights w_l . For example, the authors chose to set $w_l = 1/5$ for 'conv[1, 2, 4, 5]_1' and 0 otherwise. For the ratio α/β , they explored 1×10^{-3} and 1×10^{-4} .

Neural Conversation Model

Table of Contents Local

Written by Brandon McKinzie

[Reminder: **red text** means I need to come back and explain what is meant, once I understand it.]

Abstract. This paper presents a simple approach for conversational modeling which uses the sequence to sequence framework. It can be trained end-to-end, meaning fewer hand-crafted rules. The **lack of consistency** is a common failure of our model.

Introduction. Major advantage of the seq2seq model is it requires little feature engineering and domain specificity. Here, the model is tested on chat sessions from an IT helpdesk dataset of conversations, as well as movie subtitles.

Related Work. The authors' approach is based on the following (linked and saved) papers on seq2seq:

- Kalchbrenner & Blunsom, 2013.
- Sutskever et al., 2014. (Describes Seq2Seq model)
- Bahdanau et al., 2014.

Model. Succinctly described by the authors:

The model reads the input sequence one token at a time, and predicts the output sequence, also one token at a time. During training, the true output sequence is given to the model, so learning can be done by backpropagation. The model is trained to maximize the cross entropy of the correct sequence given its context. During inference, in which the true output sequence is not observed, we simply feed the predicted output token as input to predict the next output. This is a greedy inference approach.

Example of less greedy approach: **beam search**.



The **thought vector** is the hidden state of the model when it receives [as input] the end of sequence symbol $\langle eos \rangle$, because it stores the info of the sentence, or *thought*, “ABC”. The authors acknowledge that this model will *not* be able to “solve” the problem of modeling dialogue due to the objective function not capturing the actual objective achieved through human communication, which is typically longer term and based on exchange of information [rather than next step prediction]⁴³.

Ponder: what *would* be a reasonable objective function & model for conversation?

IT Data & Experiment.

Reminder: Check out this git repo

- **Data Description:** Customers talking to IT support, where typical interactions are 400 words long and turn-taking is clearly signaled.
- **Training Data:** 30M tokens, 3M of which are used as validation. They built a vocabulary of the most common 20K words, and introduced special tokens indicating turn-taking and actor.
- **Model:** A single-layer LSTM with 1024 memory cells.
- **Optimization:** SGD with gradient clipping.
- **Perplexity:** At convergence, achieved **perplexity** of 8, whereas an n-gram model achieved 18.

⁴³I’d imagine that, in order to model human conversation, one obvious element needed would be a *memory*. Reminds me of DeepMind’s DNC. There would need to be some online filtering & output process to capture the crucial aspects/info to store in memory for later, and also some method of retrieving them when needed later. The method for retrieval would likely be some inference process where, given a sequence of inputs, the probability of them being related to some portion of memory could be trained. This would allow for conversations that stretch arbitrarily back in the past. Also, when storing the memories, I’d imagine a reasonable architecture would be some encoder-decoder for a sparse distributed representation of memory.

NMT By Jointly Learning to Align & Translate

Table of Contents Local

Written by Brandon McKinzie

[Bahdanau et. al, 2014]. The primary motivation for me writing this is to better understand the **attention mechanism** in my sequence to sequence chatbot implementation.

Abstract. The authors claim that using a fixed-length vector [in the vanilla encoder-decoder for NMT] is a bottleneck. They propose allowing a model to (soft-)search for parts of a source sentence that are relevant to predicting a target word, without having to form these parts as a hard segment explicitly.

Learning to Align⁴⁴ and translate.

- **Decoder.** Their encoder defines the conditional output distribution as

$$p(y_i \mid y_1, \dots, y_{i-1}, \mathbf{x}) = g(y_{i-1}, s_i, c_i) \quad (134)$$

$$s_i = f(s_{i-1}, y_{i-1}, c_i) \quad (135)$$

where s_i is the RNN [decoder] hidden state at time i .

- NOTE: c_i is *not* the i th element of the standard context vector; rather, it is *itself* a distinct context vector that depends on a sequence of **annotations** (h_1, \dots, h_{T_x}). It seems that each annotation h_i is a hidden (encoder) state “that contains information about the whole input sequence with a strong focus on the parts surrounding the i -th word of the input sequence.”
- The context vector c_i is computed as follows:

$$c_i = \sum_{j=1}^{T_x} \alpha_{ij} h_j \quad (136)$$

$$\alpha_{ij} = \frac{\exp(e_{ij})}{\sum_{k=1}^{T_x} \exp(e_{ik})} \quad (137)$$

$$e_{ij} = a(s_{i-1}, h_j) \quad (138)$$

where the function e_{ij} is given by an **alignment model** which scores how well the inputs around position j and the output at position i match.

- **Encoder.** It’s just a bidirectional RNN. What they call “annotation h_j ” is literally just a concatenated vector of $h_j^{forward}$ and $h_j^{backward}$

⁴⁴By “align” the authors are referring to aligning the source-search to the relevant parts for prediction.

DETAILED MODEL ARCHITECTURE

(Appendix A). Explained with the TensorFlow user in mind.

Decoder Internals. It's just a GRU. However, it will be helpful to detail how we format the inputs (given we now have attention). Wherever we'd usually pass the previous decoder state s_{i-1} , we now pass a *concatenated* state, $[s_{i-1}, c_i]$, that also contains the i th context vector. Below I go over the flow of information from GRU input to output:

1. **Notation:** y_t is the loop-embedded output of the decoder (prediction) at time t , s_t is the internal hidden state of the decoder at time t , and c_t is the context vector at time t . \tilde{s}_t is the proposed/proposal state at time t .
2. **Gates:**

$$z_t = \sigma(W_z y_{t-1} + U_z[s_{t-1}, c_t]) \quad \text{[update gate]} \quad (139)$$

$$r_t = \sigma(W_r y_{t-1} + U_r[s_{t-1}, c_t]) \quad \text{[reset gate]} \quad (140)$$

$$(141)$$

3. **Proposal state:**

$$\tilde{s}_t = \tanh(W y_{t-1} + U[r_t \circ s_{t-1}, c_t]) \quad (142)$$

4. **Hidden state:**

$$s_t = (1 - z_t) \circ s_{t-1} + z_t \circ \tilde{s}_t \quad (143)$$

Alignment Model. All equations enumerated below are for some timestep t during the decoding process.

1. **Score:** For all $j \in [0, L_{enc} - 1]$ where L_{enc} is the number of words in the encoder sequence, compute:

$$a_j = a(s_{t-1}, h_j) = v_a^T \tanh(W_a s_{t-1} + U_a h_j) \quad (144)$$

2. **Alignments:** Feed the unnormalized alignments (scores) through a softmax so they represent a valid probability distribution.

$$a_j \leftarrow \frac{e^{a_j}}{\sum_{k=0}^{L_{enc}-1} e^{a_k}} \quad (145)$$

3. **Context:** The context vector input for our decoder at this timestep:

$$c = \sum_{j=1}^{L_{enc}} a_j h_j \quad (146)$$

Decoder Outputs. All below is for some timestep t during the decoding process. To find the probability of some (one-hot) word y [at timestep t]:

$$\Pr(y \mid s, c) \propto e^{y^T W_o u} \quad (147)$$

$$u = [\max\{\tilde{u}_{2j-1}, \tilde{u}_{2j}\}]_{j=1, \dots, \ell}^T \quad (148)$$

$$\tilde{u} = U_o[s_{t-1}, c] + V_o y_{t-1} \quad (149)$$

N.B.: From reading other (and more recent) papers, these last few equations do not appear to be the way it is usually done (thank the lord). See Luong's work for a much better approach.

Effective Approaches to Attention-Based NMT

[Luong et. al, 2015]

Attention-Based Models. For attention especially, the devil is in the details, so I’m going to go into somewhat excruciating detail here to ensure no ambiguities remain. For both global and local attention, the following information holds true:

- “At each time step t in the decoding phase, both approaches first take as input the hidden state \mathbf{h}_t at the top layer of a stacking LSTM.”
- Then, they derive [with different methods] a context vector \mathbf{c}_t to capture source-side info.
- Given \mathbf{h}_t and \mathbf{c}_t , they both compute the **attentional hidden state** as:

$$\tilde{\mathbf{h}}_t = \tanh(\mathbf{W}_c[\mathbf{c}_t; \mathbf{h}_t]) \quad (150)$$

- Finally, the predictive distribution (decoder outputs) is given by feeding this through a softmax:

$$p(y_t \mid y_{<t}, x) = \text{softmax}(\mathbf{W}_s \tilde{\mathbf{h}}_t) \quad (151)$$

Global Attention. Now I’ll describe in detail the processes involved in $\mathbf{h}_t \rightarrow \mathbf{a}_t \rightarrow \mathbf{c}_t \rightarrow \tilde{\mathbf{h}}_t$.

1. \mathbf{h}_t : Compute the hidden state \mathbf{h}_t in the normal way (not obvious if you’ve read e.g. Bahdanau’s work...)
2. \mathbf{a}_t :
 - (a) Compute the **scores** between \mathbf{h}_t and each source $\bar{\mathbf{h}}_s$, where our options are:

$$\text{score}(\mathbf{h}_t, \bar{\mathbf{h}}_s) = \begin{cases} \mathbf{h}_t^T \bar{\mathbf{h}}_s & \text{dot} \\ \mathbf{h}_t^T \mathbf{W}_a \bar{\mathbf{h}}_s & \text{general} \\ \mathbf{v}_a^T \tanh(\mathbf{W}_a[\mathbf{h}_t; \bar{\mathbf{h}}_s]) & \text{concat} \end{cases} \quad (152)$$

- (b) Compute the **alignment vector** \mathbf{a}_t of length L_{enc} (number of words in the encoder sequence):

$$\mathbf{a}_t(s) = \text{align}(\mathbf{h}_t, \bar{\mathbf{h}}_s) \quad (153)$$

$$= \frac{\exp(\text{score}(\mathbf{h}_t, \bar{\mathbf{h}}_s))}{\sum_{s'} \exp(\text{score}(\mathbf{h}_t, \bar{\mathbf{h}}_{s'}))} \quad (154)$$

3. \mathbf{c}_t : The weighted average over all source (encoder) hidden states⁴⁵:

$$\mathbf{c}_t = \sum_{i=1}^{L_{enc}} \mathbf{a}_t(i) \bar{\mathbf{h}}_i \quad (155)$$

4. $\tilde{\mathbf{h}}_t$: For convenience, I'll copy the equation for $\tilde{\mathbf{h}}_t$ again here:

$$\tilde{\mathbf{h}}_t = \tanh(\mathbf{W}_c[\mathbf{c}_t; \mathbf{h}_t]) \quad (156)$$

Input-Feeding Approach. A copy of each output $\tilde{\mathbf{h}}_t$ is sent forward and concatenated with the inputs for the next timestep, i.e. the inputs go from \mathbf{x}_{t+1} to $[\tilde{\mathbf{h}}_t; \mathbf{x}_{t+1}]$.

⁴⁵NOTE: Right after mentioning the context vector, the authors have the following cryptic footnote that may be useful to ponder: *For short sentences, we only use the top part of \mathbf{a}_t and for long sentences, we ignore words near the end.*

Using Large Vocabularies for NMT

Table of Contents Local

Written by *Brandon McKinzie*

Paper information:

- Full title: On Using Very Large Target Vocabulary for Neural Machine Translation.
- Authors: Jean, Cho, Memisevic, Bengio.
- Date: 18 Mar 2015.
- [arXiv link]

NMT Overview. Typical implementation is encoder-decoder network. Notation for inputs & encoder:

$$x = (x_1, \dots, x_T) \quad [\text{source sentence}] \quad (157)$$

$$h = (h_1, \dots, h_T) \quad [\text{encoder state seq}] \quad (158)$$

$$h_t = f(x_t, h_{t-1}) \quad (159)$$

where f is the function defined by the *cell state* (e.g. GRU/LSTM/etc.). Then the decoder generates the output sequence y , and with probability given below:

$$y = (y_1, \dots, y_T) \quad [y_i \in \mathbb{Z}] \quad (160)$$

$$\Pr[y_t \mid y_{<t}, x] \propto e^{q(y_{t-1}, z_t, c_t)} \quad (161)$$

$$z_t = g(y_{t-1}, z_{t-1}, c_t) \quad [\text{decoder hidden?}] \quad (162)$$

$$c_t = r(z_{t-1}, h_1, \dots, h_T) \quad [\text{decoder inp?}] \quad (163)$$

The functions q , g , and r are just placeholders – “some function of [inputs].”

As usual, model is jointly trained to maximize the conditional log-likelihood of correct translation. For N training sample pairs (x^n, y^n) , and denoting the length of the n -th target sentence as T_n , this can be written as,

$$\theta^* = \arg \max_{\theta} \sum_{n=1}^N \sum_{t=1}^{T_n} \log (\Pr[y_t^n \mid y_{<t}^n, x^n]) \quad (164)$$

Model Details. Above is the general structure. Here I'll summarize the specific model chosen by the authors.

- **Encoder.** Bi-directional, which just means $h_t = [h_t^{backward}, h_t^{forward}]$. The chosen cell state (the function f) is GRU.
- **Decoder.** At each timestep, computes the following:
→ **Context vector** c_t .

$$c_t = \sum_{i=1}^T \alpha_i h_i \quad (165)$$

$$\alpha_t = \frac{e^{a(h_t, z_{t-1})}}{\sum_k e^{a(h_k, z_{t-1})}} \quad (166)$$

a is a standard single-hidden-layer NN.

→ **Decoder hidden state** z_t . Also a GRU cell. Computed based on the previous hidden state z_{t-1} , the previously generated symbol y_{t-1} , and also the computed context vector c_t .

- **Next-word probability.** They model equation 161 as⁴⁶,

$$\Pr[y_t \mid y_{<t}, x] = \frac{1}{Z} e^{\mathbf{w}_t^T \phi(y_{t-1}, z_t, c_t) + b_t} \quad (167)$$

$$Z = \sum_{k: y_k \in V} e^{\mathbf{w}_k^T \phi(y_{t-1}, z_t, c_t) + b_k} \quad (168)$$

Reminder: y_i is an integer token, while \mathbf{w}_i is the target vector of length vocab size

where ϕ is affine transformation followed by a nonlinear activation, \mathbf{w}_t and b_t are the **target word vector** and bias. V is the set of all target *vocabulary*.

Approximate Learning Approach. Main idea:

“In order to avoid the growing complexity of computing the normalization constant, we propose here to use only a small subset V' of the target vocabulary at each update.”

Consider the gradient of the log-likelihood⁴⁷, written in terms of the energy \mathcal{E} .

$$\nabla \log (\Pr[y_t \mid y_{<t}, x]) = \nabla \mathcal{E}(y_t) - \sum_{k: y_k \in V} \Pr[y_k \mid y_{<t}, x] \nabla \mathcal{E}(y_k) \quad (169)$$

$$\mathcal{E}(y_j) = \mathbf{w}_j^T \phi(y_{t-1}, z_t, c_t) + b_j \quad (170)$$

⁴⁶Note: The formula for Z is correct. Notice that the only part of the RHS of $\Pr(y_t)$ with a t is as the subscript of w . To be clear, w_k is a full word vector and the sum is over all words in the output *vocabulary*, the index k has absolutely nothing to do with timestep. They use the word target but make sure not to misinterpret that as somehow meaning target words in the sentence or something.

⁴⁷**NOTE TO SELF:** After long and careful consideration, I'm concluding that the authors made a typo when defining $\mathcal{E}(y_j)$, which they choose to subscript all parts of the RHS with j , but that is in direct contradiction with a step-by-step derivation, which is why I have written it the way it is. I'm pretty sure my version is right, but I know you'll have to re-derive it yourself next time you see this. And you'll somehow prove me wrong. Actually, after reading on further, I doubt you'll prove me wrong. Challenge accepted, me. Have fun!

The crux of the approach is interpreting the second term as $\mathbb{E}_P [\nabla \mathcal{E}(y)]$, where P denotes $Pr(y \mid y_{<t}, x)$. They approximate this expectation by taking it over a subset V' of the predefined proposal distribution Q . So Q is a p.d.f. over the possible y_i , and we sample *from* Q to generate the elements of the subset V' .

$$\mathbb{E}_P [\nabla \mathcal{E}(y)] \approx \sum_{k: y_k \in V'} \frac{\omega_k}{\sum_{k': y_{k'} \in V'} \omega_{k'}} \nabla \mathcal{E}(y_k) \quad (171)$$

$$\omega_k = e^{\mathcal{E}(y_k) - \log Q(y_k)} \quad (172)$$

Here is some math I did that was illuminating to me; I'm not sure why the authors didn't point out these relationships.

$$\omega_k = \frac{e^{\mathcal{E}(y_k)}}{Q(y_k)} \quad \text{thus} \quad p(y_k \mid y_{<t}, x) = \omega_k \frac{Q(y_k)}{Z} \quad (173)$$

$$\rightarrow e^{\mathcal{E}(y_k)} = Z \cdot p(y_k \mid y_{<t}, x) = Q(y_k) \cdot \omega_k \quad (174)$$

Now check this out

Below are the exact and approximate formulas for $\mathbb{E}_P [\nabla \mathcal{E}(y)]$ written in a seductive suggestive manner. Pay careful attention to subscripts and primes.

$$\mathbb{E}_P [\nabla \mathcal{E}(y)] = \sum_{k: y_k \in V} \frac{\omega_k \cdot Q(y_k)}{\sum_{k': y_{k'} \in V} \omega_{k'} \cdot Q(y_{k'})} \nabla \mathcal{E}(y_k) \quad (175)$$

$$\mathbb{E}_P [\nabla \mathcal{E}(y)] = \sum_{k: y_k \in V'} \frac{\omega_k}{\sum_{k': y_{k'} \in V'} \omega_{k'}} \nabla \mathcal{E}(y_k) \quad (176)$$

They're almost the same! It's much easier to see why when written this way. I interpret the difference as follows: in the exact case, we explicitly attach the probabilities $Q(y_k)$ and sum over all values in V . In the second case, by sampling a subset V' from Q , we have encoded these probabilities implicitly as the relative frequency of elements y_k in V'

How to do in practice (very important).

"In practice, we partition the training corpus and define a subset V' of the target vocabulary for each partition prior to training. Before training begins, we sequentially examine each target sentence in the training corpus and accumulate unique target words until the number of unique target words reaches the predefined threshold τ . The accumulated vocabulary will be used for this partition of the corpus during training. We repeat this until the end of the training set is reached. Let us refer to the subset of target words used for the i -th partition by V'_i .

Candidate Sampling – TensorFlow

Table of Contents Local

Written by Brandon McKinzie

[\[Link to article\]](#)

What is Candidate Sampling The goal is to learn a compatibility function $F(x, y)$ which says something about the compatibility of a class y with a context x . Candidate sampling: for each training example (x_i, y_i) , only need to evaluate $F(x, y)$ for a small set of classes $\{C_i\} \subset \{L\}$, where $\{L\}$ is the set of all possible classes (vocab size number of elements). We represent $F(x, y)$ as a *layer that is trained by back-prop from/within the loss function*.

C.S. for Sampled Softmax. I'll further narrow this down to my use case of having exactly 1 target class (word) at a given time. Any other classes are referred to as **negative** classes (for that example).

Sampling algorithm. For each training example (x_i, y_i) , do:

- Sample the subset $S_i \subset L$. How? By sampling from $Q(y|x)$ which gives the probability of any particular y being included in S_i .
- Create the set of **candidates**, which is just $C_i := S_i \cup y_i$.

Training task. We are given this set C_i and want to find out which element of C_i is the target class y_i . In other words, we want the posterior probability that any of the y in C_i are the target class, given what we know about C_i and x_i . We can evaluate and rearrange as usual with Bayes' rule to get:

$$\Pr(y_i^{true} = y \mid C_i, x_i) = \frac{\Pr(y_i^{true} = y \mid x_i) \cdot \Pr(C_i \mid y_i^{true} = y, x_i)}{\Pr(C_i \mid x_i)} \quad (177)$$

$$= \frac{\Pr(y \mid x_i)}{Q(y \mid x_i)} \cdot \frac{1}{K(x_i, C_i)} \quad (178)$$

where they've just defined

$$K(x_i, C_i) \triangleq \frac{\Pr(C_i \mid x_i)}{\prod_{y' \in C_i} Q(y' \mid x_i) \prod_{y' \in (L - C_i)} (1 - Q(y' \mid x_i))} \quad (179)$$

Clarifications.

- The learning function $F(x, y)$ is the *input* to our softmax. It is our neural network, excluding the softmax function.
- After training our network, it should have learned the general form

$$F(x, y) = \log(\Pr(y | x)) + K(x) \quad (180)$$

which is the general form because

$$\text{Softmax}(\log(\Pr(y | x)) + K(x)) = \frac{e^{\log(\Pr(y|x))+K(x)}}{\sum_{y'} e^{\log(\Pr(y'|x))+K(x)}} \quad (181)$$

$$= \Pr(y | x) \quad (182)$$

Note that I've been a little sloppy here, since $\Pr(y | x)$ up until the last line actually represented the (possibly) unnormalized/*relative* probabilities.

- **[MAIN TAKEAWAY]**. Time to bring it all together. Notice that we've only trained $F(x, y)$ to include *part* of what's needed to compute the probability of any y being the target given x_i and C_i ... equation 180 doesn't take into account C_i at all! Luckily we know the form of the full equation because it just the log of equation 178. We can easily satisfy that by subtracting $\log(Q(y | x))$ from $F(x, y)$ right before feeding into the softmax.

TL;DR. Train network to learn $F(x, y)$ before softmax, but instead of feeding $F(x, y)$ to softmax directly, feed

$$\text{Softmax Input: } F(x, y) - \log(Q(y | x)) \quad (183)$$

instead. That's it.

Attention Terminology

Table of Contents Local

Written by Brandon McKinzie

Generally useful info. Seems like there are a few notations floating around, and here I'll attempt to set the record straight. The order of notes here will loosely correspond with the order that they're encountered going from encoder output to decoder output.

Jargon. The people in the attention business *love* obscure names for things that don't need names at all. Terminology:

- **Attentions keys/values:** Encoder output sequence.
- **Query:** Decoder [cell] state. Typically the most recent one.
- **Scores:** Values of e_{ij} . For the Bahdanau version, in code this would be computed via

$$e_i = v^T \tanh(\text{FC}(s_{i-1}) + \text{FC}(h)) \quad (184)$$

where we'd have FC be `tf.layers.fully_connected` with `num_outputs` equal to our attention size (up to us). Note that v is a vector.

- **Alignments:** output of the softmax layer on the attention scores.
- **Memory:** The α matrix in the equation $c_i = \sum_{j=1}^{T_x} \alpha_{ij} h_j$.

When someone lazily calls some layer output the “attention”, they are usually referring to the layer *just after* the linear combination/map of encoder hidden states. You'll often see this as some vague function of the previous decoder state, context vector, and possibly even decoder output (after project), like $f(s_{i-1}, y_{i-1}, c_i)$. In 99.9% of cases, this function is just a fully connected layer (if even needed) to map back to the state size for decoder input. That is it.

From encoder to decoder. The path of information flow from encoder outputs to decoder inputs, a non-trivial process that isn't given the *attention* (heh) it deserves⁴⁸

1. **Encoder outputs.** Tensor of shape [batch size, sequence length, state size]. The state is typically some RNNCell state.
 - Note: TensorFlow's AttentionMechanism classes will actually convert this to [batch size, L_{enc} , attention size], and refer to it as the “memory”. It is also what is returned when calling `myAttentionMech.values`.

⁴⁸For some reason, the literature favors explaining the path “backwards”, starting with the highly abstracted “decoder inputs as a weighted sum of encoder states” and then breaking down what the weights are. Unfortunately, the weights are computed via a multi-stage process so that becomes very confusing very quick.

2. **Compute the scores.** The attention scores are the computation described by Luong/Bahdanau techniques. They both take an inner product of sorts on *copies* of the encoder outputs and decoder previous state (query). The main choices are:

$$\text{score}(\mathbf{h}_t, \bar{\mathbf{h}}_s) = \begin{cases} \mathbf{h}_t^T \bar{\mathbf{h}}_s & \text{dot} \\ \mathbf{h}_t^T \mathbf{W}_a \bar{\mathbf{h}}_s & \text{general} \\ \mathbf{v}_a^T \tanh(\mathbf{W}_a [\mathbf{h}_t; \bar{\mathbf{h}}_s]) & \text{concat} \end{cases} \quad (185)$$

Synonyms:
- scores
- unnormalized alignments

where the shapes are as follows (for single timestep during decoding process):

- $\bar{\mathbf{h}}_s$: [batch size, 1, state size]
- \mathbf{h}_t : [batch size, 1, state size]
- \mathbf{W}_a : [batch size, state size, state size]
- $\text{score}(\mathbf{h}_t, \bar{\mathbf{h}}_s)$: [batch size]

3. **Softmax the scores.** In the vast majority of cases, the attention scores are next fed through a softmax to convert them into a valid probability distribution. Most papers will call this some vague probability function, when in reality they are using softmax only.

Synonyms:
- softmax outputs
- attention dist.
- alignments

$$\mathbf{a}_t(s) = \text{align}(\mathbf{h}_t, \bar{\mathbf{h}}_s) \quad (186)$$

$$= \frac{\exp(\text{score}(\mathbf{h}_t, \bar{\mathbf{h}}_s))}{\sum_{s'} \exp(\text{score}(\mathbf{h}_t, \bar{\mathbf{h}}_{s'}))} \quad (187)$$

where the alignment vector \mathbf{a}_t has shape [batch size, L_{enc}]

4. **Compute the context vector.** The inner product of the softmax outputs and the raw encoder outputs. This will have shape [batch size, attention size] in TensorFlow, where attention size is from the constructor for your AttentionMechanism.

Synonyms:
- context vector
- attention

5. **Combine context vector and decoder output:** Typically with a concat. *The result is what people mean when they say “attention”.* Luong et al. denotes this as $\tilde{\mathbf{h}}_t$, the decoder output at timestep t . This is what TensorFlow means by “Luong-style mechanisms output the attention.” And yes, these are used (at least for Luong) to compute the prediction:

$$\tilde{\mathbf{h}}_t = \tanh(\mathbf{W}_c [\mathbf{c}_t, \mathbf{h}_t]) \quad (188)$$

$$p(y_t \mid y_{<t}, x) = \text{softmax}(\mathbf{W}_s \tilde{\mathbf{h}}_t) \quad (189)$$

TextRank

Table of Contents Local

Written by Brandon McKinzie

Introduction. A graph-based ranking algorithm is a way of deciding on the importance of a vertex within a graph, by taking into account global information recursively computed from the entire graph, rather than relying only on local vertex-specific information. **TextRank** is a graph-based ranking model for graphs extracted from natural language texts. The authors investigate/evaluate TextRank on unsupervised keyword and sentence extraction.

Semantic graph: one whose structure encodes meaning between the nodes (semantic elements).

The TextRank PageRank Model. In general [graph-based ranking], a vertex can be ranked based on certain properties such as the number of vertices pointing to it (in-degree), how highly-ranked *those* vertices are, etc. Formally, the authors [of *PageRank*] define the score of a vertex V_i as follows:

$$S(V_i) = (1 - d) + d * \sum_{V_j \in \text{In}(V_i)} \frac{1}{|\text{Out}(V_j)|} S(V_j) \quad \text{where } d \in \mathbb{R}[0, 1] \quad (190)$$

The factor d is usually set to 0.85.

and the damping factor d is interpreted as the probability of jumping from a given vertex⁴⁹ to another random vertex in the graph. In practice, the algorithm is implemented through the following steps:

- (1) Initialize all vertices with arbitrary values.⁵⁰
- (2) Iterate over vertices, computing equation 4.9 until convergence [of the error rate] below a predefined threshold. The error-rate, defined as the difference between the "true score" and the score computed at iteration k , $S^k(V_i)$, is *approximated* as:

$$\text{Error}^k(V_i) \approx S^k(V_i) - S^{k-1}(V_i) \quad (191)$$

⁴⁹Note that d is a single parameter for the graph, i.e. it is the same for all vertices.

⁵⁰The authors do not specify what they mean by arbitrary. What range? What sampling distribution? Arbitrary as in uniformly random? **EDIT:** The authors claim that the vertex values upon completion are not affected by the choice of initial value. Investigate!

Weighted Graphs. In contrast with the PageRank model, here we are concerned with natural language texts, which may include multiple or partial links between the units (vertices). The authors hypothesize that modifying equation 4.9 to incorporate *weighted* connections may be useful for NLP applications.

$$WS(V_i) = (1 - d) + d * \sum_{j \in \text{In}(V_i)} \frac{w_{ji}}{\sum_{V_k \in \text{Out}(V_j)} w_{jk}} WS(V_j) \quad (192)$$

w_{ij} denotes the connection between vertices V_i and V_j .

where I've shown the modified part in green. The authors mention they set all weights to random values in the interval 0-10 (no explanation).

Text as a Graph. In general, the application of graph-based ranking algorithms to natural language texts consists of the following main steps:

- (1) Identify text units that best define the task at hand, and add them as vertices in the graph.
- (2) Identify relations that connect such text unit in order to draw edges between vertices in the graph. Edges can be directed or undirected, weighted or unweighted.
- (3) Iterate the algorithm until convergence.
- (4) Sort [in reversed order] vertices based on final score. Use the values attached to each vertex for ranking/selection decisions.

KEYWORD EXTRACTION

Graph. The authors apply TextRank to extract words/phrases that are representative for a given document. The individual graph components are defined as follows:

- **Vertex:** sequence of one or more lexical units from the text.
 - In addition, we can restrict which vertices are added to the graph with syntactic filters.
 - Best filter [for the authors]: *nouns and adjectives only*.
- **Edge:** two vertices are connected if their corresponding lexical units co-occur within a window of N words⁵¹. Typically $N \in \mathbb{Z}[2, 10]$

Procedure:

- (1) **Pre-Processing:** Tokenize and annotate [with POS] the text.
- (2) **Build the Graph:** Add all [single] words to the graph that pass the syntactic filter, and connect [undirected/unweighted] edges as defined earlier (co-occurrence).
- (3) **Run algorithm:** Initialize all scores to 1. For a convergence threshold of 0.0001, usually takes about 20-30 iterations.
- (4) **Post-Processing:**
 - (i) Keep the top T vertices (by score), where the authors chose $T = |V|/3$.⁵² Remember that vertices are still individual words.
 - (ii) From the new subset of T keywords, collapse any that were adjacent in the original text in a single lexical unit.

Evaluation. The data set used is a collection of 500 abstracts, each with a set of keywords. Results are evaluated using **precision**, **recall**, and **F-measure**⁵³. The best results were obtained with a co-occurrence window of 2 [on an undirected graph], which yielded:

Precision: 31.2% Recall: 43.1% F-measure: 36.2

The authors found that larger window size corresponded with lower precision, and that directed graphs performed worse than undirected graphs.

⁵¹That is ... simpler than expected. *Can we do better?*

⁵²Another approach is to have T be a fixed value, where typically $5 < T < 20$.

⁵³ Brief terminology review:

- **Precision:** fraction of keywords extracted that are in the "true" set of keywords.
- **Recall:** fraction of "true" keywords that are in the extracted set of keywords.
- **F-score:** combining precision and recall to get a single number for evaluation:

$$F = \frac{2pr}{p + r}$$

A PR Curve plots precision as a function of recall.

Graph. Now we move to “sentence extraction for automatic summarization.”

- **Vertex:** a vertex is added to the graph for each sentence in the text.
- **Edge:** each weighted edge represents the similarity between two sentences. The authors use the following similarity measure between two sentences S_i and S_j :

$$\text{Similarity}(S_i, S_j) = \frac{|S_i \cap S_j|}{\log(|S_i|) + \log(|S_j|)} \quad (193)$$

where the numerator is the number of words that occur in both S_i and S_j .

The **procedure** is identical to the algorithm described for keyword extraction, except we run it on full sentences.

Evaluation. The data set used is 567 news articles. For each article, TextRank generates a 100-word summary (i.e. they set $T = 100$). They evaluate with the ROUGE evaluation toolkit (Ngram statistics).

Simple Baseline for Sentence Embeddings

Table of Contents Local

Written by Brandon McKinzie

Overview. It turns out that simply taking a weighted average of word vectors and doing some PCA/SVD is a competitive way of getting unsupervised word embeddings. Apparently it *beats* supervised learning with LSTMs (?!). The authors claim the theoretical explanation for this method lies in a latent variable generative model for sentences (of course).

Discussion based on paper by Arora et al., (2017).

Algorithm.

1. Compute the weighted average of the word vectors in the sentence:

$$\frac{1}{N} \sum_i^N \frac{a}{a + p(\mathbf{w}_i)} \mathbf{w}_i \quad (194)$$

The authors call their weighted average the **Smooth Inverse Frequency (SIF)**.

where \mathbf{w}_i is the word vector for the i th word in the sentence, a is a parameter, and $p(\mathbf{w}_i)$ is the (estimated) word frequency [over the entire corpus].

2. Remove the projections of the average vectors on their first principal component (“common component removal”) (y tho?).

Algorithm 1 Sentence Embedding

Input: Word embeddings $\{v_w : w \in \mathcal{V}\}$, a set of sentences \mathcal{S} , parameter a and estimated probabilities $\{p(w) : w \in \mathcal{V}\}$ of the words.

Output: Sentence embeddings $\{v_s : s \in \mathcal{S}\}$

```

1: for all sentence  $s$  in  $\mathcal{S}$  do
2:    $v_s \leftarrow \frac{1}{|s|} \sum_{w \in s} \frac{a}{a + p(w)} v_w$ 
3: end for
4: Compute the first principal component  $u$  of  $\{v_s : s \in \mathcal{S}\}$ 
5: for all sentence  $s$  in  $\mathcal{S}$  do
6:    $v_s \leftarrow v_s - uu^\top v_s$ 
7: end for
```

Theory. Latent variable generative model. The model treats corpus generation as a dynamic process, where the t -th word is produced at time step t , driven by the random walk of a **discourse vector** $c_t \in \mathbb{R}^d$ (d is size of the embedding dimension). The discourse vector is *not* pointing to a specific word; rather, it describes what is being talked about. We can tell how related (correlation) the discourse is to any word w and corresponding vector v_w by taking the inner product $c_t \cdot v_w$. Similarly, we model the probability of observing word w at time t , w_t , as:

$$\Pr[w_t \mid c_t] \propto e^{c_t \cdot v_w} \quad (195)$$

- **The Random Walk.** If we assume that c_t doesn't change much over the words in a single sentence, we can assume it stays at some c_s . The authors claim that in their previous paper they showed that the MAP⁵⁴ estimate of c_s is – up to multiplication by a scalar – the average of the embeddings of the words in the sentence.
- **Improvements/Modifications to 195.**
 1. Additive term $\alpha p(w)$ where α is a scalar. Allows words to occur even if $c_t \cdot v_w$ is very small.
 2. Common discourse vector $c_0 \in \mathbb{R}^d$. Correction term for the most frequent discourse that is often related to syntax.
- **Model.** Given the discourse vector c_s for a sentence s , the probability that w is in the sentence (at all (?)):

$$\Pr[w \mid c_s] = \alpha p(w) + (1 - \alpha) \frac{e^{\tilde{c}_s \cdot v_w}}{Z_{\tilde{c}_s}} \quad (196)$$

$$\tilde{c}_s = \beta c_0 + (1 - \beta) c_s \quad (197)$$

with $c_0 \perp c_s$ and $Z_{\tilde{c}_s}$ is a normalization constant, taken over all $w \in V$.

Embedding Calculations. Here I'll go into more depth on how the algorithm described earlier is derived. **TODO**

⁵⁴Review of MAP:

$$\theta_{MAP} = \arg \max_{\theta} \sum_i \log(p_X(x \mid \theta) p(\theta))$$

Survey of Text Clustering Algorithms

Table of Contents Local

Written by Brandon McKinzie

Aggarwal et al., “A Survey of Text Clustering Algorithms,” (2012).

Introduction. The unique characteristics for clustering *text*, as opposed to more traditional (numeric) clustering, are (1) large dimensionality but highly sparse data, (2) words are typically highly correlated, meaning the number of principal components is much smaller than the feature space, and (3) the number of words per document can vary, so we must normalize appropriately.

Common types of clustering algorithms include agglomerative clustering algorithms, partitioning algorithms, and standard parametric modeling based methods such as the EM-algorithm.

FEATURE SELECTION AND TRANSFORMATION METHODS

Feature Selection.

- **Document Frequency-Based.** Using document frequency to filter *out* irrelevant features. Dealing with certain words, like “the”, should probably be taken a step further and simply removed (stop words).
- **Term Strength.** A more aggressive technique for stop-word removal. It’s used to measure how informative a word/term t is for identifying two related documents, x and y . Denoted $s(t)$, it is defined as:

See ref 94 of the paper for more.

$$s(t) = \Pr[t \in y \mid t \in x] \quad (198)$$

So, how do we know x and y are related to begin with? One way is a user-defined cosine similarity threshold. Say we gather a set of such *pairs* and randomly identify one of the pair as the “first” document of the pair, then we can approximate $s(t)$ as

$$s(t) = \frac{\text{Num pairs in which } t \text{ occurs in both}}{\text{Num pairs in which } t \text{ occurs in the first of the pair}} \quad (199)$$

In order to prune features, the term strength may be compared to the expected strength of a term which is randomly distributed in the training documents with the same frequency. If the term strength of t is not at least two standard deviations greater than that of the random word, then it is removed from the collection.

- **Entropy-Based Ranking.** The quality of a term is measured by the entropy reduction when it is removed [from the collection]. The entropy $E(t)$ of term t in a collection of n

documents is:

$$E(t) = - \sum_{i=1}^n \sum_{j=1}^n (S_{ij} \cdot \log(S_{ij}) + (1 - S_{ij}) \cdot \log(1 - S_{ij})) \quad (200)$$

$$S_{ij} = 2^{-d_{ij}/\bar{d}} \quad (201)$$

where

- $S_{ij} \in (0, 1)$ is the similarity between doc i and j .
- d_{ij} is the distance between i and j after the term t is removed
- \bar{d} is the average distance between the documents after the term t is removed.

LSI-based Methods. Latent Semantic Indexing is based on dimensionality reduction where the new (transformed) features are a linear combination of the originals. This helps magnify the semantic effects in the underlying data. LSI is quite similar to PCA⁵⁵, except that we use an approximation of the covariance matrix C which is appropriate for the sparse and high-dimensional nature of text data.

See ref 28. of paper for more on LSI.

Let $\mathbf{A} \in \mathbb{R}^{n \times d}$ be term-document matrix, where $\mathbf{A}_{i,j}$ is the (normalized) frequency for term j in document i . Then $\mathbf{A}^T \mathbf{A} = n \cdot \Sigma$ is the (scaled) approximation to covariance matrix⁵⁶, assuming the data is mean-centered. Quick check/reminder:

$$(\mathbf{A}^T \mathbf{A})_{ij} = \mathbf{A}_{:,i}^T \mathbf{A}_{:,j} \triangleq \mathbf{a}_i^T \mathbf{a}_j \quad (202)$$

$$\approx n \cdot \mathbb{E} [\mathbf{a}_i \mathbf{a}_j] \quad (203)$$

where the expectation is technically over the underlying data distribution, which gives e.g. $P(a_i = x)$, the probability the i th word in our vocabulary having frequency x . Apparently, since the data is sparse, we don't have to worry much about it actually being mean-centered (**why?**).

As usual, we using the eigenvectors of $\mathbf{A}^T \mathbf{A}$ with the largest variance in order to represent the text⁵⁷. In addition:

One excellent characteristic of LSI is that the truncation of the dimensions removes the noise effects of synonymy and polysemy, and the similarity computations are more closely affected by the semantic concepts in the data.

⁵⁵The difference between LSI and PCA is that PCA subtracts out the means, which destroys the sparseness of the design matrix.

⁵⁶Approximation because it is based on our training data, not on true expectations over the underlying data-distribution.

⁵⁷In typical collections, only about 300 to 400 eigenvectors are required for the representation.

Non-negative Matrix Factorization. Another latent-space method (like LSI), but particularly suitable for clustering. The main characteristics of the NMF scheme:

- In LSI, the new basis system consists of a set of orthonormal vectors. This is *not* the case for NMF.
- In NMF, the vectors in the basis system **directly correspond to cluster topics**. Therefore, the cluster membership for a document may be determined by examining the largest component of the document along any of the [basis] vectors.

Assume we want to create k clusters, using our n documents and vocabulary size d . The goal of NMF is to find matrices $\mathbf{U} \in \mathbb{R}^{n \times k}$ and $\mathbf{V} \in \mathbb{R}^{d \times k}$ that minimize:

$$J = \frac{1}{2} \|\mathbf{A} - \mathbf{UV}^T\|_F^2 \quad (204)$$

$$= \frac{1}{2} \left(\text{tr}(\mathbf{AA}^T) - 2\text{tr}(\mathbf{AVU}^T) + \text{tr}(\mathbf{UV}^T\mathbf{VU}^T) \right) \quad (205)$$

$$u_{ij} \geq 0$$

$$v_{ij} \geq 0$$

Note that the columns of \mathbf{V} provide the k basis vectors which correspond to the k different clusters. We can interpret this as trying to factorize $\mathbf{A} \approx \mathbf{UV}^T$. For each row, \mathbf{a} , of \mathbf{A} (document vector), this is

$$\mathbf{a} \approx \mathbf{u} \cdot \mathbf{V}^T \quad (206)$$

$$= \sum_{i=1}^k \mathbf{u}_i \mathbf{V}_i^T \quad (207)$$

Therefore, the document vector \mathbf{a} can be rewritten as an approximate linear (non-negative) combination of the basis vector which corresponds to the k columns of \mathbf{V}^T .

Lagrange-multiplier stuff: Our optimization problem can be solved using the Lagrange method.

- Variables to optimize: All elements of both $\mathbf{U} = [u_{ij}]$ and $\mathbf{V} = [v_{ij}]$
- Constraint: non-negativity, i.e. $\forall i, j, u_{ij} \geq 0$ and $v_{ij} \geq 0$.
- Multipliers: Denote as matrices α and β , with same dimensions as \mathbf{U} and \mathbf{V} , respectively.
- Lagrangian: I'll just show it here first, and then explain in this footnote⁵⁸:

$$\mathcal{L} = J + \text{tr}(\alpha \cdot \mathbf{U}^T) + \text{tr}(\beta \cdot \mathbf{V}^T) \quad (208)$$

$$\text{where } \text{tr}(\alpha \cdot \mathbf{U}^T) = \sum_{i=1}^n \alpha_i \cdot \mathbf{u}_i = \sum_{i=1}^n \sum_{j=1}^n \alpha_{ij} u_{ij} \quad (209)$$

Any matrix multiplication with a \cdot is just a reminder to think of the matrices as column vectors.

You should think of α as a column vector of length n , and \mathbf{U}^T as a row vector of length n . The reason we prefer \mathcal{L} over just J is because now we have an *unconstrained* optimization problem.

⁵⁸ Recall that in Lagrangian minimization, \mathcal{L} takes the form of [the-function-to-be-minimized] + λ ([constraint-function] - [expected-value-of-constraint-at-optimum]). So the second term is expected to tend toward zero (i.e. critical point) at the optimal values. In our case, since our optimal value is sort-of (?) at 0 for any value of u_{ij} and/or v_{ij} , we just have a sum over [lagrange-mult] \times [variable].

- Optimization: Set the partials of \mathcal{L} w.r.t both U and V (separately) to zero⁵⁹:

$$\frac{\partial \mathcal{L}}{\partial U} = -\mathbf{A} \cdot \mathbf{V} + \mathbf{U} \cdot \mathbf{V}^T \cdot \mathbf{V} + \boldsymbol{\alpha} = 0 \quad (210)$$

$$\frac{\partial \mathcal{L}}{\partial V} = -\mathbf{A}^T \cdot \mathbf{U} + \mathbf{V} \cdot \mathbf{U}^T \cdot \mathbf{U} + \boldsymbol{\beta} = 0 \quad (211)$$

Since, ultimately, these just say [some matrix] = 0, we can multiply both sides (element-wise) by a constant ($x \times 0 = 0$). Using⁶⁰ the **Kuhn-Tucker conditions** $\alpha_{ij} \cdot u_{ij} = 0$ and $\beta_{ij} \cdot v_{ij} = 0$, we get:

$$(\mathbf{A} \cdot \mathbf{V})_{ij} \cdot u_{ij} - (\mathbf{U} \cdot \mathbf{V}^T \cdot \mathbf{V})_{ij} \cdot u_{ij} = 0 \quad (212)$$

$$(\mathbf{A}^T \cdot \mathbf{U})_{ij} \cdot v_{ij} - (\mathbf{V} \cdot \mathbf{U}^T \cdot \mathbf{U})_{ij} \cdot v_{ij} = 0 \quad (213)$$

- Update rules:

$$u_{ij} = \frac{(\mathbf{A} \cdot \mathbf{V})_{ij} \cdot u_{ij}}{(\mathbf{U} \cdot \mathbf{V}^T \cdot \mathbf{V})_{ij}} \quad (214)$$

$$v_{ij} = \frac{(\mathbf{A}^T \cdot \mathbf{U})_{ij} \cdot v_{ij}}{(\mathbf{V} \cdot \mathbf{U}^T \cdot \mathbf{U})_{ij}} \quad (215)$$

DISTANCE-BASED CLUSTERING ALGORITHMS

One challenge in clustering short segments of text (e.g., tweets) is that exact keyword matching may not work well. One general strategy for solving this problem is to expand text representation by exploiting related text documents, which is related to smoothing of a document language model in information retrieval.

See ref. 66 in the paper for computing similarities of short text segments.

Agglomerative and Hierarchical Clustering. “Agglomerative” refers to the process of bottom-up clustering to build a tree – at the bottom are leaves (documents) and internal nodes correspond to the merged groups of clusters. The different methods for merging groups of documents for the different agglomerative methods are as follows:

- **Single Linkage Clustering.** Defines similarity between two groups (clusters) of documents as the largest similarity between any pair of documents from these two groups. First, (1) compute all similarity pairs [between documents; ignore cluster labels], then (2) sort in decreasing order, and (3) walk through the list in that order, merging clusters if the pair belong to different clusters. One drawback is *chaining*: the resulting clusters assume transitivity of similarity⁶¹.

⁵⁹Recall that the Lagrangian consists entirely of traces (re: scalars). Derivatives of traces with respect to matrices output the same dimension as that matrix, and derivatives are taken element-wise as always.

⁶⁰i.e. the equations that follow are *not* the KT conditions, they just *use/exploit* them...

⁶¹Here, transitivity of similarity means if A is similar to B , and B is similar to C , then A is similar to C . This is not guaranteed by any means for textual similarity, and so we can end up with A and Z in the same cluster, even though they aren’t similar at all.

- **Group-Average Linkage Clustering.** Similarity between two clusters is the *average* similarity over all unique pairwise combinations of documents from one cluster to the other. One way to speed up this computation with an approximation is to just compute the similarity between the mean vector of either cluster.
- **Complete Linkage Clustering.** Similarity between two clusters is the *worst-case* similarity between any pair of documents.

Distance-Based Partitioning Algorithms.

- **K-Medoid Clustering.** Use a set of points from training data as anchors (medoids) around which the clusters are built. Key idea is we are using an optimal set of representative documents *from the original corpus*. The set of k reps is successively improved via randomized inter-changes. In each iteration, we replace a randomly picked rep in the current set of medoids with a randomly picked rep from the collection, if it improves the clustering objective function. This approach is applied until convergence is achieved.
- **K-Means Clustering.** Successively (1) assigning points to the nearest cluster centroid and then (2) re-computing the centroid of each cluster. Repeat until convergence. Requires typically few iterations (about 5 for many large data sets). Disadvantage: sensitive to initial set of seeds (initial cluster centroids). One method for improving the initial set of seeds is to use some supervision - e.g. initialize with k pre-defined topic vectors (see ref. 4 in paper for more).

K-Medoid isn't great for clustering text, especially short texts.

Hybrid Approach: Scatter-Gather Method. Use a hierarchical clustering algorithm on a sample of the corpus in order to find a robust initial set of seeds. This robust set of seeds is used in conjunction with a standard k-means clustering algorithm in order to determine good clusters. **TODO:** resume note-taking; page 19/52 of PDF.

Scatter-Gather is discussed in detail in ref. 25 of the paper

PROBABILISTIC DOCUMENT CLUSTERING AND TOPIC MODELS

Overview. Primary assumptions in any topic modeling approach:

From pg. 31/52 of paper.

- The n documents in the corpus are assumed to each have a probability of belonging to one of k topics. Denote the probability of document D_i belonging to topic T_j as $\Pr [T_j | D_i]$. This allows for *soft cluster membership* in terms of probabilities.
- Each topic is associated with a probability vector, which quantifies the probability of the different terms in the lexicon for that topic. For example, consider a document that belongs completely to topic T_j . We denote the probability of term t_l occurring in that document as $\Pr [t_l | T_j]$.

The two main methods for topic modeling are **Probabilistic Latent Semantic Indexing** (PLSA) and **Latent Dirichlet Allocation** (LDA).

PLSA. We note that the two aforementioned probabilities, $\Pr [T_j | D_i]$ and $\Pr [t_l | T_j]$ allow us to calculate $\Pr [t_l | D_i]$: the probability that term t_l occurs in some document D_i :

$$\Pr [t_l | D_i] = \sum_{j=1}^k \Pr [t_l | T_j] \cdot \Pr [T_j | D_i] \quad (216)$$

which should be interpreted as a weighted average⁶². From here, we can generate a $n \times d$ matrix of probabilities.

Recall that we also have our $n \times d$ term-document matrix \mathbf{X} , where $\mathbf{X}_{i,l}$ gives the number of times term l occurred in document D_i . This allows us to do maximum likelihood! Our negative log-likelihood, J can be derived as follows:

$$J = -\log (\Pr [\mathbf{X}]) \quad (217)$$

$$= -\log \left(\prod_{i,l} \Pr [t_l | D_i]^{\mathbf{X}_{i,l}} \right) \quad (218)$$

$$= -\sum_{i,l} \mathbf{X}_{i,l} \cdot \log (\Pr [t_l | D_i]) \quad (219)$$

Interpret $\Pr [\mathbf{X}]$ as the joint probability of observing the words in our data and with their assoc. frequencies.

and we can plug-in eqn 216 to for evaluating $\Pr [t_l | D_i]$. We want to optimize the value of J , subject to the constraints:

$$(\forall T_j) : \sum_l \Pr [t_l | T_j] = 1 \quad (\forall D_i) : \sum_j \Pr [T_j | D_i] = 1 \quad (220)$$

This can be solved with a Lagrangian method, similar to the process for NMF described earlier. See page 33/52 of the paper for details.

Latent Dirichlet Allocation (LDA). The term-topic probabilities and topic-document probabilities are modeled with a *Dirichlet distribution* as a prior⁶³. Typically preferred over PLSI because PLSI more prone to overfitting.

⁶²This is actually pretty bad notation, and borderline incorrect. $\Pr [T_j | D_i]$ is NOT a conditional probability! It is our prior! It is literally $\Pr [\text{ClusterOf}(D_i) = T_j]$.

⁶³LDA is the Bayesian version of PLSI

ONLINE CLUSTERING WITH TEXT STREAMS

Reference List: [3]: A Framework for Clustering Massive Text and Categorical Data Streams; [112]: Efficient Streaming Text Clustering; [48]: Bursty feature representation for clustering text streams; [61]: Clustering Text Data Streams (Liu et al.)

Overview. Maintaining text clusters in real time. One method is the **Online Spherical K-Means Algorithm** (OSKM)⁶⁴.

See ref. 112 for more on OSKM

Condensed Droplets Algorithm. I'm calling it that because they don't call it anything – it is the algorithm in [3].

- **Fading function:** $f(t) = 2^{-\lambda \cdot t}$. A time-dependent weight for each data point (text stream). Non-monotonic decreasing; decays uniformly with time.
- **Decay rate:** $\lambda = 1/t_0$. Inverse of the half-life of the data stream.

When a cluster is created by a new point, it is allowed to remain as a trend-setting outlier for at least one half-life. During that period, if at least one more data point arrives, then the cluster becomes an active and mature cluster. If not, the trend-setting outlier is recognized as a true anomaly and is removed from the list of current clusters (*cluster death*). Specifically, this happens when the (weighted) number of points in the [single-point] cluster is 0.5. The same criterion is used to define the death of mature clusters. The statistics of the data points are referred to as **condensed droplets**, which represent the word distributions within a cluster, and can be used in order to compute the similarity of an incoming data point to the cluster. Main idea of algorithm is as follows:

1. Initialize empty set of clusters $\mathcal{C} = \{\}$. As new data points arrive, unit clusters containing individual data points are created. Once a maximum number k of such clusters have been created, we can begin the process of online cluster maintenance.
2. For a new data point X , compute its similarity to each cluster C_j , denoted as $S(X, C_j)$.
 - If $S(X, C_{best}) > \text{thresh}_{\text{outlier}}$, or if there are no inactive clusters left⁶⁵, insert X to the cluster with maximum similarity.
 - Otherwise, a new cluster is created⁶⁶ containing the solitary data point X .

⁶⁴ Authors only provide a very brief description, which I'll just copy here:

This technique divides up the incoming stream into small segments, each of which can be processed effectively in main memory. A set of k-means iterations are applied to each such data segment in order to cluster them. The advantage of using a segment-wise approach for clustering is that since each segment can be held in main memory, we can process each data point multiple times as long as it is held in main memory. In addition, the centroids from the previous segment are used in the next iteration for clustering purposes. A decay factor is introduced in order to age- out the old documents, so that the new documents are considered more important from a clustering perspective.

⁶⁵ We specify some max allowed number of clusters k .

⁶⁶ The new cluster replaces the least recently updated inactive cluster.

Misc.

- Mandatory read: reference [61]. Details phrase extraction/**topic signatures**. The use of using phrases instead of individual words is referred to as **semantic smoothing**.
- For *dynamic* (and more recent) topic modeling, see reference [107] of the paper, titled “A probabilistic model for online document clustering with application to novelty detection.”

SEMI-SUPERVISED CLUSTERING

Overview. Useful when we have any prior knowledge about the kinds of clusters available in the underlying data. Some approaches:

- Incorporate this knowledge when seeding the cluster centroids for k -means clustering.
- Iterative EM approach: unlabeled documents are assigned labels using a naive Bayes approach on the currently labeled documents. These newly labeled documents are then again used for re-training a Bayes classifier. Iterate to convergence.
- Graph-based approach: graph nodes are documents and the edges are similarities between the connected documents (nodes). We can incorporate prior knowledge by adding certain edges between nodes that we know are similar. A *normalized cut algorithm* is then applied to this graph in order to create the final clustering.

We can also use partially supervised methods in conjunction with pre-existing categorical *hierarchies*.

Deep Sentence Embedding Using LSTMs

Table of Contents Local

Written by Brandon McKinzie

Palangi et al., “Deep Sentence Embeddings Using Long Short-Term Memory Networks: Analysis and Application to Information Retrieval,” (2016).

Abstract. Sentence embeddings using LSTM cells, which automatically attenuate unimportant words and detect salient keywords. Main emphasis on applications for document retrieval (matching a query to a document⁶⁷).

Introduction. Sentence embeddings are learned using a loss function defined on *sentence pairs*. For example, the well-known Paragraph Vector⁶⁸ is learned in an unsupervised manner as a distributed representation of sentences and documents, which are then used for sentiment analysis.

The authors appear to use a dataset of their own containing examples of (search-query, clicked-title) for a search engine. Their training objective is to maximize the similarity between the two vectors mapped by the LSTM-RNN from the query and the clicked document, respectively. One very interesting claim to pay close attention to:

*We further show that different cells in the learned model indeed correspond to **different topics**, and the keywords associated with a similar topic activate the same cell unit in the model.*

Related Work. (Identified by reference number)

- [2] Good for sentiment, but doesn’t capture fine-grained sentence structure.
- [6] Unsupervised embedding method trained on the BookCorpus [7]. Not good for document retrieval task.
- [9] Semi-supervised Recursive Autoencoder (RAE) for sentiment prediction.
- [3] DSSM (uses bag-of-words) and [10] CLSM (uses bag of n-grams) models for IR and also sentence embeddings.
- [12] Dynamic CNN for sentence embeddings. Good for sentiment prediction and question type classification. In [13], a CNN is proposed for **sentence matching**⁶⁹

⁶⁷Note that this similar to topic extraction.

⁶⁸Q. V. Le and T. Mikolov, Distributed representations of sentences and documents.

⁶⁹Might want to look into this.

Basic RNN. The information flow (sequence of operations) is enumerated below.

1. Encode t th word [of the given sentence] in one-hot vector $\mathbf{x}(t)$.
2. Convert $\mathbf{x}(t)$ to a letter tri-gram vector $\mathbf{l}(t)$ using fixed hashing operator⁷⁰ \mathbf{H} :

$$\mathbf{l}(t) = \mathbf{H}\mathbf{x}(t) \quad (221)$$

3. Compute the hidden state $\mathbf{h}(t)$, which is the sentence embedding for $t = T$, the length of the sentence.

$$\mathbf{h}(t) = \tanh(\mathbf{U}\mathbf{l}(t) + \mathbf{W}\mathbf{h}(t-1) + \mathbf{b}) \quad (222)$$

where \mathbf{U} and \mathbf{W} are the usual parameter matrices for the input/recurrent paths, respectively.

LSTM. With peephole connections that expose the internal cell state s to the sigmoid computations. I'll rewrite the standard LSTM equations from my textbook notes, but with the modifications for peephole connections:

$$f_i^{(t)} = \sigma \left(b_i^f + \sum_j U_{i,j}^f x_j^{(t)} + \sum_j W_{i,j}^f h_j^{(t-1)} + \sum_j P_{i,j}^f s_j^{(t-1)} \right) \quad (223)$$

$$s_i^{(t)} = f_i^{(t)} \odot s_i^{(t-1)} + g_i^{(t)} \odot \sigma \left(b_i + \sum_j U_{i,j} x_j^{(t)} + \sum_j W_{i,j} h_j^{(t-1)} \right) \quad (224)$$

$$g_i^{(t)} = \sigma \left(b_i^g + \sum_j U_{i,j}^g x_j^{(t)} + \sum_j W_{i,j}^g h_j^{(t-1)} + \sum_j P_{i,j}^g s_j^{(t-1)} \right) \quad (225)$$

$$q_i^{(t)} = \sigma \left(b_i^o + \sum_j U_{i,j}^o x_j^{(t)} + \sum_j W_{i,j}^o h_j^{(t-1)} + \sum_j P_{i,j}^o s_j^{(t)} \right) \quad (226)$$

The final hidden state can then be computed via

$$h_i^{(t)} = \tanh(s_i^{(t)}) \odot q_i^{(t)} \quad (227)$$

⁷⁰Details aside, the hashing operator serves to lower the dimensionality of the inputs a bit. In particular we use it to convert one-hot word vectors into their letter tri-grams. For example, the word “good” gets surrounded by hashes, ‘#good#’, and then hashed from the one-hot vector to vectorized tri-grams, “#go”, “goo”, “ood”, “od#”.

Learning method. We want to maximize the likelihood of the clicked document given query, which can be formulated as the following optimization problem:

$$L(\mathbf{\Lambda}) = \min_{\mathbf{\Lambda}} \left\{ -\log \prod_{r=1}^N \Pr [D_r^+ | Q_r] \right\} = \min_{\mathbf{\Lambda}} \sum_{r=1}^N l_r(\mathbf{\Lambda}) \quad (228)$$

$$l_r(\mathbf{\Lambda}) = \log \left(1 + \sum_{j=1}^n e^{-\gamma \cdot \Delta_{r,j}} \right) \quad (229)$$

where

- N is the number of (query, clicked-doc) pairs in the corpus, while n is the number of negative samples used during training.
- D_r^+ is the clicked document for r th query.
- $\Delta_{r,j} = R(Q_r, D_r^+) - R(Q_r, D_{r,j}^-)$ (R is just cosine similarity)⁷¹.
- $\mathbf{\Lambda}$ is all the parameter matrices (and biases) in the LSTM.

The authors then describe standard BPTT updates with momentum, which need not be detailed here. See the “Algorithm 1” figure in the paper for extremely detailed pseudo-code of the training procedure.

⁷¹ Note that $\Delta_{r,j} \in [-2, 2]$. We use γ as a scaling factor so as to expand this range.

Clustering Massive Text Streams

Table of Contents Local

Written by Brandon McKinzie

Aggarwal et al., “A Framework for Clustering Massive Text and Categorical Data Streams,” (2006).

Overview. Authors present an online approach for clustering massive text and categorical data streams with the use of a statistical summarization methodology. First, we will go over the process of storing and maintaining the data structures necessary for the clustering algorithm. Then, we will discuss the differences which arise from using different kinds of data, and the empirical results.

Maintaining Cluster Statistics. The data stream consists of d -dimensional records, where each dimension corresponds to the numeric frequency of a given word in the vector space representation. Each data point is weighted by the **fading function** $f(t)$, a non-monotonic decreasing function which decays uniformly with time t . The authors define the **half-life** of a data point (e.g. a tweet) as:

$$t_0 \text{ s.t. } f(t_0) = \frac{1}{2}f(0) \quad (230)$$

and, similarly, the **decay-rate** as its inverse, $\lambda = 1/t_0$. Thus we have $f(t) = 2^{-\lambda \cdot t}$.

To achieve greater accuracy in the clustering process, we require a high level of granularity in the underlying data structures. To do this, we will use a process in which condensed clusters of data points are maintained, referred to as **cluster droplets**. We define them differently for the case of text and categorical data, beginning with categorical:

- **Categorical.** A cluster droplet $\mathcal{D}(t, \mathcal{C})$ for a set of categorical data points \mathcal{C} at time t is defined as the tuple:

$$\mathcal{D}(t, \mathcal{C}) \triangleq (D\bar{F}2, D\bar{F}1, n, w(t), l) \quad (231)$$

where

- Entry k of the vector $D\bar{F}2$ is the (weighted) number of points in cluster \mathcal{C} where the i th dimension had value x and the j dimension had value y . In other words, all pairwise combinations of values in the categorical vector⁷². $\sum_{i=1}^d \sum_{j \neq i}^d v_i v_j$ entries total⁷³.
- Similarly, $D\bar{F}1$ consists of the (weighted) counts that some dimension i took on the value x . $\sum_{i=1}^d v_i$ entries total.

⁷²This is intentionally written hand-wavy because I’m really concerned with *text* streams and don’t want to give this much space.

⁷³ v_i is the number of values the i th categorical dimension can take on.

- $w(t)$ is the sum of the weights of the data points at time t .
- l is the time stamp of the last time a data point was added to the cluster.
- **Text.** Can be viewed as an example of a sparse numeric data set. A cluster droplet $\mathcal{D}(t, \mathcal{C})$ for a set of text data points \mathcal{C} at time t is defined as the tuple:

$$\mathcal{D}(t, \mathcal{C}) \triangleq (D\bar{F}2, D\bar{F}1, n, w(t), l) \quad (232)$$

where

- $D\bar{F}2$ contains $3 \cdot wb \cdot (wb - 1)/2$ entries, where wb is the number of distinct words in the cluster \mathcal{C} .
- $D\bar{F}1$ contains $2 \cdot wb$ entries.
- n is the number of data points in the cluster \mathcal{C} .

Cluster Droplet Maintenance.

1. We first start of with k trivial clusters (the first k data points that arrived).
2. When a new point \bar{X} arrives, the cosine similarity to each cluster's $D\bar{F}1$ is computed.
3. \bar{X} is inserted into the cluster for which this is a maximum, so long as the associated $S(\bar{X}, D\bar{F}1) > \text{thresh}$, a predefined threshold. If not above the threshold *and* some **inactive cluster** (wtf does this mean??) exists, a new cluster is created containing the solitary point \bar{X} , which replaces the inactive cluster. If not above threshold but no inactive clusters, then we just insert it into the max similarity cluster anyway.
4. If \bar{X} was inserted (i.e. didn't replace an inactive cluster), then we need to:
 - (a) Update the statistics to reflect the decay of the data points at the current moment in time⁷⁴. This is done by multiplying entries in the droplet vectors by $2^{-\lambda \cdot (t-l)}$.
 - (b) Add the statistics for each newly arriving data point to the cluster statistics.

⁷⁴In other words, the statistics for a cluster do not decay, until a new point is added to it.

Supervised Universal Sentence Representations (InferSent)

Table of Contents Local

Written by Brandon McKinzie

Conneau et al., “Supervised Learning of Universal Sentence Representations from Natural Language Inference Data,” Facebook AI Research (2017).

Overview. Authors claim universal sentence representations trained using the supervised data of the Stanford Natural Language Inference (SNLI) dataset can consistently outperform unsupervised methods like SkipThought on a wide range of transfer tasks. They emphasize that training on NLI tasks in particular results in embeddings that perform well in transfer tasks. Their best encoder is a Bi-LSTM architecture with max pooling, which they claim is SOTA when trained on the SNLI data.

The Natural Language Inference Task. Also known as *Recognizing Textual Entailment* (RTE). The SNLI data consists of sentence pairs labeled as one of entailment, contradiction, or neutral. Below is a typical architecture for training on SNLI.



Note that the same sentence encoder is used for both u and v . To obtain a sentence vector from a BiLSTM encoder, they experiment with (1) the average h_t over all t (**mean pooling**), and (2) selecting the max value over each dimension of the hidden units [over all timesteps] (**max pooling**)⁷⁵.

⁷⁵Since the authors have already mentioned that BiLSTM did the best, I won't go over the other architectures they tried: self-attentive networks, hierarchical convnet, vanilla LSTM/GRU.

Dist. Rep. of Sentences from Unlabeled Data (FastSent)

Table of Contents Local

Written by Brandon McKinzie

Hill et al., “Learning Distributed Representations of Sentences from Unlabelled Data,” (2016).

Overview. A systematic comparison of models that learn distributed representations of phrases/sentences from unlabeled data. Deeper, more complex models are preferable for representations to be used in supervised systems, but shallow log-linear models work best for building representation spaces that can be decoded with simple spatial distance metrics.

Authors propose two new phrase/sentence representation learning objectives:

1. **Sequential Denoising Autoencoders** (SDAEs)
2. **FastSent**: a sentence-level log-linear BOW model.

Distributed Sentence Representations. Existing models trained on text:

- **SkipThought Vectors** (Kiros et al., 2015). Predict target sentences $S_{i\pm 1}$ given source sentence S_i . Sequence-to-sequence model.
- **ParagraphVector** (Le and Mikolov, 2014). Defines 2 log-linear models:
 1. **DBOW**: learns a vector \mathbf{s} for every sentence S in the training corpus which, together with word embeddings v_w , define a softmax distribution to predict words $w \in S$ given S .
 2. **DM**: select k-grams of consecutive words $\{w_i \cdots w_{i+k} \in S\}$ and the sentence vector \mathbf{s} to predict w_{i+k+1} .
- **Bottom-Up Methods**. Train **CBOW** and **Skip-Gram** word embeddings on the Books corpus.

The authors use **gensim** to implement ParagraphVector.

Models trained on *structured* (and freely-available) resources:

- **DictRep** (Hill et al., 2015a). Map dictionary definitions to pre-trained word embeddings, using either BOW or RNN-LSTM encoding.
- **NMT**. Consider sentence representations learned by sequence-to-sequence NMT models.

Novel Text-Based Methods.

- **Sequential (Denoising) Autoencoders.** To avoid needing coherent inter-sentence narrative, try this representation-learning objective based on DAEs. For a given sentence S and **noise function** $N(S \mid p_o, p_x)$ (where $p_o, p_x \in [0, 1]$), the approach is as follows:
 1. For each $w \in S$, N deletes w with probability p_o .
 2. For each non-overlapping bigram $w_i w_{i+1} \in S$, N swaps w_i and w_{i+1} with probability p_x .

Authors recommend
 $p_o = p_x = 0.1$

We then train the same LSTM-based encoder-decoder architecture as NMT, but with the denoising objective to predict (as target) the original source sentence S given a corrupted version $N(S \mid p_o, p_x)$ (as source).

- **FastSent.** Designed to be a more efficient/quicker to train version of SkipThought.

Latent Dirichlet Allocation

Table of Contents Local

Written by Brandon McKinzie

Blei et al., “Latent Dirichlet Allocation,” (2003).

Introduction. At minimum, one should be familiar with generative probabilistic models, mixture models, and the notion of latent variables before continuing. The “Dirichlet” in LDA of course refers to the **Dirichlet distribution**, which is a generalization of the beta distribution, B . It’s PDF is defined as⁷⁶⁷⁷:

$$\text{Dir}(\mathbf{x}; \boldsymbol{\alpha}) = \frac{1}{B(\boldsymbol{\alpha})} \prod_{i=1}^K x_i^{\alpha_i - 1} \quad \text{where} \quad B(\boldsymbol{\alpha}) = \frac{\prod_{i=1}^K \Gamma(\alpha_i)}{\Gamma(\sum_{i=1}^K \alpha_i)} \quad (233)$$

$$\sum_{i=1}^K x_i = 1$$

$(\forall i \in [1, K]) : x_i \geq 0$

Main things to remember about LDA:

- Generative probabilistic model for collections of discrete data such as text corpora.
- Three-level **hierarchical Bayesian model**. Each document is a mixture of topics, each topic is an infinite mixture over a set of topic probabilities.

Condensed comparisons/history of related models leading up to LDA:

- **TF-IDF**. Design matrix $\mathbf{X} \in \mathbb{R}^{V \times M}$, where M is the number of docs, and $\mathbf{X}_{i,j}$ gives the TF-IDF value for i th word in vocabulary and corresp. to document j .
- **LSI**:⁷⁸ Performs SVD on the TF-IDF design matrix \mathbf{X} to identify a linear subspace in the space of tf-idf features that captures most of the variance in the collection.
- **pLSI**: **TODO**

pLSI is incomplete in that it provides no probabilistic model at the level of documents. In pLSI, each document is represented as a list of numbers (the mixing proportions for topics), and there is no generative probabilistic model for these numbers.

⁷⁶ Recall that for positive integers n , $\Gamma(n) = (n-1)!$.

⁷⁷The Dirichlet distribution is **conjugate** to the multinomial distribution. TODO: Review how to interpret this.

⁷⁸Recall that LSI is basically PCA but without subtracting off the means

Model. LDA assumes the following generative process for each document (word sequence) \mathbf{w} :

1. $N \sim \text{Poisson}(\lambda)$: Sample N , the number of words (length of \mathbf{w}), from $\text{Poisson}(\lambda) = e^{-\lambda} \frac{\lambda^n}{n!}$. The parameter λ *should* represent the average number of words per document.
2. $\theta \sim \text{Dir}(\alpha)$: Sample k -dimensional vector θ from the Dirichlet distribution (eq. 233), $\text{Dir}(\alpha)$. k is the number of topics (pre-defined by us). Recall that this means θ lies in the $(k-1)$ simplex. The Dirichlet distribution thus tells us the probability density of θ over this simplex – it defines the probability of θ being at a given position on the simplex.
3. Do the following N times to generate the words for this document.
 - (a) $z_n \sim \text{Multinomial}(\theta)$. Sample a topic z_n .
 - (b) $w_n \sim \text{Pr}[w_n | z_n, \beta]$: Sample a word w_n from $\text{Pr}[w_n | z_n, \beta]$, a “multinomial probability conditioned on topic z_n .”⁷⁹ The parameter β gives the distribution of words given a topic:

$$\beta_{ij} = \text{Pr}[w_j | z_i] \quad (234)$$

In other words, we really sample $w_n \sim \beta_{i,:}$:


The defining equations for LDA are thus:

$$\text{Pr}[\theta, \mathbf{z}, \mathbf{w} | \alpha, \beta] = \text{Pr}[\theta | \alpha] \prod_{n=1}^N \text{Pr}[z_n | \theta] \text{Pr}[w_n | z_n, \beta] \quad (235)$$

$$\text{Pr}[\mathbf{w} | \alpha, \beta] = \int \text{Pr}[\theta' | \alpha] \left(\prod_{n=1}^N \sum_{z'_n} \text{Pr}[z'_n | \theta'] \text{Pr}[w_n | z'_n, \beta] \right) d\theta' \quad (236)$$

$$\text{Pr}[\mathcal{D} = \{\mathbf{w}^{(1)}, \dots, \mathbf{w}^{(M)}\} | \alpha, \beta] = \prod_{d=1}^M \text{Pr}[\mathbf{w}^{(d)} | \alpha, \beta] \quad (237)$$

Below is the plate notation for LDA, followed by an interpretation:



- **Outermost Variables:** α and β . Both represent a (Dirichlet) prior distribution: α parameterizes the probability of a given *topic*, while β a given *word*.
- Document Plate. M is the number of documents, θ_m gives the true distribution of topics for document m ⁸⁰.

⁷⁹**TODO:** interpret meaning of the multinomial distributions here. Seems a bit different than standard interp...

⁸⁰In other words, the meaning of $\theta_{m,i} = x$ is “ x percent of document m is about topic i .”

- **Topic/Word Place.** z_{mn} is the topic for word n in doc m , and w_{mn} is the word. It is shaded gray to indicate it is the only **observed variable**, while all others are **latent variables**.

Theory. I'll quickly summarize and interpret the main theoretical points. Without having read all the details, this won't be of much use (i.e. it is for someone who has read the paper already).

- **LDA and Exchangeability.** We assume that each document is a bag of words (order doesn't matter; frequency still does) *and* a bag of topics. In other words, a document of N words *is* an unordered list of words and topics. De Finetti's theorem tells us that we can model the joint probability of the words and topics as if a random parameter θ were drawn from some distribution and then the variables within w, z were **conditionally independent given** θ . LDA posits that a good distribution to sample θ from is a Dirichlet distribution.
- **Geometric Interpretation: TODO**

Inference and Parameter Estimation. As usual, we need to find a way to compute the posterior distribution of the hidden variables given a document w :

$$\Pr[\theta, z \mid w, \alpha, \beta] = \frac{\Pr[\theta, z, w \mid \alpha, \beta]}{\Pr[w \mid \alpha, \beta]} \quad (238)$$

Computing the denominator exactly is intractable. Common approximate inference algorithms for LDA include Laplace approximation, variational approximation, and Markov Chain Monte Carlo.

Conditional Random Fields

Table of Contents Local

Written by Brandon McKinzie

Lafferty et al., “Conditional Random Fields: Probabilistic Models for Segmenting and Labeling Sequence Data,” (2001).

Introduction. CRFs offer improvements to HMMs, MEMMs, and other discriminative Markov models. MEMMs and other non-generative models share a weakness called the **label bias problem**: the transitions leaving a given state compete only against each other, rather than against all other transitions in the model. The key difference between CRFs and MEMMs is that a CRF has a single exponential model for the joint probability of the entire sequence of labels given the observation sequence.

The Label Bias Problem. Recall that MEMMs are run left-to-right. One way of interpreting such a model is to consider how the probabilities (of state sequences) are distributed as we continue through the sequence of observations. The issue with MEMMs is that there’s nothing we can do if, somewhere along the way, we observe something that makes one of these state paths extremely likely/unlikely; we can’t redistribute the probability mass amongst the various allowed paths. The CRF solution:

Account for whole state sequences at once by letting some transitions “vote” more strongly than others depending on the corresponding observations. This implies that score mass will not be conserved, but instead individual transitions can “amplify” or “dampen” the mass they receive.

Conditional Random Fields. Here we formalize the model and notation. Let \mathbf{X} be a random variable over data sequences to be labeled (e.g. over all words/sentences), and let \mathbf{Y} the random variable over corresponding label sequences⁸¹. Formal definition:

Let $G = (V, E)$ be a graph such that $Y = (Y_v)_{v \in V}$, so that Y is indexed by the vertices of G . Then (X, Y) is a CRF if, when conditioned on X , the random variables Y_v obey the Markov property with respect to the graph:

$$Pr[Y_v | X, Y_w, w \neq v] = Pr[Y_v | X, Y_w, w \sim v] \quad (239)$$

where $w \sim v$ means that w and v are neighbors in G .

In human-speak, all this means is a CRF is a random field (discrete set of random-valued points in a space) where all points (i.e. globally) are conditioned on \mathbf{X} . If the graph $G = (V, E)$ of \mathbf{Y} is a tree, its cliques⁸² are the edges and vertices. By the fundamental theorem of random fields:

⁸¹We assume all components \mathbf{Y}_i can only take on values in some finite label set \mathcal{Y} .

⁸²A clique is a subset of vertices in an undirected graph such that every two distinct vertices in the clique are adjacent

$$p_{\theta}(\mathbf{y} \mid \mathbf{x}) \propto \exp \left(\sum_{e \in E, k} \lambda_k f_k(e, \mathbf{y}|_e, \mathbf{x}) + \sum_{v \in V, k} \mu_k g_k(v, \mathbf{y}|_v, \mathbf{x}) \right) \quad (240)$$

where $\mathbf{y}|_S$ is the set of components of \mathbf{y} associated with the vertices in subgraph S . We assume the K feature [functions] f_k and g_k are given and fixed.

Our estimation problem is thus to determine parameters $\theta = (\lambda_1, \lambda_2, \dots; \mu_1, \mu_2, \dots)$ from the labeled training data.

Attention Is All You Need

Table of Contents Local

Written by Brandon McKinzie

Vaswani et al., “Attention Is All You Need,” (2017)

Overview. Authors refer to sequence *transduction* models a lot – just a fancy way of referring to models that transform input sequences into output sequences. Authors propose new architecture, the **Transformer**, based solely on attention mechanisms (no recurrence!).

Model Architecture.

- **Encoder.** $N = 6$ identical layers, each with 2 sublayers: (1) a **multi-head** self-attention mechanism and (2) a position-wise FC feed-forward network. They employ a residual connection around each of the two, followed by layer normalization.
- **Decoder.** $N = 6$ with 3 sublayers each. In addition to the two sublayers described for the encoder, the decoder has a third sublayer, which performs **multi-head** attention over the output of the encoder stack. Same residual connections and layer norm.

Attention. An attention function can be described as a mapping:

$$\text{Attn}(\text{query}, \{(k_1, v_1), \dots, \}) \Rightarrow \sum_i f n(\text{query}, k_i) v_i \quad (241)$$

$$\Rightarrow \text{output} \quad (242)$$

where the query, keys, values, and output are all vectors.

- **Scaled Dot-Product Attention.**

1. **Inputs:** queries q , keys k of dimension d_k , values v of dimension d_v
2. **Dot Products:** Compute $\forall k : (q \cdot k) / \sqrt{d_k}$.
3. **Softmax:** on each dot product above. This gives the weights on the values shown earlier.

Appears that $d_q \equiv d_k$.

In practice, this is done simultaneously for all queries in a set via the following matrix equation:

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V \quad (243)$$

Note that this is identical to the standard dot-product attention mechanism, except for the *scaling* factor (hence the name) of $1/\sqrt{d_k}$. The scaling factor is motivated by the fact that additive attention outperforms dot-product attention for large d_k and the authors stipulate this is due to the softmax having small gradients in this case, due to the large dot products⁸³.

⁸³Thus, they insert the $1/\sqrt{d_k}$ to counteract this effect.

- **Multi-Head Attention.** Basically just doing some number h of parallel attention computations. Before each of these, the queries, keys, and values are linearly projected with different, learned linear projections to d_k , d_k and d_v dimensions respectively (and then fed to their respective attention function). The h outputs are then concatenated and once again projected, resulting in the final values.

$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h) W^O \quad (244)$$

$$\text{where head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V) \quad (245)$$

The authors employ $h = 8$, $d_k = d_v = d_{\text{model}}/h = 64$.

$$W_i^Q \in \mathbb{R}^{d_{\text{model}} \times d_k}$$

$$W_i^K \in \mathbb{R}^{d_{\text{model}} \times d_k}$$

$$W_i^V \in \mathbb{R}^{d_{\text{model}} \times d_v}$$

$$W^O \in \mathbb{R}^{hd_v \times d_{\text{model}}}$$

The Transformer uses multi-headed attention in 3 ways:

1. **Encoder-decoder attention:** the normal kind. Queries are previous decoder layer (state?), and memory keys and values come from output of the encoder.
2. **Encoder self-attention:** all of the keys, values, and queries come from the previous *layer* in the encoder. Each position in the encoder can attend to all positions in the previous layer of the encoder.
3. **Decoder self-attention:** Similarly, self-attention layers in the decoder allow each position in the decoder to attend to all positions in the decoder up to and including that position. Need to prevent leftward information flow in the decoder to preserve the auto-regressive property. This is implemented inside of a scaled dot-product attention by masking out (setting to $-\infty$) all values in the input of the softmax which correspond to illegal connections.

Other Components.

- **Position-wise Feed-Forward Networks (FFN):** each layer of the encoder and decoder contains a FC FFN, applied to each position separately and identically:

$$\text{FFN}(x) = \max(0, xW_1 + b_1) W_2 + b_2 \quad (246)$$

The FFN is linear \rightarrow
ReLU \rightarrow linear.

- **Embeddings and Softmax:** use learned embeddings to convert input/output tokens to vectors of dimension d_{model} , and for the pre-softmax layer at the output of the decoder⁸⁴.
- **Positional Encoding:** how the authors deal with the lack of recurrence (to make use of the sequence order). They add a sinusoid function of the position (timestep) pos and vector index i to the input embeddings for the encoder and decoder⁸⁵:

$$\text{PE}(pos, 2i) = \sin\left(\frac{pos}{10000^{2i/d_{\text{model}}}}\right) \quad (247)$$

$$\text{PE}(pos, 2i + 1) = \cos\left(\frac{pos}{10000^{2i/d_{\text{model}}}}\right) \quad (248)$$

For inputs to encoder/decoder, the embedding weights are multiplied by $\sqrt{d_{\text{model}}}$

The authors justify this choice:

⁸⁴In other words, they use the same weight matrix for all three of (1) encoder input embedding, (2) decoder input embedding, and (3) (opposite direction) from decoder output to pre-softmax.

⁸⁵Note that the positional encodings must necessarily be of dimension d_{model} to be summed with the input embeddings.

We chose this function because we hypothesized it would allow the model to easily learn to attend by relative positions, since for any fixed offset k , PE_{pos+k} can be represented as a linear function of PE_{pos} .

Hierarchical Attention Networks

Table of Contents Local

Written by Brandon McKinzie

Yang et al., “Hierarchical Attention Networks for Document Classification.”

Overview. Authors introduce the Hierarchical Attention Network (HAN) that is designed to capture insights regarding (1) the hierarchical structure of documents (words \rightarrow sentences \rightarrow documents), and (2) the context dependence between words and sentences. The latter is implemented by including two levels of attention mechanisms, one at the word level and one at the sentence level.

Hierarchical Attention Networks. Below is an illustration of the network. The first stage is familiar to sequence to sequence models - a bidirectional encoder for outputting sentence-level representations of a sequence of words. The HAN goes a step further by feeding this another bidirectional encoder for outputting document-level representations for sequences of sentences.




Figure 2: Hierarchical Attention Network.

The authors choose the GRU as their underlying RNN. For ease of reference, the defining equations of the GRU are shown below:

$$h_t = (1 - z_t) \odot h_{t-1} + z_t \odot \tilde{h}_t \quad (249)$$

$$z_t = \sigma(W_z x_t + U_z h_{t-1} + b_z) \quad (250)$$

$$\tilde{h}_t = \tanh(W_h x_t + r_t \odot (U_h h_{t-1}) + b_h) \quad (251)$$

$$r_t = \sigma(W_r x_t + U_r h_{t-1} + b_r) \quad (252)$$

Hierarchical Attention. Here I'll overview the main stages of information flow.

1. **Word Encoder.** Let the t th word in the i th sentence be denoted w_{it} . They embed the vectors with a word embedding matrix W_e , $x_{it} = W_e w_{it}$, and then feed x_{it} through a bidirectional GRU to ultimately obtain $h_{it} := [\vec{h}_{it}; \overleftarrow{h}_{it}]$.
2. **Word Attention.** Extracts words that are important to the meaning of the sentence and aggregates the representation of these informative words to form a sentence vector.

$$u_{it} = \tanh(W_w h_{it} + b_w) \quad (253)$$

$$\alpha_{it} = \frac{\exp(u_{it}^T u_w)}{\sum_t \exp(u_{it}^T u_w)} \quad (254)$$

$$s_i = \sum_t \alpha_{it} h_{it} \quad (255)$$

Note the *context vector* u_w , which is shared for all words⁸⁶ and randomly initialized and jointly learned during the training process.

3. **Sentence Encoder.** Similar to the word encoder, but uses the sentence vectors s_i as the input for the i th sentence in the document. Note that the output of this encoder, h_i contains information from the neighboring sentences too (bidirectional) but focuses on sentence i .
4. **Sentence Attention.** For rewarding sentences that are clues to correctly classify a document. Similar to before, we now use a sentence level context vector u_s to measure the importance of the sentences.

$$u_i = \tanh(W_s h_i + b_s) \quad (256)$$

$$\alpha_i = \frac{\exp(u_i^T u_s)}{\sum_i \exp(u_i^T u_s)} \quad (257)$$

$$v = \sum_i \alpha_i h_i \quad (258)$$

where v is the document vector that summarizes all the information of sentences in a document.

As usual, we convert v to a normalized probability vector by feeding through a softmax:

$$p = \text{softmax}(W_c v + b_c) \quad (259)$$

⁸⁶To emphasize, there is only a single context vector u_w in the network, period. The subscript just tells us that it is the word-level context vector, to distinguish it from the sentence-level context vector in the later stage.

Configuration and Training. Quick overview of some parameters chosen by the authors:

- **Tokenization:** Stanford CoreNLP. Vocabulary consists of words occurring more than 5 times, all others are replaced with UNK token.
- **Word Embeddings:** train word2vec on the training and validation splits. Dimension of 200.
- **GRU.** Dimension of 50 (so 100 because bidirectional).
- **Context vectors.** Both u_w and u_s have dimension of 100.
- **Training:** batch size of 64, grouping documents of similar length into a batch. SGD with momentum of 0.9.

Joint Event Extraction via RNNs

Table of Contents Local

Written by Brandon McKinzie

Nguyen, Cho, and Grishman, “Joint Event Extraction via Recurrent Neural Networks,” (2016).

Event Extraction Task. Automatic Context Extraction (ACE) evaluation. Terminology:

- **Event:** something that happens or leads to some change of state.
- **Mention:** phrase or sentence in which an event occurs, including one trigger and an arbitrary number of arguments.
- **Trigger:** main word that most clearly expresses an event occurrence.
- **Argument:** an entity mention, temporal expression, or value that serves as a participant/attribute with a specific role in an event mention.

Example:

In Baghdad, a **cameraman** **died**{*Die*} when an American tank **fired**{*Attack*} on the Palestine hotel.

TRIGGER
ARGUMENT

Each event subtype has its own set of roles to be filled by the event arguments. For example, the roles for the *Die* event subtype include *Place*, *Victim*, and *Time*.

Model.

- **Sentence Encoding.** Let w_i denote the i th token in a sentence. It is transformed into a real-valued vector x_i , defined as

$$x_i := [\text{GloVe}(w_i); \text{Embed}(\text{EntityType}(w_i)); \text{DepVec}(w_i)] \quad (260)$$

where “Embed” is an embedding we learn, and “DepVec” is the binary vector whose dimensions correspond to the possible relations between words in the dependency trees.

- **RNN.** Bidirectional LSTM on the inputs x_i .
- **Prediction.** Binary memory vector G_i^{trg} for triggers; binary memory matrices G_i^{arg} and $G_i^{arg/trg}$ for arguments (at each timestep i). At each time step i , do the following in order:
 1. Predict trigger t_i for w_i . First compute the feature representation vector R_i^{trig} , defined as:

$$R_i^{trig} := [h_i; L_i^{trg}; G_{i-1}^{trg}] \quad (261)$$

where h_i is the RNN output, L_i^{trg} is the local context vector for w_i , and G_{i-1}^{trg} is the memory vector from the previous step. $L_i^{trg} := [\text{GloVe}(w_{i-d}); \dots; \text{GloVe}(w_{i+d})]$ for

some predefined window size d . This is then fed to a fully-connected layer with softmax activation, F^{trg} , to compute the probability over possible trigger subtypes:

$$P_{i;t}^{trg} := F_t^{trg}(R_i^{trg}) \quad (262)$$

As usual, the predicted trigger type for w_i is computed as $t_i = \arg \max_t (P_{i;t}^{trg})$. If w_i is not a trigger, t_i should predict “*Other*.”

2. Argument role predictions, a_{i1}, \dots, a_{ik} , for all of the [already known] entity mentions in the sentence, e_1, \dots, e_k with respect to w_i . a_{ij} denotes the argument role of e_j with respect to [the predicted trigger of] w_i . If NOT(w_i is trigger AND e_j is one of its arguments), then a_{ij} is set to *Other*. For example, if w_i was the word “died” from our example sentence, we’d hope that its predicted trigger would be $t_i = Die$, and that **the entity associated with “cameraman” would get a predicted argument role of *Victim*.**

```
def getArgumentRoles(triggerType=t, entities=e):
    k = len(e)
    if isOther(t):
        return [Other] * k
    else:
        for e_j in e:
```

$$R_{ij}^{arg} := [h_i; h_{i_j}; L_{ij}^{arg}; B_{ij}; G_{i-1}^{arg}[j]; G_{i-1}^{arg/trg}[j]] \quad (263)$$

3. Update memory. TO BE CONTINUED...(moving onto another paper because this model is getting a *bit* too contrived for my tastes. Also not a fan of the reliance on a dependency parse.)

Event Extraction via Bidi-LSTM Tensor NNs

Table of Contents Local

Written by Brandon McKinzie

Y. Chen, S. Liu, S. He, K. Liu, and J. Zhao, “Event Extraction via Bidirectional Long Short-Term Memory Tensor Neural Networks.”

Overview. The task/goal is the event extraction task as defined in *Automatic Content Extraction* (ACE). Specifically, given a text document, our goal is to do the following in order for each sentence:

1. Identify any event triggers in the sentence.
2. If triggers found, predict their subtype. For example, given the trigger “fired,” we may classify it as having the *Attack* subtype.
3. If triggers found, identify their candidate argument(s). ACE defines an event argument as “an entity mention, temporal expression, or value that is involved in an event.”
4. For each candidate argument, predict its role: “the relationship between an argument to the event in which it participates.”

Context-aware Word Representation. Use pre-trained word embeddings for the input word tokens, the predicted trigger, and the candidate argument. Note: *we assume we already have predictions for the event trigger t and are doing a pass for one of (possibly many) candidate arguments a .*

1. Embed each word in the sentence with pre-trained embeddings. Denote the embedding for i th word as $e(w_i)$.
2. Feed each $e(w_i)$ through a bidirectional LSTM. Denote the i th output of the forward LSTM as $c_l(w_{i+1})$ and the output of the backward LSTM at the same time step as $c_r(w_{i-1})$. As usual, they take the general functional form:

$$c_l(w_i) = \overrightarrow{LSTM}(c_l(w_{i-1}), e(w_{i-1})) \quad (264)$$

$$c_r(w_i) = \overleftarrow{LSTM}(c_r(w_{i+1}), e(w_{i+1})) \quad (265)$$

$$(266)$$

3. Concatenate $e(w_i)$, $c_l(w_i)$, $c_r(w_i)$ together along with the embedding of the candidate argument $e(a)$ and predicted trigger $e(t)$. Also include the relative distance of w_i to t or (??) a , denoted as pi for position information, and the embedding of the predicted event type pe of the trigger. Denote this massive concatenation result as x_i :

$$x_i := c_l(w_i) \oplus e(w_i) \oplus c_r(w_i) \oplus pi \oplus pe \oplus e(a) \oplus e(t) \quad (267)$$

Dynamic Multi-Pooling. This is easiest shown by example. Continue with our example sentence:

In California, **Peterson** was arrested for the **murder** of his wife and unborn son.

where the colors are given for *this specific case where murder is our predicted trigger and we are considering the candidate argument Peterson*⁸⁷. Given our n outputs from the previous stage, $y^{(1)} \in \mathbb{R}^{n \times m}$, where n is the length of the sentence and m is the size of that huge concatenation given in equation 267. We split our sentence by trigger and candidate argument, then (confusingly) redefine our notation as

$$y_{1j}^{(1)} \leftarrow \begin{bmatrix} y_{1j}^{(1)} & y_{2j}^{(1)} \end{bmatrix} \quad (268)$$

$$y_{2j}^{(1)} \leftarrow \begin{bmatrix} y_{3j}^{(1)} & \cdots & y_{7j}^{(1)} \end{bmatrix} \quad (269)$$

$$y_{3j}^{(1)} \leftarrow \begin{bmatrix} y_{8j}^{(1)} & \cdots & y_{nj}^{(1)} \end{bmatrix} \quad (270)$$

Peterson is the 3rd word,
and murder is the 8th
word.

where it's important to see that, for some $1 \leq j \leq m$, each new $y_{ij}^{(1)}$ is a *vector* of length equal to the number of words in segment i . Finally, the dynamic multi-pooling layer, $y^{(2)}$, can be expressed as

$$y_{i,j}^{(2)} := \max \left(y_{i,j}^{(1)} \right) \quad 1 \leq i \leq 3, 1 \leq j \leq m \quad (271)$$

where the max is taken over each of the aforementioned vectors, leaving us with $3m$ values total. These are concatenated to form $y^{(2)} \in \mathbb{R}^{3m}$.

Output. To predict of each argument role [for the given argument candidate], $y^{(2)}$ is fed through a dense softmax layer,

$$O = W_2 y^{(2)} + b_2 \quad (272)$$

where $W_2 \in \mathbb{R}^{n_1 \times 3m}$ and n_1 is the number of possible argument roles (including "None"). The authors also use dropout on $y^{(2)}$.

⁸⁷Yes, arrested could be another predicted trigger, but the network considers each possibility at separate times/locations in the architecture.

Reasoning with Neural Tensor Networks

Table of Contents Local

Written by Brandon McKinzie

Socher et al., “Reasoning with Neural Tensor Networks for Knowledge Base Completion”

Overview. Reasoning over relationships between two entities. Goal: predict the likely truth of additional facts based on existing facts in the KB. This paper contributes (1) the new NTN and (2) a new way to represent entities in KBs. Each relation is associated with a distinct model. Inputs to a given relation’s model are pairs of database entities, and the outputs score how likely the pair has the relationship.

Neural Tensor Networks for Relation Classification. Let $e_1, e_2 \in \mathbb{R}^d$ be the vector representations of the two entities, and let R denote the relation (and thus model) of interest. The NTN computes a score of how likely it is that e_1 and e_2 are related by R via:

$$g(e_1, R, e_2) = u_R^T \tanh \left(e_1^T W_R^{[1:k]} e_2 + V_R \begin{bmatrix} e_1 \\ e_2 \end{bmatrix} + b_R \right) \quad (273)$$

$$W_R^{[1:k]} \in \mathbb{R}^{d \times d \times k}$$

$$V_R \in \mathbb{R}^{k \times 2d}$$

where the bilinear tensor product $e_1^T W_R^{[1:k]} e_2$ results in a vector $h \in \mathbb{R}^k$ with each entry computed by one slice $i = 1, \dots, k$ of the tensor.

Intuitively, we can see each slice of the tensor as being responsible for one type of entity pair or instantiation of a relation... Another way to interpret each tensor slice is that it mediates the relationship between the two entity vectors differently.

Training Objective and Derivatives. All models are trained with **contrastive max-margin objective functions** and minimize the following objective:


$$J(\mathbf{\Omega}) = \sum_{i=1}^N \sum_{c=1}^C \max \left(0, 1 - g \left(T^{(i)} \right) + g \left(T_c^{(i)} \right) \right) + \lambda \|\mathbf{\Omega}\|_2^2 \quad (274)$$

where c is for “corrupted” samples, $T_c^{(i)} := (e_1^{(i)}, R^{(i)}, e_c^{(i)})$. Notice that this function is minimized when the difference, $g(T^{(i)}) - g(T_c^{(i)})$, is maximized. The authors used minibatched **L-BFGS** for optimization.

Table of Contents Local


Written by Brandon McKinzie

Sequence-to-Tree Model. Variant of Seq2Seq that is more faithful to the compositional nature of meaning representations. It’s schematic is shown below. The authors define a “nonterminal” $< n >$ token which indicates [the root of] a subtree.



After encoding input q , the hierarchical tree decoder generates tokens at depth 1 of the subtree corresponding to parts of logical form a . If the predicted token is $\langle n \rangle$, decode the sequence by conditioning on the nonterminal's hidden vector. This process terminates when no more nonterminals are emitted.

Also note that the output posterior probability over the encoded input q is the product of subtree posteriors. For example, consider the decoding example in the figure below:



We would compute the output posterior as:

$$p(a \mid q) = p(y_1 y_2 y_3 y_4 \mid q) p(y_5 y_6 \mid y_{\leq 3}, q) \quad (275)$$

The model is trained by minimizing log-likelihood over the training data, using RMSProp for optimization. At inference time, greedy search or beam search is used to predict the most probable output sequence.

Seq2SQL: Generating Structured Queries from NL using RL

Table of Contents Local

Written by Brandon McKinzie

Zhong, Xiong, and Socher, “Seq2SQL: Generating Structured Queries From Natural Language Using Reinforcement Learning”

Overview. Deep neural network for translating natural language questions to corresponding SQL queries. Outperforms state-of-the-art semantic parser.

Seq2Tree and Pointer Baseline. Baseline model is the Seq2Tree model from the previous note on Dong & Lapata’s (2016) paper. Authors here argue their output space is unnecessarily large, and employ the idea of pointer networks with augmented inputs. The input sequence is the concatenation of (1) the column names, (2) the limited vocabulary of the SQL language such as SELECT, COUNT, etc., and (3) the question.

$$x := [< col >; x_1^c; x_2^c; \dots; x_N^c; < sql >; x^s; < question >; x^q] \quad (276)$$

$$x_j^c \in \mathbb{R}^{T_j}$$

where we also insert special (“sentinel”) tokens to demarcate the boundaries of each section. The pointer network can then produce the SQL query by selecting exclusively from the input. Let g_s denote the s th decoder [hidden] state, and y_s denote the output (index/pointer to input query token).

$$[\text{ptr net}] \ y_s = \arg \max_t \left(\alpha_s^{ptr} \right) \quad \text{where} \quad \alpha_{s,t}^{ptr} = w^{ptr} \cdot \tanh \left(U^{ptr} g_s + V^{ptr} h_t \right) \quad (277)$$

Seq2SQL.

1. **Aggregation Classifier.** Our goal here is to predict which aggregation operation to use out of COUNT, MIN, MAX, NULL, etc. This is done by projecting the attention-weighted average of encoder states, κ^{agg} , to \mathbb{R}^C where C denotes the number of unique aforementioned aggregation operations. The sequence of computations is summarized as follows:

$$\alpha_t^{inp} = w^{inp} \cdot h_t^{enc} \quad (278)$$

$$\beta^{inp} = \text{softmax} \left(\alpha^{inp} \right) \quad (279)$$

$$\kappa^{agg} = \sum_t^T \beta_t^{inp} h_t^{enc} \quad (280)$$

$$\alpha^{agg} = W^{agg} \tanh (V^{agg} \kappa^{agg} + b^{agg}) + c^{agg} \quad (281)$$

$$\beta^{agg} = \text{softmax} (\alpha^{agg}) \quad (282)$$

$$W^{agg} \in \mathbb{R}^{C \times T}$$

$$\beta^{agg} \in \mathbb{R}^C$$

where β_i^{agg} gives the probability for the i th aggregation operation. We use cross entropy loss L^{agg} for determining the aggregation operation. Note that this part isn't really a sequence-to-sequence architecture. It's nothing more than an MLP applied to an attention-weighted average of the encoder states.

2. **Get Pointer to Column.** A pointer network is used for identifying which column in the input representation should be used in the query. Recall that $x_{j,t}^c$ denotes the t th word in column j . We use the last encoder state for a given column's LSTM⁸⁸ as its representation; T_j denotes the number of words in the j th column.

$$e_j^c = h_{j,T_j}^c \quad \text{where} \quad h_{j,t}^c = \text{LSTM}\left(\text{emb}(x_{j,t}^c), h_{j,t-1}^c\right) \quad (283)$$

To construct a representation for the question, compute another input representation κ^{sel} using the same architecture (but distinct weights) as for κ^{agg} . As usual, we compute the scores for each column j via:

$$\alpha_j^{sel} = W^{sel} \tanh\left(V^{sel} \kappa^{sel} + V^c e_j^c\right) \quad (284)$$

$$\beta^{sel} = \text{softmax}\left(\alpha^{sel}\right) \quad (285)$$

Similar to the aggregation, we train the **SELECT** network using cross entropy loss L^{sel} .

3. **WHERE Clause Pointer Decoder.** Recall from equation 277 that this is a model with recurrent connections from its *outputs leading back into its inputs*, and thus a common approach is to train it with **teacher forcing**⁸⁹. However, since the boolean expressions within a **WHERE** clause can be swapped around while still yielding the same SQL query, reinforcement learning (instead of cross entropy) is used to *learn a policy to directly optimize the expected correctness of the execution result*. Note that this also implies that we will be sampling from the output distribution at decoding step s to obtain the next input for $s + 1$ [instead of teacher forcing].

⁸⁸Yes, we encode each column with an LSTM separately.

⁸⁹Teacher forcing is just a name for how we train the decoder portion of a sequence-to-sequence model, wherein we feed the ground-truth output $y^{(t)}$ as input at time $t + 1$ during training.

$$R(q(y), q_g) = \begin{cases} -2 & \text{if } q(y) \text{ is not a valid SQL query} \\ -1 & \text{if } q(y) \text{ is a valid SQL query and executes to an incorrect result} \\ +1 & \text{if } q(y) \text{ is a valid SQL query and executes to the correct result} \end{cases} \quad (286)$$

$$L^{whe} = -\mathbb{E}_y [R(q(y), q_g)] \quad (287)$$

$$\nabla L_{\Theta}^{whe} = -\nabla_{\Theta} (\mathbb{E}_{y \sim p_y} [R(q(y), q_g)]) \quad (288)$$

$$= -\mathbb{E}_{y \sim p_y} \left[R(q(y), q_g) \nabla_{\Theta} \sum_t \log p_y(y_t; \Theta) \right] \quad (289)$$

$$\approx -R(q(y), q_g) \nabla_{\Theta} \sum_t \log p_y(y_t; \Theta) \quad (290)$$

where

→ $y = [y^1, y^2, \dots, y^T]$ denotes the sequences of generated tokens in the **WHERE** clause.

→ $q(y)$ denotes the query generated by the model.

→ q_g denotes the ground truth query corresponding to the question.

and the gradient has been approximated in the last line using a single Monte-Carlo sample y .

Finally, the model is trained using gradient descent to minimize $L = L^{agg} + L^{sel} + L^{whe}$.

Speculations for Event Extraction. I want to explore using this paper’s model for the task of event extraction. Below, I’ve replaced some words (shown in green) from a sentence in the paper in order to formalize this as event extraction.

*Seq2**Event** takes as input a **sentence** and the **possible event types** of an **ontology**. It generates the corresponding **event annotation**, which, during training, is **compared** against an **event template**. The result of the **comparison** is utilized to train the reinforcement learning algorithm⁹⁰.*

⁹⁰Original: Seq2SQL takes as input a question and the columns of a table. It generates the corresponding SQL query, which, during training, is executed against a database. The result of the execution is utilized as the reward to train the reinforcement learning algorithm.

SLING: A Framework for Frame Semantic Parsing

Table of Contents Local

Written by Brandon McKinzie

M. Ringgaard, R. Gupta, F. Pereira, “SLING: A framework for frame semantic parsing” (2017)



Model.

- **Inputs.** [words; affixes; shapes]
- **Encoder.**
 1. Embed.
 2. Bidirectional LSTM.
- **Inputs to TBRU.**
 - BLSTM [forward and backward] hidden state for the current token in the parser state.
 - **Focus.** Hidden layer activations corresponding to the transition steps that evoked/brought into *focus* the top- k frames in the attention buffer.
 - **Attention.** Recall that we maintain an **attention buffer**: an ordered list of frames, where the order represents closeness to center of attention. The attention portion of inputs for the TBRU looks at the top- k frames in the attention buffer, finds the phrases in the text (if any) that evoked them. The activations from the BLSTM for the last token of each of those phrases are included as TBRU inputs⁹¹
 - **History.** Hidden layer activations from the previous k steps.
 - **Roles.** Embeddings of (s_i, r_i, t_i) , where the frame at position s_i in the attention buffer has a role (key) r_i with frame at position t_i as its value. Back-off features are added for the source roles (s_i, r_i) , target role (r_i, t_i) , and unlabeled roles (s_i, t_i) .
- **Decoder (TBRU).** Outputs a softmax over possible transitions (actions).

⁹¹Okay, how is this attention at all? Seems misleading to call it attention.

Transition System. Below is the list of possible actions. Note that, since the system is trained to predict the correct *frame graph* result, it isn't directly told what order it should take a given set of actions⁹².

- **SHIFT.** Move to next input token.
- **STOP.** Signal that we've reached end of parse.
- **EVOKE(type, num).** New frame of **type** from next **num** tokens in the input; placed at front of attention buffer.
- **REFER(frame, num).** New mention from next **num** tokens, evoking existing **frame** from attention buffer. Places at front.
- **CONNECT(source-frame, role, target-frame).** Inserts (**role, target-frame**) slot into **source-frame**, move **source-frame** to front.
- **ASSIGN(source-frame, role, value).** Same as CONNECT, but with primitive/constant **value**.
- **EMBED(target-frame, role, type).** New frame of **type**, and inserts (**role, target-frame**) slot. New frame placed to front.
- **ELABORATE(source-frame, role, type).** New frame of **type**. Inserts (**role, new-frame**) slot to **source-frame**. New frame placed at front.

Evaluation. Need some way of comparing an annotated document with its gold-standard annotation. This is done by constructing a virtual graph where the document is the start node. It is then connected to the spans (which are presumably nodes themselves), and the spans are connected to the frames they evoke. Frames that refer to other frames are given corresponding edges between them. Quality is computed by aligning the golden and predicted graphs and computing precision, recall, and F1. Specifically, these scores are computed separately for spans, frames, frame types, roles linking to other frames (referred to here as just "roles"), and roles that link to global constants (referred to here as just "labels"). Results are shown below.



⁹²This is important to keep in mind, since more than one sequence of actions can result in a given predicted frame graph.

Poincaré Embeddings for Learning Hierarchical Representations

Table of Contents Local

Written by Brandon McKinzie

M. Nickel and D. Kiela, “Poincaré Embeddings for Learning Hierarchical Representations” (2017)

Introduction. Dimensionality of embeddings can become prohibitively large when needed for complex data. Authors focus on mitigating this problem for large datasets whose objects can be organized according to a latent hierarchy⁹³. They propose to compute embeddings in a particular model of hyperbolic space, the **Poincaré ball model**, claiming it is well-suited for gradient-based optimization (they make use of **Riemannian optimization**).

Prerequisite Math. Recall that a hyperbola is a set of points, such that for any point P of the set, the absolute difference of the distances $|PF_1|$, $|PF_2|$ to two fixed points F_1 , F_2 (the foci), is constant, usually denoted by $2a$, $a > 0$. We can define a hyperbola by this set of points or by its canonical form, which are both given, respectively, as:

$$H = \{P \mid ||PF_2| - |PF_1|| = 2a\} \quad (291)$$

$$\frac{x^2}{a^2} - \frac{y^2}{b^2} = 1 \quad (292)$$

where $b^2 := c^2 - a^2$, $(\pm a, 0)$ are the two vertices, and $(\pm c, 0)$ are the two foci. Cannon et al. define n -dimensional hyperbolic space by the formula

$$H^n = \{x \in \mathbb{R}^{n+1} : x * x = -1\} \quad (293)$$

where $*$ denotes the non-euclidean inner product (subtracts last term; same as Minkowski sapce-time). Notice that this is the defining equation for a hyperboloid of two sheets, and Cannon et al. says “usually we deal only with one of the two sheets.” Hyperbolic spaces are well-suited to model hierarchical data, since both circle length and disc area grow *exponentially* with r .

⁹³This begs the question: how useful would a Poincaré embedding be for situations where this assumption isn’t valid?

Poincaré Embeddings. Let $\mathcal{B}^d = \{\mathbf{x} \in \mathbb{R}^d \mid \|\mathbf{x}\| < 1\}$ be the open d-dimensional unit ball. The **Poincaré ball** model of hyperbolic space corresponds then to the Riemannian manifold⁹⁴ $(\mathcal{B}^d, g_{\mathbf{x}})$, where

$$g_{\mathbf{x}} = \left(\frac{2}{1 - \|\mathbf{x}\|^2} \right)^2 g^E \quad (294)$$

is the **Riemannian metric tensor**, and g^E denotes the Euclidean metric tensor. The distance between two points $\mathbf{u}, \mathbf{v} \in \mathcal{B}^d$ is given as

$$d(\mathbf{u}, \mathbf{v}) = \operatorname{arccosh} \left(1 + 2 \frac{\|\mathbf{u} - \mathbf{v}\|^2}{(1 - \|\mathbf{u}\|^2)(1 - \|\mathbf{v}\|^2)} \right) \quad (295)$$

The boundary of the ball corresponds to the sphere S^{d-1} and is denoted by $\partial\mathcal{B}$. Geodesics in \mathcal{B}^d are then circles that are orthogonal to $\partial\mathcal{B}$. To compute Poincaré embeddings for a set of symbols $\mathcal{S} = \{x_i\}_{i=1}^n$, we want to find embeddings $\Theta = \{\theta_i\}_{i=1}^n$, where $\theta_i \in \mathcal{B}^d$. Given some loss function \mathcal{L} that encourages semantically similar objects to be close as defined by the Poincaré distance, our goal is to solve the optimization problem

$$\Theta' \leftarrow \arg \min_{\Theta} \mathcal{L}(\Theta) \quad s.t. \forall \theta_i \in \Theta : \|\theta_i\| < 1 \quad (296)$$

Optimization. Let $\mathcal{T}_{\theta}\mathcal{B}$ denote the **tangent space** of a point $\theta \in \mathcal{B}^d$. Let $\nabla_R \in \mathcal{T}_{\theta}\mathcal{B}$ denote the **Riemannian gradient** of $\mathcal{L}(\theta)$, and ∇_E the Euclidean gradient of $\mathcal{L}(\theta)$. Using **RS GD**, parameter updates take the form

$$\theta_{t+1} = \Re_{\theta_t}(-\eta_t \nabla_R \mathcal{L}(\theta_t)) \quad (297)$$

where \Re_{θ_t} denotes the **retraction** onto \mathcal{B} at θ and η_t denotes the learning rate at time t .

⁹⁴All five analytic models of hyperbolic geometry in Cannon et al. are differentiable manifolds with a Riemannian metric. A **Riemannian metric** ds^2 on Euclidean space \mathbb{R}^n is a function that assigns at each point $p \in \mathbb{R}^n$ a positive definite symmetric inner product on the tangent space at p , this inner product varying differentiably with the point p . If x_1, \dots, x_n are the standard coordinates in \mathbb{R}^n , then ds^2 has the form $\sum_{i,j} g_{ij} dx_i dx_j$, and the matrix (g_{ij}) depends differentiably on x and is positive definite and symmetric.

Enriching Word Vectors with Subword Information (FastText)

Table of Contents Local

Written by Brandon McKinzie

P. Bojanowski, E. Grave, A. Joulin, and T. Mikolov, “Enriching Word Vectors with Subword Information” (2017)

Overview. Based on the skipgram model, but where each word is represented as a bag of character n -grams. A vector representation is associated *each* character n -gram; words being represented as the *sum*⁹⁵ of these representations.

Skipgram with Negative Sampling. Since this is based on skipgram, recall the objective of skipgram which is to maximize:

$$\sum_{t=1}^T \sum_{c \in \mathcal{C}_t} \log \Pr [w_c \mid w_t] \quad (298)$$

for a sequence of words w_1, \dots, w_T . One way of parameterizing $\Pr [w_c \mid w_t]$ is by computing a softmax over a scoring function $s : (w_t, w_c) \mapsto \mathbb{R}$,

$$\Pr [w_c \mid w_t] = \frac{e^{s(w_t, w_c)}}{\sum_{j=1}^W e^{s(w_t, j)}} \quad (299)$$

However, this implies that, given w_t , we only predict one context word w_c . Instead, we can frame the problem as a set of independent binary classification tasks, and independently predict the presence/absence of context words. Let $\ell : x \mapsto \log(1 + e^{-x})$ denote the standard logistic negative log-likelihood. Our objective is:

$$\sum_{t=1}^T \left[\sum_{c \in \mathcal{C}_t} \ell(s(w_t, w_c)) + \sum_{n \in \mathcal{N}_{t,c}} \ell(-s(w_t, n)) \right] \quad (300)$$

where $\mathcal{N}_{t,c}$ is a set of negative examples sampled from the vocabulary. A common scoring function involves associating a distinct input vector u_w and output vector v_w for each word w . Then the score is computed as $s(w_t, w_c) = \mathbf{u}_{w_t}^T \mathbf{v}_{w_c}$.

⁹⁵It would be interesting to explore other aggregation operations than just summation.

FastText. Main contribution is a different scoring function s that utilizes subword information. Each word w is represented as a bag of character n -grams. Special symbols $<$ and $>$ delimit word boundaries, and the authors also insert the special sequence containing the full word (with the delimiters) in its bag of n -grams. The word *where* is thus represented by first building its bag of n -grams, for the choice of $n = 3$:

$$\text{where} \longrightarrow \{<wh, whe, her, ere, re>, <where>\} \quad (301)$$

Such a set of n -grams for a word w is denoted \mathcal{G}_w . Each n -gram g for a word w has its own vector \mathbf{z}_g , and the final vector representation of w is the sum of these. The scoring function becomes

$$s(w, c) = \sum_{g \in \mathcal{G}_w} \mathbf{z}_g^T \mathbf{v}_c \quad (302)$$

DeepWalk: Online Learning of Social Representations

Table of Contents Local

Written by Brandon McKinzie

B. Perozzi, R. Al-Rfou, and S. Skiena, “DeepWalk: Online Learning of Social Representations,” (2014).

Problem Definition. Classifying members of a social network into one or more categories.

Let $G = (V, E)$, where V are the members of the network, and E be its edges, $E \subseteq (V \times V)$. Given a partially labeled social network $G_L = (V, E, X, Y)$, with attributes $X \in \mathbb{R}^{|V| \times S}$ where S is the size of the feature space for each attribute vector, and $Y \in \mathbb{R}^{|V| \times |\mathcal{Y}|}$, \mathcal{Y} is the set of labels.

In other words, the elements of our training dataset, (X, Y) , are the members of the social network, and we want to label each member, represented by a vector in \mathbb{R}^S , with one or more of the $|\mathcal{Y}|$ labels. We aim to learn features that capture the graph structure *independent* of the labels’ distribution, and to do so in an unsupervised fashion.

Learning Social Representations. We want the representations to be adaptable, community-aware, low-dimensional, and continuous. The authors’ method learns representations for vertices from a stream of short random walks, optimized with techniques from language modeling.

- **Random Walks.** Denote a random walk rooted at vertex v_i as \mathcal{W}_{v_i} , where the k th visited vertex is chosen at random from the neighbors of the $(k-1)^{th}$ visited vertex, and so on. Motivation for their use here is that they’re “the foundation of a class of *output sensitive* algorithms which use them to compute local community structure information in time sublinear to the size of the input graph.”
- **Language Modeling.** Authors present a generalization of language modeling, which traditionally aims to maximize $\Pr[w_n \mid w_0, \dots, w_{n-1}]$ over all words in a training corpus. The motivation of the generalization is to explore the graph through a stream of short random walks. The walks are thought of as short sentences/phrases in a special language, and we want to estimate the probability of observing vertex v_i given all previous vertices so far in the random walk. Since we want to learn a latent social representation of each vertex, and not simply a probability distribution over node co-occurrences, we condition on the *embeddings* of visited nodes in this latent space (rather than the nodes themselves directly)

$$\Pr[v_i \mid \Phi(v_1), \Phi(v_2), \dots, \Phi(v_{i-1})] \quad (303)$$

where, in practice, the mapping function Φ is represented by a $|V| \times d$ matrix of free parameters (an embedding matrix). Since this becomes infeasible to compute as the walk length grows, the authors opt for an approach resembling CBOW: minimizing the NLL of vertices in the the context of a given vertex.

$$\min_{\Phi} -\log \Pr[v_{i-w}, \dots, v_{i-1}, v_{i+1}, \dots, v_{i+w} \mid \Phi(v_i)] \quad (304)$$

Remember that, here, v_j is the j th vertex visited in some given random walk.

DeepWalk Algorithm. Below is a conceptual summary of the procedure, followed by a figure/illustration of the formal algorithm definition.

1. **Inputs.** Graph $G(V, E)$, window size w , embedding size d , walks per vertex γ , walk length t .
2. **Random Walk.** For each vertex v_i , compute $\mathcal{W}_{v_i} := \text{RandomWalk}(G, v_i, t)$.
3. **Updates.** Upon finishing a walk, \mathcal{W}_{v_i} , run skipgram on the sequence of walked vertices to update the embedding matrix Φ .
4. **Outputs.** The embedding matrix $\Phi \in \mathbb{R}^{|V| \times d}$.

Algorithm 1 DEEPWALK(G, w, d, γ, t)

Input: graph $G(V, E)$

 window size w

 embedding size d

 walks per vertex γ

 walk length t

Output: matrix of vertex representations $\Phi \in \mathbb{R}^{|V| \times d}$

1: Initialization: Sample Φ from $\mathcal{U}^{|V| \times d}$

2: Build a binary Tree T from V

3: **for** $i = 0$ to γ **do**

4: $\mathcal{O} = \text{Shuffle}(V)$

5: **for each** $v_i \in \mathcal{O}$ **do**


6: $\mathcal{W}_{v_i} = \text{RandomWalk}(G, v_i, t)$

7: $\text{SkipGram}(\Phi, \mathcal{W}_{v_i}, w)$

8: **end for**

9: **end for**

where $\text{SkipGram}(\Phi, \mathcal{W}_{v_i}, w)$ performs SGD updates on Φ to minimize $-\log \Pr[u_k | \Phi(v_j)]$ for each visited v_j , for each u_k in the “context” of v_j . Notice that a binary tree T is build from the set of vertices V (line 2) – this is done as preparation for computing each $\Pr[u_k | \Phi(v_j)]$ via a *hierarchical softmax*, to reduce computational burden of its partition function. Finally, a visual overview of the DeepWalk algorithm is shown below. The authors use this algorithm, com-



bined with a one-vs-rest logistic regression implementation by LibLinear, for various multiclass multilabel classification tasks.

Review of Relational Machine Learning for Knowledge Graphs

Table of Contents Local

Written by Brandon McKinzie

Nickel, Murphy, Tresp, and Gabrilovich, “Review of Relational Machine Learning for Knowledge Graphs,” (2015).

Introduction. Paper discusses latent feature models such as tensor factorization and multiway neural networks, and mining observable patterns in the graph. In **Statistical Relational Learning (SRL)**, the representation of an object can contain its relationships to other objects. The main goals of SRL include:

- Prediction of missing edges (relationships between entities).
- Prediction of properties of nodes.
- Clustering nodes based on their connectivity patterns.

We’ll be reviewing how SRL techniques can be applied to large-scale **knowledge graphs** (KGs), i.e. graph structured knowledge bases (KBs) that store factual information in the form of relationships between entities.

Probabilistic Knowledge Graphs. Let $\mathcal{E} = \{e_1, \dots, e_{N_e}\}$ be the set of all entities and $\mathcal{R} = \{r_1, \dots, r_{N_r}\}$ be the set of all relation types in a KG. We model each *possible* triple $x_{ijk} = (e_i, r_k, e_j)$ as a binary random variable $y_{ijk} \in \{0, 1\}$ that indicates its existence. The full tensor $\mathbf{Y} \in \{0, 1\}^{N_e \times N_r \times N_e}$ is called the *adjacency tensor*, where each possible realization of \mathbf{Y} can be interpreted as a possible world.

Clearly, \mathbf{Y} will be large and sparse in most applications. Ideally, a relational model for large-scale KGs should scale at most linearly with the data size, i.e., linearly in the number of entities N_e , linearly in the number of relations N_r , and linearly in the number of *observed* triples $|\mathcal{D}| = N_d$.

Types of SRL Models. The presence or absence of certain triples in relational data is correlated with (i.e. predictive of) the presence or absence of certain other triples. In other words, the random variables y_{ijk} are correlated with each other. There are three main ways to model these correlations:

1. **Latent feature models:** Assume all y_{ijk} are conditionally independent given latent features associated with the subject, object and relation type and additional parameters.
2. **Graph feature models:** Assume all y_{ijk} are conditionally independent given observed graph features and additional parameters.
3. **Markov Random Fields:** Assume all y_{ijk} have local interactions.

The first two model classes predict the existence of a triple x_{ijk} via a score function $f(x_{ijk}; \Theta)$

which represents the model's confidence that a triple exists given the parameters Θ . The conditional independence assumptions can be written as

$$\Pr [\mathbf{Y} | \mathcal{D}, \Theta] = \prod_{i=1}^{N_e} \prod_{j=1}^{N_e} \prod_{k=1}^{N_r} \text{Ber}(y_{ijk} | \sigma(f(x_{ijk}; \Theta))) \quad (305)$$

where Ber is the Bernoulli distribution⁹⁶. Such models will be referred to as *probabilistic models*. We will also discuss *score-based models*, which optimize $f(\cdot)$ via maximizing the margin between existing and non-existing triples.

Latent Feature Models. We assume the variables y_{ijk} are conditionally independent given a set of global latent features and parameters. All LFMs explain triples (observable facts) via latent features of entities⁹⁷. One task of all LFMs is to infer these [latent] features automatically from the data.

- **RESCAL:** a bilinear model. Models the score of a triple x_{ijk} as

$$f_{ijk}^{\text{RESCAL}} := \mathbf{e}_i^T \mathbf{W}_k \mathbf{e}_j = \sum_{a=1}^{H_e} \sum_{b=1}^{H_e} w_{abk} e_{ia} e_{jb} \quad (307)$$

where the entity vectors $\mathbf{e}_i \in \mathbb{R}^{H_e}$ and H_e denotes the number of latent features in the model. The parameters of the model are $\Theta = \{\{\mathbf{e}_i\}_{i=1}^{N_e}, \{\mathbf{W}_k\}_{k=1}^{N_r}\}$. Note that entities have the same latent representation regardless of whether they occur as subjects or objects in a relationship (shared representation), *thus allowing the model to capture global dependencies in the data*. We can make a connection to **tensor factorization** methods by seeing that the equation above can be written compactly as

$$\mathbf{F}_k = \mathbf{E} \mathbf{W}_k \mathbf{E}^T \quad (308)$$

where $\mathbf{F}_k \in \mathbb{R}^{N_e \times N_e}$ is the matrix holding all scores for the k -th relation, and the i th row of $\mathbf{E} \in \mathbb{R}^{N_e \times H_e}$ holds the latent representation of \mathbf{e}_i .

- **Multi-layer perceptrons.** We can rewrite RESCAL as

$$f_{ijk}^{\text{RESCAL}} := \mathbf{w}_k^T \boldsymbol{\phi}_{i,j}^{\text{RESCAL}} \quad (309)$$

$$\boldsymbol{\phi}_{i,j}^{\text{RESCAL}} := \mathbf{e}_j \otimes \mathbf{e}_i \quad (310)$$

where $\mathbf{w}_k = \text{vec}(\mathbf{W}_k)$ (vector of size H_e^2 obtained by stacking columns of \mathbf{W}_k). The

⁹⁶Notation used:

$$\text{Ber}(y | p) = \begin{cases} p & \text{if } y = 1 \\ 1 - p & \text{if } y = 0 \end{cases} \quad (306)$$

⁹⁷It appears that "latent" is being used here synonymously with "not directly observed in the data".

authors extend this to what they call the E-MLP (E for entity) model:

$$f_{ijk}^{E-MLP} := \mathbf{w}_k^T \mathbf{g}(\mathbf{h}_{ijk}^a) \quad (311)$$

$$\mathbf{h}_{ijk}^a := \mathbf{A}_k^T \phi_{ij}^{E-MLP} \quad (312)$$

$$\phi_{ij}^{E-MLP} := [\mathbf{e}_i; \mathbf{e}_j] \quad (313)$$

Graph Feature Models. Here we assume that the existence of an edge can be predicted by extracting features from the observed edges in the graph. In contrast to LFMs, this kind of reasoning explains triples directly from the observed triples in the KG.

- **Similarity measures for uni-relational data.** Link prediction in graphs that consist only of a single relation (e.g. (Bob, isFriendOf, Sally) for a social network). Various **similarity indices** have been proposed to measure similarity of entities, of which there are three main classes:

1. **Local** similarity indices: Common Neighbors, Adamic-Adar index, Preferential Attachment derive entity similarities from number of common neighbors.
2. **Global** similarity indices: Katz index, Leicht-Holme-Newman index (ensembles of all paths bw entities). Hitting Time, Commute Time, PageRank (random walks).
3. **Quasi-local** similarity indices: Local Katz, Local Random Walks.

- **Path Ranking Algorithm (PRA):** extends the idea of using random walks of bounded lengths for predicting links in multi-relational KGs. Let $\pi_L(i, j, k, t)$ denote a path of length L of the form $e_i \xrightarrow{r_1} e_2 \xrightarrow{r_2} e_3 \cdots \xrightarrow{r_L} e_j$, where t represents the sequence of edge types $t = (r_1, r_2, \dots, r_L)$. We also require there to be a direct arc $e_i \xrightarrow{r_k} e_j$, representing the existence of a relationship of type k from e_i to e_j . Let $\Pi_L(i, j, k)$ represent the set of all such paths of length L , ranging over path types t .

We can compute the probability of following a given path by assuming that at each step we follow an outgoing link uniformly at random. Let $\Pr[\pi_L(i, j, k, t)]$ be the probability of the path with type t . *The key idea in PRA is to use these path probabilities as features for predicting the probabilities of missing edges.* More precisely, the feature vector and score function (logistic regression) are as follows:

$$\phi_{ijk}^{PRA} = [\Pr[\pi] : \pi \in \Pi_L(i, j, k)] \quad (314)$$

$$f_{ijk}^{PRA} := \mathbf{w}_k^T \phi_{ijk}^{PRA} \quad (315)$$

TODO: Finish...


Fast Top-K Search in Knowledge Graphs

Table of Contents Local

Written by Brandon McKinzie

S. Yang, F. Han, Y. Wu, X. Yan, “Fast Top-K Search in Knowledge Graphs.”

Introduction. Task: Given a knowledge graph G , scoring function F , and a graph query Q , top-k subgraph search over G returns k answers with the highest matching scores. An example is searching for movie makers (directors) worked with “Brad” and have won awards, illustrated below:



Clearly, it would be extremely inefficient to enumerate all possible matches and then rank them.

Preliminaries/Terminology.

- **Queries.** A query [graph] is defined as $Q = (V_Q, E_Q)$. Each query node in Q provides information/constraints about an entity, and an edge between two nodes specifies the relationship or the connectivity constraint posed on the two nodes. Q^* denotes a star-shaped query, which is basically a graph that looks like a star (central node with tree-like structure radially outward).
- **Subgraph Matching.** Given a graph query Q and a knowledge graph G , a match $\phi(Q)$ of Q in G is a subgraph of G , specified by a one-to-one matching function ϕ . It maps each node u (edge $e = (u, u')$) in Q to a node match $\phi(u)$ (edge match $\phi(e) = (\phi(u), \phi(u'))$) in $\phi(Q)$.

The matching score between query Q and its match $\phi(Q)$ is

$$F(\phi(Q)) = \sum_{v \in V_Q} F_V(v, \phi(v)) + \sum_{e \in E_Q} F_E(e, \phi(e)) \quad (316)$$

$$F_V(v, \phi(v)) = \sum_i \alpha_i f_i(v, \phi(v)) \quad (317)$$

$$F_E(e, \phi(e)) = \sum_j \beta_j f_j(e, \phi(e)) \quad (318)$$

Star-Based Top-K Matching.

1. **Query decomposition:** Given query Q , STAR decomposes Q to a set of star queries \mathcal{Q} . A star query contains a pivot node and a set of leaves as its neighbors in Q .
2. **Star querying engine:** Generate a set of top matches for each star query \mathcal{Q} .
3. **Top-k rank join.** The top matches for multiple star queries are collected and joined to produce top-k complete matches of Q .

Dynamic Recurrent Acyclic Graphical Neural Networks (DRAGNN)

Table of Contents Local

Written by Brandon McKinzie

Kong et al., “DRAGNN: A Transition-based Framework for Dynamically Connected Neural Networks,” (2017).

Transition Systems. Define a transition system $\mathcal{T} \triangleq \{\mathcal{S}, \mathcal{A}, t\}$, where

- $\mathcal{S} = \mathcal{S}(x)$ is a set of states, where that set depends on the input sequence x .
- A special start state $s^\dagger \in \mathcal{S}(x)$.
- A set of allowed decisions $\mathcal{A}(s, x) \forall s \in \mathcal{S}(x)$.
- A transition function $t(s, d, x)$ returning a new state s' for any decision $d \in \mathcal{A}(s, x)$.

The authors then define a **complete structure** as a sequence of state/decision pairs $(s_1, d_1) \dots (s_n, d_n)$ such that $s_1 = s^\dagger$, $d_i \in \mathcal{A}(s_i)$ for $i = 1, \dots, n$ and $s_{i+1} = t(s_i, d_i)$, where $n = n(x)$ is the number of decisions for input x ⁹⁸. We'll use transition systems to map inputs x into a sequence of output symbols d_1, \dots, d_n .

Transition Based Recurrent Networks. When combining transition systems with recurrent networks, we will refer to them as **Transition Based Recurrent Units (TBRU)**, which consist of:

- Transition system \mathcal{T} .
- Input function $\mathbf{m}(s) : \mathcal{S} \mapsto \mathbb{R}^K$ that maps states to some fixed-size vector representation (e.g. an embedding lookup operation).
- Recurrence function $\mathbf{r}(s) : \mathcal{S} \mapsto \mathbb{P}\{1, \dots, i-1\}$ that maps states to a set of previous time steps, where \mathbb{P} is the power set. Note that $|\mathbf{r}(s)|$ may vary with s . We use \mathbf{r} to specify state-dependent recurrent links in the unrolled computation graph.
- The RNN cell $\mathbf{h}_s \leftarrow \mathbf{RNN}(\mathbf{m}(s), \{\mathbf{h}_i \mid i \in \mathbf{r}(s)\})$.

Example: Sequential tagging RNN. Let $\mathbf{x} = \{x_1, \dots, x_n\}$ be a sequence of input tokens. Let the i th output, d_i , be a tag from some predefined set of tags \mathcal{A} . Then our model can be defined as:

- Transition system: $\mathcal{T} = \{ s_i = \mathcal{S}(x_i) = \{1, \dots, d_{i-1}\}, \mathcal{A}, t(s_i, d_i, x_i) = s_{i+1} = s_i + \{d_i\} \}$.
- Input function: $\mathbf{m}(s_i) = \text{embed}(x_i)$.
- Recurrence function: $\mathbf{r}(s_i) = \{i-1\}$ to connect the network to the previous state.
- RNN cell: $\mathbf{h}_i \leftarrow \text{LSTM}(\mathbf{m}(s_i) \mid \mathbf{r}(s_i) = \{i-1\})$.

⁹⁸The authors state that we are only concerned with complete structures that have the same number of decisions $n(x)$ for the same input x .

Example: Parsey McParseface.

- Transition system: the **arc-standard** transition system, defined in image below⁹⁹.

Initialization:	$c_s(x = x_1, \dots, x_n) = ([0], [1, \dots, n], \emptyset)$	
Terminal:	$C_t = \{c \in C \mid c = ([0], [], A)\}$	
Transitions:	$(\sigma, [i \beta], A) \Rightarrow ([\sigma i], \beta, A)$	(SHIFT)
	$([\sigma i j], B, A) \Rightarrow ([\sigma j], B, A \cup \{(j, i)\})^1$	(LEFT-ARC _l)
	$([\sigma i j], B, A) \Rightarrow ([\sigma i], B, A \cup \{(i, l, j)\})$	(RIGHT-ARC _l)
¹ Permitted only if $i \neq 0$.		

FIGURE 1. The **arc-standard** stack-based transition system for projective dependency parsing. The notation $[\sigma|i]$ (for the stack) denotes a right-headed list with head i and tail σ ; the notation $[j|\beta]$ (for the buffer) denotes a left-headed list with head j and tail β .

so the state contains all words and partially built trees (stack) as well as unseen words (buffer).

- Input function: $\mathbf{m}(s_i)$ is the concatenation of 52 feature embeddings extracted from tokens based on their positions in the stack and the buffer.
- Recurrence function: $\mathbf{r}(s_i)$ is empty, as this is a feed-forward network.
- RNN cell: a feed-forward MLP (so not an RNN...).

Inference with TBRUs. To predict the output sequence $\{d_1, \dots, d_n\}$ given input sequence $\mathbf{x} = \{x_1, \dots, x_n\}$, do:

1. Initialize $s_1 = s^\dagger$.
2. For $i = 1, \dots, n$:
 - (a) Compute $\mathbf{h}_i = \mathbf{RNN}(\mathbf{m}(s_i), \{\mathbf{h}_j \mid j \in \mathbf{r}(s_i)\})$.
 - (b) Update transition state:

$$d_i \leftarrow \arg \max_{d \in \mathcal{A}(s_i)} \mathbf{w}_d^T \mathbf{h}_i \quad (319)$$

$$s_{i+1} \leftarrow t(s_i, d_i) \quad (320)$$

⁹⁹Image taken from “Transition-Based Parsing” by Joakim Nivre. Note that “right-headed” means “goes from left to right” or “headed to the right”.

Combining multiple TBRUs. We connect multiple TBRUs with different transition systems via $\mathbf{r}(s)$.

1. We execute a list of T TBRU components sequentially, so that each TBRU advances a global step counter.
2. *Each* transition state, s^τ , from the τ 'th component has access the *terminal* states from every prior transition system, and the recurrence function $\mathbf{r}(s^\tau)$ for any given component can pull hidden activations from every prior one as well¹⁰⁰.

Example: Multi-task bi-directional tagging. Say we want to do both POS and NER tagging (indices start at 1).

- Left-to-right: $\mathcal{T} = \text{shift-only}$, $\mathbf{m}(s_i) = \mathbf{x}_i$, $\mathbf{r}(s_i) = \{i - 1\}$.
- Right-to-left: $\mathcal{T} = \text{shift-only}$, $\mathbf{m}(s_{n+i}) = \mathbf{x}_{(n-i)+1}$, $\mathbf{r}(s_{n+i}) = \{n + i - 1\}$.
- POS Tagger: $\mathcal{T}_{POS} = \text{tagger}$, $\mathbf{m}(s_{2n+i}) = \{\}$, $\mathbf{r}(s_{2n+i}) = \{i, (2n - i) + 1\}$
- NER Tagger: $\mathcal{T}_{NER} = \text{tagger}$, $\mathbf{m}(s_{3n+i}) = \{\}$, $\mathbf{r}(s_{3n+i}) = \{i, (2n - i) + 1, 2n + i\}$

which illustrates the most important aspect of the TBRU:

A TBRU can serve as both an encoder for downstream tasks and a decoder for its own task simultaneously.

For this example, the POS Tagger served as both an encoder for the NER task as well as a decoder for the POS task.

Training a DRAGNN. Assume training data consists of examples \mathbf{x} along with gold decision sequences for a given TBRU in the DRAGNN. Given decisions d_1, \dots, d_N from prior components $1, \dots, T - 1$, the log-likelihood for training the T 'th TBRU along its gold decision sequence $d_{N+1}^*, \dots, d_{N+n}^*$ is then:

$$L(\mathbf{x}, d_{N+1:N+n}^*; \theta) = \sum_i \log \Pr [d_{N+i}^* \mid d_{1:N}, d_{N+1:N+i-1}^*; \theta] \quad (321)$$


¹⁰⁰Sketchy...

Neural Architecture Search with Reinforcement Learning

Table of Contents Local

Written by Brandon McKinzie

B. Zoph and Q. Le, “Neural Architecture Search with Reinforcement Learning,” (2017).



Controller RNN. Generates architectures with a predefined number of layers, which is increased manually as training progresses. At convergence, validation accuracy of the generated network is recorded. Then, the controller parameters θ_c are optimized to maximize the expected validation accuracy over a batch of generated architectures.

Reinforcement Learning to learn the controller parameters θ_c . Let $a_{1:T}$ denote a list of actions taken by the controller¹⁰¹, which defines a generated architecture. We denote the resulting validation accuracy by R , which is the reward signal for our RL task. Concretely, we want our controller to maximize its expected reward, $J(\theta_c)$:

$$J(\theta_c) = \mathbb{E}_{P(a_{1:T}; \theta_c)} [R] \quad (322)$$

Since the quantity $\nabla_{\theta_c} R$ is non-differentiable¹⁰², we use a **policy gradient** method to iteratively update θ_c . All this means is that we instead compute gradients over the softmax outputs (the action probabilities), and use the value of R as a simple weight factor.

$$\nabla_{\theta_c} J(\theta_c) = R \sum_{t=1}^T \mathbb{E}_{P(a_{1:T}; \theta_c)} [\nabla_{\theta_c} \log P(a_t | a_{1:t-1}; \theta_c)] \quad (323)$$

$$\approx \frac{1}{m} \sum_{k=1}^m R_k \sum_{t=1}^T \nabla_{\theta_c} \log P(a_t | a_{1:t-1}; \theta_c) \quad (324)$$

where the second equation is the empirical approximation (batch-average instead of expectation) over a batch of size m , an unbiased estimator for our gradient¹⁰³. Also note that we *do*

¹⁰¹Note that T is not necessarily the number of layers, since a single generated layer can correspond to multiple actions (e.g. stride height, stride width, num filters, etc.).

¹⁰² R is a function of the action sequence $a_{1:T}$ and the parameters θ_c , and implicitly depends on the samples used for the validation set. Clearly, we do not have access to an analytical form of R , and computing numerical gradients via small perturbations of θ_c is computationally intractable.

¹⁰³It is unbiased for the same reason that any average over samples x drawn from a distribution $P(x)$ is an unbiased estimator for $\mathbb{E}_P[x]$.

have access to the distribution $P(a_{1:T}; \theta_c)$ since it is *defined* to be the joint softmax probabilities of our controller, given its parameter values θ_c (i.e. this is not a p_{data} vs p_{model} situation). The approximation 323 is an unbiased estimator for 324, but has high variance. To reduce the variance of our estimator, the authors employ a baseline function b that does not depend on the current action:

$$\frac{1}{m} \sum_{k=1}^m \sum_{t=1}^T \nabla_{\theta_c} \log P(a_t \mid a_{1:t-1}; \theta_c) (R_k - b) \quad (325)$$

Joint Extraction of Events and Entities within a Document Context

Table of Contents Local

Written by Brandon McKinzie

B. Yang and T. Mitchell, “Joint Extraction of Events and Entities within a Document Context,” (2016).

Introduction. Two main reasons that state-of-the-art event extraction systems have difficulties:

1. They extract events and entities in separate stages.
2. They extract events independently from each individual sentence, ignoring the rest of the document.

This paper proposes an approach that simultaneously extracts events and entities within a document context. They do this by first decomposing the problem into 3 tractable subproblems:

1. Learning the dependencies between a single event [trigger] and all of its potential arguments.
2. Learning the co-occurrence relations between events across the document.
3. Learning for entity extraction.

and then combine these learned models into a joint optimization framework.

Learning Within-Event Structures. For now, assume we have some document x , a set of candidate event triggers \mathcal{T} , and a set of candidate entities \mathcal{N} . Denote the set of entity candidates that are potential arguments for trigger candidate i as \mathcal{N}_i . The joint distribution over the possible trigger types, roles, and entities for those roles, is given by

$$\Pr_{\theta} [t_i, \mathbf{r}_i, \mathbf{a} \mid i, \mathcal{N}_i, x] \propto \quad (326)$$

$$\exp \left(\theta_1^T \mathbf{f}_1(t_i) + \sum_{j \in \mathcal{N}_i} \left[\theta_2^T \mathbf{f}_2(r_{ij}) + \theta_3^T \mathbf{f}_3(t_i, r_{ij}) + \theta_4^T \mathbf{f}_4(a_j) + \theta_5^T \mathbf{f}_5(r_{ij}, a_j) \right] \right) \quad (327)$$

All f_i also depend on i, x . In addition, all f_i except f_1 depend on the current j in the summation.

where each f_i is a feature function, and I’ve colored the unary feature functions green. The unary features are tabulated in Table 1 of the original paper. They use simple indicator functions $1_{t,r}$ and $1_{r,a}$ for the pairwise features. They train using maximum-likelihood estimates with L2 regularization:

$$\theta^* = \arg \max_{\theta} \mathcal{L}(\theta) - \lambda \|\theta\|_2^2 \quad (328)$$

$$\mathcal{L}(\theta) = \sum_i \log (\Pr_{\theta} [t_i, \mathbf{r}_i, \mathbf{a} \mid i, \mathcal{N}_i, x]) \quad (329)$$

and use **L-BFGS** to optimize the training objective.

Learning Event-Event Relations. A pairwise model of event-event relations in a document. Training data consists of all pair of trigger candidates that co-occur in the same sentence or are connected by a co-referent subject/object if they're in different sentences. Given a trigger candidate pair (i, i') , we estimate the probabilities for their event types $(t_i, t_{i'})$ as

$$\Pr_{\phi} [t_i, t_{i'} \mid x, i, i'] \propto \exp \left(\phi^T g(t_i, t_{i'}, x, i, i') \right) \quad (330)$$

where g is a feature function that depends on the trigger candidate pair and their context. In addition to re-using the trigger features in Table 1 of the paper, they also introduce relational trigger features:

1. whether they're connected by a conjunction dependency relation
2. whether they share a subject or an object
3. whether they have the same head word lemma
4. whether they share a semantic frame based on FrameNet.

As before, they using L-BFGS to compute the maximum-likelihood estimates of the parameters ϕ .

Entity Extraction. Trained a standard linear-chain **CRF** using the BIO scheme. Their CRF features:

1. current words and POS tags
2. context words in a window of size 2
3. word type such as all-capitalized, is-capitalized, all-digits
4. Gazetteer-based entity type if the current word matches an entry in the gazetteers collected from Wikipedia.
5. pre-trained word2vec embeddings for each word

Joint Inference. Allows information flow among the 3 local models and finds globally-optimal assignments of all variables. Define the following objective:

$$\max_{\mathbf{t}, \mathbf{r}, \mathbf{a}} \sum_{i \in \mathcal{T}} E(t_i, \mathbf{r}_i, \mathbf{a}) + \sum_{i, i' \in \mathcal{T}} R(t_i, t_{i'}) + \sum_{j \in \mathcal{N}} D(a_j) \quad (331)$$

where

- The first term is the sum of confidence scores for individual event mentions from the within-event model.

$$E(t_i, \mathbf{r}_i, \mathbf{a}) = \log p_{\theta}(t_i) + \sum_{j \in \mathcal{N}_i} [\log p_{\theta}(t_i, r_{ij}) + \log p_{\theta}(r_{ij}, a_j)] \quad (332)$$

- The second term is the sum of confidence scores for event relations based on the pairwise event model.

$$R(t_i, t_{i'}) = \log p_{\phi}(t_i, t_{i'} \mid i, i', x) \quad (333)$$

- The third term is sum of confidence scores for entity mentions, where

$$D(a_j) = \log p_\psi(a_j \mid j, x) \quad (334)$$

and $p_\psi(a_j \mid j, x)$ is the marginal probability derived from the linear-chain CRF.

The optimization is subject to agreement constraints that enforce the overlapping variables among the 3 components to agree on their values. The joint inference problem can be formulated as an **integer linear problem (ILP)**¹⁰⁴. To solve it efficiently, they find solutions for the relaxation of the problem using a **dual decomposition algorithm AD**³.

¹⁰⁴From Wikipedia: An integer linear program in canonical form: maximize $c^T x$ subject to $Ax \leq b, x \geq 0, x \in \mathbb{Z}^n$

Globally Normalized Transition-Based Neural Networks

Table of Contents Local

Written by Brandon McKinzie

D. Andor et al., “Globally Normalized Transition-Based Neural Networks,” (2016).

Introduction. Authors demonstrate that simple FF neural networks can achieve comparable or better accuracies than LSTMs, as long as they are **globally normalized**. They don’t use any recurrence, but perform **beam search** for maintaining multiple hypotheses and introduce global normalization with a **CRF** objective to overcome the label bias problem that locally normalized models suffer from.

Transition System. Given an input sequence x , define:

- Set of states $S(x)$.
- Start state $s^\dagger \in S(x)$.
- Set of decisions $\mathcal{A}(s, x)$ for all $s \in S(x)$.
- Transition function $t(s, d, x)$ returning new state s' for any decision $d \in \mathcal{A}(s, x)$.

The scoring function $\rho(s, d; \theta)$, which gives the score for decision d in state s , will be defined:

$$\rho(s, d; \theta) = \phi(s; \theta^{(l)}) \cdot \theta^{(d)} \quad (335)$$

which is just the familiar logits computation for decision d . $\theta^{(l)}$ are the parameters of the network excluding the parameters at the final layer, $\theta^{(d)}$. $\phi(s; \theta^{(l)})$ gives the representation for state s computed by the neural network under parameters $\theta^{(l)}$.

Global vs. Local Normalization.

- **Local.** Conditional probabilities $\Pr [d_j \mid s_j; \theta]$ are normalized locally over the scores for each possible action d_j from the current state s_j .

$$\Pr_L [d_{1:n}] = \prod_{j=1}^n \Pr [d_j \mid s_j; \theta] \quad (336)$$

$$= \frac{\exp \left(\sum_{j=1}^n \rho(s_j, d_j; \theta) \right)}{\prod_{j=1}^n Z_L(s_j; \theta)} \quad (337)$$

$$Z_L(s_j; \theta) = \sum_{d' \in \mathcal{A}(s_j)} \exp (\rho(s_j, d'; \theta)) \quad (338)$$

Beam search can be used to attempt to find the action sequence with highest probability.

- **Global.** In contrast, a CRF defines:

$$\Pr_G [d_{1:n}] = \frac{\exp \left(\sum_{j=1}^n \rho(s_j, d_j; \theta) \right)}{Z_G(\theta)} \quad (339)$$

$$Z_G(\theta) = \sum_{d'_{1:n} \in \mathcal{D}_n} \exp \left(\sum_{j=1}^n \rho(s'_j, d'_j; \theta) \right) \quad (340)$$

where \mathcal{D}_n is the set of all valid sequences of decisions of length n . The inference problem is now to find

$$\arg \max_{d_{1:n} \in \mathcal{D}_n} \Pr_G [d_{1:n}] = \arg \max_{d_{1:n} \in \mathcal{D}_n} \sum_{j=1}^n \rho(s_j, d_j; \theta) \quad (341)$$

and we can also use beam search to approximately find the argmax.

Training. SGD on the NLL of the data under the model. The NLL takes a different form depending on whether we choose a locally normalized model vs a globally normalized model.

- **Local.**

$$L_{local}(d_{1:n}^*; \theta) = -\ln \Pr_L [d_{1:n}^*; \theta] \quad (342)$$

$$= -\sum_{j=1}^n \rho(s_j^*, d_j^*; \theta) + \sum_{j=1}^n \ln Z_L(s_j^*; \theta) \quad (343)$$

- **Global.**

$$L_{global}(d_{1:n}^*; \theta) = -\ln \Pr_G [d_{1:n}^*; \theta] \quad (344)$$

$$= -\sum_{j=1}^n \rho(s_j^*, d_j^*; \theta) + \ln Z_G(\theta) \quad (345)$$

To make learning tractable for the globally normalized model, the authors use beam search with early updates.

The Label Bias Problem. Locally normalized models often have a very weak ability to revise earlier decisions. Here we will prove that globally normalized models are strictly more expressive than locally normalized models. Let \mathcal{P}_L denote the set of all possible distributions $p_L(d_{1:n} \mid x_{1:n})$ under the local model as the scores ρ vary. Let \mathcal{P}_G be the same, but for the global model.

Theorem 3.1 *\mathcal{P}_L is a strict subset of \mathcal{P}_G , that is $\mathcal{P}_L \subsetneq \mathcal{P}_G$.*

An Introduction to Conditional Random Fields

Table of Contents Local

Written by Brandon McKinzie

Sutton et al., “An Introduction to Conditional Random Fields,” (2012).

Graphical Modeling. Some notation. Denote factors as $\psi_a(\mathbf{y}_a)$ where $1 \leq a \leq A$ and A is the total number factors. \mathbf{y}_a is an assignment to the subset $Y_a \subseteq Y$ of variables associated with ψ_a . The value returned by ψ_a is a non-negative scalar that can be thought of as a measure of how compatible the values \mathbf{y}_a are with each other. Given a collection of subsets $\{Y_a\}_{a=1}^A$ of Y , an **undirected graphical model** is the set of all distributions that can be written as

$$p(\mathbf{y}) = \frac{1}{Z} \prod_{a=1}^A \psi_a(\mathbf{y}_a) \quad (346)$$

$$Z = \sum_{\mathbf{y}} \prod_{a=1}^A \psi_a(\mathbf{y}_a) \quad (347)$$

for any choice of **factors** $\mathcal{F} = \{\psi_a\}$ that have $\psi_a(\mathbf{y}_a) \geq 0$ for all \mathbf{y}_a . The sum for the **partition function**, Z , is over all possible assignments \mathbf{y} of the variables Y . We’ll use the term **random field** to refer to a particular distribution among those defined by an undirected model¹⁰⁵.

We can represent the factorization with a **factor graph**: a bipartite graph $G = (V, F, E)$ in which one set of nodes $V = \{1, 2, \dots, |Y|\}$ indexes the RVs in the model, and the set of nodes $F = \{1, 2, \dots, A\}$ indexes the factors. A connection between a variable node Y_s for $s \in V$ to some factor node ψ_a for $a \in F$ means that Y_s is one of the arguments of ψ_a . It is common to draw the factor nodes as squares, and the variable nodes as circles.

Generative versus Discriminative Models. Naive Bayes is generative, while logistic regression (a.k.a maximum entropy) is discriminative. Recall that Naive Bayes and logistic are defined as, respectively,

$$p(\mathbf{y}, \mathbf{x}) = p(\mathbf{y}) \prod_{k=1}^K p(x_k | \mathbf{y}) \quad (348)$$

$$p(\mathbf{y} | \mathbf{x}) = \frac{1}{Z(\mathbf{x})} \exp \left(\sum_{k=1}^K \theta_k f_k(\mathbf{y}, \mathbf{x}) \right) \quad (349)$$

where the f_k in the definition of logistic regression denote the feature functions. We could set them, for example, as $f_{y',j}(y, \mathbf{x}) = 1_{y'=y} x_j$.

¹⁰⁵i.e. a particular set of factors.

An example generative model for sequence prediction is the **HMM**. Recall that an HMM defines

$$p(\mathbf{y}, \mathbf{x}) = \prod_{t=1}^T p(y_t | y_{t-1})p(x_t | y_t) \quad (350)$$


where we are using the dummy notation of assuming an initial-initial state y_0 clamped to 0 and begins every state sequence, so we can write the initial state distribution as $p(y_1 | y_0)$.

We see that the generative models, like naive Bayes and the HMM, define a family of joint distributions that factorizes as $p(y, x) = p(y)p(x | y)$. Discriminative models, like logistic regression, define a family of conditional distributions $p(y | x)$. The main conceptual difference here is that a conditional distribution $p(y | x)$ doesn't include a model of $p(x)$. The principal advantage of discriminative modeling is that it's better suited to include rich, overlapping features. Discriminative models like CRFs make conditional independence assumptions both (1) among y and (2) about how the y can depend on x , but do *not* make conditional independence assumptions among x .

The difference between NB and LR is due *only* to the fact that NB is generative and LR is discriminative. Any LR classifier can be converted into a NB classifier with the same decision boundary, and vice versa. In other words, NB defines the same family as LR, if we interpret NB generatively as

$$p(y, \mathbf{x}) = \frac{\exp(\sum_k \theta_k f_k(y, \mathbf{x}))}{\sum_{\tilde{y}, \tilde{x}} \exp(\sum_k \theta_k f_k(\tilde{y}, \tilde{x}))} \quad (351)$$

and train it to maximize the conditional likelihood. Similarly, if the LR model is interpreted as above, and trained to maximize the joint likelihood, then we recover the same classifier as NB.



Linear-Chain CRFs. Key point: the conditional distribution $p(\mathbf{y} \mid \mathbf{x})$ that follows from the joint distribution $p(\mathbf{y}, \mathbf{x})$ of an HMM is in fact a CRF with a particular choice of feature functions. First, we rewrite the HMM joint in a form that's more amenable to generalization:

$$p(\mathbf{y}, \mathbf{x}) = \frac{1}{Z} \prod_{t=1}^T \exp \left(\sum_{i,j \in S} \theta_{i,j} 1_{\{y_t=i, y_{t-1}=j\}} + \sum_{i \in S, o \in O} \mu_{o,i} 1_{\{y_t=i, x_t=o\}} \right) \quad (352)$$

$$= \frac{1}{Z} \prod_{t=1}^T \exp \left(\sum_{k=1}^K \theta_k f_k(y_t, y_{t-1}, x_t) \right) \quad (353)$$

and the latter provides the more compact notation. We can use Bayes rule to then write $p(\mathbf{y} \mid \mathbf{x})$, which would give us a particular kind of linear-chain CRF. The general definition of linear-chain CRFs is given below:

Let Y, X be random vectors, $\theta = \{\theta_k\} \in \mathbb{R}^K$ be a parameter vector, and $\mathcal{F} = \{f_k(y, y', x_t)\}_{k=1}^K$ be a set of real-valued feature functions. Then a **linear-chain conditional random field** is a distribution $p(\mathbf{y} \mid \mathbf{x})$ that takes the form:

$$p(\mathbf{y} \mid \mathbf{x}) = \frac{1}{Z(\mathbf{x})} \prod_{t=1}^T \exp \left(\sum_k \theta_k f_k(y_t, y_{t-1}, x_t) \right) \quad (354)$$

Notice that a linear chain CRF can be described as a factor graph over \mathbf{x} and \mathbf{y} , where each local function (factor) ψ_t has the special log-linear form:

$$\psi_t(y_t, y_{t-1}, x_t) = \exp \left(\sum_k \theta_k f_k(y_t, y_{t-1}, x_t) \right) \quad (355)$$

General CRFs. Let G be a factor graph over X and Y . Then (X, Y) is a conditional random field if for any value \mathbf{x} of X , the distribution $p(\mathbf{y} \mid \mathbf{x})$ factorizes according to G . If $F = \{\psi_a\}$ is the set of A factors in G , then the conditional distribution for a CRF is

$$p(\mathbf{y} \mid \mathbf{x}) = \frac{1}{Z(\mathbf{x})} \prod_{a=1}^A \psi_a(\mathbf{y}_a, \mathbf{x}_a) \quad (356)$$

It is often useful to require that the factors be log-linear over a prespecified set of feature functions, which allows us to write the conditional distribution as

$$p(\mathbf{y} \mid \mathbf{x}) = \frac{1}{Z(\mathbf{x})} \prod_{\psi_a \in F} \exp \left(\sum_{k=1}^{K(a)} \theta_{a,k} f_{a,k}(\mathbf{y}_a, \mathbf{x}_a) \right) \quad (357)$$

In addition, most models rely extensively on **parameter tying**¹⁰⁶. To denote this, we partition the factors of G into $\mathcal{C} = \{C_1, C_2, \dots, C_P\}$, where each C_p is a **clique template**: a set of factors

¹⁰⁶Note how, for CRFs, the actual parameters θ are tightly coupled with the feature functions f .

sharing a set of feature functions $\{f_{p,k}(\mathbf{x}_c, \mathbf{y}_c)\}_{k=1}^{K(p)}$ and a corresponding set of parameters $\boldsymbol{\theta}_p \in \mathbb{R}^{K(p)}$. A CRF that uses clique templates can be written as

$$p(\mathbf{y} \mid \mathbf{x}) = \frac{1}{Z(\mathbf{x})} \prod_{C_p \in \mathcal{C}} \prod_{\psi_c \in C_p} \psi_c(\mathbf{x}_c, \mathbf{y}_c; \boldsymbol{\theta}_p) \quad (358)$$

Feature engineering.

- **Label-observation features.** When our label variables are discrete, the features $f_{p,k}$ of a clique template C_p are ordinarily chosen to have a particular form:

$$f_{p,k}(\mathbf{y}_c, \mathbf{x}_c) = 1_{\{\mathbf{y}_c = \tilde{\mathbf{y}}_c\}} q_{p,k}(\mathbf{x}_c) \quad (359)$$

and we refer to the functions $q_{p,k}(\mathbf{x}_c)$ as *observation functions*.

- **Unsupported features.** Many observation-label pairs may never occur in our training data (e.g. having the word “with” being associated with label “CITY”). Such features are called unsupported features, and can be useful since often their weights will be driven negative, which can help prevent the model from making predictions in the future that are far from what was seen in the training data.
- **Edge-Observation and Node-Observation features:** the two most common types of label-observation features. Edge-observation features are for the transition factors, while node-observation features are the form introduced for label-observation features above.
- **Feature Induction.** The model begins with a number of base features, and the training procedure adds conjunctions of those features.

There are two inference problems that arise:

1. Wanting to predict the labels of a new input \mathbf{x} using the most likely labeling $\mathbf{y}^* = \arg \max_{\mathbf{y}} p(\mathbf{y} | \mathbf{x})$.
2. Computing marginal distributions (during parameter estimation, for example) such as node marginals $p(y_t | \mathbf{x})$ and edge marginals $p(y_t, y_{t-1} | \mathbf{x})$.

For linear-chain CRFs, the **forward-backward algorithm** is used for computing marginals, and the **Viterbi algorithm** for computing the most probable assignment. We'll first derive these for the case of HMMs, and then generalize to the linear-chain CRF case.

Forward-backward algorithm (HMMs). An efficient technique for computing marginals. We begin by writing out $p(\mathbf{x})$, and using the distributive law to convert the sum of products to a product of sums:

$$p(\mathbf{x}) = \sum_{\mathbf{y}} p(\mathbf{x}, \mathbf{y}) \quad (360) \quad \text{We define the } \psi_t \text{ as } p(y_t | y_{t-1})p(x_t | y_t)$$

$$= \sum_{\mathbf{y}} \prod_{t=1}^T \psi_t(y_t, y_{t-1}, x_t) \quad (361)$$

$$= \sum_{y_1} p(y_1) p(x_1 | y_1) \sum_{y_2} p(y_2 | y_1) p(x_2 | y_2) \sum_{y_3} \cdots \sum_{y_T} p(y_T | y_{T-1}) p(x_T | y_T) \quad (362)$$

$$= \sum_{y_1} \psi_1(y_1, x_1) \sum_{y_2} \cdots \sum_{y_T} \psi_T(y_T, y_{T-1}, x_T) \quad (363)$$

$$= \sum_T \sum_{T-1} \psi_T(y_T, y_{T-1}, x_T) \sum_{y_{T-2}} \cdots \sum_{y_1} \psi_1(y_1, x_1) \quad (364)$$

We see that we can save an exponential amount of work by caching the inner sums as we go. Let M denote the number of possible states for the y variables. We define a set of T **forward variables** α_t , each of which is a vector of size M :

$$\alpha_t(j) \triangleq p(\mathbf{x}_{\langle 1 \dots t \rangle}, y_t = j) \quad (365)$$

$$= \sum_{\mathbf{y}_{\langle 1 \dots t-1 \rangle}} p(\mathbf{x}_{\langle 1 \dots t \rangle}, \mathbf{y}_{\langle 1 \dots t-1 \rangle}, y_t = j) \quad (366)$$

$$= \sum_{\mathbf{y}_{\langle 1 \dots t-1 \rangle}} \psi_t(j, y_{t-1}, x_t) \prod_{t'=1}^{t-1} \psi_{t'}(y_{t'}, y_{t'-1}, x_{t'}) \quad (367)$$

$$= \sum_{i \in S} \psi_t(j, i, x_t) \sum_{\mathbf{y}_{\langle 1 \dots t-2 \rangle}} \psi_{t-1}(y_{t-1}, y_{t-2}, x_{t-1}) \prod_{t'=1}^{t-2} \psi_{t'}(y_{t'}, y_{t'-1}, x_{t'}) \quad (368)$$

$$= \sum_{i \in S} \psi_t(j, i, x_t) \alpha_{t-1}(i) \quad (369)$$

where S is the set of M possible states. By recognizing that $p(\mathbf{x}) = \sum_{j \in S} \sum_{\mathbf{y}_{\langle 1 \dots t-1 \rangle}} p(\mathbf{x}_{\langle 1 \dots t \rangle}, \mathbf{y}_{\langle 1 \dots t-1 \rangle}, j)$, we can rewrite $p(\mathbf{x})$ as

$$p(\mathbf{x}) = \sum_{j \in S} \alpha_T(j) \quad (370)$$

Notice how in the step from equation 365 to 366, we marginalized over all possible y subsequences that could've been aligned with $\mathbf{x}_{\langle 1 \dots t \rangle}$. We will repeat this pattern to derive the backward recursion for β_t , which is the same idea except now we go from T backward until some t (instead of going from 1 *forward* until some t).

$$\beta_t(i) \triangleq p(\mathbf{x}_{\langle t+1 \dots T \rangle} \mid y_t = i) \quad (371)$$

$$= \sum_{\mathbf{y}_{\langle t+1 \dots T \rangle}} \psi_{t+1}(y_{t+1}, i, x_{t+1}) \prod_{t'=t+2}^T \psi_{t'}(y_{t'}, y_{t'-1}, x_{t'}) \quad (372)$$

We initialize $\beta_T(i) = 1$.

$$= \sum_{y_{t+1}} \psi_{t+1}(y_{t+1}, i, x_{t+1}) \beta_{t+1}(y_{t+1}) \quad (373)$$

Similar to how we obtained equation 370, we can rewrite $p(\mathbf{x})$ in terms of the β :

$$p(\mathbf{x}) = \beta_0(y_0) \triangleq \sum_{y_1} \psi_1(y_1, y_0, x_1) \beta_1(y_1) \quad (374)$$

We can then combine the definition for α and β to compute marginals of the form $p(y_{t-1}, y_t \mid \mathbf{x})$:

$$p(y_{t-1}, y_t \mid \mathbf{x}) = \frac{1}{p(\mathbf{x})} \alpha_{t-1}(y_{t-1}) \psi_t(y_t, y_{t-1}, x_t) \beta_t(y_t) \quad (375)$$

$$\text{where } p(\mathbf{x}) = \sum_{j \in S} \alpha_T(j) = \beta_0(y_0) \quad (376)$$

In summary, the forward-backward algorithm consists of the following steps:

1. Compute α_t for all t using equation 370.
2. Compute β_t for all t using equation 374.
3. Return the marginal distributions computed from equation 375.

Viterbi algorithm (HMMs). For computing $\mathbf{y}^* = \arg \max_{\mathbf{y}} p(\mathbf{y} \mid \mathbf{x})$. The derivation is nearly the same as how we derived the forward-backward algorithm, except now we've replaced the summations in equation 364 with maximization. The analog of α for viterbi are defined as:

$$\delta_t(j) = \max_{\mathbf{y}_{\langle 1 \dots t-1 \rangle}} \psi_t(j, y_{t-1}, x_t) \prod_{t'=1}^{t-1} \psi_{t'}(y_{t'}, y_{t'-1}, x_{t'}) \quad (377)$$

$$= \max_{i \in S} \psi_t(j, i, x_t) \delta_{t-1}(i) \quad (378)$$

and the maximizing assignment is computed by a backwards recursion,

$$y_T^* = \arg \max_{i \in S} \delta_T(i) \quad (379)$$

$$y_t^* = \arg \max_{i \in S} \psi_t(y_{t+1}^*, i, x_{t+1}) \delta_t(i) \text{ for } t < T \quad (380)$$

Computing the recursions for δ_t and y_t^* together is the *Viterbi algorithm*.

Forward-backward and Viterbi for linear-chain CRF. Generalizing to the linear-chain CRF, where now

$$p(\mathbf{y} \mid \mathbf{x}) = \frac{1}{Z(\mathbf{x})} \prod_{t=1}^T \psi_t(y_t, y_{t-1}, x_t) \quad (381)$$

$$\text{where } \psi_t(y_t, y_{t-1}, x_t) = \exp \left\{ \sum_k \theta_k f_k(y_t, y_{t-1}, x_t) \right\} \quad (382)$$

and the results for the forward-backward algorithm become

$$p(y_{t-1}, y_t \mid \mathbf{x}) = \frac{1}{Z(\mathbf{x})} \alpha_{t-1}(y_{t-1}) \psi_t(y_t, y_{t-1}, x_t) \beta_t(y_t) \quad (383)$$

$$\text{where } Z(\mathbf{x}) = \sum_{j \in S} \alpha_T(j) = \beta_0(y_0) \quad (384)$$

Co-sampling: Training Robust Networks for Extremely Noisy Supervision

Table of Contents Local

Written by Brandon McKinzie

Han et al., “Co-sampling: Training Robust Networks for Extremely Noisy Supervision,” (2018).

Introduction. The authors state that current methodologies [for training networks under noisy labels] involves estimating the **noise transition matrix** (which they don’t define). Patrini et al. (2017) define the matrix as follows:

Denote by $T \in [0, 1]^{c \times c}$ the noise transition matrix specifying the probability of one label being flipped to another, so that $\forall i, j$ $T_{ij} \triangleq \Pr[\tilde{y} = e^j \mid y = e^i]$. The matrix is row-stochastic¹⁰⁷ and not necessarily symmetric across the classes.

Algorithm. Authors propose a learning paradigm called **Co-sampling**. They maintain two networks f_{w_1} and f_{w_2} simultaneously. For each mini-batch data $\hat{\mathcal{D}}$, each network selects \mathcal{R}_T small-loss instances as a “clean” mini-batch $\hat{\mathcal{D}}_1$ and $\hat{\mathcal{D}}_2$, respectively. Each of the two networks then uses the clean mini-batch data to update the parameters w_2 (w_1) of its *peer* network.

- Why small-loss instances? Because deep networks tend to fit clean instances first, then noisy/harder instances progressively after.
- Why two networks? Because if we just trained a single network on clean instances, we would not be robust in extremely high-noise rates, since the training error would accumulate if the selected instances are not “fully clean.”

The Co-sampling paradigm algorithm is shown below.

```

Input:  $w_1$  and  $w_2$ ; learning rate  $\eta$ ; fixed  $\mathfrak{D}$ ; epoch  $T_k$  and  $T_{max}$ ; iteration  $N_{max}$ ;
for  $T = 1, 2, \dots, T_{max}$  do
  Shuffle: training set  $\tilde{\mathcal{D}}$  //noisy dataset
  for  $N = 1, \dots, N_{max}$  do
    Draw: mini-batch  $\hat{\mathcal{D}}$  from  $\tilde{\mathcal{D}}$ 
    Sample:  $\hat{\mathcal{D}}_1 = \arg \min_{\hat{\mathcal{D}}} \ell(f_{w_1}, \hat{\mathcal{D}}, \mathfrak{R}_T)$  //sample  $\mathfrak{R}_T$  small-loss instances
    Sample:  $\hat{\mathcal{D}}_2 = \arg \min_{\hat{\mathcal{D}}} \ell(f_{w_2}, \hat{\mathcal{D}}, \mathfrak{R}_T)$  //sample  $\mathfrak{R}_T$  small-loss instances
    Update:  $w_1 = w_1 - \eta \nabla f_{w_1}(\hat{\mathcal{D}}_2)$  //update  $w_1$  by  $\hat{\mathcal{D}}_2$ 
    Update:  $w_2 = w_2 - \eta \nabla f_{w_2}(\hat{\mathcal{D}}_1)$  //update  $w_2$  by  $\hat{\mathcal{D}}_1$ 
  end
  Update:  $\mathfrak{R}_T = 1 - \min\{\frac{T}{T_k} \mathfrak{D}, \mathfrak{D}\}$ 
end
Output:  $f_{w_1}$  and  $f_{w_2}$ 

```

¹⁰⁷Each row sums to 1.

NLP WITH DEEP LEARNING

CONTENTS

5.1	Word Vector Representations (Lec 2)	157
5.2	GloVe (Lec 3)	160

Word Vector Representations (Lec 2)

Table of Contents Local

Written by Brandon McKinzie

Meaning of a word. Common answer is to use a *taxonomy* like WordNet that has hypernyms (is-a) relationships. Problems with this discrete representation: misses nuances, e.g. the words in a set of synonyms can actually have quite different meanings/connotations. Furthermore, viewing words as atomic symbols is a bit like using one-hot vectors of words in a vocabulary space (inefficient).

Distributed representations. Want a way to encode word vectors such that two similar words have a similar structure/representation. The **distributional similarity-based**¹⁰⁸ approach represents words by means of its *neighbors* in the sentences in which it appears. You end up with a dense “vector for each word type, chosen so that it is good at predicting other words appearing in its context.”

Skip-gram prediction. Given a word w_t at position t in a sentence, learn to predict [probability of] some number of surrounding words, given w_t . Standard minimization with negative log-likelihood:

$$J(\theta) = -\frac{1}{T} \sum_{t=1}^T \sum_{-m \leq j \leq m} \log \Pr(w_{t+j} | w_t) \quad (385)$$

$$\Pr(o | c) = \frac{e^{\mathbf{u}_o^T \mathbf{v}_c}}{\sum_{w=1}^{\text{vocab size}} e^{\mathbf{u}_w^T \mathbf{v}_c}} \quad (386)$$

I cannot believe this actually works.

where

- The params θ are the vector representation of the words (they are the *only* learnable parameters here).
- m is our radius/window size.
- o and c are indices into our vocabulary (somewhat inconsistent notation).
- Yes, they are using different vector representations for \mathbf{u} (context words) and \mathbf{v} (center words). I’m assuming one reason this is done is because it makes the model architecture simpler/easier to build.

Some subtleties:

¹⁰⁸Note that this is distinct from the way “distributed” is meant in “distributed representation.” In contrast, distributional similarity-based representations refers to the notion that you can describe the meaning of words by understanding the context in which they appear.

- Looks like e.g. $Pr(w_{t+j} | w_t)$ doesn't really care what the value of j is, it is just modeling the probability that it is *somewhere* in the context window. The w_t are one-hot vectors into the vocabulary. Standard tricks for simplifying the cross-entropy loss apply.
- Equation 386 should be interpreted as the probability that the o^{th} word in our vocabulary occurs in the context window of the c^{th} word in our vocabulary.
- The model architecture is *identical* to an autoencoder. However, the (big) difference is the training procedure and interpretation of the model parameters going “in” versus the parameters going “out”.

Sentence Embeddings. It turns out that simply taking a weighted average of word vectors and doing some PCA/SVD is a competitive way of getting unsupervised word embeddings. Apparently it *beats* supervised learning with LSTMs (!). The authors claim the theoretical explanation for this method lies in a latent variable generative model for sentences (of course). Approach:

Discussion based on paper by Arora et al. (2017).

1. Compute the weighted average of the word vectors in the sentence:

$$\frac{1}{N} \sum_i^N \frac{a}{a + p(\mathbf{w}_i)} \mathbf{w}_i \quad (387)$$

The authors call their weighted average the **Smooth Inverse Frequency (SIF)**.

where \mathbf{w}_i is the word vector for the i th word in the sentence, a is a parameter, and $p(\mathbf{w}_i)$ is the (estimated) word frequency [over the entire corpus].

2. Remove the projections of the average vectors on their first principal component (“common component removal”) (y tho?).

1

Algorithm 1 Sentence Embedding

Input: Word embeddings $\{v_w : w \in \mathcal{V}\}$, a set of sentences \mathcal{S} , parameter a and estimated probabilities $\{p(w) : w \in \mathcal{V}\}$ of the words.

Output: Sentence embeddings $\{v_s : s \in \mathcal{S}\}$

```

1: for all sentence  $s$  in  $\mathcal{S}$  do
2:    $v_s \leftarrow \frac{1}{|s|} \sum_{w \in s} \frac{a}{a + p(w)} v_w$ 
3: end for
4: Compute the first principal component  $u$  of  $\{v_s : s \in \mathcal{S}\}$ 
5: for all sentence  $s$  in  $\mathcal{S}$  do
6:    $v_s \leftarrow v_s - uu^\top v_s$ 
7: end for
```

Further Reading.

- Learning representations by back-propagating errors (Rumelhard et al., 1986)
- A Neural Probabilistic Language Model (Bengio et al., 2003)
- NLP (almost) from Scratch (Collobert & Watson, 2008)
- Word2Vec (Miklov et al. 2013)

GloVe (Lec 3)

Table of Contents Local

Written by Brandon McKinzie

Skip-gram and negative sampling. Main idea:

- Split the loss function from last lecture into two (additive) terms corresponding to the numerator and denominator respectively (you’ve done this a trillion times).
- The second term is an expectation over all the words in your vocab space. That is huge, so instead we only use a subsample of size k (the negative samples).
- Interpretation: First term is maximizing $\Pr(o | c)$, the probability that the true outside word (given by index o) occurs given context (index) c . Second term is minimizing the probability of random words (the negative samples) occurring around the center (context) word given by c .

To sample the negative samples, draw from $P(w) = U(w)^{3/4} / Z$, where U is the **unigram distribution**.

GloVe (Global Vectors). Given some co-occurrence matrix we computed with previous methods, we can use the following GloVe loss function over all pairs of co-occurring words in our matrix:

$$J(\theta) = \sum_{i,j=1}^W f(P_{ij})(u_i^T v_j - \log P_{ij})^2 \quad (388)$$

where P_{ij} is computed simply from the counts of words i and j co-occurring (empirical probability) and f is some squashing function that really isn’t discussed in this lecture (**TODO**).

Evaluating word vectors.

- **Word Vector Analogies:** Basically, determining if we can do standard analogy fill-in-the-blank problems: “*man [a] is to woman [b] as king [c] is to <blank>*” (if you answered “queen”, you’d make a good AI). We can determine this using a standard cosine distance measure:

$$d = \arg \max_i \frac{(x_b - x_a + x_c)^T x_i}{\|x_b - x_a + x_c\|} \quad (389)$$

Woah that is pretty neat. The solution is $x_i = \text{queen}$. $x_b - x_a$ is the vector pointing from man to woman, which encodes the type of similarity we are looking for with the other pair. Therefore, we take the vector to “king” and *add* the aforementioned difference vector – the resultant vector should point to “queen”. Neat!

Derivation. Based on the descriptions in the original paper¹⁰⁹ We want to develop a model for learning word vectors.

1. The authors argue that “the appropriate starting point for word vector learning should be with ratios of co-occurrence probabilities rather than the probabilities themselves.” The most general such model takes the form,

$$F(w_i, w_j, \tilde{w}_k) = \frac{\Pr[\tilde{w}_k | w_i]}{\Pr[\tilde{w}_k | w_j]} \equiv \frac{P_{ik}}{P_{jk}} \quad \text{where all } w \in \mathbb{R}^d \quad (390)$$

and the tilde in \tilde{w}_k denotes that \tilde{w}_k is a *context* word vector, which are given a distinct space from the word vectors w_i and w_j being compared. We compute all P_{ik} via frequency counts over the corpus.

2. Now that we’ve specified the inputs and ratio of interest, we can start specifying some desirable constraints on the function F that we’re trying to find. The authors speculate that, since vector spaces are linear structures, we should have F encode the information of the ratio in the vector space via vector differences:

$$F(w_i - w_j, \tilde{w}_k) = \frac{P_{ik}}{P_{jk}} \quad (391)$$

which basically says “our representation of the word vectors should be s.t. the **relative** probability of some word \tilde{w}_k occurring in the context of a word w_i compared to it occurring in the context of a different word w_j can be captured by $w_i - w_j$ and \tilde{w}_k alone.”

3. Next we notice that F maps arguments in \mathbb{R}^d to a scalar in \mathbb{R} . The most straightforward way of doing so while maintaining the linear structure we are trying to capture is via a dot product:

$$F((w_i - w_j)^T \tilde{w}_k) = \frac{P_{ik}}{P_{jk}} \quad (392)$$

Note that now $F : \mathbb{R} \mapsto \mathbb{R}$.

4. We want our model to be invariant under the exchanges $w \leftrightarrow \tilde{w}$ and $X \leftrightarrow X^T$. We can restore this symmetry by first requiring that F be a **homomorphism**¹¹⁰ between the groups $(\mathbb{R}, +)$ and $(\mathbb{R}_{>0}, \times)$ (in our case, negation and division would be better symbols, but it’s equivalent). This requires that the following relation hold¹¹¹

$$F(w_i^T \tilde{w}_k - w_j^T \tilde{w}_k) = \frac{F(w_i^T \tilde{w}_k)}{F(w_j^T \tilde{w}_k)} \quad (393)$$

where I’ve grouped terms on the LHS to emphasize how this is the definition of homomorphism. The solution for this equation is that $F(\cdot) \equiv \exp(\cdot)$. Combining this realization with equations 392 and 393 yields,

$$F(w_i^T \tilde{w}_k) = e^{w_i^T \tilde{w}_k} = P_{ik} \quad (394)$$

$$\Rightarrow w_i^T \tilde{w}_k = \log(P_{ik}) = \log(X_{ik}) - \log(X_i) \quad (395)$$

¹⁰⁹Pennington et al., “GloVe: Global Vectors for Word Representation.”

¹¹⁰In more detail, $F : (\mathbb{R}, +) \mapsto (\mathbb{R}_{>0}, \times)$, which reads “the function F maps elements in \mathbb{R} and any summation of elements in \mathbb{R} to elements in \mathbb{R} that are greater than zero or any product of positive elements in \mathbb{R} .”

¹¹¹Note that we do not need to know anything about the RHS of the equations above to state this relation. We write it by definition of homomorphism.

where X_{ik} is the number of times word k appears in the context of word i , and $X_i = \sum_k X_{ik}$ is the number of times any word appears in the context of i .

5. Restore symmetry under the exchanges $w \leftrightarrow \tilde{w}$ and $X \leftrightarrow X^T$. We absorb X_i into a bias b_i for w_i since it is independent of k .

$$w_i^T \tilde{w}_k + b_i + \tilde{b}_k = \log(X_{ik}) \quad (396)$$

6. A main drawback to this model is that it weighs all co-occurrences equally, even those that happen rarely or never. The authors propose a new weighted least squares regression model, introducing a weighting function $f(X_{ij})$ into the cost function of our model:

$$J = \sum_{i,j=1}^V f(X_{ij})(w_i^T \tilde{w}_j + b_i + \tilde{b}_j - \log X_{ij})^2 \quad (397)$$

SPEECH AND LANGUAGE PROCESSING

CONTENTS

6.1	Introduction (Ch. 1 2nd Ed.)	164
6.2	Morphology (Ch. 3 2nd Ed.)	165
6.3	N-Grams (Ch. 6 2nd Ed.)	166
6.4	Naive Bayes and Sentiment (Ch. 6 3rd Ed.)	168
6.5	Hidden Markov Models (Ch. 9 3rd Ed.)	170
6.6	POS Tagging (Ch. 10 3rd Ed.)	173
6.7	Formal Grammars (Ch. 11 3rd Ed.)	176
6.8	Vector Semantics (Ch. 15)	177
6.9	Semantics with Dense Vectors (Ch. 16)	180
6.10	Information Extraction (Ch. 21 3rd Ed)	183

Introduction (Ch. 1 2nd Ed.)

Table of Contents Local

Written by Brandon McKinzie

Overview. Going to rapidly jot down what seems most important from this chapter.

- **Morphology:** captures information about the shape and behavior of words in context (Ch. 2/3).
- **Syntax:** knowledge needed to order and group words together.
- **Lexical semantics:** knowledge of the meanings of individual words.
- **Compositional semantics:** knowledge of how these components (words) combine to form larger meanings.
- **Pragmatics:** the appropriate use of polite and indirect language.
- The knowledge of language needed to engage in complex language behavior can be separated into the following 6 distinct categories:
 1. Phonetics and Phonology – The study of linguistic sounds.
 2. Morphology – The study of the meaningful components of words.
 3. Syntax – The study of the structural relationships between words.
 4. Semantics – The study of meaning.
 5. Pragmatics – The study of how language is used to accomplish goals.
 6. Discourse – The study of linguistic units larger than a single utterance.
- Methods for **resolving ambiguity:** pos-tagging, word sense disambiguation, probabilistic parsing, and speech act interpretation.
- **Models and Algorithms.** Among the most important are **state space search** and **dynamic programming** algorithms.

Morphology (Ch. 3 2nd Ed.)

Table of Contents Local

Written by Brandon McKinzie

English Morphology. Morphology is the study of the way words are built up from smaller meaning-bearing units, **morphemes**. A morpheme is often defined as the minimal meaning-bearing unit in a language¹¹². The two classes of morphemes are **stems** (the “main” morpheme of the word) and **affixes** (the “additional” meanings of various kinds).

Affixes are further divided into prefixes (precede stem), suffixes (follow stem), circumfixes (do both), and infixes (inside the stem).

Two classes of ways to form words from morphemes: **inflection** and **derivation**. Inflection is the combination of a word stem with a grammatical morpheme, usually resulting in a word of the same class as the original stem, and usually filling some syntactic function like agreement. Derivation is the combination of a word stem with a grammatical morpheme, usually resulting in a word of a different class, often with a meaning hard to predict exactly.

¹¹²Examples: “fox” is its own morpheme, while “cats” consists of the morpheme “cat” and the morpheme “-s”.

Counting Words. Most N -gram based systems deal with the *wordform*, meaning they treat words like “cats” and “cat” as distinct. However, we may want to treat the two as instances of a single abstract word, or **lemma**: a set of lexical forms having the same stem, the same major part of speech, and the same word-sense.

Simple (Unsmoothed) N-Grams. An N -gram is a $N-1$ th order Markov model (because it looks $N-1$ steps in the past). Notation: the authors use the convention that $w_1^n \triangleq w_1, w_2, \dots, w_n$ to denote a sequence of n words. Given this, we can write the general equation for the N -gram approximation for the probability of the n th word ($n > N$) in a sentence:

$$\Pr[w_n \mid w_1^{n-1}] \approx \Pr[w_n \mid w_{n-N+1}^{n-1}] \quad (398)$$

for $N \geq 2$. We can compute these probabilities by simply counting:

$$\Pr[w_n \mid w_1^{n-1}] = \frac{C(w_{n-N+1}^{n-1} w_n)}{C(w_{n-N+1}^{n-1})} \quad (399)$$

where $C(\cdot)$ is the number of times the sequence, denoted as \cdot , occurred in the corpus.

Entropy. Denote the random variable of interest as \mathbf{x} with probability function $p(\mathbf{x})$. The entropy of this random variable is:

$$H(\mathbf{x}) = - \sum_x p(\mathbf{x} = x) \log_2 p(\mathbf{x} = x) \quad (400)$$

which should be thought of as a lower bound on the number of bits it would take to encode a certain decision or piece of information in the optimal coding scheme. The value 2^H is the **perplexity**, which can be interpreted as the weighted average number of choices a random variable has to make.

Cross Entropy for Comparing Models. Useful when we don't know the actual probability distribution p that generated some data. Assume we have some model m that's an approximation of p . The cross-entropy of m on p is defined by:

$$H(p, m) = \lim_{n \rightarrow \infty} \frac{1}{n} \sum_{W \in L} p(w_1, \dots, w_n) \log m(w_1, \dots, w_n) \quad (401)$$

That is we draw sequences according to the probability distribution p , but sum the log of their probability according to m . **TODO:** continue off at pg 250/975 of textbook.

Naive Bayes and Sentiment (Ch. 6 3rd Ed.)

Table of Contents Local

Written by Brandon McKinzie

[\[3rd Ed.\]](#) [\[Quick Review\]](#)

Overview. This chapter is concerned with **text categorization**, the task of classifying an entire text by assigning it a label drawn from some set of labels. *Generative classifiers* like naive Bayes build a model of each class. Given an observation, they return the class most likely to have generated the observation. *Discriminative classifiers* like logistic regression instead learn what features from the input are most useful to discriminate between the different possible classes. Notation: we will assume we have a training set of N documents, each hand-labeled with some class: $\{(d_1, c_1), \dots, (d_N, c_N)\}$.

Discriminative systems are often more accurate and hence more commonly used.

Naive Bayes. A multinomial¹¹³ classifier with a naive assumption about how the features interact. We model a text document as a **bag of words**, meaning we store (1) the words that occurred and (2) their frequencies. It's a probabilistic classifier, meaning it estimates the label/class of a document d as

$$\hat{c} = \arg \max_{c \in C} \Pr [c \mid d] \quad (403)$$

$$= \arg \max_{c \in C} \frac{\Pr [d \mid c] \Pr [c]}{\Pr [d]} \quad (404)$$

$$= \arg \max_{c \in C} \Pr [d \mid c] \Pr [c] \quad (405)$$

Computing the class-conditional distribution (the likelihood) $\Pr [d \mid c]$ over all possible $d \in D$ is typically intractable; we must introduce some simplifying assumptions and use an approximation of it. Our assumptions in this section will be:

- **Bag-of-Words:** Assume that word position within a document is irrelevant. (Counts still matter)

¹¹³

Rapid review: **multinomial distribution.** Let $\mathbf{x} = (x_1, \dots, x_k)$ denote the result of an experiment with n independent trials ($n = \sum_i x_i$) and k possible outcomes for any given trial. i.e. x_i is the number of trials that had outcome i ($1 \leq i \leq k$). The pmf of this multinomial distribution, over all possible \mathbf{x} constrained by $n = \sum_i x_i$, is:

$$\Pr [\mathbf{x} = (x_1, \dots, x_k); n] = \frac{n!}{x_1! \cdots x_k!} p_1^{x_1} \times \cdots \times p_k^{x_k} \quad (402)$$

where p_i is the probability of outcome i for any single trial.

- **Naive Bayes Assumption:** First, recall that d is typically modeled as a (random) vector consisting of features f_1, \dots, f_n , each of which has an associated probability distribution $\Pr[f_i | c]$. The NB assumption is that the features are mutually independent given the class c :

$$\Pr[f_1, \dots, f_n | c] = \Pr[f_1 | c] \cdots \Pr[f_n | c] \quad (406)$$

$$c_{NB} = \arg \max_{c \in C} \Pr[c] \prod_{f \in F} \Pr[f | c] \quad (407)$$

where 407 is the final equation for the class chosen by the naive Bayes classifier.

In text classification we typically use the word at position i in the document as f_i , and move to log space to avoid underflow/increase speed:

$$c_{NB} = \arg \max_{c \in C} \log \Pr[c] + \sum_i^{\text{len}(d)} \log \Pr[w_i | c] \quad (408)$$

Classifiers that use a linear combination of the inputs to make a classification decision (e.g. NB, logistic regression) are called linear classifiers.

Training the NB Classifier. No real "training" in my opinion, just simple counting from the data:

$$\hat{P}[c] = \frac{N_c}{N_{docs}} \quad (409)$$

$$\hat{P}[w_i | c] = \frac{\text{count}(w_i, c) + 1}{(\sum_{w \in V} \text{count}(w, c)) + |V|} \quad (410)$$

The Laplace smoothing is added to avoid the occurrence of zero-probabilities in equation 407.

Optimizing for Sentiment Analysis.

- It often improves performance [for sentiment] to clip word counts in each document to 1 ("binary NB").
- deal with *negations* in some way.
- Use sentiment lexicons, lists of words that are pre-annotated with positive or negative sentiment.

Hidden Markov Models (Ch. 9 3rd Ed.)

Table of Contents Local

Written by Brandon McKinzie

Overview. Here we will first go over the math behind HMMs: the **Viterbi**, **Forward**, and **Baum-Welch** (EM) algorithms for unsupervised or semi-supervised learning. Recall that a HMM is defined by specifying the set of N states Q , transition matrix A , sequence of T observations O , sequence of observation likelihoods B , and the initial/final states. They can be characterized by three fundamental problems:

1. **Likelihood.** Given an HMM $\lambda = (A, B)$ and observation sequence, compute the likelihood (prob. of the observations given the model). (**Forward**)
2. **Decoding.** Given an HMM $\lambda = (A, B)$ and observation sequence, discover the best hidden state sequence. (**Viterbi**)
3. **Learning.** Given an observation sequence and the set of states in the HMM, learn the HMM parameters A and B . (**Baum-Welch/Forward-Backward/EM**)

The Forward Algorithm. For likelihood computation. We want to compute the probability of some sequence of observations O , without knowing the sequence of hidden states (that emitted the observations) Q . In general, this can be expressed by summing over all possible hidden state sequences:

$$\Pr [O] = \sum_Q \Pr [Q] \Pr [O | Q] \quad (411)$$

However, for N hidden states and T observations, this summation involves N^T terms, which becomes intractable rather quickly. Instead, we can use the $\mathcal{O}(N^2T)$ **Forward Algorithm**. The forward algorithm can be defined by initialization, recursion definition, and termination, shown respectively as follows:

$$\alpha_1(j) = a_{0j}b_j(o_1) \quad 1 \leq j \leq N \quad (412)$$

$$\alpha_t(j) = \sum_{i=1}^N \alpha_{t-1}(i)a_{ij}b_j(o_t) \quad 1 \leq j \leq N, 1 \leq t \leq T \quad (413)$$

$$\Pr [O | \lambda] = \alpha_T(q_F) = \sum_{i=1}^N \alpha_T(i)a_{iF} \quad (414)$$

Viterbi Algorithm. For decoding. Want the most likely hidden state sequence given observations. Let $v_t(j)$ represent the probability that we are in state j after t observations and passing through the most probable state sequence q_0, q_1, \dots, q_{t-1} . Similar to the forward algorithm, we show the defining equations for the Viterbi algorithm below:

$$v_1(j) = a_{0j}b_j(o_1) \quad 1 \leq j \leq N \quad (415)$$

$$v_t(j) = \max_{i=1}^N v_{t-1}(i) a_{ij} b_j(o_t) \quad (416)$$

$$P^* = v_T(q_F) = \max_{i=1}^N v_T(i) a_{iF} \quad (417)$$

N.B.: At each step, the best path up to that point can be found by taking the *argmax* instead of max.

Baum-Welch Algorithm. AKA forward-backward algorithm, a special case of the EM algorithm. Given an observation sequence O and the set of best possible states in the HMM, learn the HMM parameters A and B . First, we must define some new notation. The **backward probability** β is defined as:

$$\beta_t(i) \triangleq \Pr[o_{t+1}, \dots, o_T \mid q_t = i, \lambda] \quad (418) \quad \text{Remember } \lambda \equiv (A, B)$$

As usual, we can compute its values inductively as follows:

$$\beta_T(i) = a_{iF} \quad 1 \leq i \leq N \quad (419)$$

$$\beta_t(i) = \sum_{j=1}^N a_{ij} b_j(o_{t+1}) \beta_{t+1}(j) \quad 1 \leq i \leq N, 1 \leq t \leq T \quad (420)$$

$$\Pr[O \mid \lambda] = \alpha_T(q_F) = \beta_1(q_0) \quad (421)$$

$$= \sum_{j=1}^N a_{0j} b_j(o_1) \beta_1(j) \quad (422)$$

We can use the forward and backward probabilities α and β to compute the transition probabilities a_{ij} and observation probabilities $b_i(o_t)$ from an observation sequence. The derivation steps are as follows:

1. **Estimating \hat{a}_{ij} .** Begin by defining quantities that will prove useful:

$$\xi_t(i, j) \triangleq \Pr [q_t = i, q_{t+1} = j \mid O, \lambda] \quad (423)$$

$$\tilde{\xi}_t(i, j) \triangleq \Pr [q_t = i, q_{t+1} = j, O \mid \lambda] \quad (424)$$

$$= \alpha_t(i) a_{ij} b_j(t+1) \beta_{t+1}(j) \quad (425)$$

Remember, knowing the observation sequence does NOT give us the sequence of hidden states.

where you should be able to derive eq. 425 in your head using just logic. If you cannot, review before continuing. We can then derive $\xi_t(i, j)$ using basic definitions of conditional probability, combined with eq. 421. Finally, we estimate \hat{a}_{ij} as the expected number of transitions $q_i \rightarrow q_j$ divided by the expected number of transitions from q_i total:

$$\hat{a}_{ij} = \frac{\sum_{t=1}^{T-1} \xi_t(i, j)}{\sum_{t=1}^{T-1} \sum_{k=1}^N \xi_t(i, k)} \quad (426)$$

2. **Estimating $\hat{b}_j(v_k)$.** We define our estimate as the expected number of times we are in q_j and emit observation v_k , divided by the expected number of times we are in state j . Similar to our approach for \hat{a}_{ij} we define helper quantities for these values at a given timestep, then sum over them (all t) to obtain our estimate.

$$\gamma_t(j) \triangleq \Pr [q_t = j \mid O, \lambda] \quad (427)$$

$$= \frac{\Pr [q_t = j, O \mid \lambda]}{\Pr [O \mid \lambda]} \quad (428)$$

$$= \frac{\alpha_t(j) \beta_t(j)}{\Pr [O \mid \lambda]} \quad (429)$$

Thus, we obtain $\hat{b}_j(v_k)$ by summing over all timesteps where $o_t = v_k$, denoted as the set T_{v_k} , divided by the summation over all t regardless of o_t :

$$\hat{b}_j(v_k) = \frac{\sum_{t \in T_{v_k}} \gamma_t(j)}{\sum_{t=1}^T \gamma_t(j)} \quad (430)$$

At last we can finally define the **Forward-Backward Algorithm** as follows:

1. Initialize A and B .
2. **E-step.** Compute $\gamma_t(j)$ ($\forall t, j$), and compute $\xi_t(i, j)$ ($\forall t, i, j$).
3. **M-step.** Update all \hat{a}_{ij} and $\hat{b}_j(v_k)$.
4. Upon convergence, return A and B .

POS Tagging (Ch. 10 3rd Ed.)

Table of Contents Local

Written by Brandon McKinzie

English Parts-of-Speech. POS are traditionally defined based on syntactic and morphological function, grouping words that have similar neighboring words or take similar affixes.

	Part of Speech	Definition	Properties
Open classes:	noun	people, places, things	occur with determiners, take possessives,
	verb	actions, processes	3rd-person-sg, progressive, past participle
	adjective	properties, qualities	
	adverb	modify verbs, adverbs, verb phrases	directional, locative, degree, manner, tempo

Common nouns can be divided into *count* (e.g. goat/goats) and *mass* (e.g. snow) nouns.

Closed classes. POS with relatively fixed membership. Some of the most important in English are:

prepositions: on, under, over, near, by, at, from, to, with
determiners: a, an, the
pronouns: she, who, I, others
conjunctions: and, but, or, as, if, when
auxiliary verbs: can, may, should, are
particles: up, down, on, off, in, out, at, by
numerals: one, two, three, first, second, third

Some subtleties: the **particle** resembles a preposition or an adverb and is used in combination with a verb. An example case where “over” is a particle: “she turned the paper over.” When a verb and a particle behave as a single syntactic and/or semantic unit, we call the combination a **phrasal verb**. Phrasal verbs cause widespread problems with NLP because they often behave as a semantic unit with a noncompositional meaning – one that is not predictable from the distinct meanings of the verb and the particle. Thus, “turn down” means something like “reject”, “rule out” means “eliminate”, “find out” is “discover”, and “go on” is “continue”.

HMM POS Tagging. Since we typically train on labeled data, we need only use the Viterbi algorithm for decoding¹¹⁴. In the POS case, we wish to find the sequence of n tags, \hat{t}_1^n , given the observation sequence of n words w_1^n .

$$\hat{t}_1^n = \arg \max_{t_1^n} \Pr [t_1^n | w_1^n] = \arg \max_{t_1^n} \Pr [w_1^n | t_1^n] \Pr [t_1^n] \quad (431)$$

where we’ve dropped the denominator after using Bayes’ rule (since argmax is the same).

¹¹⁴Recall that decoding is the problem of finding the best hidden state sequence, given $\lambda = (A, B)$ and observation sequence O .

HMM taggers made two further simplifying assumptions:

$$\Pr [w_1^n | t_1^n] \approx \prod_{i=1}^n \Pr [w_i | t_i] \quad (432)$$

$$\Pr [t_1^n] \approx \prod_{i=1}^n \Pr [t_i | t_{i-1}] \quad (433)$$

We can thus plug-in these values into eq. 431 to obtain the equation for \hat{t}_1^n .

$$\hat{t}_1^n = \arg \max_{t_1^n} \prod_{i=1}^n \Pr [w_i | t_i] \Pr [t_i | t_{i-1}] \quad (434)$$


$$= \arg \max_{t_1^n} \prod_{i=1}^n b_i(w_i) a_{i-1,i} \quad (435)$$

where I've written a “translated” version on the second line using the familiar syntax from the previous chapter. In practice, we can obtain quick estimates for the two probabilities on the RHS by taking counts/averages over our tagged training data. We then run through the Viterbi algorithm to find all the argmaxes over states for the most likely hidden state sequence.

Maximum Entropy Markov Models (MEMMs). A sequence model adaptation of the MaxEnt (multinomial logistic regression) classifier¹¹⁵. Since HMMs are generative models, they decompose $\Pr [T | W]$ into $\Pr [W | T] \Pr [T]$ when computing the best tag sequence \hat{T} . Since MEMMs are discriminative, they compute/model $\Pr [T | W]$ directly:

$$\hat{T} = \arg \max_T \Pr [T | W] = \arg \max_T \prod_i \Pr [t_i | w_i, t_{i-1}] \quad (436)$$

Visually, we can think of the difference between HMMs and MEMMs via the direction of arrows, as illustrated below.



¹¹⁵Because it is based on logistic regression, the MEMM is a **discriminative sequence model**. By contrast, the HMM is a **generative sequence model**.

The top shows the HMM representation, while the bottom is MEMM.

The reason to use a discriminative sequence model is that discriminative models make it easier to incorporate a much wider variety of features.

Bidirectionality. The one problem with the MEMM and HMM models as presented is that they are exclusively run left-to-right. MEMMs¹¹⁶ have a weakness known as the **label bias problem**. Consider the tagged fragment: “*will*/NN *to*/TO *fight*/VB¹¹⁷.” Even though the word “*will*” is followed by “*to*”, which strongly suggests “*will*” is a NN, a MEMM will incorrectly label “*will*” as MD (modal verb). The culprit lies in the fact that $\Pr[\text{TO} \mid \text{to}, t_{\text{will}}]$ is essentially 1 regardless of t_{will} ; i.e. the fact that “*to*” must have the tag TO has **explained away** the presence of TO and so the model doesn’t learn the importance of the previous NN tag for predicting TO.

One way to implement bidirectionality (and thus allowing e.g. the link between TO being available when tagging the NN) is to use a **Conditional Random Field** (CRF) model. However, CRFs are much more computationally expensive than MEMMs and don’t work better for tagging.

¹¹⁶And other non-generative finite-state models based on next-state classifiers

¹¹⁷Note on the tag meanings: TO literally means “to”. MD means “modal” and refers to modal verbs such as *will*, *shall*, etc.

Constituency and CFGs. Discovering the inventory of constituents present in the language. Groups of words like *noun phrases* or *prepositional phrases* can be thought of as single units which can only be placed within certain parts of a sentence.

The most widely used formal system for modeling constituent structure in English is the **Context-Free Grammar**¹¹⁸. A CFG consists of a set of **productions** (rules), e.g.

$$\text{NP} \longrightarrow \text{Det Nominal} \quad (437)$$

$$\text{NP} \longrightarrow \text{ProperNoun} \quad (438)$$

$$\text{Nominal} \longrightarrow \text{Noun} \mid \text{Nominal Noun} \quad (439)$$

where the arrow is to be read “is composed of” or “consists of.”

The sequence of rule expansions going from left to right is called a **derivation** of the string of words, commonly represented by a **parse tree**. The formal language defined by a CFG is the set of strings that are derivable from the designated **start symbol**.

¹¹⁸Also called Phrase-Structure Grammars. Equiv formalism as Backus-Naur Form (BNF)

Vector Semantics (Ch. 15)

Table of Contents Local

Written by Brandon McKinzie

Words and Vectors. Vector models are generally based on a **co-occurrence**, an example of which is a **term-document matrix**: Each row is identified by a word, and each column a document. A given cell value is the number of times the assoc. word occurred in the assoc. document. Can also view each column as a document vector.

Information Retrieval:
task of finding document
 d , from D docs total,
that best matches a
query q .

For individual word vectors, however, it is most common to instead use a **term-term** matrix¹¹⁹, in which columns are also identified by individual words. Now, cell values are the number of times the row (target) word and the column (context) word co-occur in some context in some training corpus. The context is most commonly a window around the row/target word, meaning a cell gives the number of times the column word occurs in a window of $\pm N$ words from the row word.

- **Q:** What about the co-occurrence of a word with itself (row i , col i)?
 - **A:** It is included, yes. Source: “The size of the window ... is generally between 1 and 8 words on each side of the target word (*for a total context of 3-17 words*).”
- **Q:** Why is the size of each vector generally $|V|$ (vocab size)? Shouldn't this vary substantially with window and corpus size?
 - **A:** idk

(**TODO:** revisit end of page 5 in my pdf of this)

Pointwise Mutual Information (PMI). Motivation: raw frequencies in a co-occurrence matrix aren't that great, and words like “the” (which aren't useful and occur everywhere) can really skew things. *The best weighting or measure of association between words should tell us how much more often than chance the two words co-occur.* PMI is such a measure.

$$[\text{Mutual Information}] \quad I(X, Y) = \sum_x \sum_y P(X = x, Y = y) \text{PMI}(x, y) \quad (440)$$

$$[\text{PMI}] \quad \text{PMI}(x, y) = \ln \frac{P(x, y)}{P(x)P(y)} \quad (441)$$

which can be applied for our specific use case as $\text{PMI}(w, c) = \ln \frac{P(w, c)}{P(w)P(c)}$. The interpretation is simple: the denominator tells the joint probability of the given target word w occurring with context word c if they were independent of each other, while the numerator tells us how often we observed the two words together (assuming we compute probability by using the MLE).

¹¹⁹Also called the word-word or term-context matrix

Therefore, the ratio gives us how an estimate of how much more the target and feature co-occur than we expect by chance¹²⁰. Most people use **Positive PMI**, which is just $\max(0, \text{PMI})$. We can compute a **PPMI matrix** (to replace our co-occurrence matrix), where PPMI_{ij} gives the PPMI value of word w_i with context c_j . The authors show a few formulas which is really distracting since all we actually need is the counts $f_{ij} = \text{counts}(w_i, c_j)$, and from there we can use basic probability and Bayes rule to get the PPMI formula.

- **Q**: Explain why the following is true: very rare words tend to have very high PMI values.
 - **A**: hi
- **Q**: What is the range of α used in PPMI_α ? What is the intuition behind doing this?
 - **A**: For reference:

$$\text{PPMI}_\alpha(w, c) = \max \left(\ln \frac{P(w, c)}{P(w)P_\alpha(c)}, 0 \right) \quad (442)$$

$$P_\alpha(c) = \frac{\text{count}(c)^\alpha}{\sum_{c'} \text{count}(c')^\alpha} \quad (443)$$

Although there are better methods than PPMI for weighted co-occurrence matrices, most notably **TF-IDF**, things like tf-idf are not generally used for measuring *word similarity*. For that, PPMI and significance-testing metrics like t-test and likelihood-ratio are more common. The **t-test** statistic, like PMI, measures how much more frequent the association is than chance.

$$t = \frac{\bar{x} - \mu}{\sqrt{s^2/N}} \quad (444)$$

$$\text{t-test}(a, b) = \frac{P(a, b) - P(a)P(b)}{\sqrt{P(a)P(b)}} \quad (445)$$

where \bar{x} is the observed mean, while μ is the expected mean [under our null-hypothesis of independence].

Measuring Similarity. By far the most common similarity metric is the **cosine** of the angle between the vectors:

$$\text{cosine}(\mathbf{v}, \mathbf{w}) = \frac{\mathbf{v} \cdot \mathbf{w}}{|\mathbf{v}| |\mathbf{w}|} \quad (446)$$

Note that, since we've been defining vector elements as frequencies/PPMI values, they won't have negative elements, and thus our cosine similarities will be between 0 and 1 (not -1 and 1).

¹²⁰Computing PMI this way can be problematic for word pairs with small probability, especially if we have a small corpus. Recognize that PMI should never really be negative, but in practice this happens for such cases

Alternatives to cosine:

- **Jaccard measure:** Described as "weighted number of overlapping features, normalized", but looks like a silly hack in my opinion:

$$\text{sim}_{Jac}(\mathbf{v}, \mathbf{w}) = \frac{\sum_{i=1}^N \min(\mathbf{v}_i, \mathbf{w}_i)}{\sum_{i=1}^N \max(\mathbf{v}_i, \mathbf{w}_i)} \quad (447)$$

- **Dice measure:** Another hack. This displeases me.

$$\frac{2 \times \sum_{i=1}^N \min(\mathbf{v}_i, \mathbf{w}_i)}{\sum_{i=1}^N (\mathbf{v}_i + \mathbf{w}_i)} \quad (448)$$

- **Jensen-Shannon Divergence:** An alternative to the KL-divergence¹²¹, which represents the divergence of each distribution from the mean of the two:

$$\text{sim}_{JS}(\mathbf{v}||\mathbf{w}) = D\left(\mathbf{v} \left\| \frac{\mathbf{v} + \mathbf{w}}{2}\right.\right) + D\left(\mathbf{w} \left\| \frac{\mathbf{v} + \mathbf{w}}{2}\right.\right) \quad (450)$$

¹²¹ Idea:if two vectors, \mathbf{v} and \mathbf{w} , each express a probability distribution (their values sum to one), then they are similar to the extent that these probability distributions are similar. The basis of comparing two probability distributions P and Q is the **Kullback-Leibler** divergence or relative entropy, defined as:

$$D(P||Q) = \sum_x P(x) \log \frac{P(x)}{Q(x)} \quad (449)$$

Semantics with Dense Vectors (Ch. 16)

Table of Contents Local

Written by Brandon McKinzie

Overview. This chapter introduces three methods for generating short, dense vectors: (1) dimensionality reduction like SVD, (2) neural networks like skip-gram or CBOW, and (3) Brown clustering.

Dense Vectors via SVD. Method for finding more important dimensions of a dataset, “important” defined as dimensions wherein the data most varies. First applied (for language) for generating embeddings from term-document matrices in a model called **Latent Semantic Analysis** (LSA). LSA is just SVD on a $|V| \times c$ term-document matrix \mathbf{X} , factorized into $\mathbf{W}\mathbf{\Sigma}\mathbf{C}^T$. By using only the top $k < m$ dimensions of these three matrices, the product becomes a least-squares approx. to the original \mathbf{X} . It also gives us the reduced $|V| \times k$ matrix \mathbf{W}_k , where each row (word) is a k -dimensional vector (embedding). Voilà, we have our dense vectors!

$$\begin{aligned}\mathbf{W} &\in \mathbb{R}^{|V| \times m} \\ \mathbf{\Sigma} &\in \mathbb{R}^{m \times m} \\ \mathbf{C}^T &\in \mathbb{R}^{m \times c}\end{aligned}$$

Note that LSA implementations typically use a particular weighting of each cell in the term-document matrix called the **local** and **global** weights.

$$[\text{local}] \quad \log f(i, j) + 1 \quad (451)$$

$$[\text{global}] \quad 1 + \frac{\sum_j p(i, j) \log p(i, j)}{\log D} \quad (452)$$

$f(i, j)$ is the raw frequency of word i in context j . D is number of docs.

For the case of a word-word matrix, it is common to use PPMI weighting.

Skip-Gram and CBOW. Neural models learn an embedding by starting with a random vector and then iteratively shifting a word’s embeddings to be more like the embeddings of neighboring words, and less like the embeddings of words that don’t occur nearby¹²² Word2vec, for example, learns embeddings by training to predict neighboring words¹²³.

- **Skip-Gram:** Learns two embeddings for each word w : the **word embedding** v (within matrix \mathbf{W}) and **context embedding** c (within matrix \mathbf{C}). Visually:

$$\mathbf{W} = \begin{pmatrix} v_0^T \\ v_1^T \\ \vdots \\ v_{|V|}^T \end{pmatrix} \quad \mathbf{C} = \begin{pmatrix} c_0 & c_1 & \cdots & c_{|V|} \end{pmatrix} \quad (453)$$

¹²²Why? Why is this a sensible assumption? I see no reason a priori why it ought to be true.

¹²³Note that the prediction task is not the goal – it just happens to result in good word embeddings. Hacky.

For a context window of $L = 2$, and at a given word $\mathbf{v}^{(t)}$ inside the corpus¹²⁴, our goal is to predict the context [words] denoted as $[\mathbf{c}^{(t-2)}, \mathbf{c}^{(t-1)}, \mathbf{c}^{(t+1)}, \mathbf{c}^{(t+2)}]$.

- Example: Consider one of the context words, say $\mathbf{c}^{(t+1)} \triangleq \mathbf{c}_k$, where we also assume it's the k th word in our vocab. Also assume that our target word $\mathbf{v}^{(t)} \triangleq \mathbf{v}_j$ is the j th word in our vocab.
- Our task is to compute $\Pr[\mathbf{c}_k | \mathbf{v}_j]$. We do this with a softmax:

$$\Pr[\mathbf{c}_k | \mathbf{v}_j] = \frac{e^{\mathbf{c}_k^T \mathbf{v}_j}}{\sum_{i \in |V|} e^{\mathbf{c}_i^T \mathbf{v}_j}} \quad (454)$$

- **CBOV**: Continuous bag of words. Basically the mirror-image of skip-gram. Goal is to predict current word $\mathbf{v}^{(t)}$ from the context window of $2L$ words $[\mathbf{c}^{(t-2)}, \mathbf{c}^{(t-1)}, \mathbf{c}^{(t+1)}, \mathbf{c}^{(t+2)}]$.

As usual, the denominator of the softmax is computationally expensive, and usually we approximate it with **negative sampling**.

In the training phase, the algorithm walks through the corpus, at each target word choosing the surrounding context words as positive examples, and for each positive example also choosing k noise samples or negative samples: non-neighbor words. The goal will be to move the embeddings toward the neighbor words and away from the noise words.

Example: Suppose we come along the following window (in “[]”) ($L=2$) in our corpus:

lemon, a [tablespoon of apricot preserves or] jam

Ultimately, we want dot products, $\mathbf{c}_i \cdot \text{vector}(\text{“apricot”})$, to be high for all four of the context words \mathbf{c}_i . We do negative sampling by sampling k random noise words according to their [unigram] frequency. So here, for e.g. $k = 2$, this would amount to 8 noise words, 2 for each context word. We want the dot products between “apricot” and these noise words to be low. For a given single context-word pair (w, c) , our training objective is to maximize:

$$\log \sigma(c \cdot w) + \sum_{i=1}^k \mathbb{E}_{w_i \sim p(w)} [\log \sigma(-w_i \cdot w)] \quad (455)$$

In practice, common to use $p^{3/4}(w)$ instead of $p(w)$

Again, the above is for a single context-target word-pair and, accordingly, the summation is only over $k = 2$ (for our example). Don't try to split the expectation into a summation or anything – just view it as an expected value. To iteratively shift parameters, we use an optimizer like SGD.

¹²⁴Note that, technically, the position t of $\mathbf{v}^{(t)}$ is irrelevant for our computation; we are predicting those words based on which word $\mathbf{v}^{(t)}$ is in the vocabulary, not it's position in the corpus.

The actual model architecture is a typical neural net, progressing as follows:

$$\text{“apricot”} \rightarrow \mathbf{w}^{\text{one-hot}} = [0 \ 0 \ \dots \ 1 \ \dots \ 0] \quad (456)$$

$$\rightarrow \mathbf{h} = \mathbf{W}^T \mathbf{w}^{\text{one-hot}} \quad (457)$$

$$\rightarrow \mathbf{o} = \mathbf{C}^T \mathbf{h} = [c_0^T \mathbf{h}, c_1^T \mathbf{h}, \dots c_{|V|}^T \mathbf{h}]^T \quad (458)$$

$$\rightarrow \mathbf{y} = \text{softmax}(\mathbf{o}) = [\text{Pr}[c_0|\mathbf{h}], \text{Pr}[c_2|\mathbf{h}], \dots \text{Pr}[c_{|V|}|\mathbf{h}]]^T \quad (459)$$

$$(460)$$

Brown Clustering. An agglomerative clustering algorithm for deriving vector representations of words by clustering words based on their associations with the preceding or following words. Makes use of the **class-based language model** (CBLM), wherein each word w belongs to some class $c \in C$ via the probability $P(w | c)$. CBLMs define

$$P(w_i | w_{i-1}) = P(c_i | c_{i-1})P(w_i | c_i) \quad (461)$$

$$P(\text{corpus} | C) = \prod_{i=1}^n P(w_i | w_{i-1}) \quad (462)$$

A naive and extremely inefficient version of Brown clustering, a hierarchical clustering algorithm, is as follows:

1. Each word is initially assigned to its own cluster.
2. For each cluster pair $(c_i, c_{j \neq i})$, compute the value of eq 462 that would result from merging c_i and c_j into a single cluster. The pair whose merger results in the smallest decrease in eq 462 is merged.
3. Clustering proceeds until all words are in one big cluster.

This process builds a binary tree from the bottom-up, and the binary string corresp. to a word’s traversal from leaf-to-root is its representation.

Information Extraction (Ch. 21 3rd Ed)

Table of Contents Local

Written by Brandon McKinzie

Overview. The first step in most IE tasks is **named entity recognition** (NER). Next we can do **relation extraction**: finding and classifying semantic relations among the entities, e.g. “spouse-of.” **Event extraction** is finding the events in which the entities participate, and **event coreference** for figuring out which event mentions actually refer to the same event. It’s also common to extract dates/times (temporal expression) and perform **temporal expression normalization** to map them onto specific calendar dates. Finally, we can do **template filling**: finding recurring/stereotypical situations in documents and filling the template slots with appropriate material.

Named Entity Recognition. Standard algorithm is word-by-word sequence labeling task by a MEMM or CRF, trained to label tokens with tags indicating the presence of particular kinds of NEs. It is common to label with **BIO tagging**, for beginning, inside, and outside of entities. If we have n unique entity types, then we’d have $2n + 1$ BIO tags¹²⁵. A helpful illustration is shown below:




Figure 21.7 Named entity recognition as sequence labeling. The features available to the classifier during training and classification are those in the boxed area.

Here we see a classifier determining the label for *Corp.* with a context window of size 2 and various features shown in the boxed region. For evaluation of NER, we typically use the familiar **recall**, **precision**, and **F1 measure**.

¹²⁵ $2n$ for B-<NE> and I-<NE>, with +1 for the blanket *O* tag (not any of our NEs)

Relation Extraction. The four main algorithm classes used are (1) hand-written patterns, (2) supervised ML, (3) semi-supervised, and (4) unsupervised. Terminology:

- **Infobox:** structured tables associated with certain articles/topics/etc. For example, the Wikipedia infobox for Stanford includes structured facts like `state = 'California'`.
- **Resource Description Framework (RDF):** a metalanguage of RDF triples, tuples of (entity, relation, entity), called a subject-predicate-object expression. For example: (Golden Gate Park, location, San Francisco).
- **hypernym:** the “is-a” relation.
- **hyponym:** the “kind-of” relation. *Gelidium is a kind of red algae.*

Overview of the four algorithm classes:

1. **Patterns.** Consider a sentence that has the following form:

NP_0 such as $NP_1\{, NP_2, \dots, (and|or)NP_i\}, i \geq 1$

also known as a *lexico-syntactic pattern*, which implies $\forall NP_i, i \geq 1, \text{hyponym}(NP_i, NP_0)$ ¹²⁶. Patterns typically have high precision but low-recall.

2. **Supervised.** The general approach for finding relations in a given sequence of words is the following:
 - (a) Find all pairs of named entities in the sequence (typically a single sentence).
 - (b) For each pair, use a trained binary classifier to predict whether or not the entities in the pair are indeed related.
 - (c) If related, use a classifier trained to predict the relation given the entity-pair.

As with NER, **the most important step** in this process is to identify useful surface features that will be useful for relation classification, including word features, NER features, syntactic paths (chunk seqs, constituent paths, dependency-tree paths), and more.

3. **Semi-supervised** via bootstrapping. Suppose we have a few high-precision **seed patterns** (or seed tuples)¹²⁷. **Bootstrapping** proceeds by taking the entities in the seed pair, and then finding sentences (on the web, or whatever dataset we are using) that contain both entities. From all such sentences, we extract and generalize the context around the entities to learn new patterns.

```

function BOOTSTRAP(Relation R) returns new relation tuples

  tuples ← Gather a set of seed tuples that have relation R
  iterate
    sentences ← find sentences that contain entities in seeds
    patterns ← generalize the context between and around entities in sentences
    newpairs ← use patterns to grep for more tuples
    newpairs ← newpairs with high confidence
    tuples ← tuples + newpairs
  return tuples

```

Figure 21.14 Bootstrapping from seed entity pairs to learn relations.

¹²⁶Here, $\text{hyponym}(A, B)$ means “A is a kind-of (hyponym) of B.”

¹²⁷seed tuples are tuples of the general form (M1, M2) where M1 and M2 are each specific named entities we know have the relation of interest R.

4. **Unsupervised.** The Re Verb system extracts a relation from a sentence s in 4 steps:

1. Run a part-of-speech tagger and entity chunker over s
2. For each verb in s , find the longest sequence of words w that start with a verb and satisfy syntactic and lexical constraints, merging adjacent matches.
3. For each phrase w , find the nearest noun phrase x to the left which is not a relative pronoun, wh-word or existential “there”. Find the nearest noun phrase y to the right.
4. Assign confidence c to the relation $r = (x, w, y)$ using a confidence classifier and return it.

Event Extraction. An event mention is any expression denoting an event or state that can be assigned to a particular point, or interval, in time. Note that this is quite different than the colloquial usage of the word “event,” you should think of the two as distinct. Here, most event mentions correspond to verbs, and most verbs introduce events. Event extraction is typically modeled via ML, detecting events via sequence models with BIO tagging, and assigning event classes/attributes with multi-class classifiers.

Template Filling. The task is creation of one template for each event in the input documents, with the slots filled with text from the document. For example, an event could be “Fare-Raise Attempt” with corresponding template (slots to be filled) “(<Lead Airline>, <Amount>, <Effective Date>, <Follower>)”. This is generally modeled by training two separate supervised systems:

1. **Template recognition.** Trained to determine if template T is present in sentence S . Here, “present” means there is a sequence within the sentence that could be used to fill a slot within template T .
2. **Role-filler extraction.** Trained to detect each role (slot-name), e.g. “Lead Airline”.

PROBABILISTIC GRAPHICAL MODELS

CONTENTS

7.1	Foundations (Ch. 2)	187
7.1.1	Appendix	188
7.2	The Bayesian Network Representation (Ch. 3)	190
7.3	Undirected Graphical Models (Ch. 4)	194
7.3.1	Exercises	200
7.4	Local Probabilistic Models (Ch. 5)	201
7.5	Template-Based Representations (Ch. 6)	202
7.6	Variable Elimination (Ch. 9)	204
7.7	Clique Trees (Ch. 10)	208
7.8	Inference as Optimization (Ch. 11)	212

Foundations (Ch. 2)

Table of Contents Local

Written by Brandon McKinzie

(Brief summary of the book's notation, for ease of reference.)

Graphs. Authors denote directed graphs as \mathcal{G} and undirected graphs as \mathcal{H} .

- **Induced Subgraph.** Let $\mathcal{K} = (\mathcal{X}, \mathcal{E})$ and $\mathbf{X} \subset \mathcal{X}$. Define the *induced subgraph* $\mathcal{K}[\mathbf{X}]$ to be the graph $(\mathbf{X}, \mathcal{E}')$ where \mathcal{E}' are all edges $X \rightleftharpoons Y$ such that $X, Y \in \mathbf{X}$.
- **Complete Subgraph.** A subgraph over \mathbf{X} is *complete* if every two nodes in \mathbf{X} are connected by some edge. The set \mathbf{X} is often called a *clique*; we say that a clique \mathbf{X} is *maximal* if for any superset of nodes $\mathbf{Y} \supset \mathbf{X}$, \mathbf{Y} is not a clique.
- **Upward Closure.** We say that a subset of nodes $\mathbf{X} \in \mathcal{X}$ is *upwardly closed* in \mathcal{K} if $\forall X \in \mathbf{X}$, we have that $\text{Boundary}_X \subset \mathbf{X}$ ¹²⁸. We define the *upward closure* of \mathbf{X} to be the minimally upwardly closed subset \mathbf{Y} that contains \mathbf{X} . We define the *upwardly closed subgraph* of \mathbf{X} , denoted $\mathcal{K}^+[\mathbf{X}]$, to be the induced subgraph over \mathbf{Y} , $\mathcal{K}[\mathbf{Y}]$.

Paths and Trails. Definitions for longer-range connections in graphs. We use the notation $X_i \rightleftharpoons X_{j+1}$ to denote that X_i and X_j are connected via some edge, whether directed (in any direction) or undirected.

- **Trail/Path.** We say that X_1, \dots, X_k form a *trail* in the graph $\mathcal{K} = (\mathcal{X}, \mathcal{E})$ if $\forall i = 1, \dots, k-1$, we have that $X_i \rightleftharpoons X_{i+1}$. A *path* makes an additional restriction: either $X_i \rightarrow X_{i+1}$ or $X_i - X_{i+1}$.
- **Connected Graph.** A graph is *connected* if $\forall X_i, X_j$ there is a trail between X_i and X_j .
- **Cycle.** A *cycle* in \mathcal{K} is a directed path X_1, \dots, X_k where $X_1 = X_k$.
- **Loop.** A *loop* in \mathcal{K} is a trail where $X_1 = X_k$. A graph is *singly connected* if it contains no loops. A node in a singly connected graph is called a *leaf* if it has exactly one adjacent node.
- **Polytree/Forest.** A singly connected graph is also called a *polytree*. A singly connected undirected graph is called a *forest*; if a forest is also connected, it is called a *tree*.
 - A directed graph is a forest if each node has at most one parent. A directed forest is a tree if it is also connected.
- **Chordal Graph.** Let $X_1 - X_2 - \dots - X_k - X_1$ be a loop in the graph. A *chord* in the loop is an edge connecting X_i and X_j for two nonconsecutive nodes X_i, X_j . An undirected graph \mathcal{H} is said to be *chordal* if any loop $X_1 - X_2 - \dots - X_k - X_1$ for $k \geq 4$ has a chord.

¹²⁸ $\text{Boundary}_X \triangleq \text{Pa}_X \cup \text{Nb}_X$. For DAGs, this is simply X 's parents, and for undirected graphs X 's neighbors.

Figured this would be a good place to put some of the definitions in the Appendix, too.

Information Theory (A.1).

- **Entropy.** Let $P(X)$ be a distribution over a random variable X . The *entropy* of X is defined as.

$$H_P(X) = \mathbb{E}_P \left[\lg \frac{1}{P(X)} \right] = \sum_x P(X = x) \lg \frac{1}{P(X = x)} \quad (463)$$

We treat $0 \log \frac{1}{0} = 0$

$$0 \leq H_P(X) \leq \lg |Val(X)| \quad (464)$$

$H_P(X)$ is a lower bound for the expected number of bits required to encode instances sampled from $P(X)$. Another interpretation is that the entropy is a measure of our uncertainty about the value of X .

- **Conditional Entropy.** The *conditional entropy* of X given Y is

$$H_P(X | Y) = H_P(X, Y) - H_P(Y) = \mathbb{E}_P \left[\lg \frac{1}{P(X | Y)} \right] \quad (465)$$

$$H_P(X | Y) \leq H_P(X) \quad (466)$$

which captures the additional cost (in bits) of encoding X when we're already encoding Y .

- **Mutual Information.** The *mutual information* between X and Y is

$$I_P(X; Y) = H_P(X) - H_P(X | Y) = \mathbb{E}_P \left[\lg \frac{P(X | Y)}{P(X)} \right] \quad (467)$$

which captures how many bits we save (on average) in the encoding of X if we know the value Y .

- **Distance Metric.** A *distance metric* is any distance measure d evaluating the distance between two distributions that satisfies all of the following properties:

- **Positivity:** $d(P, Q) \geq 0$ and $d(P, Q) = 0$ if and only if $P = Q$.
- **Symmetry:** $d(P, Q) = d(Q, P)$.
- **Triangle inequality:** For any three distributions P, Q, R , we have that

$$d(P, R) \leq d(P, Q) + d(Q, R) \quad (468)$$

- **Kullback-Liebler Divergence.** Let P and Q be two distributions over random variables X_1, \dots, X_n . The relative entropy, or KL-divergence, of P and Q is

$$D(P \| Q) = \mathbb{E}_P \left[\lg \frac{P(X_1, \dots, X_n)}{Q(X_1, \dots, X_n)} \right] \quad (469)$$

Note that this only satisfies the positivity property, and is thus not a true distance metric.

Algorithms and Algorithmic Complexity (A.3).

- **Decision Problems.** A decision problem Π is a task that accepts an input (instance) ω and decides whether it satisfies a certain condition or not. The **SAT** problem accepts a formula in propositional logic and decides whether there is an assignment to the variables in the formula such that it evaluates to true. **3-SAT** restricts this to accepting only formulas in conjunctive normal form (CNF), and further restricted s.t. each clause contains at most 3 literals.
- **P.** A decision problem is in the class \mathcal{P} if there exists a deterministic algorithm that takes an instance ω and determines whether $\omega \in \mathcal{L}_\Pi$ (the set of instances for which a correct algorithm must return true), in polynomial time in the size of the input ω .
- **NP.** A *non-deterministic algorithm* takes the general form: (1) nondeterministically guess some assignment γ to the variables of ω , (2) deterministically verify whether γ satisfies the condition of the problem. The algorithm will repeat these steps until it produces a γ that satisfies the problem. A decision problem Π is in the class \mathcal{NP} if there exists a nondeterministic algorithm that accepts ω if and only if $\omega \in \mathcal{L}_\Pi$, and if the verification stage can be executed in polynomial time in the length of ω .
- **NP-hard.** Π is \mathcal{NP} -hard if for every DP $\Pi' \in \mathcal{NP}$, there is a polynomial-time transformation of inputs such that an input for Π' belongs to $\mathcal{L}_{\Pi'}$ if and only if the transformed instance belongs to \mathcal{L}_Π . Note that \mathcal{NP} -hard is a superset of \mathcal{NP} .
- **NP-complete.** A problem Π is said to be \mathcal{NP} -complete if it is both \mathcal{NP} -hard and in \mathcal{NP} .

The SAT problem is \mathcal{NP} -hard.

The Bayesian Network Representation (Ch. 3)

Table of Contents Local

Written by Brandon McKinzie

Koller and Friedman (2009). The Bayesian Network Representation.

Probabilistic Graphical Models: Principles and Techniques.

Goal: represent a joint distribution P over some set of variables $\mathcal{X} = \{X_1, \dots, X_n\}$. Consider the case where each X_i is binary-valued. A *single* joint distribution requires access to the probability for each of the 2^n possible assignments for \mathcal{X} . The set of all such possible joint distributions,

$$\{(p_1, \dots, p_{2^n}) \in \mathbb{R}^{2^n} : \sum_{i=1}^{2^n} p_i = 1\} \quad (470)$$


Understanding the exponential blowup.

is a $2^n - 1$ dimensional subspace of \mathbb{R}^{2^n} . Note that each p_i represents the probability for a unique instantiation of \mathcal{X} . Furthermore, in the general case, knowing p_i tells you nearly nothing about $p_{j \neq i}$ – i.e. you require an instantiation of $2^n - 1$ independent parameters to specify a given joint distribution.

But it would be foolish to parameterize any joint distribution in this way, since we can often take advantage of independencies. Consider the case where each X_i gives the outcome (H or T) of coin i being tossed. Then our distribution satisfies $(\mathbf{X} \perp \mathbf{Y})$ for any disjoint subsets of the variables \mathbf{X} and \mathbf{Y} . Let θ_i denote the probability that coin i lands heads. The key observation is that you only need each of the n θ_i to specify a unique joint distribution over \mathcal{X} , reducing the $2^n - 1$ dimensional subspace to an n dimensional manifold in \mathbb{R}^{2^n} ¹²⁹.

Taking advantage of independencies.

The Naive Bayes Model. Say we want to determine the intelligence of an applicant based on their grade G in some course, and their score S on the SAT. A naive bayes model can be illustrated as below.




This encodes our general assumption that¹³⁰ $P \models (S \perp G \mid I)$.

¹²⁹**TODO:** Formally define this manifold using the same notation as in 470. Edit: Not sure how to actually write it out, but intuitively it's because each of the 2^n p_i values, when going from the general case to the case of i.i.d, become *functions* of the n θ_i values. Whereas before they were independent free parameters.

¹³⁰We say that an event α is independent of event β in P with the notation $P \models (\alpha \perp \beta)$. (Def 2.2, pg 23 of book)

In general, a naive bayes model assumes that instances fall into one of a number of mutually exclusive and exhaustive *classes*, defined as the set of values that the top variable in the graph can take on¹³¹. The model also includes some number of *features* X_1, \dots, X_k , whose values are typically observed. The **naive Bayes assumption** is that the features are conditionally independent given the instance's class.

Bayesian Networks. A Bayesian network \mathcal{B} is defined by a network structure together with its set of CPDs. **Causal reasoning** (or prediction) refers to computing the downstream effects of various factors (such as intelligence). **Evidential reasoning** (or explanation) is the reverse case, where we reason from effects to causes. Finally, **intercausal reasoning** (or explaining away) is when different causes of the same effect can interact. For our student example, we could be trying to determine $\Pr [I \mid G]$, the intelligence of an a student given his/her grade in a class. In addition to intelligence being a cause for the grade, we could have another causal variable D for the difficulty of the class:



An example of intercausal reasoning would be observing D , so that we now we want $\Pr [I \mid G, d]$; the difficulty of the course can help *explain away* good/bad grades, thus changing our value for the probability of intelligence based on the grade alone.

A **Bayesian network structure** \mathcal{G} is a DAG whose nodes represent RVs X_1, \dots, X_n . Let $Pa_{X_i}^{\mathcal{G}}$ denote the parents of X_i in \mathcal{G} , and NonDesc_{X_i} denotes the variables that are NOT descendants of X_i . Then \mathcal{G} encodes the following set of **local independencies**,

$$\mathcal{I}_{\ell}(\mathcal{G}) \triangleq \{\forall X_i : (X_i \perp \text{NonDesc}_{X_i} \mid Pa_{X_i}^{\mathcal{G}})\} \quad (471)$$

¹³¹For the intelligence example, the classes are high intelligence and low intelligence.

Graphs and Distributions. Here we see that a distribution P satisfies the local independencies associated with a graph \mathcal{G} iff P is representable as a set of CPDs associated with \mathcal{G} .

- **Local independencies.** Let P be a distribution over \mathcal{X} . We define $\mathcal{I}(P)$ to be the set of **independence assertions** of the form $(\mathbf{X} \perp \mathbf{Y} \mid \mathbf{Z})$ that hold in P . The statement “ P satisfies the local independencies associated with \mathcal{G} ” can thus be succinctly written:

$$\mathcal{I}_\ell(\mathcal{G}) \subseteq \mathcal{I}(P) \quad (472)$$

and we’d say that \mathcal{G} is an **I-map** (independency map) for P .


- **I-maps.** More generally, let \mathcal{K} be *any* graph object associated with a set of independencies $\mathcal{I}(\mathcal{K})$. We say that \mathcal{K} is an **I-map** for a set of independencies \mathcal{I} if $\mathcal{I}(\mathcal{K}) \subseteq \mathcal{I}$. Note that the complete graph (every two nodes connected) is an I-map for any distribution.
- **Factorization.** Let \mathcal{G} be a BN graph over X_1, \dots, X_n . We say that a distribution P over the same space **factorizes** according to \mathcal{G} if P can be expressed as

$$P(X_1, \dots, X_n) = \prod_{i=1}^n P(X_i \mid Pa_{X_i}^{\mathcal{G}}) \quad (473)$$

For BNs, the following equivalence holds:

$$\mathcal{G} \text{ is an I-map for } P \iff P \text{ factorizes according to } \mathcal{G}$$

D-separation. We want to understand when we can *guarantee* that an independence $(\mathbf{X} \perp \mathbf{Y} \mid \mathbf{Z})$ holds in a distribution associated with a BN structure \mathcal{G} . Consider a 3-node network consisting of X, Y, and Z, where X and Y are not directly connected. There are four such cases, which I’ve drawn below.



The fourth trail is called a *v-structure*.

From left-to-right: indirect causal effect, indirect evidential effect, common cause, common effect. The first 3 satisfy $(X \perp Y \mid Z)$, but the 4th does not. Another way of saying this is that the first 3 trails are **active**¹³² IFF Z is not observed, while the 4th trail is active IFF Z (or a descendent of Z) is observed.

¹³²When influence can flow from X to Y via Z, we say that the trail $X \rightleftharpoons Y \rightleftharpoons Z$ is *active*.

General case:

Let \mathcal{G} be a BN structure, and $X \rightleftharpoons \cdots \rightleftharpoons X_n$ a trail in \mathcal{G} . Let \mathbf{Z} be a subset of observed variables. The trail is **active** given \mathbf{Z} if

- Any v-structure $X_{i-1} \rightarrow X_i \leftarrow X_{i+1}$ has X_i or one of its descendants in \mathbf{Z} .
- No other node along the trail is in \mathbf{Z} .

D-separation: (Directed separation)

Let $\mathbf{X}, \mathbf{Y}, \mathbf{Z}$ be three sets of nodes in \mathcal{G} . We say that \mathbf{X} and \mathbf{Y} are **d-separated** in \mathcal{G} , denoted $d\text{-sep}_{\mathcal{G}}(\mathbf{X}; \mathbf{Y} \mid \mathbf{Z})$, if there is no active trail between any node $X \in \mathbf{X}$ and $Y \in \mathbf{Y}$ given \mathbf{Z} . We use $\mathcal{I}(\mathcal{G})$ to denote the set of independencies that correspond to d-separation,

$$\mathcal{I}(\mathcal{G}) = \{(\mathbf{X} \perp \mathbf{Y} \mid \mathbf{Z}) : d\text{-sep}_{\mathcal{G}}(\mathbf{X}; \mathbf{Y} \mid \mathbf{Z})\} \quad (474)$$

also called the set of **global Markov independencies**.

Soundness and Completeness of d-separation as a method for determining independence.

- **Soundness** (Thm 3.3). If a distribution P factorizes according to \mathcal{G} , then $\mathcal{I}(\mathcal{G}) \subseteq \mathcal{I}(P)$.
- **Completeness**. For any distribution P that factorizes over \mathcal{G} , we have that P is **faithful**¹³³ to \mathcal{G} :

$$X \perp Y \mid \mathbf{Z} \in \mathcal{I}(P) \implies d\text{-sep}_{\mathcal{G}}(X; Y \mid \mathbf{Z}) \quad (475)$$

To see the detailed algorithm for finding nodes reachable from X given \mathbf{Z} via active trails, see Algorithm 3.1 on pg. 75 of the book.

¹³³ P is faithful to \mathcal{G} if, whenever $X \perp Y \mid \mathbf{Z} \in \mathcal{I}(P)$, then $d\text{-sep}_{\mathcal{G}}(X; Y \mid \mathbf{Z})$.

Undirected Graphical Models (Ch. 4)


Table of Contents Local

Written by Brandon McKinzie

Koller and Friedman (2009). Undirected Graphical Models.

Probabilistic Graphical Models: Principles and Techniques.

The Misconception Example. Consider a scenario where we have four students who get together in pairs to work on homework for a class. Only the following pairs meet: (Alice, Bob), (Bob, Charles), (Charles, Debbie), (Debbie, Alice). The professor misspoke in class, giving rise to a possible misconception among the students. We have four binary random variables, $\{A, B, C, D\}$, representing whether the student has the misconception (1) or not (0)¹³⁴. Intuitively, we want to model a distribution that satisfies $(A \perp C \mid \{B, D\})$, and $(B \perp D \mid \{A, C\})$, but no other independencies¹³⁵. Note that the interactions between variables seem symmetric here – students influence each other (out of the ones they have a pair with).



The nodes in the graph of a **Markov network** represent the variables, and the edges correspond to a notion of direct probabilistic interaction between the neighboring variables – an interaction that is not mediated by any other variable in the network. So, how should we parameterize our network? We want to capture the **affinities** between the related variables (e.g. Alice and Bob are more likely to agree than disagree).

Let \mathbf{D} be a set of random variables. We define a **factor** ϕ to be a function from $\text{Val}(\mathbf{D})$ to \mathbb{R} . A factor is nonnegative if all its entries are nonnegative. \mathbf{D} is called the **scope** of the factor, denoted $\text{Scope}[\phi]$.

We restrict our attention to nonnegative factors.

The factors need not be normalized. Therefore, to interpret probabilities over factors, we must

¹³⁴A student might not have the misconception if e.g. they went home and figured out the problem via reading the textbook instead.

¹³⁵These independencies cannot be naturally captured in a Bayesian (i.e. directed) network.

normalize it with what we'll call the **partition function**, Z :

$$\Pr [a, b, c, d] = \frac{1}{Z} \phi_1(a, b) \cdot \phi_2(b, c) \cdot \phi_3(c, d) \cdot \phi_4(d, a) \quad (476)$$

$$Z = \sum_{a,b,c,d} \phi_1(a, b) \cdot \phi_2(b, c) \cdot \phi_3(c, d) \cdot \phi_4(d, a) \quad (477)$$

Parameterization. Associating the graph structure with a set of parameters. We parameterize undirected graphs by associating a set of factors with it. First, we introduce the definition of **factor product**:

Let \mathbf{X} , \mathbf{Y} , and \mathbf{Z} be three disjoint sets of variables, and let $\phi_1(\mathbf{X}, \mathbf{Y})$ and $\phi_2(\mathbf{Y}, \mathbf{Z})$ be two factors. We define the factor product $\phi_1 \times \phi_2$ to be a factor $\psi : \text{Val}(\mathbf{X}, \mathbf{Y}, \mathbf{Z}) \mapsto \mathbb{R}$ as follows:

$$\psi(\mathbf{X}, \mathbf{Y}, \mathbf{Z}) = \phi_1(\mathbf{X}, \mathbf{Y}) \cdot \phi_2(\mathbf{Y}, \mathbf{Z}) \quad (478)$$

We use this to define an undirected parameterization of a distribution:

*A distribution P_{Φ} is a **Gibbs distribution** parameterized by a set of factors $\Phi = \{\phi_1(\mathbf{D}_1), \dots, \phi_K(\mathbf{D}_K)\}$ if it is defined as follows:*

$$P_{\Phi} = \frac{1}{Z} \tilde{P}_{\Phi}(X_1, \dots, X_n) \quad (479)$$

$$\tilde{P}_{\Phi}(X_1, \dots, X_n) = \phi_1(\mathbf{D}_1) \times \phi_2(\mathbf{D}_2) \times \dots \times \phi_K(\mathbf{D}_K) \quad (480)$$

where the authors (Koller and Friedman) have made a point to emphasize: **A factor is only one contribution to the overall joint distribution. The distribution as a whole has to take into consideration the contributions from all of the factors involved.** Now, we relate the parameterization of a Gibbs distribution to a graph structure.

*A distribution P_{Φ} with $\Phi = \{\phi_1(\mathbf{D}_1), \dots, \phi_K(\mathbf{D}_K)\}$ **factorizes** over a Markov network \mathcal{H} if each $\mathbf{D}_k (k = 1, \dots, K)$ is a complete subgraph¹³⁶ of \mathcal{H} .*

The factors that parameterize a Markov network are often called **clique potentials**. Although it can be used without loss of generality, the parameterization using maximal clique potentials generally obscures structure that is present in the original set of factors. Below are some useful definitions that we will use often.

¹³⁶ A subgraph is complete if every two nodes in the subgraph are connected by some edge. The set of nodes in such a subgraph is often called a **clique**. A clique \mathbf{X} is maximal if for any superset of nodes $\mathbf{Y} \supset \mathbf{X}$, \mathbf{Y} is *not* a clique.

Factor reduction:

Let $\phi(\mathbf{Y})$ be a factor, and $\mathbf{U} = \mathbf{u}$ an assignment for $\mathbf{U} \subseteq \mathbf{Y}$. Define the **reduction** of the factor ϕ to the context $\mathbf{U} = \mathbf{u}$, denoted $\phi[\mathbf{u}]$, to be a factor over the scope $\mathbf{Y}' = \mathbf{Y} - \mathbf{U}$, such that

$$\phi[\mathbf{u}](\mathbf{y}') = \phi(\mathbf{y}', \mathbf{u}) \quad (481)$$

For $\mathbf{U} \not\subseteq \mathbf{Y}$, define $\phi[\mathbf{u}]$ only for the assignments in \mathbf{u} to the variables in $\mathbf{U}' = \mathbf{U} \cap \mathbf{Y}$.

Reduced Gibbs distribution:

Let $P_{\Phi}(\mathbf{X})$ be a Gibbs distribution parameterized by $\Phi = \{\phi_1, \dots, \phi_K\}$ and let \mathbf{u} be a context. The **reduced Gibbs distribution** $P_{\Phi}[\mathbf{u}]$ is the Gibbs distribution defined by the set of factors $\Phi[\mathbf{u}] = \{\phi_1[\mathbf{u}], \dots, \phi_K[\mathbf{u}]\}$. More formally:

$$P_{\Phi}[\mathbf{u}] = P_{\Phi}(\mathbf{W} \mid \mathbf{u}) \quad \text{where} \quad \mathbf{W} = \mathbf{X} - \mathbf{U} \quad (482)$$

Reduced Markov Network:

Let \mathcal{H} be a Markov network over \mathbf{X} and $\mathbf{U} = \mathbf{u}$ a context. The **reduced Markov network** $\mathcal{H}[\mathbf{u}]$ is a Markov network over the nodes $\mathbf{W} = \mathbf{X} - \mathbf{U}$, where we have an edge $X - Y$ if there's an edge $X - Y$ in \mathcal{H} .

Note that if a Gibbs distribution $P_{\Phi}(\mathbf{X})$ factorizes over \mathcal{H} , then $P_{\Phi}[\mathbf{u}]$ factorizes over $\mathcal{H}[\mathbf{u}]$.

Markov Network Independencies. A formal presentation of the undirected graph as a representation of independence assertions.

- **Active Path.** Let \mathcal{H} be a Markov network structure, and let $X_1 - \dots - X_k$ be a path in \mathcal{H} . Let $\mathbf{Z} \subseteq \mathcal{X}$ be a set of *observed variables*. The path is *active* given \mathbf{Z} if none of the X_i 's is in \mathbf{Z} .
- **Separation.** A set of nodes \mathbf{Z} *separates* \mathbf{X} and \mathbf{Y} in \mathcal{H} , denoted $\text{sep}_{\mathcal{H}}(\mathbf{X}; \mathbf{Y} \mid \mathbf{Z})$, if there is no active path between any node $X \in \mathbf{X}$ and $Y \in \mathbf{Y}$ given \mathbf{Z} .
- **Global Independencies.** The *global independencies* associated with \mathcal{H} are defined as

$$\mathcal{I}(\mathcal{H}) = \{(\mathbf{X} \perp \mathbf{Y} \mid \mathbf{Z}) : \text{sep}_{\mathcal{H}}(\mathbf{X}; \mathbf{Y} \mid \mathbf{Z})\} \quad (483)$$

This is the *separation criterion*. Note that the definition of separation is monotonic in \mathbf{Z} : if it holds for \mathbf{Z} , then it holds for any $\mathbf{Z}' \supset \mathbf{Z}$ as well¹³⁷.

¹³⁷Which means that Markov networks are fundamentally incapable of representing nonmonotonic independence relations!

- **Soundness** of the separation criterion for detecting independence properties in distributions over \mathcal{H} . In other words, we want to prove that

$$(P \text{ factorizes over } \mathcal{H}) \implies [\text{sep}_{\mathcal{H}}(\mathbf{X}; \mathbf{Y} \mid \mathbf{Z}) \implies P \models (\mathbf{X} \perp \mathbf{Y} \mid \mathbf{Z})]$$

where the portion in brackets can equivalently be said as “ \mathcal{H} is an I-map for P ”.

Proof:

- Consider the case where $\mathbf{X} \cup \mathbf{Y} \cup \mathbf{Z} = \mathcal{X}$.
- Then, any clique in \mathcal{X} is fully contained in either $\mathbf{X} \cup \mathbf{Z}$ or $\mathbf{Y} \cup \mathbf{Z}$. In other words,

$$P(\mathbf{X}, \mathbf{Y}, \mathbf{Z}) = \frac{1}{Z} f(\mathbf{X}, \mathbf{Z}) g(\mathbf{Y}, \mathbf{Z}) \quad (484)$$

$$\text{which implies } P(\mathbf{X}, \mathbf{Y} \mid \mathbf{Z}) = \frac{f(\mathbf{X}, \mathbf{Z}) g(\mathbf{Y}, \mathbf{Z})}{\sum_{x,y} f(x, \mathbf{Z}) g(y, \mathbf{Z})} \quad (485)$$

- We can prove that $P \models (\mathbf{X} \perp \mathbf{Y} \mid \mathbf{Z})$ by showing that $P(\mathbf{X}, \mathbf{Y} \mid \mathbf{Z}) = P(\mathbf{X} \mid \mathbf{Z}) P(\mathbf{Y} \mid \mathbf{Z})$.

$$P(\mathbf{X}, \mathbf{Y} \mid \mathbf{Z}) = \frac{f(\mathbf{X}, \mathbf{Z}) g(\mathbf{Y}, \mathbf{Z})}{P(\mathbf{Z})} \quad (486)$$

$$= \frac{f(\mathbf{X}, \mathbf{Z}) g(\mathbf{Y}, \mathbf{Z})}{P(\mathbf{Z})} \frac{P(\mathbf{Z})}{P(\mathbf{Z})} \quad (487)$$

$$= \frac{f(\mathbf{X}, \mathbf{Z}) g(\mathbf{Y}, \mathbf{Z})}{P(\mathbf{Z})} \frac{\sum_x f(x, \mathbf{Z}) \sum_y g(y, \mathbf{Z})}{P(\mathbf{Z})} \quad (488)$$

$$= \frac{f(\mathbf{X}, \mathbf{Z}) \sum_y g(y, \mathbf{Z})}{P(\mathbf{Z})} \frac{g(\mathbf{Y}, \mathbf{Z}) \sum_x f(x, \mathbf{Z})}{P(\mathbf{Z})} \quad (489)$$

$$= P(\mathbf{X} \mid \mathbf{Z}) P(\mathbf{Y} \mid \mathbf{Z}) \quad (490)$$

- For the general case where $\mathbf{X} \cup \mathbf{Y} \cup \mathbf{Z} = \mathcal{X}$, let $\mathbf{U} = \mathcal{X} - (\mathbf{X} \cup \mathbf{Y} \cup \mathbf{Z})$. Since we know that $\text{sep}_{\mathcal{H}}(\mathbf{X}; \mathbf{Y} \mid \mathbf{Z})$, we can partition \mathbf{U} into two disjoint sets \mathbf{U}_1 and \mathbf{U}_2 such that $\text{sep}_{\mathcal{H}}(\mathbf{X} \cup \mathbf{U}_1; \mathbf{Y} \cup \mathbf{U}_2 \mid \mathbf{Z})$. Combining the previous result with the decomposition property¹³⁸ give us the desired result that $P \models (\mathbf{X} \perp \mathbf{Y} \mid \mathbf{Z})$.

Pairwise Independencies:

Let \mathcal{H} be a Markov network. We define the **pairwise independencies** assoc. with \mathcal{H} :

$$\mathcal{I}_p(\mathcal{H}) = \{(X \perp Y \mid \mathcal{X} - \{X, Y\}) : X - Y \notin \mathcal{H}\} \quad (491)$$

which just says “ X is indep. of Y given everything else if there’s no edge between X and Y .”

¹³⁸The decomposition property:

$$(\mathbf{X} \perp (\mathbf{Y}, \mathbf{W}) \mid \mathbf{Z}) \implies (\mathbf{X} \perp \mathbf{Y} \mid \mathbf{Z})$$

Local Independencies:

For a given graph \mathcal{H} , define the **Markov blanket** of X in \mathcal{H} , denoted $\text{MB}_{\mathcal{H}}(X)$, to be the neighbors of X in \mathcal{H} . We define the **local independencies** associated with \mathcal{H} :

$$\mathcal{I}_{\ell}(\mathcal{H}) = \{(X \perp \mathcal{X} - \{X\} - \text{MB}_{\mathcal{H}}(X) \mid \text{MB}_{\mathcal{H}}(X)) : X \in \mathcal{X}\} \quad (492)$$

which just says “ X is indep. of the rest of the nodes in the graph given its immediate neighbors.”

Log-Linear Models. Certain patterns involving particular values of variables for a given factor can often be more easily seen by converting factors into log-space. More precisely, we can rewrite a factor $\phi(\mathbf{D})$ as

$$\phi(\mathbf{D}) = e^{-\epsilon(\mathbf{D})} \quad (493)$$

$$\epsilon(\mathbf{D}) \triangleq -\ln \phi(\mathbf{D}) \quad (494)$$

$$\Pr[X_1, \dots, X_n] \propto e^{-\sum \epsilon_i(\mathbf{D}_i)} \quad (495)$$

where $\epsilon(\mathbf{D})$ is often called an **energy function**. Note how ϵ can take on any value along the real line (i.e. removes our nonnegativity constraint).

This motivates introducing the notion of a **feature**, which is just a factor without the nonnegativity requirement. A popular type of feature is the **indicator feature** that takes on value 1 for some values $\mathbf{y} \in \text{Val}(\mathbf{D})$ and 0 otherwise. We can now provide a more general definition for our notion of log-linear models:

*A distribution P is a **log-linear model** over a Markov network \mathcal{H} if it is associated with:*


- *A set of features $\mathcal{F} = \{f_1(\mathbf{D}_1), \dots, f_k(\mathbf{D}_k)\}$ where each \mathbf{D}_i is a complete subgraph (i.e. a clique) in \mathcal{H} .*
- *A set of weights w_1, \dots, w_k .*

such that

$$\Pr[X_1, \dots, X_n] = \frac{1}{Z} \exp \left[-\sum_{i=1}^k w_i f_i(\mathbf{D}_i) \right] \quad (496)$$

The log-linear model provides a much more compact representation for many distributions, especially in situations where variables have large domains (such as text).

Conditional Random Fields. So far, we’ve only described Markov network representation as encoding a joint distribution over \mathcal{X} . The same undirected graph representation and parameterization can also be used to encode a *conditional distribution* $\Pr[\mathbf{Y} | \mathbf{X}]$, where \mathbf{Y} is a set of *target variables* and \mathbf{X} is a (disjoint) set of *observed variables*.



Note how there are no connection between any of the X s

Formally, a CRF is an undirected graph \mathcal{H} whose nodes correspond to $\mathbf{Y} \cup \mathbf{X}$. Since we want to avoid representing¹³⁹ a probabilistic model over \mathbf{X} , we disallow potentials that involve only variables in \mathbf{X} ; our set of factors is $\phi_1(\mathbf{D}_1), \dots, \phi_m(\mathbf{D}_m)$, such that each $\mathbf{D}_i \not\subseteq \mathbf{X}$. The network encodes a conditional distribution as follows:

$$P(\mathbf{Y} | \mathbf{X}) = \frac{1}{Z(\mathbf{X})} \tilde{P}(\mathbf{Y}, \mathbf{X}) \quad (497)$$

$$\tilde{P}(\mathbf{Y}, \mathbf{X}) = \prod_{i=1}^m \phi_i(\mathbf{D}_i) \quad (498)$$


$$Z(\mathbf{X}) = \sum_{\mathbf{Y}} \tilde{P}(\mathbf{Y}, \mathbf{X}) \quad (499)$$

where now the partition function is a function of the assignment \mathbf{x} to \mathbf{X} .

¹³⁹Also note how we never have to deal with a summation over all possible \mathbf{X} , due to restricting ourselves to $Z(\mathbf{X})$.

EXERCISES

Exercise 4.1. Let \mathcal{H} be the graph of binary variables below. Show that P does not factorize over \mathcal{H} (Hint: proof by contradiction).



- Example 4.4 recap: P satisfies the global independencies w.r.t \mathcal{H} . They showed this by manually checking the two global indeps of \mathcal{H} , $(X_1 \perp X_3 \mid X_2, X_4)$ and $(X_2 \perp X_4 \mid X_1, X_3)$, against the tabulated list of possible assignments for P (given in example). Nothing fancy.
- P factorizes over \mathcal{H} if it can be written as a product of clique potentials.
- My proof:
 1. Assume that P does factorize over \mathcal{H} .
 2. Then P can be written as

$$P(X_1, X_2, X_3, X_4) = \frac{1}{Z} \phi_1(X_1, X_2) \phi_2(X_2, X_3) \phi_3(X_3, X_4) \phi_4(X_4, X_1) \quad (500)$$

Let the above assertion be denoted as C , and the statement that P factorizes according to \mathcal{H} be denoted simply as $P_{\mathcal{H}}$. Since $P_{\mathcal{H}} \iff C$, if we can prove that C does not hold, then we've found our contradiction, meaning $P_{\mathcal{H}}$ also must not hold.

3. I know that the proof must take advantage of the fact that we know P is zero for certain assignments to \mathcal{X} . For example $P(0100) = 0$. Furthermore, by looking at the assignments where P is *not* zero, I can see that all possible combinations of (X_1, X_2) are present, which means ϕ_1 never evaluates to zero.
4. From the example, we know that $P(1100) = 1/8 \neq 0$. However, since

$$0 = \frac{P(0100)}{P(1100)} = \frac{\phi_1(0, 1) \phi_4(0, 0)}{\phi_1(1, 1) \phi_4(0, 1)} \quad (501)$$

$$= \frac{\phi_4(0, 0)}{\phi_4(0, 1)} \quad (\phi_1 > 0) \quad (502)$$

and we also know that both the numerator and denominator of eq. 502 are positive, and thus we have a contradiction.

Local Probabilistic Models (Ch. 5)

Table of Contents Local


Written by Brandon McKinzie

Koller and Friedman (2009). Local Probabilistic Models.
Probabilistic Graphical Models: Principles and Techniques.

Deterministic CPDs. When X is a deterministic function of its parents Pa_X :

$$P(x \mid Pa_X) = \begin{cases} 1 & x = f(Pa_X) \\ 0 & \text{otherwise.} \end{cases} \quad (503)$$

Consider the example below, where the double-line notation on C means that C is a deterministic function of A and B. What new conditional dependencies do we have?



Answer: $(D \perp E \mid A, B)$, which would not be true by d-separation alone. It only holds because C is a deterministic function of A and B.

Template-Based Representations (Ch. 6)


Table of Contents Local

Written by Brandon McKinzie

Koller and Friedman (2009). Template-Based Representations.

Probabilistic Graphical Models: Principles and Techniques.

In what follows, we will build on an example where a vehicle tries to track its true location (L) using various sensor readings: velocity (V), weather (W), failure of sensor (F), observed location (O) of the noisy sensor.



Temporal Models. We discretize time into slices of interval Δ , and denote the ground random variables at time $t \cdot \Delta$ by $\mathcal{X}^{(t)}$. We can simplify our formulation considerably by assuming a **Markovian system**: a dynamic system over template variables \mathcal{X} that satisfies the Markov assumption:

$$(\mathcal{X}^{(t+1)} \perp \mathcal{X}^{(0:(t-1))} \mid \mathcal{X}^{(t)}) \quad (504)$$

which allows us to define a more compact representation of the joint distribution from time 0 to T:

$$P(\mathcal{X}^{(0:T)}) = P(\mathcal{X}^{(0)}) \prod_{t=0}^{T-1} P(\mathcal{X}^{(t+1)} \mid \mathcal{X}^{(t)}) \quad (505)$$

One last simplifying assumption, to avoid having unique transition probabilities for each time t , is to assume a **stationary**¹⁴⁰ Markovian dynamic system, defined s.t. $P(\mathcal{X}^{(t+1)} \mid \mathcal{X}^{(t)})$ is the same for all t .

¹⁴⁰Also called time invariant or homogeneous.

Dynamic Bayesian Networks. Above, I've drawn the **2-time-slice Bayesian network** (2-TBN) for our location example. A 2-TBN is a conditional BN over \mathcal{X}' given \mathcal{X}_I , where $\mathcal{X}_I \subseteq \mathcal{X}$ is a set of **interface variables**¹⁴¹. For each template variable X_i , the CPD $P(X'_i \mid Pa_{X'_i})$ is a **template factor**.

¹⁴¹Interface variables are those variables whose values at time t can have a direct effect on the variables at time $t + 1$.

Variable Elimination (Ch. 9)

Table of Contents Local

Written by Brandon McKinzie

Koller and Friedman (2009). Variable Elimination.

Probabilistic Graphical Models: Principles and Techniques.

Analysis of Exact Inference. The focus of this chapter is the conditional probability query,

$$\Pr[\mathbf{Y} \mid \mathbf{E} = \mathbf{e}] = \frac{\Pr[\mathbf{Y}, \mathbf{e}]}{\Pr[\mathbf{e}]} \quad (9.1)$$

Ideally, we want to obtain all instantiations $\Pr[\mathbf{y} \mid \mathbf{e}]$ of equation 9.1. Let $\mathbf{W} = \mathcal{X} - \mathbf{Y} - \mathbf{E}$ be the RVs that are neither query nor evidence. Then

$$\Pr[\mathbf{y} \mid \mathbf{e}] = \frac{\sum_{\mathbf{w}} \Pr[\mathbf{y}, \mathbf{e}, \mathbf{w}]}{\sum_{\mathbf{y}} \Pr[\mathbf{y}, \mathbf{e}]} \quad (9.3)$$

and note that, by computing all instantiations for the numerator first, we can reuse them to obtain the denominator. We can formulate the inference problem as a decision problem, which we will call *BNPrDP*, defined as follows:

Given: Bayesian network \mathcal{B} over \mathcal{X} , a variable $X \in \mathcal{X}$, and a value $x \in \text{Val}(X)$.

Decide: whether $\Pr_{\mathcal{B}}[X = x] > 0$.

BNPrDP

Thm 9.1: The decision problem *BNPrDP* is \mathcal{NP} -complete. **Proof:**

1. **BNPrDP is in \mathcal{NP} :** Guess a full assignment ξ to the network¹⁴². If the guess is successful, where success is defined as $(X = x) \in \xi$ and $P(\xi) > 0$, then we know that $P(X = x) > 0$ ¹⁴³. Computing $P(\xi)$ is linear in the number of factors for a BN, since we just multiply them together.
2. **BNPrDP is \mathcal{NP} -hard.** We show this by proving that we can solve 3-SAT (which is \mathcal{NP} -hard) by transforming inputs to 3-SAT to inputs of BNPrDP in polynomial time. Given any 3-SAT formula ϕ , we can create a BN B_{ϕ} with some special variable X s.t. ϕ is satisfiable IFF $P_{B_{\phi}}(X = x^1) > 0$. You can easily build such a network by having a node Q_i for each binary RV q_i , and a node C_i for each of the clauses that's a deterministic function of its parents (up to 3 Q nodes). Then, the node X is a deterministic function of its parents, which are chains of AND gates along the C_i . Since each node has at most 3 parents, we can ensure that construction is bounded by polynomial time in the length of ϕ .

¹⁴² Apparently the time it takes to generate a guess is irrelevant.

¹⁴³ This is true because $P(x)$ can be decomposed as $P(\xi) + \sum P(\dots) \geq P(\xi)$.

Analysis of Approximate Inference. Consider a specific query $P(\mathbf{y} \mid \mathbf{e})$, where we focus on a particular assignment \mathbf{y} . Let ρ denote some approximate answer, whose accuracy we wish to evaluate relative to the correct probability. We can use the **relative error** to estimate the quality of the approximation: *An estimate ρ has relative error ϵ if:*

$$\frac{\rho}{1 + \epsilon} \leq P(\mathbf{y} \mid \mathbf{e}) \leq \rho(1 + \epsilon) \quad (506)$$

Unfortunately, the task of finding some approximation ρ with relative error ϵ *is also \mathcal{NP} -hard*. Furthermore, even if we relax this metric by using absolute error instead, we end up finding that **in the case where we have evidence, approximate inference is no easier than exact inference, in the worst case.**

Variable Elimination. Basically, a dynamic programming approach for performing exact inference. Consider the simple BN $X_1 \rightarrow \dots \rightarrow X_n$, where each variable can take on k possible values. The dynamic programming approach for computing $P(X_n)$ involves computing

$$P(X_{i+1}) = \sum_{x_i} P(X_{i+1} \mid x_i) P(x_i) \quad (507)$$

$n - 1$ times, starting with $i = 1$, all the way up to $i = n - 1$, reusing the previous computation at each step, with total cost $\mathcal{O}(nk^2)$. So for this simple network, even though the size of the joint is k^n (exponential in n), we can do inference in linear time.

First, we formalize some basic concepts before defining the algorithm.

Factor marginalization:

*Let \mathbf{X} be a set of variables, and $Y \notin \mathbf{X}$ a variable. Let $\phi(\mathbf{X}, Y)$ be a factor. Define the **factor marginalization** of Y in ϕ , denoted $\sum_Y \phi$, to be a factor ψ over \mathbf{X} such that:*

$$\psi(\mathbf{X}) = \sum_Y \phi(\mathbf{X}, Y)$$

The key observation that's easy to miss is that *we're only summing entries in the table where the values of \mathbf{X} match up*. One useful rule for exchanging factor product and summation: If $X \notin \text{Scope}[\phi_1]$, then

$$\sum_X (\phi_1 \cdot \phi_2) = \phi_1 \cdot \sum_X \phi_2 \quad (9.6)$$

So, when computing some marginal probability, the main idea is to group factors together and compute expressions of the form

$$\sum_{\mathbf{Z}} \prod_{\phi \in \Phi} \phi \quad (508)$$

where Φ is the set of all factors ϕ for which $\mathbf{Z} \in \text{Scope}[\phi]$. This is commonly called the **sum-product** inference task. The full algorithm for sum-product variable elimination, which is an instantiation of the sum-product inference task, is illustrated below.

```


Procedure Sum-Product-VE (
     $\Phi$ , // Set of factors
     $\mathbf{Z}$ , // Set of variables to be eliminated
     $\prec$  // Ordering on  $\mathbf{Z}$ 
)
    Let  $Z_1, \dots, Z_k$  be an ordering of  $\mathbf{Z}$  such that
         $Z_i \prec Z_j$  if and only if  $i < j$ 
    for  $i = 1, \dots, k$ 
         $\Phi \leftarrow \text{Sum-Product-Eliminate-Var}(\Phi, Z_i)$ 
     $\phi^* \leftarrow \prod_{\phi \in \Phi} \phi$ 
    return  $\phi^*$ 

Procedure Sum-Product-Eliminate-Var (
     $\Phi$ , // Set of factors
     $Z$  // Variable to be eliminated
)
     $\Phi' \leftarrow \{\phi \in \Phi : Z \in \text{Scope}[\phi]\}$ 
     $\Phi'' \leftarrow \Phi - \Phi'$ 
     $\psi \leftarrow \prod_{\phi \in \Phi'} \phi$ 
     $\tau \leftarrow \sum_Z \psi$ 
    return  $\Phi'' \cup \{\tau\}$ 

```

This is what we use to compute the marginal probability $P(\mathbf{X})$ where $\mathbf{X} = \mathcal{X} - \mathbf{Z}$. To compute conditional queries of the form $P(\mathbf{Y} \mid \mathbf{E} = \mathbf{e})$, simply replace all factors whose scope overlaps with \mathbf{E} with their reduced factor (see chapter 4 notes for definition) to get the unnormalized $\phi^*(\mathbf{Y})$ (the numerator of $P(\mathbf{Y} \mid \mathbf{e})$). Then divide by $\sum_{\mathbf{y}} \phi^*$ to obtain the final result.

Example. We will work through computing $P(\text{Job})$ for the BN below.



Due to Happy being a child of Job, $P(J)$ actually requires using all factors in the graph. Below shows how VE with elimination ordering C, D, I, H, G, S, L progressively simplifies the equation for computing $P(J)$.

$$P(J) = \sum P(J | s, \ell) P(s | i) P(\ell | g) P(g | d, i) P(d | c) P(h | J, g) P(c) P(i) \quad (509)$$

$$= \sum_{\substack{c, d, i, h, g, s, \ell}} (\textcolor{red}{P(c)P(d | c)}) \cdot P(g | d, i) P(i) P(s | i) P(h | J, g) P(\ell | g) P(J | s, \ell) \quad (510)$$

$$= \sum_{\substack{d, i, h, g, s, \ell}} (\textcolor{red}{\tau_1(d)P(g | d, i)}) \cdot P(i) P(s | i) P(h | J, g) P(\ell | g) P(J | s, \ell) \quad (511)$$

$$= \sum_{\substack{i, h, g, s, \ell}} (\textcolor{red}{\tau_2(g, i)P(i)P(s | i)}) \cdot P(h | J, g) P(\ell | g) P(J | s, \ell) \quad (512)$$

$$= \sum_{\substack{h, g, s, \ell}} (\textcolor{red}{P(h | J, g)}) \cdot \tau_3(g, s) P(\ell | g) P(J | s, \ell) \quad (513)$$

$$= \sum_{\substack{g, s, \ell}} (\textcolor{red}{\tau_4(g, J)\tau_3(g, s)P(\ell | g)}) \cdot P(J | s, \ell) \quad (514)$$

$$= \sum_{\substack{s, \ell}} \textcolor{red}{\tau_5(J, \ell, s)} \cdot P(J | s, \ell) \quad (515)$$

$$= \sum_{\ell} \textcolor{red}{\tau_6(J, \ell)} \quad (516)$$

where red indicates the focus of the given step in the VE algorithm.

Clique Trees (Ch. 10)

Table of Contents Local

Written by Brandon McKinzie

Koller and Friedman (2009). Clique Trees.


Probabilistic Graphical Models: Principles and Techniques.

Cluster Graphs. A graphical flowchart of the factor-manipulation process that will be relevant when we discuss message passing. Each node is a *cluster*, which is associated with a subset of variables. Formally,

A **cluster graph** \mathcal{U} for a set of factors Φ over \mathcal{X} is an undirected graph. Each node i is associated with a subset $C_i \subseteq \mathcal{X}$. Each factor $\phi \in \Phi$ must be associated with a cluster C_i , denoted $\alpha(\phi)$, such that $\text{Scope}[\phi] \subseteq C_i$ (**family-preserving**). Each edge between a pair of clusters C_i and C_j is associated with a **sepset** $S_{i,j} \subseteq C_i \cap C_j$.

Recall that each step of variable elimination involves creating a factor ψ_i by multiplying a group of factors¹⁴⁴. Then, denoting the variable we are eliminating at this step as Z , we obtain another factor τ_i that's the factor marginalization of Z in ψ_i (denoted $\sum_Z \psi_i$). An execution of variable elimination defines a cluster graph: we have a cluster for each of the ψ_i , defined as $C_i = \text{Scope}[\psi_i]$. We draw an edge between C_i and C_j if the **message** τ_i is used in the computation of τ_j .

Consider when we applied variable elimination to the student graph network below, to compute $P(J)$. Elimination ordering C, D, I, H, G, S, L .



¹⁴⁴All whose scope contains the variable we are currently trying to eliminate.

Clique Trees. Since VE uses each intermediate τ_i at most once, the cluster graph induced by an execution VE is necessarily a *tree*, and it also defines a directionality: the direction of the message passing (left-to-right in the above illustration). All the messages flow toward a single cluster where the final result is computed – the **root** of the tree; we say the messages “flow up” to the root. Furthermore, for cluster trees induced by VE, the scope of each message (edge) τ_i is *exactly* $C_i \cap C_j$, not just a subset¹⁴⁵.

*Let Φ be a set of factors over \mathcal{X} . A cluster tree over Φ that satisfies the running intersection property is called a **clique tree**. For clique trees, the clusters are also called **cliques**.*

Message Passing: Sum Product. Previously we saw how an execution of VE can be illustrated with a clique tree. We now go the other direction – given a clique tree, we show how it can be used for variable elimination. Given a clique tree representation of some BN, we can use it to guide us along an execution of VE to compute any marginal we’d like. First, before any run, we generate the set of **initial potentials** ψ_i associated with each clique C_i in the tree, defined as just the multiplication of the initial factors associated with the clique. We define the root of the tree as any clique containing the variable whose marginal we want to compute (we pick arbitrarily). Starting from the leaves and moving toward the root, we pass messages along from clique to clique. A clique is *ready* to send a message when it has received a message from all of its downstream neighbors. The message from C_i to [a neighbor] C_j is computed using the **sum-product message passing** computation:

$$\delta_{i \rightarrow j} = \sum_{C_i - S_{i,j}} \psi_i \cdot \prod_{k \in (Nb_i - \{j\})} \delta_{k \rightarrow i} \quad (10.2)$$

where the summation is simply over the variables in C_i that aren’t passed along to C_j , and the product is over all messages that C_i received. Stated even simpler, we multiply all the incoming messages by our initial potential, then sum out all variables except those in $S_{i,j}$. When the root clique has received all messages, it multiplies them with its own initial potential, resulting in a factor called the **beliefs**, $\beta_r(C_r)$. It represents

$$\tilde{P}_\Phi(C_r) = \sum_{\mathcal{X} - C_r} \prod_{\phi} \phi \quad (517)$$

where, to be clear, the product is over all ϕ in the graph.

¹⁴⁵This follows from the **running intersection property**, which is satisfied by any cluster tree that’s defined by variable elimination. It’s defined as, if any variable X is in both cluster C_i and C_j , then X is also in every cluster in the (unique) path in the tree between C_i and C_j .

Below is a more compact summary of all of this, showing the procedure for computing *all* final factors (belief) β_i for some marginal probability query on the variables in \mathbf{C}_r *asynchronously*.

Algorithm 10.2: Sum-Product Belief Propagation

1. For each clique \mathbf{C}_i , compute its initial potential:

$$\psi_i(\mathbf{C}_i) \leftarrow \prod_{\phi_j: \alpha(\phi_j)=i} [\phi_j]$$

2. While $\exists i, j$ such that i is ready to transmit to j , compute:

$$\delta_{i \rightarrow j} \leftarrow \sum_{\mathbf{C}_i - \mathbf{S}_{i,j}} \psi_i \cdot \prod_{k \in (Nb_i - \{j\})} \delta_{k \rightarrow i}$$

3. Then, compute each belief factors β_i by multiplying the initial potential ψ_i by the incoming messages to \mathbf{C}_i :

$$\beta_i \leftarrow \psi_i \cdot \prod_{k \in Nb_{\mathbf{C}_i}} \delta_{k \rightarrow i}$$

4. Return the set of beliefs $\{\beta_i\}$, where

$$\beta_i = \sum_{\mathcal{X} - \mathbf{C}_i} \tilde{P}_{\Phi}(\mathcal{X}) = \tilde{P}_{\Phi}(\mathbf{C}_i) \quad (518)$$

The SP Belief Propagation algorithm above is also called **clique tree calibration**. A clique tree \mathcal{T} is *calibrated* if all pairs of adjacent cliques are calibrated. A calibrated clique tree satisfies the following property for what we'll now call the **clique beliefs**, β_i , and the **sepset beliefs**, $\mu_{i,j}$ over $\mathbf{S}_{i,j}$:

$$\mu_{i,j}(\mathbf{S}_{i,j}) \triangleq \sum_{\mathbf{C}_i - \mathbf{S}_{i,j}} \beta_i = \sum_{\mathbf{C}_j - \mathbf{S}_{i,j}} \beta_j \quad (519)$$

$$\mu_{i,j} = \tilde{P}_{\Phi}(\mathbf{S}_{i,j})$$

The main advantage of the clique tree algorithm is that it computes the posterior probability of all variables in a graphical model using only twice the computation¹⁴⁶ of the upward pass in the same tree.

We can also show that $\mu_{i,j} = \delta_{j \rightarrow i} \delta_{i \rightarrow j}$, which then allows us to derive:

$$\tilde{P}_{\Phi}(\mathcal{X}) = \frac{\prod_{i \in \mathcal{V}_{\mathcal{T}}} \beta_i}{\prod_{ij \in \mathcal{E}_{\mathcal{T}}} \mu_{i,j}} \quad (10.10)$$

In other words, the clique and sepset beliefs provide a **reparameterization** of the unnormalized measure, a property called the **clique tree invariant**.

¹⁴⁶The algorithm is equivalent to doing one upward pass, one downward pass.

Message Passing: Belief Update. We now discuss an alternative message passing approach that is mathematically equivalent but intuitively different than the sum-product approach. First, we introduce some new definitions.

Factor Division:

Let \mathbf{X} and \mathbf{Y} be disjoint sets of variables, and let $\phi_1(\mathbf{X}, \mathbf{Y})$ and $\phi_2(\mathbf{Y})$ be two factors. We define the division of ϕ_1 and ϕ_2 as a factor ψ with scope \mathbf{X}, \mathbf{Y} as follows:


Define $0/0 = 0$

$$\psi(\mathbf{X}, \mathbf{Y}) \triangleq \frac{\phi_1(\mathbf{X}, \mathbf{Y})}{\phi_2(\mathbf{Y})}$$

Looking back at equation 10.2, we can now see that another way to write $\delta_{i \rightarrow j}$ is

$$\delta_{i \rightarrow j} = \frac{\sum_{\mathbf{C}_{i-S_{i,j}}} \beta_i}{\delta_{j \rightarrow i}} \quad (10.13)$$

Now, consider the clique tree below for the simple Markov network A-B-C-D:



If we assigned \mathbf{C}_2 as the root, then our previous approach would compute $\delta_{2 \rightarrow 1}$ as $\sum_C \psi_2 \cdot \delta_{3 \rightarrow 2}$. Alternatively, we can use equation 10.13 to realize this is equivalent to dividing β_2 by $\delta_{1 \rightarrow 2}$ and marginalizing out C . This observation motivates the algorithm below, which allows us to execute message passing in terms of the clique and sepset beliefs, without having to remember the initial potentials ψ_i or explicitly compute the messages $\delta_{i \rightarrow j}$.

Algorithm 10.3: Belief-Update Message Passing

1. For each clique \mathbf{C}_i , set its initial belief β_i to its initial potential ψ_i . For each edge in $\mathcal{E}_{\mathcal{T}}$, set $\mu_{i,j} = \mathbf{1}$.
2. While there exists an uninformed¹⁴⁷ clique in \mathcal{T} , select any edge in $\mathcal{E}_{\mathcal{T}}$, and compute

$$\sigma_{i \rightarrow j} \leftarrow \sum_{\mathbf{C}_i - \mathbf{S}_{i,j}} [\beta_i] \quad (520)$$

$$\beta_j \leftarrow \beta_j \cdot \frac{\sigma_{i \rightarrow j}}{\mu_{i,j}} \quad (521)$$

$$\mu_{i,j} \leftarrow \sigma_{i \rightarrow j} \quad (522)$$

3. Return the resulting set of informed beliefs $\{\beta_i\}$.

At convergence, $\sigma_{i \rightarrow j} = \mu_{i,j} = \sigma_{j \rightarrow i}$.

¹⁴⁷A clique is informed once it has received informed messages from all of its neighbors. An informed message is one that has been sent by taking into account information from all of the sending cliques' neighbors (aside from the receiving clique of that message, of course).

Inference as Optimization (Ch. 11)


Table of Contents Local

Written by Brandon McKinzie

Koller and Friedman (2009). Inference as Optimization.

Probabilistic Graphical Models: Principles and Techniques.

Propagation-Based Approximation. We can use a general-purpose cluster graph rather than the more restrictive clique tree (needed to guarantee exact inference) for approximate inference methods. Consider the simple Markov network below on the left.



The clique tree for this network, which can be used for exact inference, has two cliques ABD and BCD and messages are passed between them consisting of $\tau(B, D)$. Suppose that, instead, we set up 4 clusters corresponding to each of the initial potentials, shown as the cluster graph above on the right. We can still apply belief propagation here, but due to it now having loops (as opposed to before when we only had trees), *the process may not converge*.

INFORMATION THEORY, INFERENCE, AND LEARNING ALGORITHMS

CONTENTS

8.1	Introduction to Information Theory (Ch. 1)	214
8.2	Probability, Entropy, and Inference (Ch. 2)	216
8.2.1	More About Inference (Ch. 3 Summary)	217
8.3	The Source Coding Theorem (Ch. 4)	219
8.3.1	Data Compression and Typicality	221
8.3.2	Further Analysis and Q&A	223
8.4	Monte Carlo Methods (Ch. 29)	225
8.5	Variational Methods (Ch. 33)	227

Introduction to Information Theory (Ch. 1)

Table of Contents Local

Written by Brandon McKinzie

[Note: Skipping most of this chapter since it's mostly introductory material.]

Preface. For ease of reference, some common quantities we will frequently be using:

- **Binomial distribution.** Let r denote the number of successful trials out of N total trials. Let f denote the probability of success for a single trial.

$$\Pr[r \mid f, N] = \binom{N}{r} f^r (1-f)^{N-r} \quad \mathbb{E}[r] = Nf \quad \text{Var}[r] = Nf(1-f) \quad (523)$$

- **Stirling's Approximation.**

$$x! \simeq x^x e^{-x} \sqrt{2\pi x} \quad \Leftrightarrow \quad \ln x! \simeq x \ln x - x + \frac{1}{2} \ln 2\pi x \quad (524)$$

Recall that
 $\log_b x = \frac{\log_a x}{\log_a b}$

$$\ln \binom{N}{r} \simeq r \ln \frac{N}{r} + (N-r) \ln \frac{N}{N-r} \quad (525)$$

- **Binary Entropy Function** and its relation with Stirling's approximation.

$$H_2(x) \triangleq x \lg \frac{1}{x} + (1-x) \lg \frac{1}{1-x} \quad (526)$$

$$\lg \binom{N}{r} \simeq NH_2(r/N) \quad (527)$$

Perfect communication over an imperfect, noisy communication channel (1.1). We want to make an encoder-decoder architecture, of the general form in the figure below, to achieve reliable communication over a noisy channel.

Information theory is concerned with the theoretical limitations and potentials of such systems. Let's explore some examples for the case of the **binary symmetric channel**¹⁴⁸:


- **Repetition codes.** Let R_N denote the repetition code that repeats each bit in the message N times¹⁴⁹. We model the channel as “adding”¹⁵⁰ a sparse noise vector \mathbf{n} to the encoded message \mathbf{t} , so $\mathbf{r} := \mathbf{n} + \mathbf{t}$. What is the optimal way of decoding r ?

$$\hat{s}_i \leftarrow \arg \max_{s_i} \Pr[s_i \mid \mathbf{r}_{i:i+N}] = \arg \max_{s_i} \Pr[\mathbf{r}_{i:i+N} \mid s_i] \Pr[s_i] \quad (528)$$

¹⁴⁸A binary symmetric channel transmits each bit correctly with probability $(1-f)$ and incorrectly with probability f , where f is assumed to be small.

¹⁴⁹So R_2 would encode 101 as 110011.

¹⁵⁰We add in modulo 2, which is NOT the same as binary arithmetic (no carry). Addition modulo 2 is the same as doing XOR.




We see that we must make assumptions about the prior probability $\Pr[s_i]$. It is common to assume all possible values of s_i (0 or 1 in this case) are equally probable. It is useful to observe the **likelihood ratio**,

$$\frac{\Pr[r_{i:i+N} | s_i = 1]}{\Pr[r_{i:i+N} | s_i = 0]} = \prod_{n=i}^{i+N-1} \frac{\Pr[r_n | t_n = 1]}{\Pr[r_n | t_n = 0]} = \prod_{n=i}^{i+N-1} \begin{cases} \gamma & \text{if } r_n = 1 \\ \gamma^{-1} & \text{if } r_n = 0 \end{cases} \quad (529)$$

where we've defined $\gamma := (1 - f)/f$, with f being the probability of a bit getting flipped by the channel. We want to assign \hat{s}_i to the most likely **hypothesis** out of the possible s_i . If the likelihood ratio [for the two hypotheses] is greater than 1, we choose $\hat{s}_i = 1$, else we choose $\hat{s}_i = 0$.

- **Block Codes - the (7, 4) Hamming Code.** Although, by increasing the number of repetitions per bit N for our R_N repetition code can decrease the error-per-bit probability p_b , we incur a substantial decrease in the *rate of information transfer* – a factor of $1/N$. The **(7, 4) Hamming Code** tries to improve this by encoding *blocks* of bits at a time (instead of per-bit). It is a *linear block code* – it encodes each 4-bit block into a 7-bit block, where the additional 3 bits are linear functions of the original $K = 4$ bits. The (7, 4) Hamming code has each of the extra 3 bits act as parity checks. It's easiest to show this with the illustration below:



which encodes 1000 as 1000101.

Probability, Entropy, and Inference (Ch. 2)

Table of Contents Local

Written by Brandon McKinzie

[Note: Skipping most of this chapter since it's mostly introductory material.]

Notation. Some of the notation this author seems to use a lot.

- **Ensemble X :** a triple $(x, \mathcal{A}_X, \mathcal{P}_X)$, where the outcome x is the value of a R.V. which can take on one of a set of possible values (“alphabet”), $\mathcal{A}_X = \{a_1, \dots, a_i, \dots, a_I\}$, having probabilities $\mathcal{P}_X = \{p_1, \dots, p_i, \dots, p_I\}$. Note that this doesn't appear technically consistent with how the author actually *uses* the term ensemble – in practice, he actually means “ X is the set of all possible triples $(x, \mathcal{A}_X, \mathcal{P}_X)$, $\forall x \in \mathcal{A}_X$, or something like that. He uses it to casually refer to the space of possibilities.

Forward Probabilities and Inverse Probabilities. Both of these involve a **generative model** of the data. In a *forward* probability problem, we want find the PDF, expectation, or some other function of a quantity that depends on the data/is *produced* by the generative process. For example, we can model a series of N coin flips as “producing” the quantity n_H , denoting the number of heads. In an *inverse* probability problem, we want to compute conditional probabilities on one or more of the *unobserved variables* in the process, *given* the observed variables.

The Likelihood Principle. For a generative model on data d given parameters θ , $\Pr[d | \theta]$, and having observed a particular outcome d_1 , all inferences and predictions should depend only on the function $\Pr[d_1 | \theta]$.

Entropy and Related Quantities.

- **Shannon information content** of an outcome x :

$$h(x) \triangleq \lg \left(\frac{1}{P(x)} \right) \quad (530)$$

- **Entropy of Ensemble X .** Defined to be the average Shannon information content of an outcome.

$$H(X) \triangleq \sum_{x \in \mathcal{A}_X} \Pr[x] \lg \left(\frac{1}{\Pr[x]} \right) \quad (531)$$

$$H(X) \leq \lg(|\mathcal{A}_X|) \quad \text{with equality iff} \quad p_i = \frac{1}{|\mathcal{A}_X|} \quad \forall i \quad (532)$$

$$0 \times \lg \left(\frac{1}{0} \right) \triangleq 0$$

- **Decomposability of the Entropy.** For any probability distribution $\mathbf{p} = \{p_1, p_2, \dots, p_I\}$ and m (where $1 \leq m \leq I$):

$$H(\mathbf{p}) = H(\Sigma_{1:m}, \Sigma_{m+1:I}) \quad (533)$$

$$+ \Sigma_{1:m} H\left(\frac{p_1}{\Sigma_{1:m}}, \dots, \frac{p_m}{\Sigma_{1:m}}\right) \quad (534)$$

$$+ \Sigma_{m+1:I} H\left(\frac{p_{m+1}}{\Sigma_{m+1:I}}, \dots, \frac{p_I}{\Sigma_{m+1:I}}\right) \quad (535)$$

where I've let $\Sigma_{1:m} := \sum_{i=1}^m p_i$.

- **Kullback-Leibler Divergence** between two probability distributions $P(X)$ and $Q(X)$ that are defined over the same alphabet \mathcal{A}_X :

$$D_{KL}(P||Q) = \sum_x P(x) \lg \frac{P(x)}{Q(x)} \quad (536)$$

$$D_{KL}(P||Q) \geq 0 \quad \text{[Gibb's Inequality]} \quad (537)$$

where, in the words of the author, “**Gibb's inequality is probably the most important inequality in this book**”.

- **Convex functions and Jensen's Inequality.** A function $f(x)$ is convex over the interval $[x = a, x = b]$ if every chord of the function lies above the function. That is, $\forall x_1, x_2 \in [a, b]$ and $0 \leq \lambda \leq 1$:

$$f(\lambda x_1 + (1 - \lambda)x_2) \leq \lambda f(x_1) + (1 - \lambda)f(x_2) \quad (538)$$

$$f(\mathbb{E}[x]) \leq \mathbb{E}[f(x)] \quad \text{[Jensen's Inequality]} \quad (539)$$

and we say f is **strictly convex** if, $\forall x_1, x_2 \in [a, b]$, we get equality for only $\lambda = 0$ and $\lambda = 1$.

MORE ABOUT INFERENCE (CH. 3 SUMMARY)

Because this is too short to have as its own chapter...

A first inference problem. Author explores a particle decay problem of finding $\Pr[\lambda | \{x\}]$ where λ is the characteristic decay length and $\{x\}$ is our collection of observed decay distances. Plotting the likelihood $\Pr[\{x\} | \lambda]$ as a function of λ for any given $x \in \{x\}$ shows each has a *peak* value. The kicker is the interpretation: if each measurement $x \in \{x\}$ is independent, the total likelihood is the product of all the individual likelihoods, which can be interpreted as updating/narrowing the interval $[\lambda_a, \lambda_b]$ within which $\Pr[\{x\} | \lambda]$ (as a function of λ) peaks. In the words of the author's mentor:

what you know about λ after the data arrive is what you knew before ($\Pr[\lambda]$) and what the data told you ($\Pr[\{x\} | \lambda]$)

We update our beliefs regarding the distribution of λ as we collect data.

Lessons learned from problems.

- (3.8) The classic Monty Hall problem. **Be careful defining probabilities after collecting data.** My blunder: when using Bayes’theorem to get the probability of the prize being behind door $i \in \{1, 2\}$ after the host opens door 3, I failed to take into account that we chose door 1 while I was computing the evidence (the marginal probability that the host opened door 3)¹⁵¹. Quite embarrassing.
- (3.9) Monty Hall problem, but an earthquake opens door 3. Although I correctly answered that, in this case, both hypotheses (the prize being behind door 1 or door 2) are equiprobable, I still failed to account for the subtle fact that the earthquake could’ve opened *multiple* doors. The lesson here is **always write down the probability of everything**, which just so happens to be suggested by the solution for this problem, too.

So, why did I still get the answer correct? The reason is because enumeration of probabilities wasn’t necessary at all, you just needed to realize that the likelihood for the two remaining hypotheses (\mathcal{H}_1 and \mathcal{H}_2) were the *same* – the probability of observing the earthquake open door 3 and the prize not being revealed was the same for the case of the prize being behind door 1 or door 2. So maybe the real lesson here is **determine whether calculations are even needed in order to solve the given problem**, which luckily I’ve had drilled in my head for years from studying physics.

- (3.15) Another biased coin variant. One of the best examples I’ve seen for favoring Bayesian methods over frequentist methods. Also, made use of the beta function:

$$\int_0^1 p^x (1-p)^y dp = \frac{\Gamma(x+1)\Gamma(y+1)}{\Gamma(x+y+2)} = B(x+1, y+1) \quad (540)$$

where B is the **beta function**, which is defined by the LHS.

¹⁵¹Note that, while important to recognize and understand, I could’ve avoided this pitfall entirely by just ignoring the evidence during calculations and normalizing after, since the evidence can be determined solely by the normalization constraint.

The Source Coding Theorem (Ch. 4)

Table of Contents Local

Written by Brandon McKinzie

Overview. We will be examining the two following assertions:

1. The Shannon information content is a sensible measure of the information content of a given outcome $x = a_i$:

$$h(x = a_i) \triangleq \lg \frac{1}{p_i} \quad (541)$$

2. The entropy of an ensemble X is a sensible measure of the ensemble's average information content.

$$H(X) = \sum_i p_i \lg \frac{1}{p_i} \quad (542)$$

The weighing problem. We are instructed to ponder the following problem.

You're given 12 balls, all equal in weight except for one that is either heavier or lighter. You're also given a two-pan balance. Your task is to determine which ball is the odd ball, and in as few uses of the balance as possible. Note: each use of the balance must place the same number of balls on either side.

An interesting observation is to consider the number of possible outcomes of the weighing process. Each outcome can be one of three possibilities: equal weight, left heavier, or left lighter. After N such weighings, the number of unique possible weighing result sequences is 3^N . Note that there are $12 \times 2 = 24$ unique final answers for our task (identifying which is the odd ball, and whether it is heavier or lighter). Therefore, since we are seeking a procedure to identify which of the 24 options is the correct option with 100% accuracy, we require our weighing procedure to take on *at least* 24 unique possible results. Since $N = 3$ weighings corresponds to $3^3 = 27$ possible outcomes, $N = 3$ is a *lower bound* on the number of weighings our approach will involve. It is impossible to guarantee a correct answer for $N < 3$ weighings¹⁵².

Things I didn't consider until reading the solution:

- It's actually *not* optimal, upon observing both sides equal, to subsequently use only the balls not involved in that measurement. My initial reaction to this was "why? we already know the oddball is not any of the balls just measured, since the outcome was equal."

¹⁵²Finally, it should be clear that, regardless of our approach, the final weighing will involve 2 balls, since we have to identify which is the oddball AND whether it is heavier/lighter AND the number of balls on the left of the scale must be the same as the right of the scale for every weighing.

The response to this reaction is: “yes, *exactly*, and we must use that information to be able to discern in the future whether, e.g., a measurement of “left side heavier” means the oddball is on the left and heavy, or if it’s on the right and light – *it’s useful to know that a given side of the scale does not contain the oddball before a measurement*.

- More generally, it’s also not optimal to greedily search for solutions that eliminate the highest number of possibilities *in any given single step*. Another way of thinking about this is that it’s undesirable for the i th measurement outcome to cause any of the 3 possible measurement outcomes to be impossible at the next stage.
- I focused a disproportionate amount of thought on handling the equal-weight measurement outcome, for whatever reason. I probably would’ve arrived at the solution faster if I’d actually thought about how my strategies would’ve handled some outcome being “left heavier” and *then considered what that strategy would put on the scale at the next step*, where the italics denote what would’ve illuminated the fatal flaw in all my approaches.

Guessing Games. What’s the smallest number of yes/no questions needed to identify an integer x between 0 and 63? Although it was obvious to me that the solution is to successively halve the possible values of x , I found it interesting that **you can write down the list of questions independent of the answers at each step** using a basic application of modular arithmetic. In other words, you can specify the full decision tree of N nodes with just $\lg N$ questions. Nice. Also, recognize that the Shannon information content for any single outcome is $\lg \frac{1}{0.5} = 1$ bit, and thus the total Shannon information content (for our predefined 6 questions) is 6 bits, which is not-coincidentally the number of possible values that x could be before we ask any questions.

In general, if an outcome x has Shannon information content $h(x)$ number of bits, I like to interpret that as “learning the result x eliminates $2^{h(x)}$ possibilities for the final result.” The battleship example follows this interpretation well. Stated another way (in the author’s words):

The Shannon information content can be intimately connected to the size of a file that encodes the outcomes of a random experiment.

Data Compression. A **lossy compressor** compresses some files, but maps some files to the *same* encoding. We introduce a parameter δ that describes the risk (aggressiveness of our compression) we are taking with a given compression method: δ is the probability that there will be no name for an outcome¹⁵³ x .

The smallest δ -sufficient subset

If S_δ is the smallest subset of \mathcal{A}_X satisfying

$$\Pr [x \in S_\delta] \geq 1 - \delta \quad (543)$$

then S_δ is the smallest δ -sufficient subset. It can be constructed by ranking the elements of \mathcal{A}_X in order of decreasing probability and adding successive elements starting from the most probable elements until the total probability is $\geq (1 - \delta)$.

- **Raw bit content** of X : $H_0(X) \triangleq \lg |\mathcal{A}_X|$. A lower bound for the number of binary questions that are always *guaranteed* to identify an outcome from the ensemble X – it simply maps each outcome to a constant-length binary string.
- **Essential bit content** of X : $H_\delta(X) \triangleq \lg |S_\delta|$. A compression code can be made by assigning a binary string of H_δ bits to each element of the smallest sufficient subset.

Finally, we can now state **Shannon’s source coding theorem**: Let X be an ensemble for the random variable x with entropy $H(X) = H$ bits, and let X^N denote a sequence of identically distributed (but not necessarily independent¹⁵⁴) of random variables/ensembles, (X_1, X_2, \dots, X_N) .

$$(\exists N_0 \in \mathbb{Z}^+)(\forall N > N_0) : \quad \left| \frac{1}{N} H_\delta(X^N) - H \right| < \epsilon \quad (0 < \delta < 1) \ (\epsilon > 0) \quad (544)$$

which, in English, can be read: *N i.i.d. random variables each with entropy $H(X)$ can be compressed into more than $NH(X)$ bits with negligible risk of information loss, as $N \rightarrow \infty$; conversely if they’re compressed into fewer than $NH(X)$ bits it is virtually certain that information will be lost.*

¹⁵³More specifically, if there is some subset, $\{a\}$, of unique values that x can take on but our compression method discards/ignores, then we say $\delta = \sum_i p(x = a_i)$.

¹⁵⁴Actually, before the actual theorem statement, the author mentions we are now concerned with “string of N i.i.d. random variables from a single ensemble X .” It’s probably fair to assume this is true for the quantities in the theorem, but I’m leaving this note here as a reminder.

Typicality. The reason that large N in equation 544 corresponds to larger potential for better compression is that the subset of likely results for a string of outcomes becomes more and more concentrated relative to the number of possible sequences as N increases¹⁵⁵. I just realized this is for the same fundamental reasons that entropy exists in thermodynamics – *there are just more ways to exist in a high entropy state than otherwise*. The author showed $\binom{N}{r}$ as a function of r (the number of 1s in the N -bit string). For large N , this becomes almost comically concentrated near the center (like a delta function at $N/2$) – see footnote for more details¹⁵⁶.

This motivates the notion of **typicality** for [a string of length N from] an arbitrary ensemble X with alphabet \mathcal{A}_X . For large N , we expect to find $p_i N$ occurrences of the outcome $x = a_i$. Hence the probability of such a string, and its information content, is roughly¹⁵⁷

$$\Pr[\mathbf{x}]_{typ} = \Pr[x_1] \Pr[x_2] \cdots \Pr[x_N] \simeq p_1^{(p_1 N)} p_2^{(p_2 N)} \cdots p_I^{(p_I N)} \quad (545)$$

$$h(\mathbf{x})_{typ} = \lg \frac{1}{\Pr[\mathbf{x}]_{typ}} \simeq N \sum_{i=1}^I p_i \lg \frac{1}{p_i} = NH(X) \quad (546)$$

Accordingly, we define the typical elements (strings of length N) of \mathcal{A}_X^N to be those elements that have probability close to 2^{-NH} . We introduce a parameter β that defines what we mean by “close,” and define the set of typical elements as the **typical set** $T_{N\beta}$:

$$T_{N\beta} \triangleq \left\{ \mathbf{x} \in \mathcal{A}_X^N : \left| \frac{1}{N} \lg \frac{1}{P(\mathbf{x})} - H \right| < \beta \right\} \quad (547)$$

It turns out that whatever value of β we choose, the $T_{N\beta}$ contains almost all the probability as N increases.

¹⁵⁵The author gave an example for a sequence of bits with probability of any given bit being 1 as 0.1. He showed how, although the *average* number of 1s in a sequence of N bits grew as $\mathcal{O}(N)$, the standard deviation of that average only grew as \sqrt{N} .

¹⁵⁶The probability of getting a string with r 1s follows a binomial distribution with mean Np_1 and standard deviation $\sqrt{Np_1(1-p_1)}$. This results in an increasingly narrower distribution $P(r)$ for larger N .

¹⁵⁷We appear to be assuming that each outcome x in the string \mathbf{x} are i.i.d. (CONFIRMED)

Proving the Source Coding Theorem.

- **Setup.** We will make use of the following:
 - **Chebyshev’s Inequalities:**

$$\Pr [x \geq \alpha] \leq \frac{\mathbb{E}[x]}{\alpha} \quad \text{and} \quad \Pr [(x - \mathbb{E}[x])^2 \geq \alpha] \leq \frac{\text{Var}[x]}{\alpha} \quad (548)$$

where α is a positive real number, and x is assumed non-negative in the first inequality¹⁵⁸.

- **Weak Law of Large Numbers (WLLN):** Consider a sample h_1, \dots, h_N of N independent RVs all with common mean \bar{h} and common variance σ_h^2 . Let $x = \frac{1}{N} \sum_{n=1}^N h_n$ be their average. Then

$$\Pr [(x - \bar{h})^2 \geq \alpha] \leq \frac{\sigma_h^2}{\alpha N} \quad (549)$$

which can be easily derived from Chebyshev’s inequalities.

- **Proving ‘asymptotic equipartition’ principle,** i.e. that an outcome \mathbf{x} is almost certain to belong to the typical set, approaching probability 1 for large enough N . It is a simple application of the WLLN to the random variable

$$\frac{1}{N} \lg \frac{1}{\Pr[\mathbf{x}]} = \frac{1}{N} \sum_{n=1}^N \lg \frac{1}{x_n} = \frac{1}{N} \sum_{n=1}^N h(x_n) \quad (550)$$

where $\mathbb{E}[h(x_n)] = H(X)$ for all terms in the summation. Observe, then, that the definition of the typical set given in equation 547 (squaring both sides) has the same form as the definition for the WLLN. Plugging in and rearranging yields

$$\Pr [\mathbf{x} \in T_{N\beta}] \geq 1 - \frac{\sigma^2}{\beta^2 N} \quad (551)$$

where $\sigma^2 \equiv \text{Var} \left[\lg \frac{1}{P(x_n)} \right]$. This proves the asymptotic equipartition principle. It will also be useful to recognize that for any \mathbf{x} in the typical set, we can rearrange equation 547 to obtain

$$2^{-N(H+\beta)} < \Pr[\mathbf{x}] < 2^{-N(H-\beta)} \quad (552)$$

- **Proof of SCT Part I.** Want to show that $\frac{1}{N} H_\delta(X^N) < H + \epsilon$. **TODO**
- **Proof of SCT Part II.** Want to show that $\frac{1}{N} H_\delta(X^N) > H - \epsilon$. **TODO**

¹⁵⁸Notice how the two inequalities are technically the same.

Questions & Answers. Collection of my questions as I read through the chapter, which I answered upon completing it.

- **Q:** Why isn't the essential bit content of a string of N i.i.d. variables NH when $\delta = 0$?
 - **A:** I'm not entirely sure how to answer this still, but it seems the question is confused. First off, the essential bit content approaches the raw bit content as δ decreases to 0: $H_\delta \rightarrow H_0$ as $\delta \rightarrow 0$. It's important to notice that both H_δ and H_0 define an entropy where all members (S_δ^N for H_δ ; \mathcal{A}_X^N for H_0) are *equiprobable*. I remember asking this question wondering "what is the significance of $H_\delta(X^N)$ approaching $NH(X)$ (not a typo!) for tiny δ ". The answer: for larger N , more of the probability mass is concentrated in a relatively smaller region, *with elements of that region being roughly equiprobable*. The last part is what I didn't initially realize – that **allowing for tiny δ combined with large N essentially makes it so that $S_\delta^N \approx T_{N\beta}$** .
- **Q:** Why aren't the most probable elements necessarily in the typical set?
 - **A:** In the limit of $N \rightarrow \infty$, *they are*, since in that limit, all elements are in the typical set and they're equiprobable. However, in essentially any real case, we can imagine that some elements will be too unlikely to be found within the typical set, which necessarily requires that there exist elements with probability too high to be in the typical set. Remember that the typical set is basically *defined* such that all elements have probability within the range given in equation 552.

Monte Carlo Methods (Ch. 29)

Table of Contents Local

Written by Brandon McKinzie

Overview. The aims of Monte Carlo methods are to solve one or both of the following:

1. Generate samples $\{\mathbf{x}^{(r)}\}_{r=1}^R$ from a given probability distribution $\Pr[\mathbf{x}]$.
2. Estimate expectations of function under $\Pr[\mathbf{x}]$, for example

$$\Phi = \mathbb{E}_{\mathbf{x} \sim \Pr[\mathbf{x}]} [\phi(\mathbf{x})] \equiv \int d^N \mathbf{x} \Pr[\mathbf{x}] \phi(\mathbf{x}) \quad (553)$$

where it's assumed that $\Pr[\mathbf{x}]$ is sufficiently complex that we can't evaluate such expectations by exact methods.

Note that we can concentrate on the first problem (sampling), since we can use it to solve the second problem (estimating an expectation) by using the random samples $\{\mathbf{x}^{(r)}\}_{r=1}^R$ to give the estimator

$$\hat{\Phi} \equiv \frac{1}{R} \sum_r \phi(\mathbf{x}^{(r)}) \quad (554)$$

where

$$\mathbb{E}[\hat{\Phi}] = \frac{1}{R} \sum_r \mathbb{E}_{\mathbf{x}^{(r)} \sim \Pr[\mathbf{x}]} [\phi(\mathbf{x}^{(r)})] = \Phi \quad (555)$$

$$\lim_{R \rightarrow \infty} \text{Var}[\hat{\Phi}] = \lim_{R \rightarrow \infty} \frac{\sum_r (\phi(\mathbf{x}^{(r)}) - \Phi)^2}{R-1} = \frac{\sigma^2}{R} \equiv \frac{1}{R} \int d^N \mathbf{x} \Pr[\mathbf{x}] (\phi(\mathbf{x}) - \Phi)^2 \quad (556)$$

Importance Sampling. A generalization of [naively] uniformly sampling \mathbf{x} in order to approximate equation 553. We assume henceforth that we are able to evaluate (for now, a 1-D) $\Pr[\mathbf{x}]$ at any point x at least within a multiplicative constant; thus we can evaluate a function $P^*(x)$ such that $P(x) = P^*(x)/Z$. We assume we have some simpler $Q(x)$, called the **sampler density**, from which we *can* generate samples from and evaluate up to a multiplicative constant, $Q(x) = Q^*(x)/Z_Q$. We construct an approximation for our estimator in equation 554 via sampling from $Q(x)$ and computing:

$$\hat{\Phi} = \frac{\sum_r w_r \phi(x^{(r)})}{\sum_r w_r} \quad \text{where} \quad w_r \equiv \frac{P^*(x^{(r)})}{Q^*(x^{(r)})} \quad (557)$$

Rejection Sampling. In addition to the assumptions in importance sampling, we now also assume that we know the value of a constant c such that

$$\forall x : cQ^*(x) > P^*(x) \quad (558)$$

1. Generate sample x from proposal density $Q(x)$, and evaluate $cQ^*(x)$.
2. Generate a uniformly distributed random variable u from the interval $[0, cQ^*(x)]$.
3. If $u > P^*(x)$, then reject x ; else, accept x and add it to our set of samples $\{x^{(r)}\}$.

Metropolis-Hastings Method. Proposal density Q now depends on the current state $x^{(t)}$.

1. Sample tentative next state x' from proposal density $Q(x'; x^{(t)})$.
2. Compute:

$$a = \frac{P^*(x')}{P^*(x^{(t)})} \frac{Q(x^{(t)}; x')}{Q(x'; x^{(t)})} \quad (559)$$

3. If $a \geq 1$, the new state is accepted. Otherwise, the new state is accepted with probability a .
4. If accepted, we set $x^{(t+1)} = x'$. If rejected, we set $x^{(t+1)} = x^{(t)}$.

Variational Methods (Ch. 33)

Table of Contents Local

Written by Brandon McKinzie

Probability Distributions in Statistical Physics¹⁵⁹. Consider the common situation below, where the state vector $\mathbf{x} \in \{-1, +1\}^N$:

$$\Pr[\mathbf{x} \mid \beta, \mathbf{J}] = \frac{e^{-\beta E(\mathbf{x}; \mathbf{J})}}{Z(\beta, \mathbf{J})} \quad (560)$$

$$E(\mathbf{x}; \mathbf{J}) = -\frac{1}{2} \sum_{m,n} J_{mn} x_m x_n - \sum_n h_n x_n \quad (561)$$

$$Z(\beta, \mathbf{J}) = \sum_{\mathbf{x}} e^{-\beta E(\mathbf{x}; \mathbf{J})} \quad (562)$$

Note that evaluating $E(\mathbf{x}; \mathbf{J})$ for a given \mathbf{x} takes polynomial time in the number of spins N , and evaluating Z is $\mathcal{O}(2^N)$. Variational free energy minimization is a method for approximating the complex distribution $P(\mathbf{x})$ by a simpler ensemble $Q(\mathbf{x}; \boldsymbol{\theta})$ parameterized by adjustable $\boldsymbol{\theta}$.

Variational Free Energy. How do we find/evaluate Q ? The objective function chosen to measure the quality of the approximation is the **variational free energy**, $\tilde{F}(\boldsymbol{\theta})$:

$$\beta \tilde{F}(\boldsymbol{\theta}) = \sum_{\mathbf{x}} Q(\mathbf{x}; \boldsymbol{\theta}) \ln \frac{Q(\mathbf{x}; \boldsymbol{\theta})}{e^{-\beta E(\mathbf{x}; \mathbf{J})}} \quad (563)$$

$$= \beta \mathbb{E}_{\mathbf{x} \sim Q} [E(\mathbf{x}; \mathbf{J})] - H(Q) \quad (564)$$

$$= D_{KL}(Q \parallel P) + \beta F \quad (565)$$

where $F \triangleq -\ln Z(\beta, \mathbf{J})$ is the true free energy. It's not immediately clear why this approximation Q is more tractable than P – for that we turn to an example below.

¹⁵⁹Yay!

Variational Free Energy Minimization for Spin Systems. For the system with energy function given in equation 561, consider the *separable* approximating distribution,

$$Q(\mathbf{x}; \mathbf{a}) = \frac{e^{\sum_n a_n x_n}}{Z_Q} = \frac{\prod_n e^{a_n x_n}}{\sum_{\mathbf{x}'} \prod_{n'} e^{a_{n'} x'_{n'}}} = \frac{\prod_n e^{a_n x_n}}{\sum_{x'_1} e^{a_1 x'_1} \sum_{x'_2} e^{a_2 x'_2} \dots \sum_{x'_n} e^{a_n x'_n}} \quad (566)$$

To compute \tilde{F} , we must compute the mean energy and entropy under Q .

→ **Mean Energy.** Given our definition of Q , and the fact that each $x_n = \pm 1$, the mean value of any x_n is $\bar{x} = \tanh(a_n) = 2q_n - 1$, where $q_n \equiv Q(x_n = 1)$.

$$\mathbb{E}_{\mathbf{x} \sim Q} [E(\mathbf{x}; \mathbf{J})] = \sum_{\mathbf{x}} Q(\mathbf{x}; \mathbf{a}) \left[-\frac{1}{2} \sum_{m,n} J_{mn} x_m x_n - \sum_n h_n x_n \right] \quad (567)$$

$$= -\frac{1}{2} \sum_{m,n} J_{mn} \bar{x}_m \bar{x}_n - \sum_n h_n \bar{x}_n \quad (568)$$

→ **Entropy.** Recall that if $Q(\mathbf{x}; \mathbf{a}) = \prod_n Q(x_n; a_n)$ (i.e. Q is separable), that $H(\mathbf{x}) = \sum_n H(x_n)$, so we have

$$H_Q(\mathbf{x}) = \sum_{\mathbf{x}} Q(\mathbf{x}; \mathbf{a}) \ln \frac{1}{Q(\mathbf{x}; \mathbf{a})} \quad (569)$$

$$= \sum_n Q(x_n = 1; a_n) \ln \frac{1}{Q(x_n = 1; a_n)} + Q(x_n = 0; a_n) \ln \frac{1}{Q(x_n = 0; a_n)} \quad (570)$$

$$= \sum_n q_n \ln \frac{1}{q_n} + (1 - q_n) \ln \frac{1}{1 - q_n} \quad (571)$$

So the variational free energy is given by

$$\beta \tilde{F}(\mathbf{a}) = \beta \mathbb{E}_{\mathbf{x} \sim Q} [E(\mathbf{x}; \mathbf{J})] - H_Q(\mathbf{x}) \quad (572)$$

$$= \beta \left(-\frac{1}{2} \sum_{m,n} J_{mn} \bar{x}_m \bar{x}_n - \sum_n h_n \bar{x}_n \right) - \left(\sum_n q_n \ln \frac{1}{q_n} + (1 - q_n) \ln \frac{1}{1 - q_n} \right) \quad (573)$$

Remember, our goal is to find parameters \mathbf{a} that minimize $\tilde{F}(\mathbf{a})$:

$$\beta \frac{\partial \tilde{F}}{\partial a_m} = 2 \left(\frac{\partial q_m}{\partial a_m} \right) \left[-\beta \left(\sum_n J_{mn} \bar{x}_n + h_m \right) + a_m \right] \quad (574)$$

which, when set to zero, yields

$$a_m \leftarrow \beta \left(\sum_n J_{mn} \bar{x}_n + h_m \right) \quad (575)$$

$$\bar{x}_n = \tanh(a_n) \quad (576)$$

BAYESIAN DATA ANALYSIS

CONTENTS

9.1	Probability and Inference (Ch. 1)	230
9.2	Single-Parameter Models (Ch. 2)	231
9.3	Asymptotics and Connections to Non-Bayesian Approaches (Ch. 4)	234

Note: only including some useful tidbits for quick reference.

Means and Variances of Conditional Distributions.

$$\mathbb{E}[u] = \int dv \Pr[v] \int du \Pr[u | v] \quad (577)$$

$$= \mathbb{E}_v [\mathbb{E}_u [u | v]] \quad (578)$$

$$\text{Var}[u] = \mathbb{E}[u^2] - \mathbb{E}[u]^2 \quad (579)$$

$$= \mathbb{E}_v [\mathbb{E}_u [u^2 | v]] - \mathbb{E}_v [\mathbb{E}_u [u | v]]^2 \quad (580)$$

$$= \mathbb{E}_v [\mathbb{E}_u [u^2 | v] - \mathbb{E}_u [u | v]^2] + \mathbb{E}_v [\mathbb{E}_u [u | v]^2] - \mathbb{E}_v [\mathbb{E}_u [u | v]]^2 \quad (581)$$

$$= \mathbb{E}_v [\text{Var}[u | v]] + \text{Var}[\mathbb{E}_u [u | v]] \quad (582)$$

Transformation of Variables.

$$\Pr_v[v] = |J| \Pr_u[f^{-1}(v)] \quad (583)$$

$$J_{i,j} = \frac{\partial u}{\partial v} = \frac{\partial f^{-1}(v)}{\partial v}$$

where u and v have the same dimensionality, and $|J|$ is the determinant of the Jacobian of the transformation $u = f^{-1}(v)$. When working with parameters defined on the open unit interval, $(0, 1)$, we often use the logistic transformation:

$$\text{logit}(u) = \log\left(\frac{u}{1-u}\right) \quad (584)$$

$$\text{logit}^{-1}(v) = \log\left(\frac{e^v}{1+e^{-v}}\right) \quad (585)$$

Single-Parameter Models (Ch. 2)

Table of Contents Local

Written by Brandon McKinzie

Posterior as compromise between data and prior information. Intuitively, the prior and posterior distributions over θ should have some general relationship showing how the process of observing data y updates our distribution on θ . We can use the identities from the previous chapter to see that

$$\mathbb{E}[\theta] = \mathbb{E}_y[\mathbb{E}_\theta[\theta | y]] \quad (586)$$

$$\text{Var}[\theta] = \mathbb{E}_y[\text{Var}[\theta | y]] + \text{Var}[\mathbb{E}_\theta[\theta | y]] \quad (587)$$

where:

- Equation 586 states the obvious: our prior expectation for θ is the expectation, taken over the distribution of possible data, of the posterior expectation.
- Equation 587 states: *the posterior variance is on average smaller than the prior variance*, by an amount that depends on the variation [in posterior means] over the distribution of possible data. Stated another way: we can reduce our uncertainty with regard to θ by larger amounts with models whose [expected] posteriors are strongly informed by the data.

Informative Priors. Let x denote the number of successes in n trials. We want to find the posterior distribution over θ , the probability of success in a single trial. Then the likelihood is given by

$$\Pr[x | \theta] = \binom{n}{x} \theta^x (1 - \theta)^{n-x} \quad (588)$$

If we define our prior, $\Pr[\theta]$, to have the same parametric form (form of θ -dependence), then our posterior density $\Pr[\theta | x]$ will *also* have this form. For our example, we can define the prior and derive the posterior as follows:

$$\Pr[\theta] \propto \theta^{\alpha-1} (1 - \theta)^{\beta-1} \quad (589)$$

$$\Pr[\theta | x] \propto \theta^{x+\alpha-1} (1 - \theta)^{n-x+\beta-1} \quad (590)$$

The property that the posterior distribution follows the same parametric form as the prior distribution is called **conjugacy**.

Conjugacy, families, and sufficient statistics. Formally, if \mathcal{F} is a class of sampling distributions $\{\Pr_i[y | \theta]\}$, and \mathcal{P} is a class of prior distributions, $\{\Pr_j[\theta]\}$, then the class \mathcal{P} is **conjugate** for \mathcal{F} if

$$\forall \Pr[\cdot | \theta] \in \mathcal{F}, \Pr[\cdot] \in \mathcal{P} : \Pr[\theta | y] \in \mathcal{P} \quad (591)$$

Probability distributions that belong to an **exponential family** have natural conjugate prior distributions. The class \mathcal{F} is an exponential family if all its members have the form,

$$\Pr[y_i | \theta] = f(y_i)g(\theta)e^{\phi(\theta)^T u(y_i)} \quad (592)$$

$$\Pr[y | \theta] = \left(\prod_{i=1}^n f(y_i) \right) g(\theta)^n \exp \left(\phi(\theta)^T \sum_{i=1}^n u(y_i) \right) \quad (593)$$

$$\propto g(\theta)^n \exp \left(\phi(\theta)^T t(y) \right) \quad (594)$$

where

- $y = (y_1, \dots, y_n)$ denotes n iid observations.
- $\phi(\theta)$ is called the **natural parameter** of \mathcal{F} .
- $t(y) = \sum_{i=1}^n u(y_i)$ is said to be a **sufficient statistic** for θ , because the likelihood for θ depends on the data y only through the value of $t(y)$.

Normal distribution with known variance¹⁶⁰. The family of conjugate prior densities for the Gaussian likelihood looks like

$$\Pr[\theta] \propto e^{A\theta^2 + B\theta + C} \quad (596)$$

We parameterize this family as

$$\Pr[\theta] \propto e^{-\frac{1}{2\tau_0^2}(\theta - \mu_0)^2} \quad (597)$$

After some basic arithmetic/substitutions, we can find that this gives

$$\Pr[\theta | y] \propto \exp \left(-\frac{1}{2\tau_1^2}(\theta - \mu_1)^2 \right) \quad (598)$$

$$\mu_1 = \frac{\frac{1}{\tau_0^2}\mu_0 + \frac{1}{\sigma^2}y}{\frac{1}{\tau_0^2} + \frac{1}{\sigma^2}} \quad \text{and} \quad \frac{1}{\tau_1^2} = \frac{1}{\tau_0^2} + \frac{1}{\sigma^2} \quad (599)$$

where we see the interesting result that the posterior **precision** (inverse of variance) equals the prior precision plus the data precision. We can see the posterior mean μ_1 expressed as a weighted average of the prior mean and the observed value¹⁶¹ y , with weights proportional to the precisions.

¹⁶⁰The following will be useful to remember:

$$\int_{-\infty}^{\infty} e^{-ax^2 + bx + c} dx = \frac{\pi}{a} e^{\frac{b^2}{4a} + c} \quad (595)$$

¹⁶¹For now, we're considering the single data point case.

Normal distribution with unknown variance. Now, we assume the mean θ is known, and the variance σ^2 is unknown. The likelihood for a vector $y = (y_1, \dots, y_n)$ of n iid observations is

$$\Pr [y \mid \sigma^2] \propto (\sigma^2)^{-n/2} \exp \left(-\frac{n}{2\sigma^2} v \right) \quad (600)$$

$$v := \frac{1}{n} \sum_{i=1}^n (y_i - \theta)^2 \quad (601)$$

where v is the sufficient statistic. The corresponding conjugate prior density is the inverse-gamma,

$$\Pr [\sigma^2] \propto (\sigma^2)^{-(\alpha+1)} e^{-\beta/\sigma^2} \quad (602)$$

Jeffrey's Invariance Principle. An approach for defining noninformative prior distributions. Let $\phi = h(\theta)$, where the function h is one-to-one. By transformation of variables,

$$p(\phi) = p(\theta) \left| \frac{d\theta}{d\phi} \right| = p(\theta) |h'(\theta)|^{-1} \quad (603)$$

Jeffrey's principle is to basically take the above equation as a true equivalence – there is no difference between finding $p(\theta)$ and applying the equation above [to get $p(\phi)$] and directly finding $p(\phi)$.

Let $J(\theta)$ denote the **Fisher information** for θ , defined as

$$J(\theta) = \mathbb{E} \left[\left(\frac{d \log \Pr [y \mid \theta]}{d\theta} \right)^2 \mid \theta \right] = -\mathbb{E} \left[\frac{d^2 \log \Pr [y \mid \theta]}{d\theta^2} \mid \theta \right] \quad (604)$$

Jeffrey's prior model defines the noninformative prior density as $\Pr [\theta] \propto [J(\theta)]^{1/2}$. We can work out that this model is indeed invariant to parameterization¹⁶².

¹⁶²Evaluate $J(\phi)$ at $\theta = h^{-1}(\phi)$. You should find that $J(\phi)^{1/2} = J(\theta)^{1/2} \left| \frac{d\theta}{d\phi} \right|$

Asymptotics and Connections to Non-Bayesian Approaches (Ch. 4)

Table of Contents Local

Written by Brandon McKinzie

Normal Approximations to the Posterior Distribution. If the posterior $\Pr[\theta | y]$ is unimodal and roughly symmetric, it can be convenient to approximate it by a normal distribution. Here we'll consider a quadratic approximation via the Taylor series expansion up to second-order,

$$\log \Pr[\theta | y] = \log \Pr[\hat{\theta} | y] + \frac{1}{2}(\theta - \hat{\theta})^T \left[\frac{d^2}{d\theta^2} \log \Pr[\theta | y] \right]_{\theta=\hat{\theta}} (\theta - \hat{\theta}) \quad (605)$$

where $\hat{\theta}$ is the posterior mode. The remainder terms of higher order fade in importance relative to the quadratic term when θ is close to $\hat{\theta}$ and n is large. We'd like to cast this into a normal distribution. First, let

$$I(\theta) \triangleq -\frac{d^2}{d\theta^2} \log \Pr[\theta | y] \quad (606)$$

which we will refer to as the **observed information**. We can then rewrite our approximation as¹⁶³

$$\Pr[\theta | y] \approx \mathcal{N}(\hat{\theta}, [I(\hat{\theta})]^{-1}) \quad (609)$$

Under the normal approximation, the posterior distribution is summarized by its mode, $\hat{\theta}$, and the curvature of the posterior density, $I(\hat{\theta})$; that is, asymptotically, these are sufficient statistics.

¹⁶³I also found it helpful to explicitly write the substitution after raising eq 605 by power of e (all logs are assumed natural logs)

$$\Pr[\theta | y] = e^{\log \Pr[\hat{\theta} | y] - \frac{1}{2}(\theta - \hat{\theta})^T [I(\hat{\theta})]^{-1} (\theta - \hat{\theta})} \quad (607)$$

$$= \Pr[\hat{\theta} | y] e^{-\frac{1}{2}(\theta - \hat{\theta})^T [I(\hat{\theta})]^{-1} (\theta - \hat{\theta})} \quad (608)$$

Example. Let y_1, \dots, y_n be independent observations from $\mathcal{N}(\mu, \sigma^2)$. Define $\theta := (\mu, \log \sigma)$ as the parameters of interest, and assume a uniform prior¹⁶⁴ $\Pr[\theta]$. Recall that (equation 3.2 in textbook)

$$\Pr[\theta = (\mu, \log \sigma) \mid y] \propto \sigma^{-(n+2)} \exp\left(-\frac{1}{2\sigma^2} \sum_{i=1}^n (y_i - \mu)^2\right) \quad (611)$$

Note that $\sum_{i=1}^n (y_i - \bar{y}) = 0$

$$= \sigma^{-(n+2)} \exp\left(-\frac{1}{2\sigma^2} \left[n(\bar{y} - \mu)^2 + \sum_{i=1}^n (y_i - \bar{y})^2 \right]\right) \quad (612)$$

$$= \sigma^{-(n+2)} \exp\left(-\frac{1}{2\sigma^2} \left[n(\bar{y} - \mu)^2 + (n-1)s^2 \right]\right) \quad (613)$$

where $s^2 = \frac{1}{n-1} \sum_{i=1}^n (y_i - \bar{y})^2$ is the sample variance of the y_i 's. The sufficient statistics are \bar{y} and s^2 . To construct the approximation, we need the second derivatives of the log posterior density¹⁶⁵,

$$\log \Pr[\mu, \log \sigma \mid y] = \text{const} - n \log \sigma - \frac{1}{2\sigma^2} \left(n(\bar{y} - \mu)^2 + (n-1)s^2 \right) \quad (614)$$

in order to compute $I(\hat{\theta})$. After computing first derivatives, we find that the posterior mode is

$$\hat{\theta} = (\hat{\mu}, \log \hat{\sigma}) = \left(\bar{y}, \log \left(\sqrt{\frac{n-1}{n}} s \right) \right) \quad (615)$$

We then compute second derivatives and evaluate at $\theta = \hat{\theta}$ to obtain $I(\hat{\theta})$. Combining all this into the final result:

$$\Pr[\mu, \log \sigma \mid y] \approx \mathcal{N}(\hat{\theta}, [I(\hat{\theta})]^{-1}) \quad (616)$$

$$= \mathcal{N}\left(\begin{pmatrix} \bar{y} \\ \log \hat{\sigma} \end{pmatrix}, \begin{pmatrix} \hat{\sigma}^2/n & 0 \\ 0 & 1/(2n) \end{pmatrix}\right) \quad (617)$$

where $\hat{\sigma}^2/n$ is the variance along the μ dimension, and $1/(2n)$ is the variance along the $\log \sigma$ direction. This example was just meant to illustrate, with a simple case, how we work through constructing the approximate normal distribution.

¹⁶⁴Recall from Ch 3.2 that the uniform prior on $\mu, \log \sigma$ is

$$\Pr[\mu, \sigma^2] \propto (\sigma^2)^{-1} \quad (610)$$

where we continue to assume μ is uniform in $[0, 1]$ for some reason.

¹⁶⁵I'm not sure why $n+2$ has seemingly turned into n .

Large-Sample Theory. Asymptotic normality of the posterior distribution: as more data arrives from the same underlying distribution $f(y)$, the posterior distribution of the *parameter vector* θ approaches multivariate normality, even if the true distribution of the data is not within the parametric family under consideration.

Suppose the data are modeled by a parametric family, $\Pr[y \mid \theta]$, with a prior distribution $\Pr[\theta]$. If the true distribution, $f(y)$, is included in the parametric family (i.e. if $\exists \theta_0 : f(y) = \Pr[y \mid \theta_0]$), then it's also true that **consistency** holds¹⁶⁶: $\Pr[\theta \mid y]$ converges to a point mass at the true parameter value, θ_0 as $n \rightarrow \infty$.

¹⁶⁶So, what if the true $f(y)$ is *not* included in the parametric family? In that case, there is no longer a true value θ_0 , but its role in the theoretical result is replaced by a value θ_0 that makes the model distribution $\Pr[y \mid \theta]$ closest to the true distribution $f(y)$, in a technical sense involving the **Kullback-Leibler divergence**.

BLOGS

CONTENTS


10.1 Conv Nets: A Modular Perspective	238
10.2 Understanding Convolutions	239
10.3 Deep Reinforcement Learning	241
10.4 Deep Learning for Chatbots (WildML)	243
10.5 Attentional Interfaces – Neural Perspective	245

Conv Nets: A Modular Perspective

[Table of Contents](#) [Local](#)*Written by Brandon McKinzie*

From this post on Colah's Blog.

The title is inspired by the following figure. Colah mentions how groups of neurons, like A , that appear in multiple places are sometimes called **modules**, and networks that use them are sometimes called modular neural networks. You can feed the output of one convolutional layer into another. With each layer, the network can detect higher-level, more abstract features.



- Function of the A neurons: compute certain *features*.
- Max pooling layers: kind of “zoom out”. They allow later convolutional layers to work on larger sections of the data. They also make us invariant to some very small transformations of the data.


Understanding Convolutions

[Table of Contents](#) [Local](#)
Written by *Brandon McKinzie*

From Colah's Blog.

Ball-Dropping Example. The posed problem:

Imagine we drop a ball from some height onto the ground, where it only has one dimension of motion. How likely is it that a ball will go a distance c if you drop it and then drop it again from above the point at which it landed?



From basic probability, we know the result is a sum over possible outcomes, constrained by $a + b = c$. It turns out this is actually the definition of the convolution of f and g .

$$\Pr(a + b = c) = \sum_{a+b=c} f(a) \cdot g(b) \quad (618)$$


$$(f * g)(c) = \sum_{a+b=c} f(a) \cdot g(b) \quad (619)$$

$$= \sum_a f(a) \cdot g(c - a) \quad (620)$$

Visualizing Convolutions. Keeping the same example in the back of our heads, consider a few interesting facts.

- **Flipping directions.** If $f(x)$ yields the probability of landing a distance x away from where it was dropped, what about the probability that it was dropped a distance x from where it *landed*? Apparently¹⁶⁷ it is $f(-x)$.
- Above is a visualization of one term in the summation of $(f * g)(c)$. It is meant to show how we can move the bottom around to think about evaluating the convolution for different c values.

¹⁶⁷Not entirely sold on the generalization of this, or even how true it is here.



We can relate these ideas to image recognition. Below are two common kernels used to convolve images with.

0	0	0	0	0
0	1/9	1/9	1/9	0
0	1/9	1/9	1/9	0
0	1/9	1/9	1/9	0
0	0	0	0	0

0	0	0	0	0
0	0	0	0	0
0	-1	1	0	0
0	0	0	0	0
0	0	0	0	0

On the left is a kernel for *blurring* images, accomplished by taking simple averages. On the right is a kernel for *edge detection*, accomplished by taking the difference between two pixels, which will be largest at edges, and essentially zero for similar pixels.

Deep Reinforcement Learning

Table of Contents Local

Written by Brandon McKinzie

Link to tutorial – Part I of “Demystifying deep reinforcement learning.”

Reinforcement Learning. Vulnerable to the *credit assignment problem* - i.e. unsure which of the preceding actions was responsible for getting some reward and to what extent. Also need to address the famous *explore-exploit dilemma* when deciding what strategies to use.

Markov Decision Process. Most common method for representing a reinforcement problem. MDPs consist of states, actions, and rewards. Total reward is sum of current (includes previous) and *discounted* future rewards:

$$R_t = r_t \gamma (r_{t+1} + \gamma (r_{t+2} + \dots)) = r_t + \gamma R_{t+1} \quad (621)$$

Q - learning. Define function $Q(s, a)$ to be best possible score at end of game after performing action a in state s ; the “quality” of an action from a given state. The recursive definition of Q (for one transition) is given below in the *Bellman equation*.

$$Q(s, a) = r + \gamma \max_{a'} Q(s', a')$$

and updates are computed with a learning rate α as

$$Q(s_t, a_t) = (1 - \alpha) \cdot Q(s_{t-1}, a_{t-1}) + \alpha \cdot (r + \gamma \max_{a'} Q(s'_{t+1}, a'_{t+1}))$$

Deep Q Network. Deep learning can take deal with issues related to prohibitively large state spaces. The implementation chosen by DeepMind was to represent the Q-function with a neural network, with the states (pixels) as the input and Q-values as output, where the number of output neurons is the number of possible actions from the input state. We can optimize with simple squared loss:

$$L = \frac{1}{2} \left[\underbrace{r + \max_{a'} Q(s', a')}_{\text{target}} - \underbrace{Q(s, a)}_{\text{prediction}} \right]^2$$

and our algorithm from some state s becomes

1. **First forward pass** from s to get all predicted Q-values for each possible action. Choose action corresponding to max output, leading to next s' .

2. **Second forward pass** from s' and again compute $\max_{a'} Q(s', a')$.
3. **Set target output** for each action a' from s' . For the action corresponding to max (from step 2) set its target as $r + \gamma \max_{a'} Q(s', a')$, and for all other actions set target to same as originally returned from step 1, making the error 0 for those outputs. (Interpret as update to our guess for the best Q-value, and keep the others the same.)
4. **Update weights** using backprop.

Experience Replay. This the most important trick for helping convergence of Q-values when approximating with non-linear functions. During gameplay all the experience $\langle s, a, r, s' \rangle$ are stored in a replay memory. When training the network, random minibatches from the replay memory are used instead of the most recent transition.

Exploration. One could say that initializing the Q-values randomly and then picking the max is essentially a form of exploitation. However, this type of exploration is *greedy*, which can be tamed/fixed with **ϵ -greedy exploration**. This incorporates a degree of randomness when choosing next action at *all* time-steps, determined by probability ϵ that we choose the next action randomly. For example, DeepMind decreases ϵ over time from 1 to 0.1.

Deep Q-Learning Algorithm.

```

initialize replay memory D
initialize action-value function Q with random weights
observe initial state s
repeat
    select an action a
        with probability  $\epsilon$  select a random action
        otherwise select  $a = \operatorname{argmax}_{a'} Q(s, a')$ 
    carry out action a
    observe reward r and new state  $s'$ 
    store experience  $\langle s, a, r, s' \rangle$  in replay memory D

    sample random transitions  $\langle ss, aa, rr, ss' \rangle$  from replay memory D
    calculate target for each minibatch transition
        if  $ss'$  is terminal state then  $tt = rr$ 
        otherwise  $tt = rr + \gamma \max_{a'} Q(ss', aa')$ 
    train the Q network using  $(tt - Q(ss, aa))^2$  as loss

     $s = s'$ 
until terminated

```

Deep Learning for Chatbots (WildML)

Table of Contents Local


Written by Brandon McKinzie

Overview.

- **Model.** Implementing a retrieval-based model. Input: conversation/context c . Output: response r .
- **Data.** Ubuntu Dialog Corpus (UDC). 1 million examples of form (context, utterance, label). The label can be 1 (utterance was actual response to the context) or a 0 (utterance chosen randomly). Using NLTK, the data has been . . .
 - **Tokenized.** dividing strings into lists of substrings.
 - **Stemmed.** **IDK**
 - **Lemmatized.** **IDK**

The test/validation set consists (context, ground-truth utterance, [9 distractors (incorrect utterances)]). The distractors are picked at random¹⁶⁸

Dual-Encoder LSTM.



1. **Inputs.** Both the context and the response text are split by words, and each word is embedded into a vector and fed into the same RNN.
2. **Prediction.** Multiply the [vector representation ("meaning")] c with param matrix M to predict some response r' .
3. **Evaluation.** Measure similarity of predicted r' to actual r via simple dot product. Feed this into sigmoid to obtain a probability [of r' being the correct response]. Use (binary) cross-entropy for loss function:

$$L = -y \cdot \ln(y') - (1 - y) \cdot \ln(1 - y') \quad (622)$$

where y' is the predicted probability that r' is correct response r , and $y \in \{0, 1\}$ is the true label for the context-response pair (c, r) .

¹⁶⁸Better example/approach: Google's Smart Reply uses clustering techniques to come up with a set of possible responses.

Data Pre-Processing. Courtesy of WildML, we are given 3 files after preprocessing: `train.tfrecords`, `validation.tfrecords`, and `test.tfrecords`, which use TensorFlow's 'Example' format. Each Example consists of . . .

- `context`: Sequence of word ids.
- `context_len`: length of the aforementioned sequence.
- `utterance`: seq of word ids representing utterance (response).
- `utterance_len`.
- `label`: only in training data. 0 or 1.
- `distractor__[N]`: Only in test/validation. N ranges from 0 to 8. Seq of word ids reppin the distractor utterance.
- `distractor__[N]_len`.


Attentional Interfaces – Neural Perspective

Table of Contents Local

Written by Brandon McKinzie

[\[Link to article\]](#)

Attention Mechanism. Below is a close-up view/diagram of an attention layer. Technically, it only corresponds to a single time step i ; we are using the previous decoder state s_{i-1} to compute the i th context vector c_i which will be fed as an input to the decoder for step i .



For convenience, I'll rewrite the familiar equations for computing quantities at some step i .

$$\text{[decoder state]} \quad s_i = f(s_{i-1}, y_{i-1}, c_i) \quad (623)$$

$$\text{[context vect]} \quad c_i = \sum_{j=1}^{T_x} \alpha_{ij} h_j \quad (624)$$

$$\text{[annotation weights]} \quad \alpha_{ij} = \frac{\exp(e_{ij})}{\sum_{k=1}^{T_x} \exp(e_{ik})} \quad (625)$$

$$e_{ij} = a(s_{i-1}, h_j) \quad (626)$$

Now we can see just how simple this really is. Recall that Bahdanau *et al.*, 2015 use the wording: “ e_{ij} is an alignment model which scores how well the inputs around position j and the output at position i match.” But we can see an example implementation of an alignment model above: the \tanh function (that’s it).

APPENDIX A - COMMON DISTRIBUTIONS

Distribution	Density Function	Notation
Gamma	$p(\theta) = \frac{\beta^\alpha}{\Gamma(\alpha)} \theta^{\alpha-1} e^{-\beta\theta}, \quad \theta > 0$	$\theta \sim \text{Gamma}(\alpha, \beta)$
Inverse-gamma	$p(\theta) = \frac{\beta^\alpha}{\Gamma(\alpha)} \theta^{-\alpha-1} e^{-\beta/\theta}, \quad \theta > 0$	$\theta \sim \text{Inv-gamma}(\alpha, \beta)$
Chi-Square	$p(\theta) = \frac{2^{-\nu/2}}{\Gamma(\nu/2)} \theta^{\nu/2-1} e^{-\theta/2}, \quad \theta > 0$	$\theta \sim \chi_\nu^2$
Inverse-chi-square	$p(\theta) = \frac{2^{-\nu/2}}{\Gamma(\nu/2)} \theta^{-\nu/2-1} e^{-\frac{1}{2\theta}}, \quad \theta > 0$	$\theta \sim \text{Inv-}\chi_\nu^2$

APPENDIX B - QUESTIONS AND STUFF I ALWAYS FORGET

Questions:

- **Q:** In general, how can one tell if a matrix \mathbf{A} has an eigenvalue decomposition? [insert more conceptual matrix-related questions here . . .]
- **Q:** Let \mathbf{A} be real-symmetric. What can we say about \mathbf{A} ?
 - Proof that eigendecomposition $\mathbf{A} = \mathbf{Q}\mathbf{\Lambda}\mathbf{Q}^T$ exists: Wow this is apparently quite hard to prove according to many online sources. Guess I don't feel so bad now that it wasn't (and still isn't) obvious.
 - Eigendecomposition not unique. This is apparently because two or more eigenvectors may have same eigenvalue.

This is the principal axis theorem: if \mathbf{A} symmetric, then orthonorm basis of e-vects exists.

Stuff I Forget:

- Existence of eigenvalues/eigenvectors. Let $\mathbf{A} \in \mathbb{R}^{n \times n}$.
 - λ is an eigenvalue of \mathbf{A} iff it satisfies $\det(\lambda\mathbf{I} - \mathbf{A}) = 0$. Why? Because it is an equivalent statement as requiring that $(\lambda\mathbf{I} - \mathbf{A})\mathbf{x} = 0$ has a nonzero solution for \mathbf{x} .
 - The following statements are equivalent:
 - * \mathbf{A} is diagonalizable.
 - * \mathbf{A} has n linearly independent eigenvectors.
 - The **eigenspace** of \mathbf{A} corresponding to λ is the solution space of the homogeneous system $(\lambda\mathbf{I} - \mathbf{A})\mathbf{x} = 0$.
 - \mathbf{A} has at most n distinct eigenvalues.
- Diagonalizability notes from 5.2 of advanced linear alg. book (261). Recall that \mathbf{A} is defined to be diagonalizable if and only if there exists an ordered basis β for the space consisting of eigenvectors of \mathbf{A} .
 - If the standard way of finding eigenvalues leads to k distinct λ_i , then the corresponding set of k eigenvectors v_i are guaranteed to be linearly independent (but might not span the full space).
 - If \mathbf{A} has n linearly independent eigenvectors, then \mathbf{A} is diagonalizable.
 - The characteristic polynomial of any diagonalizable linear operator splits (can be factored into product of linear factors). The **algebraic multiplicity** of an eigenvalue λ is the largest positive integer k for which $(t - \lambda)^k$ is a factor of $f(t)$.
- Logistic regression is called a *linear model* because it classifies based on the linear $\boldsymbol{\theta}^T \mathbf{x}$ before feeding that to the logistic sigmoid function.

Most info here comes from chapter 5 of your "Elementary Linear Algebra" textbook (around pg305)

Recall that a linear operator is a special case of a linear map where the input space is the same as the output space.

Coreference Resolution. The task of finding all expressions that refer to the same entity in a text. It is an important step for a lot of higher level NLP tasks that involve natural language understanding such as document summarization, question answering, and information extraction.

Conditional Random Field (CRF). A type of discriminative undirected probabilistic graphical model¹⁶⁹. The CRF defines the probability, $p(\ell | s)$, of a given sentence s being described with the [vector] label ℓ as follows:

See this post by Edwin Chen for a great CRF intro.

$$\text{score}(\ell | s) = \sum_{j=1}^m \sum_{i=1}^n \lambda_j f_j(s, i, \ell_i, \ell_{i-1}) \quad (627)$$

$$p(\ell | s) = \frac{e^{\text{score}(\ell|s)}}{\sum_{\ell'} e^{\text{score}(\ell'|s)}} \quad (628)$$

A label element, ℓ_i , could be the POS label for the i th word. The sum in the denominator of (2) is over all possible label combinations for s .

where m is the number of features (e.g. different POS labels) and n is the number of words/tokens in sentence s . So, where's the graph part, you may ask? It can be inferred from the above equation. I like to imagine a picture where I have my set of label nodes \mathbf{Y} and set of word nodes \mathbf{X} for the sentence s . Note that all nodes are modeled as random variables. The built-in assumption of the CRF is that the value of \mathbf{Y}_i , the label for the i th word in s , is conditionally independent of the other labels \mathbf{Y}_j *given* the previous label \mathbf{Y}_{i-1} .

- Often used for labeling or parsing of sequential data, e.g. POS tagging, NER.

Question/Answer.

- **Q:** How is this different from logistic regression?
 - **A:** Indeed, this would be identical to logistic regression if we viewed all possible label [vector] combinations ℓ as black boxes with no internal structure (i.e. replace their vector representations with integer ids). The difference is that we are exploiting the sequential interdependencies of the structure of ℓ – we are explicitly modeling the labels as *sequential*.
- **Q:** How is this different from HMMs?
 - **A:** Basically, we can build a CRF equivalent to any HMM, but CRFs are able to define more features and they can have arbitrary weights. Let's look into this in more detail (TODO).
- **Q:** How is this different from basically any neural network-based text classifier? They all accept sequences and output a label prediction.
 - **A:** .

¹⁶⁹From browsing the literature, it appears that *by far* the most important CRF for sequential labels with NLP is the **linear-chain CRF**, so much so that often authors will say CRF when they really mean linear-chain CRF (and henceforth so will I)