

CONTENTS

1	Lectures	2
1.1	Machine Learning II	3
1.2	Machine Learning III	5
1.3	Search I	7
1.4	Search II	11
1.5	Markov Decision Processes	14
1.6	Reinforcement Learning	17
1.7	Games I	20
1.7.1	Adversarial Search (AIMA Ch. 5)	23
1.8	Games II	26
1.9	CSPs I	30
1.9.1	Constraint Propagation: Inference in CSPs (AIMA Ch. 6.2)	32
1.9.2	Event Scheduling (Section 5/10/2019)	33
1.10	CSPs II	34
1.11	Bayesian Networks I	37
1.12	Bayesian Networks II	39
1.13	Bayesian Networks III	41
2	Review	43
2.1	Discrete Math and Probability	44
2.2	Course Synthesis	45
3	Learning from Mistakes	48
3.1	Homework 1: Foundations	49

LECTURES

CONTENTS

1.1	Machine Learning II	3
1.2	Machine Learning III	5
1.3	Search I	7
1.4	Search II	11
1.5	Markov Decision Processes	14
1.6	Reinforcement Learning	17
1.7	Games I	20
1.7.1	Adversarial Search (AIMA Ch. 5)	23
1.8	Games II	26
1.9	CSPs I	30
1.9.1	Constraint Propagation: Inference in CSPs (AIMA Ch. 6.2)	32
1.9.2	Event Scheduling (Section 5/10/2019)	33
1.10	CSPs II	34
1.11	Bayesian Networks I	37
1.12	Bayesian Networks II	39
1.13	Bayesian Networks III	41

Machine Learning II

Roadmap.

- Features.
- Neural Networks.
- Gradients without tears.
- Nearest Neighbors.

Linear Classifiers. Discussion of feature maps $\phi(x)$ that can be used for making a non-linear decision boundary [52:00].

Neural Networks.

Predicting Car Collision [58:00]

Input: position of two oncoming cars $x = [x_1, x_2]$.

Output: Whether safe ($y = +1$) or collide $y = -1$.

We are told that the true function is

$$y^* = \text{sign}(|x_1 - x_2| - 1)$$

- **Insight:** In (x_1, x_2) space, there is a “band” (diagonal to the right centered about origin) representing “safe” [58:30]. Recognize that these can be thought of as two linear decision boundaries.
- **Decompose** into subproblems^a.

$$h_1 = \mathbb{1}[x_1 - x_2 \geq 1] \tag{1}$$

$$h_2 = \mathbb{1}[x_2 - x_1 \geq 1] \tag{2}$$

$$y = \text{sign}(h_1 + h_2) \tag{3}$$

- **Learning strategy.** Define $\phi(x) \triangleq [1, x_1, x_2]$.

$$h_1 = \mathbb{1}[\mathbf{v}_1 \cdot \phi(x) \geq 0] \quad \mathbf{v}_1 \triangleq [-1, +1, -1] \tag{4}$$

$$h_2 = \mathbb{1}[\mathbf{v}_2 \cdot \phi(x) \geq 0] \quad \mathbf{v}_2 \triangleq [-1, -1, +1] \tag{5}$$

$$f_{\mathbf{V}, \mathbf{w}}(x) = \text{sign}(\mathbf{w} \cdot \mathbf{h}) \quad \mathbf{w} \triangleq [1, 1] \tag{6}$$


Key idea: joint learning [1:03:42]. Learn both hidden subproblems $\mathbf{V} = (\mathbf{v}_1, \mathbf{v}_2)$ and combination weights $\mathbf{w} = [w_1, w_2]$.

^aHe is defining $\text{sign}(0) \equiv -1$.


Gradients. (Continuing off example) Our parameters are $\{v_1, v_2, w\}$. To learn these parameters, we may consider gradient descent on y (and therefore on h_1 and h_2). Notice that

$$\nabla_{v_1} h_1 = 0 \quad (7)$$

Solution: Redefine $h_1 \triangleq \sigma(v_1 \cdot \phi(x))$.




Gradients without tears [1:14:00]. Gradients of common function building blocks:



Binary classification with hinge loss [1:17:01].

$$L(x, y, w) = \max\{1 - w \cdot \phi(x)y, 0\} \quad (8)$$



Definition: Forward/backward values

Forward: f_i is value for subexpression rooted at i

Backward: $g_i = \frac{\partial \text{out}}{\partial f_i}$ is how f_i influences output

Machine Learning III


Review. Some feature extractor $\phi(\mathbf{x})$.

- Linear predictor score: $\mathbf{w} \cdot \phi(\mathbf{x})$.
- Neural network score: $\sum_{j=1}^k w_j \sigma(\mathbf{v}_j \cdot \phi(\mathbf{x}))$.

Generalization. Goal is to minimize error on unseen future examples.

- Suppose there is some space of all possible predictors f and, within this space, there exists the optimal predictor f^* .
- When we decide how we will do our feature extraction, we constrain the possible predictors we can get to be inside some **hypothesis class** \mathcal{F} [21:11]. Note that there exists a predictor g that is optimal, constrained to being inside \mathcal{F} .
- **Approximation Error:** How good is our hypothesis class \mathcal{F} ?
- **Estimation Error:** How good is our learned predictor \hat{f} *relative to* our hypothesis class \mathcal{F} ?

$$\text{Err}(\hat{f}) - \text{Err}(f^*) = \underbrace{\text{Err}(\hat{f}) - \text{Err}(g)}_{\text{Estimation}} + \underbrace{\text{Err}(g) - \text{Err}(f^*)}_{\text{Approximation}} \quad (9)$$



For linear predictors with weights $\mathbf{w} \in \mathbb{R}^d$, we can **control the size of \mathcal{F}** by ...

- Keeping dimensionality d small.
 - Add/remove features.
 - Forward selection.
 - Boosting.
 - L_1 regularization.
- Keeping the norm $\|\mathbf{w}\|$ small.

Note that **SVMs** = hinge loss + regularization.

How do we choose **hyperparameters** [43:39]? One way is to use a **validation set**.

Unsupervised Learning [1:06:00]. Data has rich **latent structures** that we'd like to learn automatically. Two types of unsupervised learning: **clustering** (e.g. K-means) and **dimensionality reduction** (e.g. PCA).

Clustering

Input: Training set of inputs $\mathcal{D}_{train} = \{x_1, \dots, x_n\}$.

Output: Assignment of each point to a cluster $[z_1, \dots, z_n]$ where $z_i \in \{1, \dots, K\}$.

K-Means Clustering [1:10:35]

Setup: Each cluster $k = 1, \dots, K$ is represented by a **centroid** $\mu_k \in \mathbb{R}^d$. Want each $\phi(x_i)$ close to its assigned centroid μ_{z_i} .

Objective: MSE.

$$L(z, \mu) = \sum_{i=1}^n \|\phi(x_i) - \mu_{z_i}\|^2 \quad (10)$$

Algorithm:

1. For each point x_i , assign:

$$z_i := \arg \min_{k=1, \dots, K} \|\phi(x_i) - \mu_k\|^2 \quad (11)$$

2. For each cluster $k = 1, \dots, K$, assign:

$$\mu_k := \frac{1}{|\{i : z_i = k\}|} \sum_{i: z_i = k} \phi(x_i) \quad (12)$$

Search I

Table of Contents Local

Written by Brandon McKinzie

Running example: farmer has a boat. Needs to get a cabbage, goat, and wolf across river. Goat and cabbage cannot be alone. Wolf and goat cannot be alone. Boat can hold at most two things (including farmer). How he do it tho?

Beyond Reflex [9:00]. Classifiers (reflex) map x to single action y . In **search**, we map x to an *action sequence* (a_1, a_2, \dots) . **Key**: Need to consider consequences of future actions.

Tree search. Root node is start state. Edges are actions. Nodes are states resulting from said action. Branches represent a sequence of actions.

- S_{start} is "FCGW|" (farmer, cabbage, goat, wolf all on left of river).
- Actions: farmer can go to/from river. Actions enumerate what he takes with him and what direction he goes.
- $Cost(s, a)$: ?
- $Succ(s, a)$: returns the state resulting from taking action a from state s .
- $IsEnd(s)$: ?

Transportation example [17:40]

Streets with blocks numbered 1 to n . Walking from s to $s+1$ takes 1 minute. Taking a magic tram from s to $2s$ takes 2 minutes. How to travel from 1 to n in the least time?

You should read this and then try formalizing it as a search problem.

- Actions: either take 1 step forward or take the tram.
- The costs are 1 minute and 2 minutes for the aforementioned actions.

She starts coding it up at [19:30]. Takeaways from coding:

- Separate modeling from inference.
- Modeling: She defines a class with methods `startState`, `isEnd(state)`, `succAndCost(state)`.

Backtracking search [23:10]. Also uses a search tree. Define branching factor b and depth D of tree¹. BS goes down all the way leftmost to leaf. Sounds like **depth-first search** tbh²

- Memory: $\mathcal{O}(D)$ (small). Why? because we only need to remember our past history, which for BS is basically a straight line from leaf back up to root.
- Time: $\mathcal{O}(b^D)$ (huge). In other words, all nodes. Why? because worst-case is that it has to visit every single node.

¹Distinguishing between uppercase D and lowercase d will be important in this lecture.

²At [34:00] she address this. Yes, it is basically DFS, but DFS will *stop* as soon as it hits a leaf (a solution). Also, DFS implicitly defines all costs as zero.

She starts coding for BS at [27:28]. Note that she is able to reuse her code from the previous part, where she just defined the *model*. Now, she implements BS which is the *inference*. Below is the BS *algorithm*³:

```
def backtracking_search(s, path):
    if is_end(s):
        minimum_cost_path = min(minimum_cost_path, path)

    for a in actions(s):
        path.append([succ(s, a), cost(s, a)])
        backtracking_search(Succ(s, a), path)

    return minimum_cost_path
```

Breadth-first Search [37:00]. Requires that cost for all actions is the same constant $c \geq 0$. Explores level-by-level. Like DFS, terminates as soon as it finds a solution⁴. Important: let d (small d) denote the depth of the *solution*.

- Memory: $\mathcal{O}(b^d)$ (a lot worse!).
- Time: $\mathcal{O}(b^d)$ (better, depends on d , not D).

Brief overview of **DFS with iterative deepening** at [41:00]. Hybrid of DFS with BFS. Call DFS at max depths 1, 2, ... etc. $\mathcal{O}(d)$ space and $\mathcal{O}(b^d)$ time. Nice summary slide of tree search algorithms at [43:11].

Dynamic programming [43:30]. Observation: the cost of going from s to end state can be decomposed into cost of $(s, a) \rightarrow s'$ plus the cost of going from s' to end state.

- Denote “future cost” of state s as $F(s)$.

$$F(s) \triangleq \begin{cases} 0 & \text{IsEnd}(s) \\ \min_{a \in \text{Actions}(s)} C(s, a) + F(\text{Succ}(s, a)) & \text{otherwise} \end{cases} \quad (13)$$

- **Key observation:** *we can cache intermediate results.*
- **Assumption:** State graph defined by $\text{Actions}(s)$ and $\text{Succ}(s, a)$ is acyclic.
- State should be sufficient to encode full history, so that $\text{futureCost}(\text{state})$ can be computed.

³Which is slightly different in form than her code.

⁴This is a consequence of the fixed global cost $c \geq 0$.

Route finding [47:38]

Find minimum cost path from city 1 to n , only moving forward. It costs c_{ij} to go from i to j .

Clarification: Yeah, you can go from any i to $i + j$ for $0 \leq j \leq n - i$. **Observation:** future cost only depends on the current city (since we can only move forward). Cache, cache, cache.

What if we add the **constraint** that you can't visit three odd cities in a row. How does this change our definition of a state? Well, now we need to encode history somehow.


- Naive: redefine our state as (prev city, current city). This increases our state space size from n to n^2 .
- More compact: redefine our state as (prevWasOdd, current city). Now our state space is $2n$.

Coding for dynamic programming:

```
def dynamic_programming(problem):
    cache = {} # state -> future_cost(state)
    def future_cost(state):
        if problem.is_end(state):
            return 0
        if state in cache:
            return cache[state]
        result = min(cost + future_cost(new_state) for _, new_state, cost in
                     problem.succ_and_cost(state))
        cache[state] = result
        return result
    return future_cost(problem.start_state(), [])
```

Uniform cost search [1:05:00]. Can handle cycles.

- **Assumption:** $\forall(s, a) : C(s, a) \geq 0$.
- **Observation:** prefixes of optimal path are optimal.
- **Key idea:** Enumerate states in order of increasing past cost.



Goal: Use UCS to get from A to D.

- Maintain 3 sets of states: explored, frontier, and unexplored.
- We start with all states in the unexplored sets.

- Pop A into the frontier. What states can we reach from A?
- Pop B and C into frontier and the cost to get to them. A is now in the explored set.

It seems we don't explore states we've already been to? UCS is *correct*: when any state s is moved from frontier to explored, its priority of $PastCost(s)$ is the minimum cost to s (discussion around [1:18:00]).

Coding for UCS begins at [1:15:00].

```
def uniform_cost_search(problem):
    frontier = PriorityQueue()
    frontier.update(problem.start_state(), 0)

    while True:
        # Move from frontier to explored.
        state, past_cost = frontier.remove_min()
        if problem.is_end(state):
            return past_cost, []
        # Push out onto frontier.
        for a, new_state, cost in problem.succ_and_cost(state):
            frontier.update(new_state, past_cost + cost)
```



Algorithm: uniform cost search [Dijkstra, 1956]

```
Add  $s_{start}$  to frontier (priority queue)
Repeat until frontier is empty:
    Remove  $s$  with smallest priority  $p$  from frontier
    If IsEnd( $s$ ): return solution
    Add  $s$  to explored
    For each action  $a \in \text{Actions}(s)$ :
        Get successor  $s' \leftarrow \text{Succ}(s, a)$ 
        If  $s'$  already in explored: continue
        Update frontier with  $s'$  and priority  $p + \text{Cost}(s, a)$ 
```

Search II

Discussion of **Final Project** up until [15:00].

Review. Key idea: a **state** is a summary of all the past actions sufficient to choose future actions **optimally**. The class paradigm includes modeling, inference, and learning. Today will talk more about learning.

Learning costs [25:00]. What if we don't have access to the costs? Goal: develop a learning algorithm that outputs the costs of actions.

$$\text{Forward problem (Search)} \quad \text{Cost}(s, a) \longrightarrow (a_1, \dots, a_k) \quad (14)$$

$$\text{Inverse problem (Learning)} \quad (a_1, \dots, a_k) \longrightarrow \text{Cost}(s, a) \quad (15)$$

Prediction (inference) problem:

- Inputs (x): search problem without costs.
- Outputs (y): sequence of optimal actions.

So our training data is search problem inputs and the labeled solutions (action sequence) for each. We need to learn that mapping from search problem to optimal action sequence. The idea is that we'll learn the costs by doing this. A simple way of modeling this⁵:

$$\text{Cost}(s, a_i) \triangleq w[a_i] \triangleq w_i \quad (16)$$

Structured Perceptron (simplified) [32:00]

1. For each action a , initialize $w[a] \leftarrow 0$.
2. For each iteration $t = 1, \dots, T$:
 - (a) For each $(x, y) \in \mathcal{D}_{train}$:
 - i. Compute minimum cost path y' given w .
 - ii. $\forall a \in y : w[a] \leftarrow w[a] - 1$.
 - iii. $\forall a \in y' : w[a] \leftarrow w[a] + 1$.

Coding for this starts at [38:00]. It just uses a modified version of `TransportationProblem` (model) that incorporates weights, and uses the `dynamicProgramming` function (inference) unchanged. Then it updates weights as defined above.

⁵Note that w does not care about the state we are taking the action from.

Generalization to features [46:10]. Costs are parameterized by [incorporating] a feature vector:

$$Cost(s, a) = \mathbf{w} \cdot \phi(s, a) \quad (17)$$

$$Cost(y) = \mathbf{w} \cdot \sum_{(s,a) \in y} \phi(s, a) \quad (18)$$

A* search [46:30]. Improving upon UCS:

- **Problem:** UCS orders states by cost from s_{start} to s .
- **Goal:** take into account cost from s to s_{end} .


Since we obviously don't know $FutureCost(s)$, A* will explore in order $PastCost(s) + h(s)$, using a **heuristic function** that estimates future cost.

$$\underbrace{Cost'(s, a)}_{A^*} = \underbrace{Cost(s, a)}_{UCS} + h(Succ(s, a)) - h(s) \quad (19)$$

- **Consistency of h .** A heuristic h is **consistent** if both of the following hold:

$$Cost'(s, a) \geq 0 \quad (20)$$

$$h(s_{end}) = 0 \quad (21)$$



Proposition: Correctness of A* [56:40]

If h is consistent, A returns the minimum cost path⁶.*

- **Admissibility of h [1:02:20]**. A heuristic h is admissible if

$$h(s) \leq FutureCost(s) \quad (22)$$

Intuition: admissible heuristics are *optimistic*. Theorem: $isConsistent(h) \Rightarrow isAdmissible(h)$.

⁶Proof hint: show that A* ultimately ends up returning UCS + a constant after you sum everything up.

Efficiency.


- UCS efficiency: UCS explores all states s satisfying

$$PastCost(s) \leq PastCost(s_{end})$$

- **A* efficiency: A* explores all states s satisfying**

$$PastCost(s) \leq PastCost(s_{end}) - h(s)$$

⁷



- If $h(s) = 0$, then A* is same as UCS.
- If $h(s) = \text{FutureCost}(s)$, then A* only explores nodes on a minimum cost path.
- Usually $h(s)$ is somewhere in between.

Key Idea: **distortion** [1:02:00]

A distorts edge costs to favor end states.*

Relaxation [1:04:00]. How we find heuristics. Basically, removing constraints.

⁷IDGI. Plz explain here when smarter.

Markov Decision Processes

Table of Contents Local

Written by Brandon McKinzie

Motivation. (Building off search problems) There is **uncertainty** in the world. Volcano crossing example starts at [7:20].

Dice Game [13:12]

For each round $r = 1, 2, \dots$

- You choose to **stay** or **quit**.
- If **quit**, you get \$10 and end the game.
- If **stay**, get \$4 and then roll a 6-sided dice.
 - If result is 1 or 2 \rightarrow end game.
 - Otherwise, continue to next round.

Train of thought should roughly go as follows:

- If I follow policy **stay**, I can compute my expected reward as

$$\frac{1}{3}(4) + \frac{2}{3}\frac{1}{3}(8) + \frac{2}{3}\frac{2}{3}\frac{1}{3}(12) + \dots = 12 \quad (23)$$

where the $\frac{1}{3}$ represent the probability that my roll ends the game.

Using MDP lingo, we revisit this at [44:30]. We have States = $\{in, end\}$ and Actions = $\{stay, quit\}$. Our policy is the simple $\pi(in) := stay$. We can compute our expected reward (assuming $\gamma = 1$ discount) as:

$$V_{\pi}(in) = \frac{1}{3}(4 + V_{\pi}(end)) + \frac{2}{3}(4 + V_{\pi}(in)) = 12 \quad (24)$$

Markov Decision Processes. We can represent an MDP as a graph containing both state nodes and *chance* nodes. Edges from state nodes to chance nodes represent an action. Edges from chance nodes to state nodes represent the random outcomes of the given action.

**Definition: Markov decision process**

States: the set of states

$s_{start} \in \text{States}$: starting state

Actions(s): possible actions from state s

$T(s, a, s')$: probability of s' if take action a in state s

Reward(s, a, s'): reward for the transition (s, a, s')

IsEnd(s): whether at end of game

$0 \leq \gamma \leq 1$: discount factor (default: 1)

- $\forall(s, a) : \sum_{s' \in \text{States}} T(s, a, s') = 1$
- Successors are defined as all s' s.t. $T(s, a, s') > 0$.

Revisit transportation example (magic tram) at [22:00], with addition that tram fails with

$p = 0.5$. A *solution* to an MDP is a **policy** π : a mapping from each state $s \in \text{States}$ to an action $a \in \text{Actions}(s)$.

Policy Evaluation [30:50]. Following a policy yields a *random path*. Policy terminology:

- **Utility** U is discounted $\sum R(s, a, s')$ on the path (this is a random quantity).
- **Value** $V_\pi(s)$ of policy π from state s is the $\mathbb{E}[U]$ received by following policy π .
- **Q-value** $Q_\pi(s, a)$ is $\mathbb{E}[U]$ of taking action a from state s , and then following policy π .

We can compute these quantities via:

$$V_\pi(s) = \begin{cases} 0 & \text{if IsEnd}(s) \\ Q_\pi(s, \pi(s)) & \text{otherwise} \end{cases} \quad (25)$$

$$Q_\pi(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_\pi(s')] \quad (26)$$

Policy Evaluation [45:00]

Goal: automatically compute value of any policy π .

1. Initialize $V_\pi^{(0)}(s) \leftarrow 0$ for all states s .
2. For each iteration $t = 1, \dots, t_{PE}$ and for each state s :

$$V_\pi^{(t)}(s) \leftarrow \underbrace{\sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_\pi^{(t-1)}(s')]}_{Q^{(t-1)}(s, \pi(s))} \quad (27)$$

where it's important to note that we use the values $V_\pi^{(t-1)}(s')$ from the previous iteration when computing the values $V_\pi^{(t)}(s)$ for the current iteration. We are iteratively updating our estimates of V . We want to choose t_{PE} such that

$$\max_{s \in \text{States}} |V_\pi^{(t)}(s) - V_\pi^{(t-1)}(s)| \leq \epsilon \quad (28)$$

Complexity of policy evaluation:

- Time: $\mathcal{O}(t_{PE}SS')$

Value iteration [51:05]. Given an MDP, how do we find a good policy π ?

Value Iteration [57:00]

1. Initialize $V_\pi^{(0)}(s) \leftarrow 0$ for all states s .
2. For each iteration $t = 1, \dots, t_{VI}$ and for each state s :

$$V_{opt}^{(t)}(s) \leftarrow \max_{a \in \text{Actions}(s)} \underbrace{\sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_{opt}^{(t-1)}(s')]}_{Q_{opt}^{(t-1)}(s, a)} \quad (29)$$

We want to choose t_{PE} such that

$$\max_{s \in \text{States}} |V_\pi^{(t)}(s) - V_\pi^{(t-1)}(s)| \leq \epsilon \quad (30)$$

Complexity of value iteration:

- Time: $\mathcal{O}(t_{VISAS'})$

Coding of VI starts around [1:00:00]. Convergence discussion around [1:08:00].

Reinforcement Learning

Table of Contents Local

Written by Brandon McKinzie

MDPs	RL
Mental model of world	Don't know how world works
Find policy to collect max R	Perform actions in world to find out and collect R

Monte Carlo Methods [18:53]. Our data consists of paths of the form $s_0; a_1, r_1, s_1; a_2, r_2, s_2; \dots$, which in RL we also call an **episode**⁸.

Model-Based [Monte Carlo] Learning

Estimate the MDP parameters $T(s, a, s')$ and $R(s, a, s')$.

$$\hat{T}(s, a, s') = \frac{\text{Count}(s \rightarrow a \rightarrow s')}{\text{Count}(s \rightarrow a)} \quad (31)$$

$$\hat{R}(s, a, s') = r \text{ in } (s, a, r, s') \quad (32)$$

We can use these parameter estimates to compute an estimate for $\hat{Q}_{opt}(s, a)$.

Problem: might not see all of the state/action space. Solution: Need π to *explore* explicitly. Next, consider that all we really need for prediction is (estimate of) $Q_{opt}(s, a)$:

$$\hat{Q}_{opt}(s, a) = \sum_{s'} \hat{T}(s, a, s') \left[\hat{R}(s, a, s') + \gamma \hat{V}_{opt}(s') \right] \quad (33)$$

Model-Free [Monte Carlo] Learning

Estimate $Q_{opt}(s, a)$ directly (without learning parameters). Recall that we define the utility at time t as

$$u_t \triangleq r_t + \gamma \cdot r_{t+1} + \gamma^2 \cdot r_{t+2} + \dots \quad (34)$$

Collect all episodes in our data for which there exists transition⁹ $s_{t-1}=s \rightarrow a_t=a$. Then¹⁰ $\hat{Q}_\pi(s, a) := \text{average}(\{u_t\})$ for those episodes.

Model-Free Monte Carlo is an example of an **on-policy** algorithm: estimating the value of data-generating policy. There are also **off-policy** algorithms, ones that estimate the value of a different policy than seen in the data. On-Policy vs Off-Policy

⁸Defined as beginning with a start state, and ending with an end state.

⁹And s, a doesn't occur in s_0, \dots, s_{t-2} (i.e. only care about first occurrence).

¹⁰Assume for now that we've decided on using some policy π .

Model-Free Equivalences [39:40].

- **Original:** $\hat{Q}_\pi(s, a) := \text{average}(\{u_t\})$
- **Convex combination:** At each occurrence of u_t , we assign $\hat{Q}_\pi(s, a)$ to the convex combination of itself (our current estimate) and u_t :

$$\hat{Q}_\pi(s, a) \leftarrow (1 - \eta)\hat{Q}_\pi(s, a) + \eta u \quad (35)$$

$$\text{where } \eta := \frac{1}{1 + \text{NumUpdates}(s, a)} \quad (36)$$

Where "NumUpdates" for a given (s, a) is the number of times we've updated our estimate of $\hat{Q}_\pi(s, a)$ so far¹¹.

- **Stochastic gradient [42:23]:**

$$\hat{Q}_\pi(s, a) \leftarrow \hat{Q}_\pi(s, a) - \eta \left[\underbrace{\hat{Q}_\pi(s, a)}_{\text{prediction}} - \underbrace{u}_{\text{target}} \right] \quad (37)$$

where the implicit objective is least squares regression $(Q - u)^2$.

Bootstrapping methods: SARSA [50:20]

Motivation: the values of u we observe in the data often have high variance^a.

^aThis is because we are defining u as the discounted sum of rewards from whenever (s, a) first happened *until the end of the episode*. A lot of random stuff can happen in that sequence of actions.

On each (s, a, r, s', a') :

$$\hat{Q}_\pi(s, a) \leftarrow (1 - \eta)\hat{Q}_\pi(s, a) + \eta \left[\underbrace{r}_{\text{data}} + \gamma \underbrace{\hat{Q}_\pi(s', a')}_{\text{estimate}} \right] \quad (38)$$

Comparison with model-free MC (colors are unrelated to colors in eq above):

- **biased^a**
- **small variance**
- **immediate updates** (don't have to wait until end of episode).

^aTODO: should probably know how to prove this.

Bootstrapping methods: Q-Learning [59:50]

On each (s, a, r, s') :

$$\hat{Q}_{opt}(s, a) \leftarrow (1 - \eta)\underbrace{\hat{Q}_{opt}(s, a)}_{\text{prediction}} + \eta \underbrace{\left(r + \gamma \hat{V}_{opt}(s') \right)}_{\text{target}} \quad (39)$$

$$\hat{V}_{opt}(s') = \max_{a' \in \text{Actions}(s')} \hat{Q}_{opt}(s', a') \quad (40)$$

¹¹So e.g. if our data only contains u corresponding to the same (s, a) , then η_i (for iteration i) would be $\frac{1}{1+i}$ (we start at $i = 0$).

Covering the unknown. Note that we still have not talked about how to decide what actions to take. There are a few approaches:

$$\text{[Exploit Only]} \quad \pi_{exploit}(s) := \arg \max_{a \in \text{Actions}(s)} \hat{Q}_{opt}(s, a) \quad (41)$$

$$\text{[Explore Only]} \quad \pi_{explore}(s) := \text{GetRandom}(\text{Actions}(s)) \quad (42)$$

$$\text{[Epsilon-Greedy]} \quad \pi_{eps} := \begin{cases} \pi_{exploit} & \text{probability } 1 - \epsilon \\ \pi_{explore} & \text{probability } \epsilon \end{cases} \quad (43)$$

- Exploit Only: $\hat{Q}(s, a)$ inaccurate, too greedy.
- Explore Only: average utility is low because exploration is not guided. Our Q-values are actually quite good, though.

Another problem is **generalization** [1:07:00], which is particularly relevant when our state space is *large*. Also note that we've not yet taken advantage of similarities between states, we've basically been treating states as individual black boxes. We can't generalize to unseen states/actions.

Q-learning with function approximation

Define features $\phi(s, a)$ and weights \mathbf{w} .

$$\hat{Q}_{opt}(s, a; \mathbf{w}) := \mathbf{w} \cdot \phi(s, a) \quad (44)$$

On each (s, a, r, s') :

$$\mathbf{w} \leftarrow \mathbf{w} - \eta \left[\underbrace{\hat{Q}_{opt}(s, a)}_{\text{prediction}} - \underbrace{(r + \gamma \hat{V}_{opt}(s'))}_{\text{target}} \right] \phi(s, a) \quad (45)$$

with implied objective $(\hat{Q} - (r + \gamma \hat{V}))^2$.

Games I

Code for the $\lfloor \frac{N}{2} \rfloor$ game at [13:00].

Two-Player Zero-Sum Game

- ▶ s_{start} : starting state
- ▶ $Actions(s)$
- ▶ $Succ(s, a) \rightarrow s'$
- ▶ $IsEnd(s)$
- ▶ $Utility(s)$: agent's utility for **end state**¹² s
- ▶ $Player(s)$: player who controls state s

Policies of Player p

- ▶ **Deterministic**: $\pi_p(s) \in Actions(s)$
- ▶ **Stochastic**: $\pi_p(s, a) \in [0, 1]$

Evaluation [22:00]. Let $V_{eval}(s)$ denote the expected value of state s . We also say that it's a recurrence for expected utility.

$$V_{eval}(s) = \begin{cases} Utility(s) & \text{if } IsEnd(s) \\ \sum_{a \in Actions(s)} \pi_{Player(s)}(s, a) V_{eval}(Succ(s, a)) & \text{otherwise} \end{cases} \quad (46)$$


Around [26:00], she introduces the idea of V_{expmax} (expectimax¹³) that uses $\max_{a \in Actions(s)}$ in the above equation, which corresponds to a player that always maximizes their expected utility. Note, however, that there is no need to clutter our brains¹⁴ with new definitions, since this is still the same definition as above with a player that has $\pi(s, a)$ equal to 1 for the action that maximizes expected utility, and zero elsewhere.

¹²**TODO**: So all non-terminal states have zero utility? Answer: No, there is no utility for non-terminal states. It's just not a thing.

¹³It seems that we say "expectimax" when we assume our opponent plays ~~randomly~~ deterministically, and "minimax" when we assume our opponent is trying to minimize our utility.

¹⁴We can, however, clutter this footer with them... [29:31]

$$V_{minmax}(s) = \begin{cases} Utility(s) & \text{if } IsEnd(s) \\ \max_{a \in Actions(s)} V_{minmax}(Succ(s, a)) & \text{Player}(s)=agent \\ \min_{a \in Actions(s)} V_{minmax}(Succ(s, a)) & \text{Player}(s)=opp \end{cases} \quad (47)$$



Minimax Property 1 [41:49]. If opponent is minimizing, agent should maximize.

$$V(\pi_{agent}, \pi_{min}) \geq V(\pi_{agent}, \pi_{min}) \quad \forall \pi_{agent} \quad (48)$$

However, if e.g. your opponent plays randomly, then agent should use expectimax (lesson: maximizing not *always* optimal).

Computation.

- **Tree Search:** $\mathcal{O}(d)$ space, $\mathcal{O}(b^{2d})$ time¹⁵.

Evaluation Functions [53:25]. We can use **Depth-limited [tree] search**: stop at some maximum depth d_{max} . This can be defined by making a slight modification to eq. 46:

$$V_{minmax}(s, d) = \begin{cases} \text{Utility}(s) & \text{if IsEnd}(s) \\ \mathbf{Eval}(s) & \mathbf{d = 0} \\ \max_{a \in \text{Actions}(s)} V_{minmax}(\text{Succ}(s, a), d) & \text{Player}(s)=agent \\ \min_{a \in \text{Actions}(s)} V_{minmax}(\text{Succ}(s, a), d-1) & \text{Player}(s)=opp \end{cases} \quad (49)$$

where $\text{Eval}(s)$ is an estimate¹⁶ of $V_{minmax}(s)$. For example, in chess we could define $\text{Eval}(s) := \text{sum}(\text{material}, \text{mobility}, \text{king-safety})$. Pacman code example at [57:28].

¹⁵She defines d really confusingly ([52:00]), but wikipedia says d should be the number of moves a given player makes, while plies (2d here since 2 players) is the total number of moves made.

¹⁶No guarantees on error of estimate, unlike A*.

Alpha-Beta Pruning [59:19]. Key Idea: branch and bound. Maintain lower and upper bounds on values. If intervals don't overlap, we can choose optimally without considering the non-optimal interval choices. Specifically, when exploring tree, we can keep track of...

- a_s : lower bound on value of max node s (we know value is *at least* (\geq) a_s)
- b_s : upper bound on value of min node s (we know value is *at most* (\leq) b_s).
- $\alpha_s \triangleq \max_{s' \preceq s} a_{s'}$, where the max is over ancestors of s .
- $\beta_s \triangleq \min_{s' \preceq s} b_{s'}$.


which clearly only makes sense in a minimax context¹⁷. Detailed walkthrough example at [1:05:00]. The values of a_s and b_s can really be thought of locally, in the sense that they are determined by either a max (a_s) or min (b_s) over the values presented to them by their immediate successors (values are propagated “up” the tree when we are computing these quantities).

Clearly, **ordering** (when traversing/considering states) makes a big difference with pruning. In practice, ordering by $\text{Eval}(s)$ can be a good strategy:

- Max nodes: order successors by decreasing $\text{Eval}(s)$.
- Min nodes: order successors by increasing $\text{Eval}(s)$.

¹⁷**TODO:** think about situations where we aren't doing minimax and want to prune. This was an exam question in CS188.

First, I think it's crucial to understand the game tree in depth. Below is a 2-ply game tree¹⁸, which translates for game problems to 2-player game tree. The leaf nodes show the *utility values as perceived by the root node*. **IMPORTANT:** the only reason it is valid for us to focus on the utility of the root is because this is a *two-player zero-sum game* (the MIN utilities are defined as the negation of the MAX utilities). If this were a multiplayer game, we'd have to track a *vector* of utilities for each player.



Define the **minimax value** $V_{minmax}(s)$ of state s to be the utility¹⁹ of being in state s , assuming both players play optimally from there to the end of the game.

$$V_{minmax}(s) = \begin{cases} \text{Utility}(s) & \text{if IsEnd}(s) \\ \max_{a \in \text{Actions}(s)} V_{minmax}(\text{Succ}(s, a)) & \text{Player}(s) = \text{agent} \\ \min_{a \in \text{Actions}(s)} V_{minmax}(\text{Succ}(s, a)) & \text{Player}(s) = \text{opp} \end{cases} \quad (50)$$

¹⁸This tree has depth 1, with 2 plys.

¹⁹as perceived by root

```

def minimax_decision(root: State) -> Action:
    """We assume the root is a MAX node."""
    best_action = None
    best_value = -float("inf")
    for action in root.actions():
        value = min_value(succ(root, action))
        if value > best_value:
            best_value = value
            best_action = action
    return best_action

def min_value(state: State) -> float:
    """Returns V_minmax from perspective of a MIN node."""
    assert isMinNode(state)
    if isEnd(state):
        return state.utility()
    v = float("inf")
    for action in state.actions():
        v = min(v, max_value(succ(state, action)))
    return v

def max_value(state: State) -> float:
    """Returns V_minmax from perspective of a MAX node."""
    assert isMaxNode(state)
    if isEnd(state):
        return state.utility()
    v = -float("inf")
    for action in state.actions():
        v = max(v, min_value(succ(state, action)))
    return v

```

Alpha-Beta Pruning (5.3). General principle:

- Consider some node n somewhere in the tree, such that *agent* has a choice of moving to that node (i.e. n is an *immediate descendent* of an *agent* (MAX) node²⁰).

Conceptual algorithm for the two-ply tree we've been using.

1. Make your way down the left-most path of the tree until you hit the MIN node just before the leafs.
 - (a) Collect each leaf utility. You can update the interval of possible values that the parent MIN node can have as you go, but it doesn't seem like we really do anything with those intermediate values (**right?**).
 - (b) You now know that the given MIN node will choose the smallest value. Let's call that value m_1 (the first min node value we've determined).
2. The parent of MIN (which is the root MAX node for our tree) sees that MIN has set its value/action choice as m_1 . This means the root's value will be *at least* m_1 .

²⁰This also implies (for a two-player game) that n is either a MIN node or a leaf.

Key points to remember:

- When we speak about “ranges” or intervals, we are talking about the range of possible values $V(s)$ could be, given the information we’ve collected thus far.

Games II

Review: Evaluation Function. New idea: can we *learn* the eval function? We can redefine $\text{Eval}(s) = V(s; \mathbf{w})$. Some possible models might be a linear model, a simple neural network:

$$V(s; \mathbf{w}) = \mathbf{w} \cdot \phi(s) \quad (51)$$

$$V(s; \mathbf{w}, \mathbf{v}_{1:k}) = \sum_{j=1}^k w_j \sigma(\mathbf{v}_j \cdot \phi(s)) \quad (52)$$

Temporal Difference Learning (TD-Learning). In the learning scenario, we treat our estimate of V as the *prediction*, and the value of $r + \gamma V(s')$ as the *label*. We then define an objective and perform gradient descent. If we use MSE objective, then TD-learning is defined as performing, for each (s, a, r, s') [25:00].

$$\mathbf{w} \leftarrow \mathbf{w} - \eta [V(s; \mathbf{w}) - (r + \gamma V(s'; \mathbf{w}))] \nabla_{\mathbf{w}} V(s; \mathbf{w}) \quad (53)$$

**Algorithm: TD learning**

On each (s, a, r, s') :

$$\mathbf{w} \leftarrow \mathbf{w} - \eta [\underbrace{\hat{V}_{\pi}(s; \mathbf{w})}_{\text{prediction}} - \underbrace{(r + \gamma \hat{V}_{\pi}(s'; \mathbf{w}))}_{\text{target}}] \nabla_{\mathbf{w}} \hat{V}_{\pi}(s; \mathbf{w})$$

**Algorithm: Q-learning**

On each (s, a, r, s') :

$$\mathbf{w} \leftarrow \mathbf{w} - \eta [\underbrace{\hat{Q}_{\text{opt}}(s, a; \mathbf{w})}_{\text{prediction}} - \underbrace{(r + \gamma \max_{a' \in \text{Actions}(s)} \hat{Q}_{\text{opt}}(s', a'; \mathbf{w}))}_{\text{target}}] \nabla_{\mathbf{w}} \hat{Q}_{\text{opt}}(s, a; \mathbf{w})$$

Q-Learning	TD Learning
$\hat{Q}_{\text{opt}}(s, a; \mathbf{w})$	$\hat{V}_{\pi}(s; \mathbf{w})$
Off-policy	On-policy
Don't need $T(s, a, s')$	Need $\text{Succ}(s, a)$

Two-Finger Morra [40:00]

Players **A** and **B** each show 1 or 2 fingers.

- If both show 1, **B** gives **A** 2 dollars.
- If both show 2, **B** gives **A** 4 dollars.
- Otherwise, **A** gives **B** 3 dollars

Single-Move Simultaneous Game [42:00]

- $Players = \{A, B\}$
- $V(a, b)$: **A's utility** if A chooses action a , B chooses b . We call V the **payoff matrix**.
- *Evaluation: the value of the game if A follows π_A and B follows π_B :*

$$V(\pi_A, \pi_B) = \sum_{a,b} \pi_A(a) \pi_B(b) V(a, b) \quad (54)$$

where we seem to be using a horrifying overload of notation. In the above, $\pi : a \mapsto [0, 1]$.

Here we'll consider **pure strategies** (a single action) and **mixed strategies** (probability distribution of action). Let's pretend for a moment that the players *do* move sequentially (non-simultaneous). For **pure strategies**, who should go first?²¹ Proposition [52:34]: going second is preferable to going first (pure strategies only).

$$\underbrace{\max_a \min_b V(a, b)}_{\text{A goes 1st}} \leq \underbrace{\min_b \max_a V(a, b)}_{\text{A goes 2nd}} \quad (55)$$

Note that the above is not specific to two-finger Morra. What if A is playing a mixed strategy [and tells B what it is]? Proposition: B should always choose a pure strategy. For any mixed strategy π_A , [56:56]

$$\min_{\pi_B} V(\pi_A, \pi_B) \quad (56)$$

can be attained by a pure strategy.

²¹It is crucial to remember that $V(a, b)$ is *defined* to be the utility from player A 's perspective. In light of this, the proposition should not be surprising at all.

What if A only tells us that they're using some mixed strategy $\pi_A := [p, 1 - p]$ (still on two-finger Morra example)? Let's also pretend that player A is going first. We should enumerate the values for all possible choices of B .

A plays mixed; pretend
 A goes first

$$\text{[B plays 1]} \quad p \cdot (2) + (1 - p) \cdot (-3) = 5p - 3 \quad (57)$$

$$\text{[B plays 2]} \quad p \cdot (-3) + (1 - p) \cdot (4) = -7p + 4 \quad (58)$$

remember that the numerical values $(2, -3, 4)$ above are still from the perspective of A (negative is good for B). B will want to choose action $b \in \text{Actions}(b)$ such that

$$b = \arg \min \{5p - 3, -7p + 4\}$$

Notice that both options are linear functions of p , and that whatever the value of p , it is always the same for both of the linear functions²². If we plot these linear functions, we can find the possible line segments (rays, technically) where this minimum lies²³. Since B is trying to minimize, we (re: player A) know that B will choose the smallest of these two options. Therefore, if player A is going first (and wants to *maximize*), while knowing that B will minimize, it should choose p that results in the largest value *in the minimum region* of our two linear functions – *this turns out to be the point where the functions intersect*. Therefore, we can compute the optimal choice of p for player A by setting both linear equations equal to each other and solving for p .

How does the aforementioned analysis change if we pretend B goes first [1:02:00]? I guess it doesn't change anything.

A plays mixed; pretend
 B goes first

Minimax Theorem [von Neumann, 1928]

For every simultaneous two-player zero-sum game with a finite number of actions:

$$\max_{\pi_A} \min_{\pi_B} V(\pi_A, \pi_B) = \min_{\pi_B} \max_{\pi_A} V(\pi_A, \pi_B) \quad (59)$$

where π_A, π_B range over *mixed strategies*.

²²There is only one p . It is “shared”. Don't overthink it.

²³Red marker area at [1:00:47]

Non Zero-Sum Games [1:06:00]. Difference of utility functions:

- (Zero-Sum) Competitive games: minimax (linear programming)
- Collaborative games: pure maximization (plain search)

Example zero-sum game is *Prisoner's dilemma* [1:07:00]. One difference is that now we need to denote which player's perspective to use with V (before we always were using A 's perspective). Now, we use $V_p(\cdot)$ to denote the value as perceived by player p .

Nash Equilibrium

A **Nash equilibrium** is (π_A^*, π_B^*) such that no player has an incentive to change his/her strategy:

$$V_A(\pi_A^*, \pi_B^*) \geq V_A(\pi_A, \pi_B^*) \quad \forall \pi_A \quad (60)$$

$$V_B(\pi_A^*, \pi_B^*) \geq V_B(\pi_A^*, \pi_B) \quad \forall \pi_B \quad (61)$$

Nash's existence theorem: In any finite-player game with finite number of actions, there exists **at least one** Nash equilibrium.

CSPs I

Map coloring problem starts at [13:00]. Formulate as a search problem where states nodes represent individual countries and edges denote that they are touching (neighbors). A *state* is a partial assignment of color to the set of countries (nodes). Actions are assigning a particular country a color.

Variable-Based Models [20:48]. AKA graphical models (yay). Here, solutions are assignments to [sets of] variables (**modeling**), while decisions about variable ordering, etc., are chosen via **inference**.

Factor Graphs [22:40]. You know all this already. Circles represent variables. They are connected to squares, which denote factors. Scope of factor is set of variables it depends on. Arity of factor is number of variables in its scope. Each assignment of variables has a weight equal to the product of factors over the assignment (weird wording but ok) [33:00].

Constraint Satisfaction Problem [34:43]

A CSP is a factor graph where all factors are **constraints**:

$$f_j(x) \in \{0, 1\} \quad \forall j = 1, \dots, m \quad (62)$$

The constraint is satisfied iff $f_j(x) = 1$.

An assignment x is **consistent** iff $\text{Weight}(x) = 1$ (all constraints satisfied).

Dynamic Ordering [37:40]. Every time we assign a value to a variable, there are additional factors we can evaluate. Formally, let x be some partial assignment, and X_i be a variable we are considering assigning a value to. We denote the **dependent factors** as $D(x, X_i)$: the set of factors depending on X_i and x but not on unassigned variables [42:10].

Backtracking Search [43:00]

Inputs: partial assignment x , weight w (not sure why we can't compute this ourselves given x), and Domains set of possible values to assign to the unassigned variables.

1. If x is complete assignment \rightarrow update best^a and return
2. **Choose** unassigned variable X_i .
3. **Order** the possible values in Domains _{i} for chosen X_i .
4. For each value v in that order:
 - Compute

$$\delta \leftarrow \prod_{f_j \in D(x, X_i)} f_j(x \cup \{X_i : v\}) \quad (63)$$

- If $\delta = 0$, continue
- Domains' \leftarrow Domains via **lookahead**
- Recurse: Backtrack($x \cup \{X_i : v\}, w \cdot \delta, \text{Domains}'$).

^a**TODO:** how do we define “best”?

Guaranteed to find maximum-weight solution.

Forward Checking (one-step lookahead)[49:53]

- After assigning X_i , eliminate inconsistent values from domains of X_i 's neighbors.
- If any domain becomes empty, don't recurse.
- When unassign X_i restore neighbors' domains [to their most recent state].

Variable Ordering

When deciding which variable to assign next, choose the variable that has fewest consistent values. This is the **most constrained variable** (MCV).

Value Ordering [55:08]

Choose value that gives you maximum flexibility thereafter. This is the **least constrained value** (LCV).

When do these heuristics help?

- MCV useful when at least *some* factors are constraints (indicator functions).
- LCV useful when *all* factors are constraints.

Arc Consistency [1:00:22]

A variable X_i is **arc consistent** w.r.t. X_j if

$$(\forall f : X_i, X_j \in \text{Scope}[f]) (\forall x_i \in \text{Val}(X_i)) (\exists x_j \in \text{Val}(X_j)) : f[x_i, x_j](\cdot) \neq 0 \quad (64)$$

We want an algorithm for *enforcing* arc consistency: remove values from $\text{Val}(X_i)$ to make X_i AC wrt X_j .

AC-3 [1:12:39]

1. Add X_j to set
2. While set is non-empty:
 - (a) Remove some X_k from the set however you want.
 - (b) For all $X_l \in \text{MB}(X_k)^a$:
 - i. EnforceArcConsistency(X_l, X_k)
 - ii. If $\text{Val}(X_l)$ changed, add X_l to the set.

^aRecall that the markov blanket of X_k , denoted $\text{MB}(X_k)$, is the set of neighbors of X_k .

Limitations:

- Only checks *pairwise* constraints.

1.9.1 CONSTRAINT PROPAGATION: INFERENCE IN CSPs (AIMA Ch. 6.2)

AC-3

Inputs: binary CSP with components (X, D, C)

Returns: false if inconsistency found, else true.

AC-3(csp):

```
queue ← map(SCOPEOF, csp.binaryFactors())
while queue is not empty do
  ( $X_i, X_j$ ) ← REMOVEFIRST(queue)
  if REVISE( $csp, X_i, X_j$ ) then
    if  $D[i].empty()$  then
      return false
    end if
    for  $X_k$  in  $X_i.neighbors - \{X_j\}$  do
      queue.push(( $X_k, X_i$ ))
    end for
  end if
end while
return true
```

REVISE(csp, X_i, X_j):

```
revised ← false
for  $xi$  in  $D[i]$  do
  if ISEMPTY(map(CONSTRAINTSATISFIED( $xi, \cdot$ ),  $D[j]$ )) then
     $D[i].remove(xi)$ 
    revised ← true
  end if
end for
return revised
```


Setup:


- Have E events, T time slots, and list A containing tuples (e, t) .
- Each event e must be put in exactly one time slot.
- Each event t can have at most one event.
- Event e only allowed at time slot t if $(e, t) \in A$.

There are 2 ways of formulating this as a CSP. My attempt/formulation:

- Variables: events.
- Values/Domains: valid time slots determined by A .
- Constraints: events cannot have the same value (time slot).

You can also formulate with time slots as variables and values as events (plus None). My formulation is the better of the two, because it doesn't require large N-ary constraints. Yay.

Auxiliary Variables [23:59]. Let's say we are dumb and choose the 2nd formulation (contains N-ary constraints). We'll end up with E total factors f_e (for each event) of the form:




Want to convert N-ary constraints to binary constraints. Can do this with **auxiliary variables** which hold intermediate computation. Let Y_i denote the i th time slot.

$$B_1[0] = 0 \tag{65}$$

$$B_i[1] = \min(B_i[0] + \mathbb{1}[Y_i=e], 2) \tag{66}$$

$$B_T[1] = 1 \tag{67}$$

$$B_i[0] = B_{i-1}[1] \tag{68}$$



CSPs II

Review. Time complexity of backtracking search [10:20].

- Vanilla: $\mathcal{O}(|Domain|^n)$
- Lookahead: forward checking, AC3: $\mathcal{O}(|Domain|^n)$
- Dynamic ordering: MCV, LCV: $\mathcal{O}(|Domain|^n)$

Example of predicting location of object given some noisy readings [12:30].

Greedy Search. Make assignment [of a given variable X_i] that maximizes the weight [over the total resulting partial assignment] [18:34].

Beam Search [19:40]. Maintain a set of K **beams**, which are partial assignments. Out of all possible $b \cdot K$ next assignments, choose the top K , which determines the updated beam.

- $\mathcal{O}(nbK \log(bK))$ time with branching factor $b = |Domain|$ [24:25].
- $K = \infty$ is BFS, $\mathcal{O}(b^n)$.

Local Search [26:00]. While backtracking/beam search *extend partial* assignments, local search will *modify complete* assignments. **Iterated Conditional Modes (ICM)** take a complete assignment and explores how the total weight changes if it changes a given single variable (keeping all others fixed). **Key idea:** only needs to consider dependent factors of the variable you're modifying at any given time.

Iterated Conditional Modes (ICM) [29:42]

1. Initialize x to a random complete assignment.
2. Loop through $i = [1..n]$ until convergence:
 - (a) Compute Weight ($x_v \triangleq x \cup \{X_i : v\}$)
 - (b) Assign $x \leftarrow x_v$ with highest weight.

Some obvious properties at [31:30].


Gibbs Sampling [33:00]. Local search with a stochastic component. Key idea: randomness. Sample an assignment [for the current variable [I think?]] with probability proportional to its weight. ~~Will eventually converge to global optimum (might take age of the universe though)~~ just kidding?

Gibbs Sampling [40:00]


1. Initialize x to a random complete assignment.
2. Loop through $i = [1..n]$ until convergence:
 - (a) Compute Weight ($x_v \triangleq x \cup \{X_i : v\}$)
 - (b) Assign $x \leftarrow x_v$ with probability proportional to its weight.

Conditioning [45:30]. Defines independence of variables in a factor graph. **Conditioning** involves assigning a variable and then removing edges/dependencies given our assignment, and replacing them with new factors with a reduced scope. Formal definition shown at [54:28]. We can use conditioning to find maximum weight assignments [59:20].


Conditioning vs elimination discussion starts at [1:01:33]. Elimination (max) example at [1:04:30].



Conditioning



Elimination



Variable Elimination (Max)

First, define the subroutine `max_eliminate(X_i)`: [1:09:39]

1. Collect all factors $\{f_j : X_i \in \text{Scope}[f_j]\}$.
2. Remove X_i and all the aforementioned factors f_j from the graph.
3. Replace those factors with a single new factor over $X := \text{MB}(X_i)$:

$$f_{\text{new}}(x) = \max_{x_i} \prod_{j=1}^k f_j[x_i](x) \quad (69)$$

Variable Elimination (max):

1. For $i = [1..n]$: `max_eliminate(X_i)`
2. For $i = [n..1]$: $X_i \leftarrow \arg \max_{x_i} f_{\text{new}}^{(i)}(x_i, x)$

Let **max-arity** be the maximum arity of any $f_{new}^{(i)}$. Then max variable elimination is $\mathcal{O}(n \cdot |Domain|^{max-arity+1})$ running time. This motivates needing a good strategy for **variable ordering** [1:15:20]. Intuitively, we want to eliminate variables with the fewest neighbors first, since these result in f_{new} with small scope (low arity).

Treewidth [1:15:38]

The **treewidth** of a factor graph is the maximum arity of any factor created by VE with the best variable ordering.


Structure	Treewidth
Chain	1
Tree	1
Cycle	2
$m \times n$ Grid	$\min(m, n)$

Bayesian Networks I

Consistency of Sub-Bayesian Networks [[~38:00](#)]

- **Marginalization** of a leaf node yields a Bayesian network without the node.
- **Local conditional distributions** (e.g. $p(x \mid a, b, c)$) are the true conditional distributions (e.g. $\mathbb{P}(X = x \mid A = a, B = b, C = c)$).

Probabilistic Programs [[46:18](#)]. Make it easier to write down complex BNs. Write a program to generate an assignment (rather than specifying the probability of an assignment). Consider the following BN²⁴.



A possible **probabilistic program** for this could be:

$$B \sim \text{Bernoulli}(\epsilon) \tag{70}$$

$$E \sim \text{Bernoulli}(\epsilon) \tag{71}$$

$$A = B \vee E \tag{72}$$

Discussion of HMMs around [[52:00](#)].

²⁴For the problem of finding probability of an alarm (A) going off given whether an earthquake (E) or burglary (B) occurred

Probabilistic Inference[56:14]

► *Input:*

‣ **BN** $\mathbb{P}(X_1=x_1, \dots, X_n=x_n)$

‣ **Evidence** $E=e$ where $E \subseteq X$

‣ **Query** $Q \subseteq X$

► *Output:*


$$\mathbb{P}(Q=q \mid E=e) \quad \forall q \in Q \quad (73)$$

General Probabilistic Inference Strategy [1:03:20]

1. Remove (marginalize) variables that are not ancestors of Q or E .
2. Convert BN to factor graph.
3. Condition on $E=e$ (shade nodes + disconnect).
4. Remove (marginalize) nodes disconnected from Q .
5. Run probabilistic inference algorithm (manual, VE, Gibbs, particle filtering).

Bayesian Networks II

Forward-Backward [6:30]. Exact algorithm for HMMs. Consider the following BN we've been using for object tracking.



where $H_t \in \{1, \dots, K\}$ is (true) location of object at time t , and $E_t \in \{1, \dots, K\}$ is (noisy) sensor reading at time t . The BN encodes the distribution

$$\Pr [H=\mathbf{h}, E=\mathbf{e}] = p(h_1) \underbrace{\prod_{t=1}^4 p(e_t \mid h_t)}_{\text{emission}} \underbrace{\prod_{t=2}^4 p(h_t \mid h_{t-1})}_{\text{transition}} \quad (74)$$

Two types of queries:

- **Filtering**: probability of hidden given past (and current) evidence, e.g. $\Pr [H_3 \mid E_1, E_2, E_3]$.
- **Smoothing**: probability of hidden given past and future evidence $\Pr [H_3 \mid E_1, \dots, E_5]$.

Lattice representation [14:00] encodes all possible paths h_1, \dots, h_T . For any given path (and observed evidence e) we can compute $\Pr [h_1, \dots, h_T \mid E=e]$ by multiplying the associated factors along that path.

Forward-Backward Algorithm (HMMs only) [31:58]

Goal: Compute smoothed marginals $\Pr [H_t | E] \propto S_t(h_t)$. Key insight: we can use a DP approach that stores probability sums before/after query nodes in our lattice.

Define/initialize $F(h_0)=1$. Define $\psi_t(h_t, h_{t-1}, e_t) \triangleq \Pr [h_t | h_{t-1}] \Pr [e_t | h_t]$.

$$F(h_t) \triangleq \sum_{h_{t-1}} F(h_{t-1}) \psi_t(h_t, h_{t-1}, e_t) \quad (75)$$

$$B(h_t) \triangleq \sum_{h_{t+1}} B(h_{t+1}) \psi_{t+1}(h_{t+1}, h_t, e_{t+1}) \quad (76)$$

$$S(h_t) \triangleq F(h_t) B(h_t) \quad (77)$$

Note that $F(h_t) = \Pr [h_t, \mathbf{e}_{\langle 1 \dots t \rangle}]$ and $B(h_t) = \Pr [\mathbf{e}_{\langle t+1 \dots T \rangle} | h_t]$.

Algorithm:

1. Compute F_1, \dots, F_T
2. Compute B_T, \dots, B_1
3. Compute S_t for each t and normalize.

Return $\Pr [H_t | E] = \text{Normalized}(S(H_t))$.

Runtime is $\mathcal{O}(TK^2)$

Particle-Based Approximation [34:15]. Use a small set of assignments (**particles**) to represent a large probability distribution. Compute inference queries as a function of our set of particles.

Gibbs Sampling [38:00]

1. Init x to random complete assignment.
2. While notConverged, do: For $i = [1..n]$:
 - (a) Compute weight of $x \cup \{X_i : v\} \forall v$
 - (b) Choose $x \cup \{X_i : v\}$ with probability proportional to its weight. In other words, set $X_i = v$ with probability $\Pr [X_i = v | X_{-i} = x_{-i}]$.

stopped around [42:00]

Bayesian Networks III

Today's lecture is about parameter estimation for BNs.

Supervised Learning. We have \mathcal{D}_{train} consisting of full assignments. We want to learn the local distributions of the BN (the parameters of a BN are probabilities). Everything up to [25:00] is literally just counting values of variables to estimate probabilities. Around [26:00] they go over some trivial examples of parameter sharing.

Maximum Likelihood for BNs [40:02]

Input: training examples \mathcal{D}_{train} of full assignments x .

Output: parameters $\theta = \{p_d\}$ for each of the local distributions defined by the graph.

1. **Count:** For each $x \in \mathcal{D}_{train}$: for each $x_i \in x$: $count_{d_i}(x_i, Pa_{x_i}) + +$.
2. **Normalize:** ($\forall d \in D$) and for all assignments to some parents Pa_{x_i} :

$$p_d(x_i \mid Pa_{x_i}) \propto count_d(x_i, Pa_{x_i}) \quad (78)$$

Laplace Smoothing [47:11]

1. For each local distribution d and partial assignment (x_i, Pa_{x_i}) , add λ to $count_d(x_i, Pa_{x_i})$.
2. Normalize to get probability estimates.

Stated informally: hallucinate λ occurrences of each local assignment.

Maximum Marginal Likelihood [51:00]

Let H be hidden variable(s) and let $E = e$ be observed variable(s). The *maximum marginal likelihood objective* is

$$\max_{\theta} \prod_{e \in \mathcal{D}_{train}} Pr[E=e; \theta] \quad (79)$$

$$= \max_{\theta} \prod_{e \in \mathcal{D}_{train}} \sum_h Pr[H=h, E=e; \theta] \quad (80)$$

Expectation Maximization

(Step 0: initialize the BN params somehow lol)

E-step:


- $\forall h : q(h) := \Pr[H=h \mid E=e; \theta]$
- Create weighted points: (h, e) with weight $q(h)$. (ok so like repeatedly make new points for all values of h and all observations of e)

M-step:

- Compute maximum likelihood (just count and normalize) to get θ . (so weighted counts or??)

EM example [56:50]

We are given the following graph and training data.



$$\mathcal{D}_{train} = \{(\textcolor{red}{?}, 2, 2), (\textcolor{red}{?}, 1, 2)\}$$

$$p(G \mid R_1, R_2) \propto p(G, R_1, R_2)$$

We also are given some initialization of the parameters to start with^a.

E-step:

1. For each (r_1, r_2) in \mathcal{D}_{train} , and [then] for each possible value of G (c or d), we compute $p(g, r_1, r_2)$ (4 values total). We compute these solely by using our existing parameter estimates (not counts or anything like that).
2. Then, for each (r_1, r_2) , we compute $q(g)$ by normalizing:

$$q(g) \triangleq p(g \mid r_1, r_2) = \frac{p(g, r_1, r_2)}{\sum_{g'} p(g', r_1, r_2)} \quad (81)$$

This results in $2 \times 2 = 4$ values, which are our weighted points.

M-step:

1. Update our estimates for p_G : Our maximum likelihood estimates for $p_G(g)$ are computed by (for some given g) summing over weights of our points from the previous step that had $G = g$. Then we normalize as usual.
2. Update our estimates for p_R : This is where things get a little weird. For a given value of $G=g$, we collect all of our weighted points that contained $G = g$. For example, for $G=c$ here, this is $\{(c, r_1=2, r_2=2), (c, r_1=1, r_2=2)\}$ with weights 0.69 and 0.5, respectively. Given that the parameters are shared ($p(r_1 \mid g) = p(r_2 \mid g) = p_R(r \mid g)$), how do we estimate $p_R(r \mid g)$? We basically just flatten things out and say that, e.g., $(c, r_1=1, r_2=2)$, which had weight 0.5, counts as both $(c, r=1)$ and $(c, r=2)$ each with the same weight as originally (0.5). The final result for, e.g. the unnormalized $p_R(r=2 \mid g=c)$ is $0.5 + 0.69 + 0.69$ [1:01:00].

^aThe initialized parameters are the probability tables for $p_G(g)$ and $p_R(r \mid g)$. R_1 and R_2 share the conditional parameters to G .

ok

REVIEW

CONTENTS

2.1	Discrete Math and Probability	44
2.2	Course Synthesis	45

Discrete Math and Probability

Table of Contents Local

Written by Brandon McKinzie

Sequences and Summations.

$$\sum_{k=0}^n ar^k \ (r \neq 0) = \frac{ar^{n+1} - a}{r - 1}, \ r \neq 1 \quad (82)$$

$$\sum_{k=1}^{\infty} kx^{k-1}, \ |x| < 1 = \frac{1}{(1-x)^2} \quad (83)$$

Recurrence Relations.

Bit Strings of Length Five

Find number of bit strings of length n , denoted as a_n , that do NOT have two consecutive 0s.

Suggested thought process:

1. I can probably define a_n as a **recurrence relation**.
2. I can think about this problem in terms of whether the last bit is a 1 or a 0:
 - (a) Last bit is 1: there are a_{n-1} such bit strings satisfying the question.
 - (b) Last bit is a 0: well then the $n-2$ bit can't be a zero, since that would violate the question. Therefore, there are a_{n-2} such bit strings.
3. $\therefore a_n = a_{n-1} + a_{n-2}$.

Themes:

- Splitting the problem into two sets corresponding to 0 and 1 somehow.
- Using the **sum rule**.

Patterns/Themes.

- **Bit Strings.**
 - Think in terms of cases: (1) the last bit is one, and (2) the last bit is zero.
- **Number of times until something happens.**
 - Probably follows a **geometric distribution**:

$$p(X=k) = (1-p)^{k-1}p \quad \text{for } k = 1, 2, 3, \dots \quad (84)$$

$$\mathbb{E}[X] = \frac{1}{p} \quad (85)$$

Course Synthesis

A lot of the course topics have strange overlap and overloaded notation. Let's simplify.

Markov Decision Process

- **Definition:**
 - ▢ $States \triangleq Set[State]$
 - * $s_{start} \in States.$
 - * $IsEnd(s)$
 - ▢ $Actions(s) \mapsto Set[Action]$
 - ▢ $T(s, a, s') \mapsto [0, 1]$
 - ▢ $Reward(s, a, s')$
 - ▢ $0 \leq \gamma \leq 1.$
- **Representation:** graph with state nodes s and chance nodes c .
 - ▢ $s \rightarrow c$ edges are actions.
 - ▢ $c \rightarrow s$ edges are random outcomes.
- **Solution:** a policy π .
- **Evaluation** of a policy π :

$$U(s_1, a_1, s_2, \dots, a_{T-1}, s_T) \triangleq \sum_{t=1}^{T-1} \gamma^{t-1} R(s_t, a_t, s_{t+1}) \quad (86)$$

$$V_{\pi}(s) \triangleq \begin{cases} 0 & \text{if } IsEnd(s) \\ Q_{\pi}(s, \pi(s)) & \text{otherwise} \end{cases} \quad (87)$$

$$Q_{\pi}(s, a) \triangleq \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_{\pi}(s')] \quad (88)$$

Now we need some way of *solving* the MDP: finding the optimal policy π_{opt} that maximizes our expected utility.

Solving MDPs

- **Value Iteration:** converges to correct answer (π_{opt}), provided that either $\gamma < 1$ or MDP graph acyclic.

Value Iteration

1. Initialize $V_{\pi}^{(0)}(s) \leftarrow 0$ for all states s .
2. For each iteration $t = 1, \dots, t_{VI}$ and for each state s :

$$V_{opt}^{(t)}(s) \leftarrow \max_{a \in \text{Actions}(s)} \underbrace{\sum_{s'} T(s, a, s') \left[R(s, a, s') + \gamma V_{opt}^{(t-1)}(s') \right]}_{Q_{opt}^{(t-1)}(s, a)} \quad (89)$$

We want to choose t_{PE} such that

$$\max_{s \in \text{States}} \left| V_{\pi}^{(t)}(s) - V_{\pi}^{(t-1)}(s) \right| \leq \epsilon \quad (90)$$

Q-learning with function approximation

Define features $\phi(s, a)$ and weights \mathbf{w} .

$$\hat{Q}_{opt}(s, a; \mathbf{w}) := \mathbf{w} \cdot \phi(s, a) \quad (91)$$

On each (s, a, r, s') :

$$\mathbf{w} \leftarrow \mathbf{w} - \eta \left[\underbrace{\hat{Q}_{opt}(s, a)}_{\text{prediction}} - \underbrace{\left(r + \gamma \hat{V}_{opt}(s') \right)}_{\text{target}} \right] \phi(s, a) \quad (92)$$

with implied objective $(\hat{Q} - (r + \gamma \hat{V}))^2$.

Key Hints/Facts/etc. These are things that, if seen on an exam, should narrow down the scope of what you're thinking about.

- Randomness \implies not a search problem.
- All N agents are all trying to minimize the same objective \implies search problem.
- Non-zero sum games generally don't have optimal policies, but merely Nash equilibria.

Models. Mega-list of all *models*.

- **Reflex.**
 - Linear predictors. Score $s(x, \mathbf{w}) = \mathbf{w} \cdot \phi(x)$.
 - Linear classifiers. $f_{\mathbf{w}}(x) = \text{sign}(s(x, \mathbf{w}))$.
 - KNN. Large $k \iff$ high bias. Low $k \iff$ high variance.
 - Neural networks.
 - K-means.
- **State-based.**
 - Search problem.
 - MDPs (probabilistic search problem).
 - Game: two-player zero-sum, single-move simultaneous, non-zero-sum.
- **Variable-based.**
 - Factor graphs.
 - CSPs.

Algorithms.

- **Reflex.**
 - SGD.
 - PCA.
- **State-based.**
 - Tree search: backtracking, DFS, BFS, iterative deepening.
 - Graph search: DP (if acyclic), UCS, A*.
 - **Policy evaluation** ($\pi \mapsto V_\pi$), **value iteration** (V_{opt}, π_{opt})
 - Model-based MC (estimate T and R)
 - Model-free MC or SARSA (estimate Q_π).
 - Q-learning (estimate Q_{opt}).
 - Game-related: expectimax (analog of VI), minimax, alpha-beta pruning, TD learning (T and R unknown).
- **Variable-based.**
 - Backtracking search.

LEARNING FROM MISTAKES

CONTENTS

3.1	Homework 1: Foundations	49
-----	-----------------------------------	----

Homework 1: Foundations

Table of Contents Local

*Written by Brandon McKinzie***Problem 1: Optimization and Probability.**

- (a) **Mistake:** I didn't show the 2nd derivative was positive. **Lesson:** Always check 2nd deriv is positive when finding a minimum.
- (a) **Mistake:** my interpretation of the case where w_i not guaranteed positive was incorrect (?) **Lesson:** No idea (TODO: figure out why I was wrong).
- (b) **Mistake:** I got it right, but my approach was way more complicated than theirs. They just pulled the max to the left of the summation and refactored both f and g to maximize over s_1, \dots, s_d and used a more elegant proof by intuition. I basically brute-forced it. **Lesson:** get more comfortable moving around summations with maxes, and understand that when comparing two quantities (like f and g), you'll make it easier on yourself if you get them in the same form.
- (c) **Mistake:** I got it right but, again, their solution was way simpler. Apparently, MDPs are relevant for solving these their way. They just used a simple expectation recursion relation, and I went full-on infinite series. **Lesson:** I need to get way more comfortable with recursive expectation value problems.